



```
fn main() -> i32 {  
    println("Hello, Amiga!")  
    return 0  
}
```

NOVUS

Programmer's Reference Guide

A modern language for the Amiga

Version 0.1 | December 2025

For Amiga 68020+ | AmigaOS 2.0+

Novus Programmer's Reference Guide

Version 0.1 — December 2025

Copyright © 2025, 2026 Barry Walker.
All Rights Reserved.

No part of this publication may be reproduced, distributed, or transmitted in any form or by any means, without the prior written permission of the author, except in the case of brief quotations embodied in critical reviews and certain other noncommercial uses permitted by copyright law.

Novus is an open source project.
Visit <https://github.com/barryw/Novus> for source code and updates.

Amiga is a trademark of Amiga Corporation.
All other trademarks are the property of their respective owners.

Typeset in Latin Modern using L^AT_EX.

Contents

Foreword	xxi
1 Introduction	1
1.1 What is Novus?	1
1.1.1 The Compilation Pipeline	1
1.1.2 Key Characteristics	2
1.1.3 A First Comparison	2
1.2 Why Novus Exists	3
1.2.1 The Limits of C	3
1.2.2 Real Bugs from Real Software	3
1.2.3 How Novus Prevents These Bugs	4
1.2.4 Modern Solutions, Amiga Constraints	5
1.3 Design Philosophy	5
1.3.1 Explicit Over Implicit	5
1.3.2 Predictable Performance	6
1.3.3 Respect the Machine	6
1.3.4 Fail Fast, Fail Clearly	7
1.4 A Taste of Novus	7
1.5 What Novus Doesn't Do	9
1.5.1 No Garbage Collection	9
1.5.2 No Exceptions	9
1.5.3 No Hidden Threads	9
1.5.4 68020 Minimum	9
1.6 Target Audience	9
1.7 How to Read This Guide	10
1.8 Showcase: Copper Rainbow	10
1.9 Language Features at a Glance	12
1.9.1 Type System	12
1.9.2 Memory Safety	12
1.9.3 Error Handling	12
1.9.4 Amiga-Specific	12
1.10 Comparison with Other Languages	13
1.10.1 Novus vs C	13
1.10.2 Novus vs Rust	13
1.10.3 Novus vs AMOS/Blitz BASIC	13
1.11 A Day in the Life: Novus Development	13
1.11.1 Morning: Starting a New Feature	14
1.11.2 Midday: Debugging	14
1.11.3 Afternoon: Refactoring	15
1.11.4 Evening: Testing	15
1.11.5 Key Takeaways	16
1.12 What's Different from Rust?	16
1.12.1 No Lifetimes	16
1.12.2 Mutable References are Marked Differently	16

1.12.3	Simpler Pattern Matching	16
1.12.4	No Closures	17
1.12.5	Smaller Standard Library	17
1.13	Community and Resources	17
1.13.1	Getting Help	17
1.13.2	Contributing	17
1.13.3	Reporting Issues	18
1.13.4	Learning Path	18
1.14	Summary	19
2	Getting Started	21
2.1	Quick Start	21
2.2	Prerequisites	21
2.3	Installation	22
2.3.1	Platform-Specific Instructions	22
2.3.2	Verifying Installation	23
2.3.3	Troubleshooting	23
2.4	Your First Program	23
2.4.1	Line-by-Line Explanation	23
2.4.2	Compiling	24
2.4.3	What Happens During Compilation	24
2.4.4	Running on the Amiga	25
2.5	A More Complete Example	25
2.6	The Compilation Pipeline	26
2.6.1	Stage 1: Parsing and Analysis	26
2.6.2	Stage 2: C Code Generation	26
2.6.3	Stage 3: Object File Compilation	27
2.6.4	Stage 4: Linking	27
2.7	Development Workflow	27
2.7.1	Fast Feedback with novus check	27
2.7.2	Code Formatting	27
2.7.3	Editor Support	27
2.8	Running on Real Hardware	28
2.8.1	Shared Folders (Emulator)	28
2.8.2	Network Transfer	28
2.8.3	CF/SD Card	28
2.8.4	Floppy Disk Images	28
2.9	Project Structure	28
2.9.1	The project.toml File	29
2.9.2	Build Settings Reference	29
2.10	Compiler Commands Reference	29
2.10.1	Common Options	30
2.10.2	Debug vs Release Builds	30
2.11	The Standard Library	30
2.12	Showcase: System Information	31
2.13	Running Tests	32
2.14	Next Steps	32
2.15	Emulator Setup	33
2.15.1	FS-UAE Setup	33
2.15.2	WinUAE Setup (Windows)	33
2.15.3	Amiberry Setup (Raspberry Pi)	34
2.16	Working with the Amiga Shell	34

2.16.1	Basic Navigation	34
2.16.2	Running Programs	34
2.16.3	Checking Exit Codes	35
2.16.4	Useful Shell Variables	35
2.16.5	Environment and Assigns	35
2.17	Common Issues and Solutions	36
2.17.1	Compilation Errors	36
2.17.2	Runtime Errors on Amiga	36
2.17.3	Linking Errors	37
2.17.4	Transfer and Execution Issues	37
2.18	Workspaces for Multi-Project Development	37
2.18.1	Creating a Workspace	37
2.18.2	Workspace Structure	38
2.18.3	Building the Workspace	38
2.18.4	Project Dependencies	38
2.19	Debugging Techniques	39
2.19.1	Print Debugging	39
2.19.2	Assert for Invariants	40
2.19.3	Memory Tracking	40
2.19.4	Using Amiga Debuggers	41
2.20	Build System Integration	41
2.20.1	Using Make with Novus	41
2.20.2	CI/CD Integration	41
2.21	Summary	42
3	Language Basics	43
3.1	Comments	43
3.1.1	Single-Line Comments	43
3.1.2	Multi-Line Comments	43
3.2	Variables and Mutability	43
3.2.1	Mutable Variables with <code>var</code>	44
3.2.2	Immutable Bindings with <code>let</code>	44
3.2.3	Compile-Time Constants with <code>const</code>	44
3.2.4	Comparison	44
3.2.5	Type Annotations	45
3.3	Primitive Types	45
3.3.1	Integer Types	45
3.3.2	Floating-Point Types	45
3.3.3	Fixed-Point Types	45
3.3.4	Boolean Type	47
3.4	Literals	47
3.4.1	Integer Literals	47
3.4.2	Floating-Point Literals	48
3.4.3	String Literals	48
3.4.4	F-Strings (Formatted Strings)	49
3.4.5	Character Literals	49
3.5	Operators	49
3.5.1	Arithmetic Operators	49
3.5.2	Comparison Operators	49
3.5.3	Logical Operators	50
3.5.4	Bitwise Operators	50
3.5.5	Compound Assignment Operators	51

3.5.6	Increment and Decrement Operators	52
3.5.7	Range Operators	52
3.5.8	Operator Precedence	52
3.6	Type Conversions	53
3.6.1	C-Style Casts	53
3.6.2	The <code>as</code> Keyword	54
3.7	References and Pointers	54
3.7.1	Immutable References	54
3.7.2	Mutable References	54
3.7.3	Raw Pointers	54
3.7.4	Dereferencing	55
3.7.5	Pointer Arithmetic	55
3.7.6	Auto-Dereferencing	55
3.7.7	Address-of Operator	56
3.8	Arrays	56
3.8.1	Fixed-Size Arrays	56
3.8.2	Repeat Syntax	56
3.8.3	Indexing	56
3.9	Tuples	56
3.9.1	Tuple Types	56
3.9.2	Tuple Literals	57
3.9.3	Destructuring	57
3.9.4	Returning Multiple Values	57
3.10	Expressions and Statements	57
3.10.1	Expression vs Statement	57
3.10.2	Expression Contexts	58
3.11	Intrinsics	58
3.11.1	Available Intrinsics	58
3.11.2	Common Uses	58
3.12	Best Practices	59
3.12.1	Variable Declarations	59
3.12.2	Type Annotations	59
3.12.3	Numeric Types	59
3.13	Common Pitfalls	59
3.13.1	Integer Overflow	59
3.13.2	Signed vs Unsigned Comparisons	60
3.13.3	Precedence Surprises	60
3.13.4	Reference vs Copy	60
3.14	Showcase: Temperature Converter	60
3.15	Summary	63
4	Types	65
4.1	Type System Overview	65
4.2	Primitive Types	65
4.3	Structs	65
4.3.1	Declaring Structs	66
4.3.2	Struct Visibility	66
4.3.3	Creating Struct Instances	67
4.3.4	Accessing Fields	67
4.3.5	Struct Update Syntax	67
4.3.6	Generic Structs	67
4.3.7	Const Generic Parameters	68

4.3.8	Where Clauses	68
4.4	Implementation Blocks	68
4.4.1	Self Parameter Variants	69
4.4.2	Associated Functions	69
4.4.3	Generic Implementations	70
4.4.4	Multiple Impl Blocks	70
4.5	Enums	71
4.5.1	Unit Variants	71
4.5.2	Tuple Variants	71
4.5.3	Generic Enums	71
4.5.4	Enum Limitations	72
4.6	Traits	72
4.6.1	Defining Traits	72
4.6.2	Default Implementations	73
4.6.3	Implementing Traits	73
4.6.4	Built-in Traits	74
4.6.5	Trait Bounds	75
4.7	Pattern Matching	75
4.7.1	Pattern Types	75
4.7.2	Match Expressions	76
4.7.3	If Let Expressions	78
4.8	Option and Result	78
4.8.1	Option<T>	78
4.8.2	Result<T, E>	80
4.8.3	The ? Operator	81
4.9	Generics	82
4.9.1	Generic Functions	82
4.9.2	Generic Structs	82
4.9.3	Const Generic Parameters	83
4.9.4	Where Clauses	83
4.9.5	Turbofish Syntax	84
4.9.6	Generic Constraints in Practice	84
4.10	Type Coercions and Casts	85
4.10.1	Integer Widening	85
4.10.2	Integer Narrowing	85
4.10.3	Reference to Pointer	85
4.10.4	Array to Pointer	86
4.10.5	Explicit Casts	86
4.11	Type Inference	86
4.11.1	Local Variable Inference	86
4.11.2	Function Return Type Inference	86
4.11.3	Generic Type Inference	87
4.12	Summary	87
5	Memory Management	89
5.1	The Amiga Memory Model	89
5.1.1	Memory Types	89
5.1.2	Memory Flag Reference	89
5.1.3	No Size Parameter on Free	89
5.2	Ownership and Move Semantics	90
5.2.1	Move by Default	90
5.2.2	Borrowing to Avoid Moves	91

5.2.3	The Borrow Checker	91
5.3	References	91
5.3.1	Immutable References	91
5.3.2	Mutable References	92
5.3.3	Self Parameter Forms	92
5.3.4	References vs Raw Pointers	92
5.3.5	Reference Lifetimes	93
5.4	The Drop Trait	95
5.4.1	Implementing Drop	95
5.4.2	Drop Execution Order	96
5.4.3	Move Prevents Double Drop	96
5.4.4	Drop for Enums	96
5.5	Defer Blocks	97
5.5.1	Basic Defer	97
5.5.2	LIFO Execution Order	97
5.5.3	Return Value Computed Before Defer	97
5.5.4	Defer vs Drop	98
5.6	The Using Statement	98
5.7	RAII: Resource Acquisition Is Initialization	99
5.7.1	Why RAII Matters on the Amiga	99
5.7.2	The Three RAII Patterns	99
5.7.3	Implementing Your Own RAII Types	101
5.7.4	RAII with Error Handling	102
5.7.5	When NOT to Use RAII	103
5.8	Standard Library Memory Types	104
5.8.1	MemoryBlock	104
5.8.2	MemHandle	104
5.8.3	Allocation<T>	105
5.8.4	Box<T>	106
5.8.5	Slice<T>	107
5.9	Copy and Clone Traits	108
5.9.1	The Copy Trait	108
5.9.2	The Clone Trait	108
5.10	Common Patterns	109
5.10.1	Resource Acquisition with Error Handling	109
5.10.2	Passing Arrays to Functions	109
5.10.3	Moving Out of a Container	110
5.11	Memory Safety Guarantees	110
5.11.1	What Is NOT Guaranteed	110
5.12	Summary	111
6	Control Flow	113
6.1	Conditional Statements	113
6.1.1	if and else	113
6.1.2	if as an Expression	114
6.1.3	Short-Circuit Evaluation	114
6.1.4	if let	115
6.1.5	let-else	115
6.2	Loops	116
6.2.1	while Loops	116
6.2.2	forever Loops	117
6.2.3	C-Style for Loops	119

6.2.4	Range-Based for Loops	119
6.2.5	for-in Loops with Collections	120
6.3	Loop Control Statements	120
6.3.1	break and continue	120
6.3.2	Labeled Loops	121
6.3.3	Postfix Conditions	122
6.4	Pattern Matching with match	122
6.4.1	Basic match Expressions	122
6.4.2	Matching on Enums	123
6.4.3	Match with Generic Enums	123
6.4.4	Match with Blocks	124
6.4.5	Match with Guards	124
6.4.6	Matching Integer Literals	125
6.4.7	Exhaustiveness Checking	125
6.5	Error Propagation with the Try Operator	125
6.5.1	Basic Usage	126
6.5.2	Try Operator in Expressions	126
6.5.3	Chaining Operations	126
6.5.4	Combining with defer	127
6.6	The defer Statement	127
6.6.1	Basic Usage	127
6.6.2	Multiple Defers	127
6.6.3	Defer with AmigaOS Resources	128
6.7	The using Statement	128
6.7.1	using vs defer	129
6.8	Common Control Flow Patterns	130
6.8.1	Early Exit Pattern	130
6.8.2	Resource Acquisition Pattern	131
6.8.3	Retry Pattern	131
6.8.4	State Machine Pattern	132
6.8.5	Search Pattern with Multiple Collections	133
6.9	Showcase: Menu System	133
6.10	Summary	137
7	Functions and Methods	139
7.1	Function Basics	139
7.1.1	Defining Functions	139
7.1.2	Implicit Returns	140
7.1.3	Parameters	140
7.1.4	Reference Parameters	140
7.1.5	The <code>consuming</code> Keyword	141
7.1.6	Early Returns	142
7.1.7	Returning Multiple Values	142
7.2	Visibility	143
7.2.1	Public Functions	143
7.2.2	Internal Visibility	143
7.3	Traits	143
7.3.1	Defining Traits	143
7.3.2	Implementing Traits	144
7.3.3	Common Standard Library Traits	144
7.4	Methods and Impl Blocks	145
7.4.1	Defining Methods	145

7.4.2	Self Parameter Types	145
7.4.3	Associated Functions	146
7.4.4	Calling Methods	146
7.5	Generic Functions	147
7.5.1	Type Parameters	147
7.5.2	Trait Bounds	147
7.5.3	Turbofish Syntax	148
7.5.4	Generic Trait Implementations	148
7.6	Function Overloading	148
7.6.1	Overloading Rules	148
7.6.2	Overloading Restrictions	149
7.7	Const Functions	149
7.7.1	Defining Const Functions	149
7.7.2	Using Const Functions	150
7.7.3	Const Function Restrictions	150
7.8	Closures	150
7.8.1	Closure Syntax	150
7.8.2	Capturing Variables	151
7.9	External Functions	151
7.9.1	Extern Declarations	151
7.9.2	Calling External Functions	152
7.9.3	AmigaOS Library Bindings	152
7.9.4	Variadic Functions	152
7.10	Best Practices	152
7.10.1	Parameter Passing Guidelines	153
7.10.2	Method Receiver Guidelines	153
7.10.3	AmigaOS FFI Best Practices	153
7.10.4	Naming Conventions	154
7.11	Common Function Patterns	154
7.11.1	The Builder Pattern	154
7.11.2	The Factory Pattern	155
7.11.3	The Callback Pattern	156
7.11.4	The RAII Guard Pattern	157
7.11.5	The Visitor Pattern	157
7.12	Showcase: Animation System	159
7.13	Summary	164
8	Modules and Project Organization	167
8.1	Module Basics	167
8.1.1	Module Naming	167
8.2	Imports	167
8.2.1	Basic Imports	168
8.2.2	Wildcard Imports	168
8.2.3	Pattern Imports	169
8.2.4	Import Aliases	169
8.2.5	Multiple Imports	169
8.3	Visibility Modifiers	170
8.3.1	Private (Default)	170
8.3.2	Public (pub)	170
8.3.3	Internal (internal)	171
8.3.4	Extern (extern)	171
8.4	Re-exports	172

8.4.1	The <code>pub use</code> Syntax	172
8.4.2	Multiple Re-exports	172
8.5	Constants and Static Variables	172
8.5.1	Constants (<code>const</code>)	172
8.5.2	Static Variables (<code>static</code>)	173
8.5.3	Mutable Statics	173
8.6	Project Structure	174
8.6.1	Single Project	174
8.6.2	Workspaces	174
8.6.3	Project Types	175
8.6.4	Creating Projects	175
8.6.5	Building Projects	175
8.7	Standard Library Organization	176
8.7.1	Core Modules	176
8.7.2	Data Structures	176
8.7.3	Memory Management	176
8.7.4	Strings	176
8.7.5	Error Handling	177
8.7.6	I/O and Networking	177
8.7.7	Concurrency	177
8.7.8	Graphics and UI	177
8.7.9	AmigaOS FFI	177
8.7.10	Other Modules	178
8.8	The Prelude	178
8.9	Best Practices	179
8.9.1	Visibility Guidelines	179
8.9.2	Import Guidelines	179
8.9.3	Project Organization	179
8.9.4	Constants vs Statics	179
8.10	Common Module Patterns	179
8.10.1	The Facade Pattern	179
8.10.2	The Error Module Pattern	180
8.10.3	The Configuration Module Pattern	181
8.10.4	The Plugin Architecture Pattern	182
8.10.5	The Module Index Pattern	183
8.11	Showcase: Modular Game Framework	183
8.12	Summary	189
9	Attributes	191
9.1	Understanding Attributes	191
9.2	Attribute Syntax	191
9.2.1	At-Sign Syntax (Recommended)	191
9.2.2	Bracket Syntax	192
9.2.3	Attribute Arguments	192
9.2.4	Multiple Attributes	192
9.3	Testing Attributes	193
9.3.1	@test	193
9.3.2	@bench / @benchmark	193
9.4	Optimization Attributes	194
9.4.1	@inline	194
9.4.2	@noinline	194
9.5	Layout Attributes	195

9.5.1	@packed	195
9.6	Trait Derivation	195
9.6.1	@derive / #[derive]	195
9.7	Method Chaining	196
9.7.1	@chain	196
9.8	Interoperability Attributes	197
9.8.1	@export	197
9.8.2	@extern_type	197
9.9	Synchronization Attributes	198
9.9.1	@atomic	198
9.9.2	@no_interrupts	199
9.10	Interrupt Handling Attributes	199
9.10.1	@interrupt	199
9.10.2	@interrupt_safe	200
9.11	Diagnostic Attributes	200
9.11.1	@suppress	200
9.12	Program Configuration	201
9.12.1	#[stack_size]	201
9.13	Library Attributes	201
9.13.1	@library	201
9.13.2	@libfunc	202
9.13.3	Library Lifecycle Hooks	202
9.14	Device Attributes	203
9.14.1	@device	203
9.14.2	@devicecmd	203
9.14.3	@beginio / @abortio	203
9.14.4	Device Lifecycle Hooks	204
9.15	Common Attribute Patterns	204
9.15.1	Interrupt-Safe State Updates	204
9.15.2	Builder Pattern with Validation	204
9.16	Showcase: Interrupt-Driven Music Player	205
9.17	Attribute Reference	207
9.18	Unknown Attributes	207
9.19	Future Attributes	208
10	Error Handling	209
10.1	Philosophy: Errors as Values	209
10.1.1	Why Not Exceptions?	209
10.1.2	Why Not errno?	209
10.2	The Result Type	210
10.2.1	Result Methods Reference	211
10.3	The Option Type	211
10.3.1	When to Use Option vs Result	211
10.3.2	Option Methods Reference	212
10.3.3	Converting Between Option and Result	212
10.4	The ? Operator In Depth	212
10.4.1	What ? Actually Does	212
10.4.2	Chaining Multiple ? Calls	213
10.4.3	Using ? in main()	213
10.4.4	Limitations of ?	214
10.5	Standard Error Types	214
10.5.1	DosError	214

10.5.2	ExecError	215
10.5.3	IntuitionError	215
10.5.4	GraphicsError	215
10.6	The NovusError Wrapper	216
10.7	Error Conversion with From	216
10.7.1	Manual Conversion	216
10.7.2	Getting Raw Error Codes	217
10.8	Error Messages	217
10.9	Error Handling Patterns	217
10.9.1	Early Return Pattern	217
10.9.2	Fallback Chains	218
10.9.3	Cleanup on Error with defer	218
10.9.4	Error Context	219
10.9.5	Retry Pattern	219
10.10	Panic: When Errors Are Bugs	220
10.10.1	Panic vs Error	220
10.10.2	assert!() for Debug Checks	220
10.11	Creating Custom Errors	220
10.11.1	Implementing message()	221
10.11.2	Extracting Error Details	221
10.12	Showcase: Config Loader	222
10.13	Summary	224
10.14	Advanced Error Patterns	224
10.14.1	Collecting Multiple Errors	224
10.14.2	Error Recovery Strategies	225
10.14.3	Error Logging Pattern	226
10.14.4	Error Wrapping for Context	227
10.15	Testing Error Conditions	228
10.15.1	Triggering Specific Errors	229
10.15.2	Error Injection with Function Parameters	229
10.15.3	Testing Error Messages	230
10.16	Performance Considerations	230
10.16.1	Result is Zero-Cost	230
10.16.2	The Cost of the ? Operator	231
10.16.3	Avoid Unwrap in Hot Paths	231
10.16.4	When to Use unwrap_or	231
10.17	Real-World Error Handling	232
10.17.1	File Processing Pipeline	232
10.17.2	GUI Error Feedback	233
10.17.3	Multi-Step Initialization	234
10.18	Error Handling Best Practices	235
10.18.1	Do's	235
10.18.2	Don'ts	236
10.18.3	Error Design Guidelines	236
11	Advanced Traits and Generics	237
11.1	Operator Overloading	237
11.1.1	Arithmetic Operator Traits	237
11.1.2	Implementing Arithmetic Operators	238
11.1.3	Fixed-Point Arithmetic Example	239
11.1.4	Common Mistakes	239
11.1.5	Index Operators	240

11.2	Comparison and Equality Traits	242
11.2.1	The Equality Contract	242
11.2.2	PartialOrd — Comparison Operators	242
11.2.3	Ord — Total Ordering	243
11.2.4	When to Use PartialOrd vs Ord	244
11.2.5	Deriving Comparison Traits	244
11.2.6	Common Mistakes	244
11.3	Conversion Traits	245
11.3.1	The Conversion Trait Family	245
11.3.2	From and Into	245
11.3.3	TryFrom and TryInto	245
11.3.4	Designing Error Types for Conversions	246
11.3.5	Choosing the Right Conversion Trait	247
11.4	The Iterator System	247
11.4.1	The Iterator Trait	247
11.4.2	Creating a Custom Iterator	247
11.4.3	Iterating Collections	248
11.4.4	Iterator for a Linked List	249
11.4.5	Performance Considerations	249
11.4.6	Common Mistakes	250
11.5	Numeric Traits	251
11.5.1	Zero and One	251
11.5.2	Writing Generic Numeric Functions	251
11.5.3	Bounded — Min and Max Values	252
11.5.4	Practical Use: Saturating Arithmetic	252
11.5.5	Fixed-Point and Bounded	252
11.6	Reference Traits	253
11.6.1	AsRef — Cheap Reference Conversion	253
11.6.2	Designing Flexible APIs	253
11.6.3	AsMut — Mutable View Conversion	253
11.6.4	Standard Library Usage	254
11.6.5	When Not to Use AsRef	254
11.7	Advanced Generics	254
11.7.1	Multiple Trait Bounds	254
11.7.2	Multiple Type Parameters	254
11.7.3	Where Clauses on Structs	255
11.7.4	Generic Implementations	255
11.7.5	Bounds on Individual Methods	256
11.7.6	Const Generic Parameters	256
11.7.7	Monomorphization	257
11.8	Common Patterns and Idioms	257
11.8.1	The Newtype Pattern	257
11.8.2	Extension Traits	258
11.8.3	Blanket Implementations	259
11.8.4	Marker Traits	259
11.8.5	The Builder Pattern	260
11.8.6	The Handle Pattern	261
11.9	Summary	261

12 Standard Library Overview	263
12.1 Library Organization	263
12.1.1 Module Hierarchy	263
12.1.2 Physical Layout	264
12.2 The Prelude	264
12.2.1 Selective Imports	264
12.3 Import Conventions	265
12.3.1 Recommended Import Order	265
12.3.2 Wildcard Imports	266
12.3.3 Module Aliases	266
12.4 Core Types	266
12.4.1 Option<T>	266
12.4.2 Result<T, E>	267
12.4.3 Core Traits	268
12.5 Error Types	269
12.5.1 Error Type Hierarchy	269
12.5.2 DosError	269
12.5.3 ExecError	269
12.5.4 NovusError	270
12.5.5 Converting Between Error Types	270
12.6 Memory Allocation Conventions	270
12.6.1 Fallible Allocation	270
12.6.2 Memory Type Flags	271
12.6.3 RAII Memory Management	271
12.7 AmigaOS FFI Organization	272
12.7.1 Structure	272
12.7.2 Using FFI Bindings	272
12.7.3 Safe Wrappers	272
12.8 Collections Overview	273
12.8.1 Available Collections	273
12.8.2 Common Patterns	273
12.9 Strings Overview	274
12.9.1 String Types	274
12.9.2 Usage Examples	274
12.10 I/O Overview	275
12.10.1 Console Output	275
12.10.2 File Operations	275
12.11 Design Principles	275
12.11.1 Amiga First	276
12.11.2 Zero-Cost Abstractions	276
12.11.3 Explicit Allocation	276
12.11.4 Safe Defaults, Unsafe Escape Hatch	276
12.12 What's Covered in Part II	276
12.13 Quick Reference	277
12.13.1 Common Imports	277
12.13.2 Error Handling Pattern	277
12.13.3 Memory Allocation Pattern	277
12.14 Summary	278

13 Core Types and Traits	279
13.1 Option<T>	279
13.1.1 When to Use Option	279
13.1.2 Creating Options	279
13.1.3 Option Methods	280
13.1.4 Pattern Matching with Option	280
13.1.5 Combining Options	280
13.2 Result<T, E>	281
13.2.1 When to Use Result	281
13.2.2 Result Methods	281
13.2.3 The ? Operator	281
13.2.4 Error Type Conversion with ?	282
13.2.5 Pattern Matching with Result	282
13.2.6 Converting Between Option and Result	283
13.3 LibraryVersion	283
13.3.1 LibraryVersion Methods	283
13.4 Resource Management Traits	283
13.4.1 Drop — Automatic Cleanup	283
13.4.2 Copy — Implicit Bitwise Copy	284
13.4.3 Clone — Explicit Deep Copy	285
13.5 Type Conversion Traits	286
13.5.1 From<T> — Infallible Conversion	286
13.5.2 Into<T> — Consuming Conversion	286
13.5.3 TryFrom and TryInto — Fallible Conversion	286
13.6 Comparison Traits	287
13.6.1 Eq — Equality	287
13.6.2 Ord — Total Ordering	287
13.6.3 PartialOrd — Comparison Operators	288
13.7 Hash Trait	288
13.8 Operator Traits	289
13.8.1 Arithmetic Operators	289
13.8.2 Index Operators	289
13.9 Iterator Traits	290
13.9.1 Iterable — For-In Loop Support	290
13.9.2 Iterator — Stateful Iteration	290
13.9.3 Collection Building Traits	291
13.10 Numeric Traits	291
13.10.1 Default — Default Values	291
13.10.2 Zero and One — Identity Elements	291
13.10.3 Bounded — Min/Max Values	292
13.10.4 Primitive Trait Implementations Summary	292
13.11 View Conversion Traits	292
13.11.1 AsRef and AsMut	292
13.11.2 Write — Unified Output	293
13.12 Error Trait	293
13.13 Function Traits	293
13.14 Summary	294

14 Collections	295
14.1 Overview	295
14.1.1 Choosing the Right Collection	295
14.2 Vec<T> — Dynamic Array	296
14.2.1 Creating Vectors	296
14.2.2 Core Operations	296
14.2.3 Growth Strategy	296
14.2.4 Iteration	297
14.2.5 Memory and RAII	297
14.3 VecDeque<T> — Double-Ended Queue	298
14.3.1 Ring Buffer Implementation	298
14.3.2 Core Operations	298
14.3.3 Use Cases	298
14.4 SmallVec<T, N> — Small Vector Optimization	299
14.4.1 When to Use SmallVec	299
14.4.2 Core Operations	299
14.5 HashMap<K, V> — Hash Map	300
14.5.1 Implementation Details	300
14.5.2 Core Operations	301
14.5.3 Iteration	301
14.5.4 Key Requirements	302
14.6 HashSet<T> — Hash Set	302
14.6.1 Set Operations	303
14.7 BitSet — Compact Bit Storage	303
14.7.1 Bitwise Operations	304
14.7.2 Core Operations	304
14.8 RingBuffer<T> — Fixed-Size Circular Buffer	304
14.8.1 Power-of-2 Optimization	305
14.8.2 Interrupt Safety	305
14.8.3 Amiga Use Cases	306
14.9 IntrusiveList<T> — AmigaOS-Compatible Lists	306
14.9.1 Key Constraints	307
14.9.2 PriorityList<T>	307
14.10 SlotMap<T> — Generational Arena	308
14.10.1 How Generations Work	308
14.10.2 SlotKey	309
14.10.3 Use Cases	309
14.10.4 SlotMapFixed<T>	309
14.11 ObjectPool<T> — Pool Allocator	310
14.11.1 Interrupt Safety	310
14.11.2 ObjectPoolFixed<T>	311
14.12 Performance Summary	312
15 Strings	313
15.1 Overview	313
15.2 Str — Borrowed String Slice	313
15.2.1 Creating Str Values	314
15.2.2 Basic Operations	314
15.2.3 Slicing	314
15.2.4 Comparison	314
15.2.5 Searching	315
15.2.6 Trimming	316

15.3	String — Owned Heap String	316
15.3.1	Creating Strings	316
15.3.2	Core Operations	316
15.3.3	Memory Management	316
15.3.4	Converting Between Types	317
15.4	Fixed-Size Strings	317
15.4.1	NSString (32 bytes)	317
15.4.2	GadgetString (128 bytes)	318
15.4.3	PathString (256 bytes)	318
15.5	StringBuilder — Efficient Construction	318
15.5.1	Formatting Methods	319
15.6	BStr — BCPL Strings for AmigaDOS	319
15.6.1	BCPL String Format	320
15.7	String Parsing	320
15.7.1	Integer Parsing	320
15.7.2	Boolean Parsing	321
15.7.3	Parse Errors	321
15.8	Splitting and Iteration	321
15.8.1	Splitting by Delimiter	321
15.8.2	Splitting into Lines	322
15.9	String Transformations	322
15.9.1	Case Conversion	322
15.9.2	Search and Replace	322
15.9.3	Joining Strings	323
15.10	AmigaDOS Path Manipulation	323
15.11	Display and Debug Traits	324
15.11.1	Implementing Display	324
15.11.2	Primitive Implementations	324
15.11.3	StackFormatter	325
15.12	Error Types	325
15.12.1	StringError	325
15.12.2	ParseError	325
15.13	C FFI Functions	325
15.14	Trait Implementations	326
15.14.1	Clone	326
15.14.2	Write	326
15.14.3	Ord (Ordering)	327
15.14.4	Iterator	327
15.14.5	AsRef	327
15.15	Performance Summary	328
16	Memory Management	329
16.1	Amiga Memory Architecture	329
16.1.1	Chip RAM vs Fast RAM	329
16.1.2	Memory Flags	330
16.2	Core Memory Types	330
16.3	MemoryBlock — Raw Memory with Size Tracking	330
16.3.1	Key Points	331
16.3.2	The try_alloc Variant	331
16.3.3	Resizing Memory	331
16.3.4	Taking Ownership with into_raw()	331
16.4	Allocation<T> — Type-Safe Bulk Allocation	332

16.4.1	Overflow Protection	333
16.5	Box<T> — Single Heap Values	333
16.6	MemHandle — Simple RAII Wrapper	333
16.7	Slice<T> — Fat Pointers for Safety	334
16.8	Chip RAM Management	335
16.8.1	ChipBuffer — Direct Allocation	335
16.8.2	Checking Memory Type	335
16.8.3	Cache-Safe Copying (68040/68060)	336
16.9	ChipPool — Pooled Allocation	336
16.9.1	Size Classes	336
16.10	ChipCache — LRU Caching	337
16.10.1	Pinning Assets	337
16.10.2	Priority Levels	338
16.10.3	Cache Statistics	338
16.10.4	Advanced Cache Operations	338
16.11	ResourceBank — Proactive Management	339
16.11.1	Bank Features	339
16.11.2	Pinning Banks for DMA Safety	340
16.11.3	Bank State Checking	340
16.11.4	Memory Budgeting	340
16.12	Streaming Large Assets	341
16.12.1	The Streamable Trait	341
16.12.2	ChipStream — High-Level Streaming	341
16.12.3	Buffer Strategy	342
16.13	DMA Cache Management (68040/68060)	342
16.13.1	Critical: WaitBlit() Before Freeing	343
16.14	Common Patterns	343
16.14.1	Pattern 1: Load-Use-Free	343
16.14.2	Pattern 2: defer for Complex Logic	344
16.14.3	Pattern 3: Pre-allocation for Performance	344
16.14.4	Pattern 4: Graceful Degradation	344
16.15	Memory Safety Checklist	345
16.16	Performance Considerations	346
17	Input/Output	347
17.1	Console Output	347
17.1.1	Format Specifiers	347
17.1.2	Array-Based Formatting	348
17.2	Console Input	348
17.3	ANSI Console Control	348
17.3.1	Basic Functions	348
17.3.2	Pen Colors	349
17.3.3	Cursor Movement	350
17.3.4	Screen and Line Clearing	350
17.3.5	The Ansi Builder	350
17.4	File I/O	351
17.4.1	OwnedFileHandle — RAII File Access	351
17.4.2	File Open Modes	351
17.4.3	Reading and Writing	351
17.4.4	Reading Lines	352
17.4.5	Seeking	353
17.4.6	Ownership Transfer	353

17.5	Directory Locks	354
17.6	File System Helpers	354
17.7	Command-Line Arguments	355
17.7.1	Basic Argument Parsing	355
17.7.2	Template Syntax	356
17.7.3	Argument Types	356
17.7.4	ArgParser Builder	357
17.7.5	Argument Errors	357
17.8	Workbench Support	358
17.8.1	Detecting Launch Mode	358
17.8.2	Workbench Startup Message	359
17.8.3	WBArgument Structure	360
17.8.4	Tool Window	360
17.9	Error Handling	361
17.9.1	Common DOS Errors	361
17.10	I/O Patterns	361
17.10.1	Pattern 1: Read Entire File	361
17.10.2	Pattern 2: Write with Backup	362
17.10.3	Pattern 3: Process File Line by Line	362
17.11	I/O Best Practices	363
A	Guide for C Programmers	365
A.1	Philosophy Differences	365
A.2	Variable Declarations	365
A.3	Type System	365
A.3.1	Primitive Types	366
A.3.2	Booleans	366
A.3.3	Structs	366
A.4	Pointers and References	367
A.5	Functions	367
A.5.1	Declaration Syntax	367
A.5.2	Methods	368
A.6	Control Flow	368
A.6.1	If Statements	368
A.6.2	Switch vs Match	369
A.6.3	Loops	369
A.7	Memory Management	370
A.7.1	The goto cleanup Pattern	370
A.7.2	RAII vs Manual Management	370
A.8	Error Handling	370
A.9	Headers and Modules	371
A.10	Visibility	372
A.11	Common Intrinsics	372
A.12	AmigaOS NDK Access	372
A.12.1	FFI Modules	372
A.12.2	Calling Raw NDK Functions	372
A.12.3	When to Use Raw vs Wrappers	373
A.12.4	Example: Raw Intuition Window	373
A.12.5	BCPL Pointers (DOS Library)	374
A.12.6	Comparison: Raw vs Wrapper	374
A.12.7	Hardware Access	375
A.13	Quick Reference Table	375

A.14 Migration Tips	376
A.15 What Stays the Same	376
B Quick Reference	377
B.1 Variables and Constants	377
B.2 Primitive Types	377
B.3 Literals	377
B.4 Operators	378
B.5 Control Flow	378
B.6 Functions	379
B.7 Structs and Enums	380
B.8 Option and Result	381
B.9 Memory and References	381
B.10 RAII Patterns	382
B.11 Arrays and Collections	382
B.12 Modules and Imports	383
B.13 Traits	383
B.14 Common Attributes	383
B.15 Ininsics	383
B.16 Compiler Commands	384
B.17 C to Novus Quick Reference	384
B.18 Escape Sequences	384

Foreword

The Amiga was ahead of its time. In 1985, while other personal computers offered beeps and static screens, the Amiga delivered pre-emptive multitasking, dedicated graphics and audio hardware, and a sophisticated operating system that wouldn't be matched by mainstream competitors for nearly a decade.

For many of us, the Amiga wasn't just a computer—it was where we learned to code, to create, to push hardware to its limits. The demoscene flourished. Games achieved visual and audio fidelity that seemed impossible. A generation of programmers cut their teeth on 68000 assembly, crafting tight loops that squeezed every cycle from the machine.

Then the platform faded from the mainstream. But it never truly died.

Today, a vibrant community keeps the Amiga alive. Original hardware is lovingly maintained and expanded. FPGA recreations achieve cycle-perfect accuracy. Emulation lets new generations experience what made these machines special. And remarkably, people still write software for them—not out of obligation, but out of genuine passion.

Yet the tools available to modern Amiga developers are largely the same ones we used thirty years ago. C compilers like VBCC and GCC produce excellent code, but the languages themselves haven't evolved. The patterns and safety mechanisms that modern systems programming takes for granted—sum types, pattern matching, resource management without garbage collection—remain out of reach.

Novus aims to change that.

This is not an attempt to modernize the Amiga by hiding what makes it special. Quite the opposite. Novus is designed to *embrace* the Amiga's architecture: its custom chips, its message-passing OS, its elegant hardware. The goal is to give developers modern ergonomics and safety while preserving—even enhancing—the direct hardware access that makes Amiga programming rewarding.

Whether you're a veteran who remembers waiting for Workbench to load from floppy, or a newcomer curious about this legendary platform, we hope Novus helps you write code that's safer, clearer, and more enjoyable—without sacrificing an ounce of the control that Amiga development demands.

New code for classic machines. Welcome to Novus.

December 2025

Chapter 1

Introduction

“The best way to predict the future is to invent it.” — Alan Kay

This guide teaches you Novus, a modern systems programming language designed exclusively for the Amiga. Whether you’re building games, demos, productivity software, or exploring what these remarkable machines can do, Novus gives you the tools to write safe, efficient code that fully leverages the platform’s unique capabilities.


1.1 What is Novus?

Novus is a systems programming language designed specifically for the Amiga 68k platform. It transpiles to C99, which VBCC then compiles to native 68020+ machine code. The result is executables that integrate deeply with AmigaOS and run on real hardware and accurate emulators alike.

The name comes from the Latin word for “new”—reflecting the project’s mission to bring modern language design to classic machines. Novus is not a port of an existing language, nor is it a lowest-common-denominator cross-platform tool. Every design decision is made with the Amiga in mind.

1.1.1 The Compilation Pipeline

Understanding how Novus code becomes an Amiga executable helps you debug problems and optimize performance. The pipeline has four stages:



1. **Novus Compiler:** Parses your source, performs type checking, generates intermediate representation (IR)
2. **Code Generation:** Converts IR to clean, readable C99 code
3. **VBCC:** Compiles C to optimized 68k assembly
4. **VASM/VLINK:** Assembles and links into an AmigaOS HUNK executable

You can examine any stage:

```
# Emit IR to see the compiler's internal representation
novus compile myfile.novus -o out --emit-ir

# Emit assembly to see exactly what instructions are generated
novus compile myfile.novus -o out --emit-asm
```

1.1.2 Key Characteristics

Native Performance Novus transpiles to C99, leveraging VBCC’s mature 68k code generator to produce Amiga executables. There is no interpreter, no virtual machine, no runtime overhead beyond what you explicitly request.

Modern Safety The type system catches errors at compile time. `Result` and `Option` types make error handling explicit. Resource cleanup is deterministic via the `Drop` trait and `defer` blocks. You get safety without garbage collection.

Direct Hardware Access Novus doesn’t hide the Amiga’s hardware behind abstractions you can’t escape. Custom chip registers, Exec library calls, copper lists, blitter operations—all are first-class citizens with typed, safe interfaces.

Readable Output The generated C code is clean and readable. When you need to understand what the compiler is doing—and on the Amiga, you often do—the output won’t fight you.

Amiga-Native Async The `async/await` model maps directly to AmigaOS primitives: Exec signals and message ports. No hidden threads, no runtime scheduler—just the OS facilities the Amiga provides.

1.1.3 A First Comparison

Here’s a simple function that opens a file and reads its size. First, the C version:

```
// C: Opening a file and getting its size
LONG get_file_size(const char *path) {
    BPTR lock = Lock(path, ACCESS_READ);
    if (!lock) {
        return -1; // Caller must check return value!
    }

    struct FileInfoBlock *fib = AllocMem(sizeof(struct
        FileInfoBlock),
                                         MEMF_PUBLIC);

    if (!fib) {
        UnLock(lock); // Must remember to clean up!
        return -1;
    }

    LONG size = -1;
    if (Examine(lock, fib)) {
        size = fib->fib_Size;
    }

    FreeMem(fib, sizeof(struct FileInfoBlock)); // Must match
        size!
    UnLock(lock); // Must remember cleanup order!
    return size;
}
```

Now the same function in Novus:

```
// Novus: Opening a file and getting its size
fn get_file_size(path: &str) -> Result<i32, DosError> {
    let lock = Lock::new(path, ACCESS_READ)? // ? propagates
        errors
}
```

```

let fib = FileInfoBlock::alloc()?

lock.examine(&fib)?
return Result::Ok(fib.size())
} // lock and fib automatically cleaned up via Drop

```

The Novus version is shorter, but more importantly:

- The return type `Result<i32, DosError>` makes failure explicit—callers *must* handle it
- The `?` operator propagates errors automatically, no forgotten checks
- Resources are cleaned up automatically when they go out of scope
- No size mismatches possible—`FileInfoBlock` knows its own size

Coming from C

In C, every function that can fail returns a value that might be an error. You must remember to check it. You must remember to clean up. You must match every allocation size.

In Novus, the type system enforces these rules. A `Result` *must* be handled. Resources with `Drop` *will* be cleaned up. The compiler tracks sizes for you.

1.2 Why Novus Exists

The Amiga platform has capable C compilers. VBCC, in particular, generates excellent code and remains actively maintained. So why create a new language?

1.2.1 The Limits of C

C was revolutionary in 1972. It remains useful today precisely because it's minimal—a portable assembler that gets out of your way. But that minimalism has costs:

- **No sum types.** Representing “this value is either an error or a result” requires conventions that the compiler can't enforce.
- **No built-in error handling.** Return codes are easily ignored. Exceptions don't exist. Every function call is an opportunity for silent failure.
- **Manual resource management.** Matching every `AllocMem` with a `FreeMem`, every `OpenLibrary` with a `CloseLibrary`, every `Lock` with an `UnLock`—across all control flow paths, including error cases.
- **Primitive module system.** Header files and `#include` are textual substitution. Name collisions, include order dependencies, and compile time explosion are facts of life.
- **No metaprogramming.** The preprocessor is powerful but crude. Anything beyond simple macros quickly becomes write-only code.

1.2.2 Real Bugs from Real Software

These aren't theoretical concerns. They're the source of real bugs in real Amiga software:

The Forgotten UnLock:

```
BOOL check_file(const char *path) {
    BPTR lock = Lock(path, ACCESS_READ);
    if (!lock) return FALSE;

    struct FileInfoBlock fib;
    if (!Examine(lock, &fib)) {
        return FALSE; // BUG: lock leaked!
    }

    UnLock(lock);
    return TRUE;
}
```

Every early return is an opportunity to forget cleanup. In a large function with multiple resources, the combinations explode.

The Ignored Return Value:

```
void process_data(void) {
    Open("output.dat", MODE_NEWFILE); // BUG: return value
    // ignored!
    Write(file, data, size);           // file is uninitialized
}
```

C happily compiles this. The result: random corruption.

The Mismatched Size:

```
struct MyData *data = AllocMem(sizeof(struct MyData),
    MEMF_PUBLIC);
// ... 500 lines later, after the struct changed ...
FreeMem(data, sizeof(struct OtherData)); // BUG: wrong size!
```

Memory corruption. Guru meditation. Hours of debugging.

1.2.3 How Novus Prevents These Bugs

Each of these bugs is impossible in Novus:

The Forgotten UnLock:

```
fn check_file(path: &str) -> Result<bool, DosError> {
    let lock = Lock::new(path, ACCESS_READ)?

    let fib = FileInfoBlock::alloc()?
    lock.examine(&fib)?

    return Result::Ok(true)
} // lock.drop() called automatically on ALL paths
```

The Ignored Return Value:

```
fn process_data() -> Result<(), DosError> {
    // This won't compile - Result must be handled:
    // let _ = open("output.dat", MODE_NEWFILE) // Error: unused
    // Result

    let file = open("output.dat", MODE_NEWFILE)? // Must use ?
    // or match
}
```

```

    write(&file, data, size)?
    return Result::Ok(())
}

```

The Mismatched Size:

```

fn allocate_data() -> Result<Allocation<MyData>, ExecError> {
    let data = Allocation::<MyData>::alloc(MEMF_PUBLIC)?
    return Result::Ok(data)
} // data.drop() frees with correct size automatically

```

1.2.4 Modern Solutions, Amiga Constraints

Languages like Rust address these problems elegantly. But Rust targets modern systems with megabytes of RAM, gigahertz CPUs, and operating systems that provide virtual memory and threads. Its runtime assumptions don’t map to the Amiga.

Novus takes a different approach: provide modern ergonomics within Amiga constraints.

- `Result<T, E>` and `Option<T>` are zero-cost at runtime—they’re just tagged unions that the compiler understands.
- The `Drop` trait compiles to cleanup code inserted at scope exit. No hidden allocations, no unwinding tables.
- Pattern matching compiles to efficient jump tables or comparison chains as appropriate.
- The module system is compile-time only. No runtime symbol resolution, no dynamic linking overhead.
- Fixed-point math uses inline 68020+ instructions, not floating-point emulation.

The result is a language that *feels* modern but *runs* like hand-tuned assembly.

1.3 Design Philosophy

Novus follows several core principles that inform every language decision.

1.3.1 Explicit Over Implicit

If something allocates memory, the code should say so. If something can fail, the type should reflect it. If a function has side effects, they should be visible at the call site.

```

// Allocation is explicit - you know this allocates
var buffer = MemoryBlock::alloc(1024, MEMF_CHIP)?

// Failure is explicit - Result type, can't ignore
let file = open("data.dat", MODE_OLDFILE)?

// Mutation is explicit at the call site
process(&var player) // Clearly modifying player

```

This doesn’t mean verbose—Novus has type inference, sensible defaults, and concise syntax. But it means no hidden costs, no action at a distance, no “magic” that obscures what the machine is actually doing.

Coming from C

In C, `process(player)` might modify `player` or it might not—you have to read the function signature.

In Novus, `process(&player)` is read-only. `process(&var player)` mutates. The call site tells you.

1.3.2 Predictable Performance

On a 7 MHz 68000, every cycle matters. On a 50 MHz 68060, you have more headroom—but you still can't afford hidden allocations in a tight loop or cache-hostile memory access patterns.

Novus gives you control. The compiler won't secretly box values, won't insert invisible runtime checks in release builds, won't generate code that's impossible to reason about.

Consider this function:

```
fn sum_array(data: &[i32]) -> i32 {
    var total: i32 = 0
    for value in data {
        total = total + value
    }
    return total
}
```

The generated 68k assembly is exactly what you'd expect:

```
sum_array:
    moveq    #0,d0                ; total = 0
    move.l   d2,-(sp)             ; save d2
    move.l   (a0),d2              ; d2 = slice length
    addq.l   #4,a0               ; a0 = data pointer
.loop:
    subq.l   #1,d2               ; decrement counter
    bmi.s    .done              ; done if negative
    add.l    (a0)+,d0            ; total += *data++
    bra.s    .loop
.done:
    move.l   (sp)+,d2            ; restore d2
    rts
```

No hidden function calls. No allocations. Just a tight loop.

1.3.3 Respect the Machine

The Amiga isn't a Unix workstation. It has custom chips that do things no other hardware can do. Its operating system uses message passing and signals, not files and processes. Its memory model is flat and physical.

Novus doesn't pretend otherwise. The standard library wraps AmigaOS idiomatically—`Result` types for operations that can fail, handles for resources that need cleanup—but it doesn't try to be POSIX. Amiga code should look like Amiga code.

```
// Copper lists, typed and safe
var builder = CopperListBuilder::new()
builder.wait(100, 0)?
builder.move_color(0, rgb12(15, 0, 0))? // Set background red
let copper_list = builder.build()?
copper_list.activate()
```



```
// Blitter operations
blitter.blit_fill(&var bitmap, x, y, width, height, color)?

// Message ports for IPC
let port = MsgPort::create("myport")?
let signal = port.signal_bit()
let sigs = Wait(1 << signal)
```

1.3.4 Fail Fast, Fail Clearly

When something goes wrong, you should know immediately and know exactly what happened. Compile-time errors are better than runtime errors. Runtime errors with clear messages are better than silent corruption.

In debug builds, Novus checks array bounds, validates pointers, and panics with meaningful diagnostics:

```
PANIC at src/player.novus:42
  Array index out of bounds: index 10, length 8

Stack trace:
  update_score (src/player.novus:42)
  game_loop (src/main.novus:156)
  main (src/main.novus:12)
```

In release builds, those checks can be removed via `-safety-level`—but the type system still prevents the errors that checks would catch.

1.4 A Taste of Novus

Before diving into the full language, let's look at a complete program. This opens an Intuition window, draws a message, and waits for the user to close it:

```
// hello_window.novus - A complete Novus program
from std::core import Result
from std::ffi::intuition import OpenWindowTags, CloseWindow
from std::ffi::graphics import SetAPen, Move, Text
from std::ffi::exec import Wait
from std::ffi::amiga_consts import (
    WA_Width, WA_Height, WA_Title, WA_Flags, WA_IDCMP,
    WFLG_CLOSEGADGET, WFLG_DRAGBAR, WFLG_ACTIVATE,
    IDCMP_CLOSEWINDOW, TAG_DONE
)

pub fn main() -> i32 {
    // Open a window with the close gadget
    var window: *Window = null

    unsafe {
        window = OpenWindowTags(null,
            WA_Width, 320,
            WA_Height, 100,
            WA_Title, "Hello from Novus!",
            WA_Flags, WFLG_CLOSEGADGET | WFLG_DRAGBAR |
                WFLG_ACTIVATE,
```

```
        WA_IDCMP, IDCMP_CLOSEWINDOW,
        TAG_DONE)
    }

    if window == null {
        return 1 // Window open failed
    }

    // Draw text in the window
    unsafe {
        let rp = window.RPort
        SetAPen(rp, 1) // Color 1 (usually white)
        Move(rp, 20, 50) // Position cursor
        Text(rp, "Welcome to Novus!", 17)
    }

    // Wait for close gadget
    unsafe {
        let signal = 1u32 << window.UserPort.mp_SigBit
        Wait(signal)
    }

    // Clean up
    unsafe {
        CloseWindow(window)
    }

    return 0
}
```

This is deliberately written in “raw” style, calling AmigaOS functions directly with `unsafe` blocks. Later chapters show the safer `stdlib` wrappers, but this demonstrates that Novus gives you full access to the platform.

Coming from C

Notice what’s different from C:

- No `#include` files—`from...import` brings in what you need
- No semicolons required (they’re optional)
- `null` instead of `NULL`
- `unsafe {}` blocks mark where you’re bypassing safety checks
- Return type comes after the function name: `fn main() -> i32`

What’s the same:

- Same AmigaOS functions: `OpenWindowTags`, `Wait`, etc.
- Same structures: `Window`, `RastPort`
- Same constants: `WA_Width`, `IDCMP_CLOSEWINDOW`
- Same calling convention and memory layout

1.5 What Novus Doesn't Do

Honesty about limitations helps you make informed decisions. Novus is not a general-purpose language. It's designed for Amiga development, and it makes trade-offs accordingly.

1.5.1 No Garbage Collection

Novus uses deterministic resource management via RAII (the `Drop` trait). When a value goes out of scope, its destructor runs immediately. This is predictable and has zero overhead—but it means you must think about ownership.

Cyclic data structures require extra care. Reference counting (via `Rc<T>`) is available when needed, but it's explicit, not automatic.

1.5.2 No Exceptions

Error handling uses `Result<T, E>` and the `?` operator, not exceptions. This means:

- No invisible control flow jumps
- No unwinding tables taking up code space
- Errors are always visible in function signatures

But it also means you can't `catch` an error from deep in a call stack without propagating it explicitly through each level.

1.5.3 No Hidden Threads

The Amiga's multitasking is cooperative and explicit. Novus's `async/await` compiles to state machines driven by Exec signals—there's no hidden thread pool, no automatic parallelism.

If you want parallel execution, you create Exec tasks explicitly. Novus gives you tools to do this safely, but it doesn't do it behind your back.

1.5.4 68020 Minimum

Novus targets 68020+ processors. The 68000/68010 lacks instructions that the compiler relies on (32-bit multiply/divide, bitfield operations). If you're targeting a stock A500 or A1000, Novus isn't the right choice—use C with VBCC directly.

This isn't a significant limitation in practice. Most Amiga users have accelerators, and the A1200/A4000 shipped with 68020/68040.

1.6 Target Audience

This guide assumes you're a working programmer. You should be comfortable with:

- Systems programming concepts: pointers, memory allocation, stack vs. heap
- Basic 68k architecture: registers, addressing modes, the supervisor/user split
- AmigaOS fundamentals: libraries, Exec, DOS, the message-passing model

Prior experience with Rust, Swift, or other modern systems languages will make Novus feel familiar, but it's not required. If you've written C for the Amiga, you have all the background you need.

1.7 How to Read This Guide

The guide is organized in layers:

Part I: Language Fundamentals covers the core language—syntax, types, control flow, functions, and modules. Read this first.

Part II: Standard Library documents the `std` packages: collections, strings, memory management, and Amiga OS bindings.

Part III: Advanced Topics covers async programming, inline assembly, FFI with existing C code, and performance optimization.

Appendices provide quick references, grammar summaries, and migration guides from C.

Code examples are designed to be complete and runnable. When a snippet requires context, we provide it. When we omit boilerplate, we say so.

1.8 Showcase: Copper Rainbow

To close this introduction, here’s a program that demonstrates what makes Novus special for Amiga development. It uses the copper coprocessor to create a smooth rainbow gradient across the screen—something that would be impossible on any other platform.

```
// copper_rainbow.novus - A visual showcase
from std::core import Result
from std::graphics::copper import CopperListBuilder,
    CopperListHandle, CopperError
from std::hardware::registers import rgb12
from std::os::timer import delay, Duration
from std::ffi::graphics import WaitTOF

/// Interpolate between two colors
fn lerp_color(c1: u16, c2: u16, t: u16) -> u16 {
    let r1 = (c1 >> 8) & $F
    let r2 = (c2 >> 8) & $F
    let g1 = (c1 >> 4) & $F
    let g2 = (c2 >> 4) & $F
    let b1 = c1 & $F
    let b2 = c2 & $F

    let r = r1 + ((r2 - r1) * t) / 256
    let g = g1 + ((g2 - g1) * t) / 256
    let b = b1 + ((b2 - b1) * t) / 256

    return rgb12((u8)r, (u8)g, (u8)b)
}

/// Get rainbow color for position 0-255
fn rainbow(pos: u8) -> u16 {
    let p = (u16)pos
    let segment = p / 43
    let offset = (p % 43) * 6

    match segment {
        0 => lerp_color(rgb12(15,0,0), rgb12(15,15,0), offset),
        1 => lerp_color(rgb12(15,15,0), rgb12(0,15,0), offset),
    }
}
```

```

        2 => lerp_color(rgb12(0,15,0), rgb12(0,15,15), offset),
        3 => lerp_color(rgb12(0,15,15), rgb12(0,0,15), offset),
        4 => lerp_color(rgb12(0,0,15), rgb12(15,0,15), offset),
        _ => lerp_color(rgb12(15,0,15), rgb12(15,0,0), offset),
    }
}

/// Build copper list with rainbow gradient
fn build_rainbow() -> Result<CopperListHandle, CopperError> {
    var builder = CopperListBuilder::new()

    var line: u16 = 0
    while line < 200 {
        let scanline = 44 + line
        let color_pos = (line * 256) / 200
        let color = rainbow((u8)color_pos)

        builder.wait(scanline, 0)?
        builder.move_color(0, color)?

        line = line + 2
    }

    builder.wait(244, 0)?
    builder.move_color(0, rgb12(0, 0, 0))?

    return builder.build()
}

pub fn main() -> i32 {
    unsafe { WaitTOF() }

    match build_rainbow() {
        Result::Ok(copper_list) => {
            unsafe { copper_list.activate() }
            let _ = delay(Duration::secs(5))
            return 0 // Copper list cleaned up by Drop
        },
        Result::Err(_) => return 1
    }
}

```

This 70-line program:

- Uses `Result` for error handling throughout
- Relies on RAII—the copper list restores the display when dropped
- Talks directly to the copper hardware through safe abstractions
- Compiles to tight, efficient 68k code

Run it on real hardware or in FS-UAE to see a smooth, colorful gradient scroll across your screen. This is what Novus is for: making the Amiga’s unique capabilities accessible with modern safety and ergonomics.

1.9 Language Features at a Glance

Before diving deep, here's a summary of what Novus offers:

1.9.1 Type System

Static Typing Every variable has a type known at compile time. No runtime type checks.

Type Inference The compiler deduces types where possible: `let x = 42` infers `i32`.

Generics Write type-safe code that works with any type: `Vec<T>`, `Result<T, E>`.

Traits Define shared behavior across types, similar to interfaces or typeclasses.

Sum Types Enums can contain data: `Option<T>` is either `Some(value)` or `None`.

Fixed-Point Native `fixed16` and `fixed32` types for efficient math without FPU.

1.9.2 Memory Safety

Ownership Each value has one owner. When the owner goes out of scope, the value is freed.

References Borrow values without taking ownership: `&T` (read) or `&var T` (write).

No Null Pointers (`*T`) can be null; references (`&T`) cannot.

Bounds Checking Array accesses are checked in debug builds.

RAII Resources are released automatically via the `Drop` trait.

1.9.3 Error Handling

Result Type Functions that can fail return `Result<T, E>`.

? Operator Propagate errors with `?`: `let file = open(path)?`

No Exceptions All error paths are explicit—no hidden control flow.

Pattern Matching Handle all cases with `match` expressions.

1.9.4 Amiga-Specific

Direct FFI Call any AmigaOS function with zero overhead.

Async/Await Map to Exec signals and message ports, not threads.

Inline Assembly Embed 68k assembly in your code when needed.

Copper/Blitter First-class support for custom chip programming.

Library/Device Build shared libraries and device drivers.

1.10 Comparison with Other Languages

1.10.1 Novus vs C

C is the traditional choice for Amiga development. Here's how Novus compares:

Aspect	C	Novus
Error handling	Return codes, <code>errno</code>	<code>Result<T, E></code> , <code>?</code> operator
Resource cleanup	Manual	Automatic via <code>Drop</code>
Null safety	None	<code>*T</code> vs <code>&T</code> distinction
Generics	Macros, <code>void*</code>	Real type parameters
Modules	<code>#include</code>	<code>from...import</code>
Compile speed	Fast	Similar (via C backend)
Runtime overhead	Zero	Zero
Tooling	Mature	Growing

When to use C: When targeting 68000 (Novus requires 68020+), when you need existing C libraries without wrappers, or when your team is more comfortable with C.

When to use Novus: When you want modern safety guarantees, cleaner code organization, and are targeting 68020 or later.

1.10.2 Novus vs Rust

Rust is the primary inspiration for Novus's safety features. Key differences:

Aspect	Rust	Novus
Borrow checker	Full lifetime analysis	Simpler scope-based
Target platforms	Many	Amiga only
Runtime size	Can be large	Minimal
Async model	<code>tokio</code> , <code>std futures</code>	Exec signals/ports
Standard library	Large	Amiga-focused
<code>unsafe</code> usage	Rarely needed	Common for FFI

Rust's borrow checker provides stronger guarantees but is more complex. Novus trades some of that safety for simplicity and Amiga-specific integration.

1.10.3 Novus vs AMOS/Blitz BASIC

AMOS and Blitz BASIC made Amiga programming accessible. Novus aims for similar accessibility with more power:

Aspect	BASIC	Novus
Learning curve	Easy	Moderate
Performance	Good	Excellent
Type safety	Limited	Strong
OS access	Limited	Full
Memory management	Automatic (GC)	Explicit (RAII)
Modern features	No	Yes

If you learned programming on AMOS or Blitz, Novus is a natural next step when you need more control.

1.11 A Day in the Life: Novus Development

Here's what developing with Novus looks like in practice:

1.11.1 Morning: Starting a New Feature

You're adding joystick support to your game. First, check the existing code:

```
# See what input-related code exists
novus search "input" --file-pattern "*.novus"
```

Create a new module for joystick input:

```
// src/input/joystick.novus
from std::ffi::exec import custom_chips
from std::ffi::amiga_consts import JOYDAT0, JOYDAT1

pub struct JoystickState {
    left: bool,
    right: bool,
    up: bool,
    down: bool,
    fire: bool,
}

pub fn read_joystick(port: u8) -> JoystickState {
    let data = unsafe {
        if port == 0 {
            custom_chips().joy0dat
        } else {
            custom_chips().joy1dat
        }
    }

    // Decode joystick register bits
    JoystickState {
        left: (data & $0002) != 0,
        right: (data & $0200) != 0,
        up: ((data >> 1) ^ data) & $0001 != 0,
        down: ((data >> 1) ^ data) & $0100 != 0,
        fire: false, // Read from CIA
    }
}
```

1.11.2 Midday: Debugging

Something's wrong—the joystick reads backwards. Add debug output:

```
pub fn read_joystick_debug(port: u8) -> JoystickState {
    let data = unsafe {
        if port == 0 { custom_chips().joy0dat }
        else { custom_chips().joy1dat }
    }

    // Debug: print raw register value
    print("JOY DAT: $")
    print_hex(data)
    println("")

    // ... decoding ...
}
```


Compile and test:

```
novus compile src/main.novus -o build/game
# Transfer to Amiga shared folder
cp build/game ~/amiga_share/
```

Run in the emulator, spot the bug: bit positions were wrong.

1.11.3 Afternoon: Refactoring

The input code works, but it's not reusable. Create an abstraction:

```
pub trait InputDevice {
  fn poll(&var self)
  fn left(&self) -> bool
  fn right(&self) -> bool
  fn up(&self) -> bool
  fn down(&self) -> bool
  fn fire(&self) -> bool
}

impl InputDevice for JoystickState {
  fn poll(&var self) {
    *self = read_joystick(0)
  }
  fn left(&self) -> bool { self.left }
  fn right(&self) -> bool { self.right }
  fn up(&self) -> bool { self.up }
  fn down(&self) -> bool { self.down }
  fn fire(&self) -> bool { self.fire }
}
```

Now the game can accept any input device without code changes.

1.11.4 Evening: Testing

Write unit tests for the input decoding:

```
#[test]
fn test_joystick_decode() {
  // Simulate joystick left
  let state = decode_joydat($0002)
  assert!(state.left, "left should be true")
  assert!(!state.right, "right should be false")
}

#[test]
fn test_joystick_up() {
  // Simulate joystick up (XOR logic)
  let state = decode_joydat($0100)
  assert!(state.up, "up should be true")
}
```

Run the tests:

```
novus test tests/input_test.novus -o build/test_runner
# Transfer and run on Amiga
```

1.11.5 Key Takeaways

This workflow demonstrates:

- **Rapid iteration** — Changes compile quickly, transfer is simple
- **Safe FFI** — Direct hardware access with `unsafe` blocks
- **Abstraction** — Traits let you write reusable code
- **Testing** — Unit tests catch bugs before they reach hardware

1.12 What's Different from Rust?

If you're coming from Rust, Novus will feel familiar but simpler. Here are the key differences:

1.12.1 No Lifetimes

Rust's lifetime system (`'a`, `'static`, etc.) doesn't exist in Novus. References are valid for their lexical scope only:

```
// Rust: explicit lifetime
fn longest<'a>(x: &'a str, y: &'a str) -> &'a str { ... }

// Novus: no lifetimes, simpler rules
fn longest(x: &str, y: &str) -> String {
    // Must return owned value, not reference
    if x.len() > y.len() {
        return x.to_string()
    }
    return y.to_string()
}
```

This means some patterns that work in Rust require copying in Novus. The trade-off is a much simpler mental model.

1.12.2 Mutable References are Marked Differently

```
// Rust
fn modify(x: &mut i32) { *x = 42; }
modify(&mut value);

// Novus
fn modify(x: &var i32) { *x = 42 }
modify(&var value)
```

Novus uses `var` consistently: variables declared with `var` are mutable, references with `&var` are mutable.

1.12.3 Simpler Pattern Matching

Novus pattern matching is powerful but doesn't include all Rust features:

```
// Novus patterns
match value {
    0 => println("zero"),
```

```
1..10 => println("small"),
n if n > 100 => println("large"),
_ => println("other"),
}
```

No slice patterns, no @ bindings, no or-patterns (yet).

1.12.4 No Closures

Novus uses function pointers and explicit context, not closures:

```
// Rust
let add_x = |y| x + y;

// Novus - explicit approach
fn add_with(x: i32, y: i32) -> i32 { x + y }

// Or use a context struct
struct Adder { x: i32 }
impl Adder {
    fn add(&self, y: i32) -> i32 { self.x + y }
}
```

This keeps the runtime simple and code explicit.

1.12.5 Smaller Standard Library

Novus's stdlib is focused on Amiga development. There's no networking, no filesystem abstraction layers, no cross-platform compatibility code. Instead, you get:

- Complete AmigaOS FFI bindings
- Amiga-optimized collections
- Fixed-point math
- Custom chip abstractions

1.13 Community and Resources

Novus is an open-source project. Here's how to get involved:

1.13.1 Getting Help

GitHub Issues Report bugs or request features at the repository

Discord Join the Amiga development community for real-time chat

Forums The English Amiga Board has an active development section

1.13.2 Contributing

Contributions are welcome:

- Bug fixes and improvements to the compiler
- Standard library additions

- Documentation improvements
- Example programs and tutorials
- Testing on different Amiga configurations

1.13.3 Reporting Issues

When reporting bugs, include:

1. A minimal reproducing example
2. Novus version (`novus -version`)
3. Target configuration (CPU, FPU, OS version)
4. Actual vs expected behavior

1.13.4 Learning Path

Depending on your background, here's a suggested reading order:

If you're new to systems programming:

1. Chapters 1–3: Basics and syntax
2. Chapter 5: Memory management (important!)
3. Chapter 10: Error handling
4. Then explore based on your project needs

If you're an experienced C programmer:

1. Skim Chapter 2 for setup
2. Read Chapter 4 (Types) for the type system
3. Read Chapter 5 (Memory) for RAII and references
4. Chapter 10 (Errors) for Result-based handling
5. Appendix: Guide for C Programmers

If you're coming from Rust:

1. Skim this chapter for Amiga-specific differences
2. Read Chapter 5 for the simpler ownership model
3. Jump to Chapter 11 for traits and generics
4. Explore FFI chapters for AmigaOS integration

1.14 Summary

Novus brings modern language design to the Amiga:

- **Safety without overhead** — Result types, RAII, bounds checking
- **Expressiveness** — Generics, traits, pattern matching
- **Amiga integration** — Direct FFI, copper/blitter support, async via Exec
- **Familiar performance** — Compiles to tight 68k code via VBCC

The following chapters teach you the language systematically, from basic syntax through advanced topics. Code examples are complete and runnable. When in doubt, compile and run—that’s the best way to learn.

Let’s begin.

Chapter 2

Getting Started

“The journey of a thousand miles begins with a single step.”
—Lao Tzu

This chapter walks you through installing the Novus compiler, writing your first program, and understanding the compilation pipeline. By the end, you’ll have a working executable running on your Amiga (or emulator).

2.1 Quick Start

If you already have Novus installed, here’s the fastest path to “Hello, Amiga!”:

1. Create `hello.novus`:

```
from std::io::file import println

pub fn main() -> i32 {
    println("Hello, Amiga!")
    return 0
}
```

2. Compile:

```
novus compile hello.novus -o hello
```

3. Copy to your Amiga and run:

```
1.Workbench:> hello
Hello, Amiga!
```

That’s it! The rest of this chapter explains the prerequisites, installation, and toolchain in detail.

2.2 Prerequisites

Before installing Novus, ensure you have:

- **.NET 9.0 or later** — The Novus compiler is written in C# and requires the .NET runtime. Download from <https://dotnet.microsoft.com>
- **VBCC toolchain** — Novus uses VBCC’s assembler (`vasmm68k_mot`) and linker (`vlink`) to produce Amiga executables. The toolchain is vendored in the repository under `vendor/vbcc`, or you can specify a custom installation via the `-vbcc-path` option.

- **An Amiga or emulator** — To run your compiled programs. FS-UAE, WinUAE, or real hardware all work. A 68020+ configuration with AmigaOS 2.0 or later is required.

2.3 Installation

2.3.1 Platform-Specific Instructions

macOS

1. Install .NET SDK using Homebrew:

```
brew install dotnet
```

2. Clone and build Novus:

```
git clone https://github.com/barryw/Novus.git
cd Novus
dotnet build
```

3. Add to your PATH (optional):

```
export PATH="$PATH:$(pwd)/Novus/bin/Debug/net9.0"
```

Linux

1. Install .NET SDK:

```
# Ubuntu/Debian
sudo apt install dotnet-sdk-9.0

# Fedora
sudo dnf install dotnet-sdk-9.0
```

2. Clone and build:

```
git clone https://github.com/barryw/Novus.git
cd Novus
dotnet build
```

Windows

1. Download and install .NET SDK from <https://dotnet.microsoft.com>
2. Clone the repository (using Git Bash, PowerShell, or WSL):

```
git clone https://github.com/barryw/Novus.git
cd Novus
dotnet build
```

WSL Note: If you prefer Linux tools, Windows Subsystem for Linux works well. Install the Linux .NET SDK inside WSL and compile there.

2.3.2 Verifying Installation

After building, verify everything works:

```
# Check the Novus compiler
./Novus/bin/Debug/net9.0/novus --version

# Check VBCC tools (vendored)
./vendor/vbcc/bin/vasm --version 2>&1 | head -1
./vendor/vbcc/bin/vlink --version 2>&1 | head -1
```

You should see version information for each tool.

2.3.3 Troubleshooting

Problem: `dotnet: command not found`

Solution: .NET SDK is not installed or not in PATH. Reinstall and ensure the install location is in your PATH.

Problem: VBCC tools not found

Solution: The vendored VBCC is under `vendor/vbcc`. If using a custom VBCC installation, pass `-vbcc-path /path/to/vbcc` to the compiler.

Problem: Build fails with missing dependencies

Solution: Run `dotnet restore` to fetch NuGet packages.

Coming from C

Setting up a C cross-compiler for Amiga traditionally involves:

1. Downloading VBCC, vasm, and vlink separately
2. Setting VBCC environment variable
3. Creating `vc.config` with include/library paths
4. Finding and installing the AmigaOS NDK
5. Setting up include paths for system headers

Novus bundles everything in the repository—no environment variables or configuration files needed. Just `dotnet build` and compile.

2.4 Your First Program

Create a file named `hello.novus`:

```
from std::io::file import println

pub fn main() -> i32 {
    println("Hello, Amiga!")
    return 0
}
```

2.4.1 Line-by-Line Explanation

Line 1: `from std::io::file import println`

This imports the `println` function from the standard library's I/O module. Unlike some languages where printing is built-in, Novus requires explicit imports. This keeps the language core small and makes dependencies visible.

`std::io::file` is the module path:

- `std` — the standard library
- `io` — the I/O subsystem
- `file` — file and console operations

Line 3: `pub fn main() -> i32`

This declares the program entry point. Every Novus program must have exactly one `main` function that:

- Is marked `pub` (public) so the linker can find it
- Takes no parameters
- Returns `i32` — the exit code passed back to AmigaOS

Line 4: `println("Hello, Amiga!")`

Calls the imported `println` function with a string literal. String literals in Novus are null-terminated C-style strings (`*u8`).

Line 5: `return 0`

Returns the exit code. Zero means success; non-zero indicates an error. AmigaOS and the Shell use this to detect failures.

2.4.2 Compiling

Compile to an Amiga executable:

```
novus compile hello.novus -o hello
```

This produces an executable named `hello` in Amiga HUNK format. Without the `-o` flag, the output defaults to `a.out`.

2.4.3 What Happens During Compilation

When you compile, Novus performs these steps:

1. **Parse** — Read `hello.novus` and build an AST
2. **Analyze** — Check types, resolve names, build IR
3. **Import** — Load `std::io::file` and its dependencies
4. **Generate** — Emit C99 code for your program and used `stdlib` functions
5. **Compile** — VBCC compiles C to 68k object files
6. **Link** — `vlink` combines objects into a HUNK executable

Use `-v` (verbose) to watch this process:

```
novus compile hello.novus -o hello -v
```

2.4.4 Running on the Amiga

Transfer the executable to your Amiga and run it from the Shell:

```
1.Workbench:> hello
Hello, Amiga!
```

Congratulations—you’ve run your first Novus program!

2.5 A More Complete Example

Let’s look at a program that does more than print a message. This example queries and displays system memory information:

```
// memory_info.novus - Display memory statistics
from std::io::file import print, println
from std::ffi::exec import AvailMem
from std::ffi::amiga_consts import MEMF_CHIP, MEMF_FAST,
    MEMF_TOTAL

pub fn main() -> i32 {
    println("Memory Information")
    println("=====")

    // Query chip RAM (DMA-accessible memory)
    let chip_mem = unsafe { AvailMem(MEMF_CHIP | MEMF_TOTAL) }

    // Query fast RAM (CPU-only memory)
    let fast_mem = unsafe { AvailMem(MEMF_FAST | MEMF_TOTAL) }

    // Display results
    print("Chip RAM: ")
    print_kb(chip_mem)
    println(" KB")

    print("Fast RAM: ")
    print_kb(fast_mem)
    println(" KB")

    print("Total: ")
    print_kb(chip_mem + fast_mem)
    println(" KB")

    return 0
}

fn print_kb(bytes: u32) {
    let kb = bytes / 1024u32
    print_number(kb)
}

fn print_number(n: u32) {
    if n == 0 { return }
    if n >= 10 { print_number(n / 10u32) }
    let digit: u8 = (u8)(n % 10u32) + 48u8
    var buf = [0u8; 2]
    buf[0] = digit
}
```

```
    print(&buf[0])  
}
```

This example demonstrates:

- **Multiple imports** from different modules
- **Unsafe FFI calls** to AmigaOS via `AvailMem`
- **Constants** like `MEMF_CHIP` from the FFI layer
- **Helper functions** for formatting output
- **Arithmetic and type casts**

Coming from C


The equivalent C code would look similar, but you'd need to:

- Include `<proto/exec.h>` or set up function prototypes
- Link with `amiga.lib` or use direct library calls
- Deal with library bases (`SysBase`)

In Novus, the FFI layer handles library opening/closing automatically. Just import and call.

2.6 The Compilation Pipeline

Understanding the compilation pipeline helps with debugging and optimization. Here's what happens when you compile:



2.6.1 Stage 1: Parsing and Analysis

The Novus compiler reads your source file, parses it into an Abstract Syntax Tree (AST), then performs semantic analysis:

- Type checking
- Name resolution
- Trait resolution
- Building the Intermediate Representation (IR)

Use `-emit-ir` to see the IR:

```
novus compile hello.novus --emit-ir
```

2.6.2 Stage 2: C Code Generation

The IR is lowered to C99 code. Each function becomes a C function; each type becomes a C struct or typedef. The generated C uses VBCC-specific extensions for 68k calling conventions.

2.6.3 Stage 3: Object File Compilation

VBCC's C compiler (`vbccm68k`) compiles the C code to 68k object files. Each function is compiled separately for incremental builds.

Use `-emit-asm` to see the generated assembly:

```
novus compile hello.novus --emit-asm
```

This is useful for understanding performance characteristics or debugging codegen issues.

2.6.4 Stage 4: Linking

`vlink` combines all object files—your code, the runtime, and `stdlib`—into a final executable in Amiga HUNK format. Dead code elimination removes unused functions to keep binaries small.

2.7 Development Workflow

2.7.1 Fast Feedback with `novus check`

For quick syntax and type checking without full compilation:

```
novus check myfile.novus
```

This skips code generation and linking, giving faster feedback during development.

2.7.2 Code Formatting

Keep your code consistent with `novus fmt`:

```
novus fmt myfile.novus
```

This reformats the file in place according to the Novus style guide.

2.7.3 Editor Support

Novus provides a Language Server Protocol (LSP) implementation for IDE integration:

- **VS Code** — Install the Novus extension from the marketplace
- **Other editors** — Configure your LSP client to use `novus lsp`

Features include:

- Syntax highlighting
- Error diagnostics as you type
- Go to definition
- Hover information
- Code completion

2.8 Running on Real Hardware

The ultimate goal is running on real Amiga hardware. Here are several ways to transfer executables.

2.8.1 Shared Folders (Emulator)

If using FS-UAE, WinUAE, or Amiberry, configure a shared folder:

1. Create a directory on your host system (e.g., `~/amiga_share`)
2. Configure the emulator to mount it as a device (e.g., `DEV:`)
3. Compile directly to the shared folder:

```
novus compile hello.novus -o ~/amiga_share/hello
```

4. Access from the Amiga: `DEV:hello`

2.8.2 Network Transfer

With networked Amigas (via Plipbox, A2065, or similar):

```
# Using scp (requires AmiSSH or similar on the Amiga)
scp hello amiga.local:Work/

# Using ftp
ftp amiga.local
put hello
```

2.8.3 CF/SD Card

For accelerator cards with CF/SD storage:

1. Mount the card on your development machine
2. Copy the executable to the card
3. Unmount and return the card to the Amiga

2.8.4 Floppy Disk Images

Create an ADF image with your executable:

```
# Create a new ADF
xdftool myfiles.adf format "MyDisk" + write hello
```

Then use the ADF in your emulator or write it to a real floppy.

2.9 Project Structure

For anything beyond a single file, create a project with `novus new`:

```
novus new myproject --type cli
cd myproject
```

This creates a standard project structure:

```
myproject/
  project.toml      # Project configuration
  src/
    main.novus      # Entry point
```

2.9.1 The project.toml File

The `project.toml` file defines project metadata and build settings:

```
[package]
name = "myproject"
version = "0.1.0"
authors = ["Your Name"]
description = "My first Novus project"

[build]
target_cpu = "68020"
fpu = "auto"
optimization_level = 0
output = "build"
```

2.9.2 Build Settings Reference

Setting	Default	Description
<code>target_cpu</code>	<code>auto</code>	Target CPU: 68020, 68030, 68040, 68060, or <code>auto</code>
<code>fpu</code>	<code>auto</code>	FPU mode: <code>soft</code> , 68881, 68040, or <code>auto</code>
<code>optimization_level</code>	<code>0</code>	VBCC optimization level (0–3)
<code>output</code>	<code>build</code>	Output directory for compiled files
<code>stack_size</code>	<code>4096</code>	Stack size in bytes

Build the project with:

```
novus build
```

2.10 Compiler Commands Reference

novus compile <file> Compile a single source file to an executable

novus build Build a project defined by `project.toml`

novus check <file> Type-check without generating code

novus test <file> Compile test runner

novus new <name> Create a new project from a template

novus fmt <file> Format source code

novus clean Remove cached build artifacts

novus lsp Start the language server

2.10.1 Common Options

Options for `novus compile`:

- `-o, -output` Output file name (default: `a.out`)
- `-cpu` Target CPU: 68020, 68030, 68040, 68060, or `auto`
- `-fpu` FPU mode: `soft`, 68881, 68040, or `auto`
- `-release` Enable optimizations and disable debug checks
- `-O` Optimization level (0–2)
- `-emit-asm` Emit assembly output only
- `-emit-ir` Emit IR to stdout
- `-v, -verbose` Show detailed progress

2.10.2 Debug vs Release Builds

By default, Novus compiles in debug mode:

- Array bounds checking enabled
- Debug symbols included
- Minimal optimizations

For production, use `-release`:

```
novus compile hello.novus -o hello --release
```

2.11 The Standard Library

Novus includes a comprehensive standard library:

- `std::core` Fundamental types: `Option`, `Result`, core traits
- `std::collections` Data structures: `Vec`, `HashMap`, `HashSet`
- `std::strings` String handling: `String`, `Str`, `StringBuilder`
- `std::memory` Memory management, slices
- `std::math` Fixed-point arithmetic, trigonometry
- `std::io` File and console I/O
- `std::ffi` AmigaOS bindings (`Exec`, `DOS`, `Graphics`, `Intuition`)
- `std::hardware` CPU/chipset detection, memory info

Import with `from`:

```
from std::collections::vec import Vec
from std::io::file import println
```


2.12 Showcase: System Information

This program demonstrates the concepts from this chapter by querying and displaying system information:

```
// sysinfo.novus - System Information Display
from std::core import Option
from std::io::file import print, println
from std::ffi::exec import AvailMem
from std::ffi::amiga_consts import MEMF_FAST, MEMF_CHIP
from std::ffi::amiga_consts import MEMF_TOTAL, MEMF_LARGEST

fn print_memory(label: *u8, bytes: u32) {
  print(" ")
  print(label)
  print(": ")
  let kb: u32 = bytes / 1024u32
  if kb == 0 {
    print("0")
  } else {
    print_number(kb)
  }
  println(" KB")
}

fn print_number(n: u32) {
  if n == 0 { return }
  if n >= 10 { print_number(n / 10u32) }
  let digit: u8 = (u8)(n % 10u32) + 48u8
  var buf = [0u8; 2]
  buf[0] = digit
  print(&buf[0])
}

fn show_memory_info() {
  println("")
  println("MEMORY")
  println("-----")

  let chip_total = unsafe { AvailMem(MEMF_CHIP | MEMF_TOTAL) }
  let chip_largest = unsafe { AvailMem(MEMF_CHIP |
    MEMF_LARGEST) }
  print_memory("Chip RAM Total", chip_total)
  print_memory("Chip RAM Largest Block", chip_largest)

  println("")

  let fast_total = unsafe { AvailMem(MEMF_FAST | MEMF_TOTAL) }
  if fast_total > 0 {
    let fast_largest = unsafe { AvailMem(MEMF_FAST |
      MEMF_LARGEST) }
    print_memory("Fast RAM Total", fast_total)
    print_memory("Fast RAM Largest Block", fast_largest)
  } else {
    println(" Fast RAM: Not installed")
  }
}
```

```
pub fn main() -> i32 {
    println("=====")
    println("          NOVUS SYSTEM INFORMATION")
    println("=====")

    show_memory_info()

    println("")
    println("=====")
    return 0
}
```

This program demonstrates:

- **Imports** from multiple stdlib modules
- **Unsafe FFI calls** to AvailMem
- **Constants** from `amiga_consts`
- **Helper functions** for output formatting
- **Control flow** with conditionals

The full source code is available in `Novus.Tests/Examples/sysinfo_showcase.novus`.

2.13 Running Tests

Novus has built-in test support. Mark test functions with the `@test` attribute:

```
from std::test::test import expect_eq

@test
fn test_addition() {
    expect_eq(2 + 2, 4, "basic addition")
}

@test
fn test_multiplication() {
    expect_eq(3 * 7, 21, "basic multiplication")
}
```

Compile the test runner:

```
novus test mytest.novus
```

Important: The test command builds the test runner but does not execute it. Transfer the executable to your Amiga to run the tests.

2.14 Next Steps

You now have a working Novus installation and understand the basic workflow. The following chapters cover:

- **Chapter 3: Language Basics** — Variables, types, and expressions
- **Chapter 4: Types** — The type system in depth

- **Chapter 5: Memory Management** — Safe and explicit memory handling
- **Chapter 6: Control Flow** — Conditionals, loops, and pattern matching

When you're ready, turn the page and start learning the language itself.

2.15 Emulator Setup

If you don't have real Amiga hardware, emulators provide an excellent development environment. This section covers setup for the most popular options.

2.15.1 FS-UAE Setup

FS-UAE is a cross-platform Amiga emulator with excellent accuracy:

1. Install FS-UAE from <https://fs-uae.net> or your package manager
2. Download Kickstart ROMs (required for legality, obtain from your real Amiga or Amiga Forever)
3. Configure a basic A4000 setup in `~/.fs-uae/Configurations/dev.fs-uae`:

```
[fs-uae]
amiga_model = A4000
kickstart_file = ~/Kickstart/kick40063.A4000.rom
chip_memory = 2048
fast_memory = 8192
jit_compiler = 1
hard_drive_0 = ~/Amiga/System
hard_drive_0_type = dir

# Shared folder for development
hard_drive_1 = ~/amiga_share
hard_drive_1_type = dir
hard_drive_1_label = DEV
```

4. Launch with: `fs-uae dev.fs-uae`

Performance Tip: Enable JIT compilation (`jit_compiler = 1`) for significantly faster 68040/68060 emulation.

2.15.2 WinUAE Setup (Windows)

WinUAE offers the most configuration options:

1. Download from <https://www.winuae.net>
2. Create a new configuration:
 - Select “A4000” as the base configuration
 - Set CPU to 68040 or 68060
 - Enable 8MB Fast RAM
 - Add a directory hardfile pointing to your shared folder
3. Under “Expansions”, ensure “BSDSocket” is enabled for network access

2.15.3 Amiberry Setup (Raspberry Pi)

Amiberry is optimized for ARM single-board computers:

```
# Install on Raspberry Pi OS
sudo apt install amiberry

# Copy your Kickstart ROMs to ~/.amiberry/Kickstarts/
```

Configuration is similar to FS-UAE. Amiberry’s built-in GUI makes setup straightforward.

Coming from C

When developing C for Amiga, you typically cross-compile on a powerful host machine and test in an emulator. The same workflow applies to Novus:

1. Write code on your host (macOS/Linux/Windows)
2. Compile with `novus compile`
3. Transfer via shared folder
4. Test in emulator or on real hardware
5. Debug, fix, repeat

The main difference: Novus gives you immediate feedback on type errors before you even try to run the program. No more “undefined reference” linker errors or segfaults from type mismatches.

2.16 Working with the Amiga Shell

When testing your programs, you’ll interact with the Amiga Shell. Here are essential commands:

2.16.1 Basic Navigation

```
# Change directory
CD Work:Projects

# List files (like ls)
DIR

# Show current directory
CD

# Create a directory
MAKEDIR NewFolder
```

2.16.2 Running Programs

```
# Run from current directory
myprogram
```

```
# Run with full path
Work:Projects/myprogram

# Run and capture output
myprogram >T:output.txt

# Run with input redirection
myprogram <T:input.txt

# Run in background
RUN myprogram
```

2.16.3 Checking Exit Codes

```
# Run your program
myprogram

# Check the return code
IF WARN
  ECHO "Program returned error (exit code >= 5)"
ENDIF

# Or use conditional execution
myprogram && ECHO "Success" || ECHO "Failed"
```

2.16.4 Useful Shell Variables

```
# Get the return code of last command
ECHO $RC

# Current directory
ECHO $CD

# Date and time
ECHO $DATE $TIME
```

2.16.5 Environment and Assigns

```
# Create an assign (like a symbolic link)
ASSIGN MYPROJECT: Work:Novus/MyProject

# Now you can use it
CD MYPROJECT:
TYPE MYPROJECT:main.novus

# Set an environment variable
SETENV MYVAR "Hello"
```

2.17 Common Issues and Solutions

This section addresses frequent issues developers encounter.

2.17.1 Compilation Errors

“Module not found: std::io::file”

The standard library isn’t being found. Ensure you’re running from the Novus repository directory, or set the `NOVUS_HOME` environment variable:

```
export NOVUS_HOME=/path/to/Novus
novus compile hello.novus
```

“Type mismatch: expected i32, found u32”

Novus requires explicit type conversions. Add a cast:

```
// Wrong
let x: i32 = some_u32_value

// Right
let x: i32 = (i32)some_u32_value
```

“Cannot borrow as mutable”

You’re trying to modify something through an immutable reference:

```
// Wrong
fn modify(x: &i32) {
    *x = 42 // Error: cannot modify through &
}

// Right
fn modify(x: &var i32) {
    *x = 42 // OK: &var is mutable reference
}
```

2.17.2 Runtime Errors on Amiga

“Error: Not enough memory”

Your program ran out of chip or fast RAM. Solutions:

- Reduce buffer sizes
- Free unused allocations earlier
- Check for memory leaks (unmatched `AllocMem/FreeMem`)
- Increase emulator memory configuration

“Guru Meditation”

A crash occurred. Common causes:

- Null pointer dereference
- Stack overflow (increase stack size with `STACK 10000` before running)
- Invalid memory access

- Incorrect FFI calling convention

Enable debug builds (the default) for bounds checking that catches many issues.

“Program just exits with no output”

Possible causes:

- Program returned non-zero exit code (check with `ECHO $RC`)
- Output buffering issue—try `print` followed by explicit newline
- Early `return` before print statements

2.17.3 Linking Errors

“Undefined reference to `__main`”

The entry point wasn’t found. Ensure:

- You have a `pub fn main() -> i32` function
- It’s marked `pub`
- The file was included in compilation

“Multiple definitions of ...”

Two files define the same symbol. Common when accidentally including the same code twice. Check your imports and ensure you’re not compiling object files multiple times.

2.17.4 Transfer and Execution Issues

Executable won’t run (“Command not found”)

Check that:

- The file has the correct protection bits (`PROTECT myprogram +e`)
- You’re using the correct path
- The file wasn’t corrupted during transfer (check file size)

“Wrong CPU type”

The executable was compiled for a CPU your Amiga doesn’t have:

```
# Compile for base 68020 (works on all supported systems)
novus compile hello.novus --cpu 68020
```

2.18 Workspaces for Multi-Project Development

For larger development efforts with multiple related projects, use workspaces.

2.18.1 Creating a Workspace

```
mkdir myworkspace
cd myworkspace

# Create workspace.toml
cat > workspace.toml << 'EOF'
[workspace]
```

```
members = [  
    "app",  
    "library",  
    "tests"  
]  
  
[build]  
target_cpu = "68020"  
optimization_level = 1  
EOF  
  
# Create member projects  
novus new app --type cli  
novus new library --type library  
novus new tests --type cli
```

2.18.2 Workspace Structure

```
myworkspace/  
  workspace.toml      # Workspace configuration  
  app/                # Main application  
    project.toml  
    src/  
      main.novus  
  library/            # Shared library  
    project.toml  
    src/  
      lib.novus  
  tests/              # Test suite  
    project.toml  
    src/  
      main.novus
```

2.18.3 Building the Workspace

```
# Build all projects  
novus build  
  
# Build specific project  
novus build app  
  
# Build with release optimizations  
novus build --release
```

2.18.4 Project Dependencies

Projects can depend on other workspace members:

```
# app/project.toml  
[package]  
name = "app"  
  
[dependencies]
```



```
library = { path = "../library" }
```

Then import from the dependency:

```
// app/src/main.novus
from library::utils import helper_function

pub fn main() -> i32 {
    helper_function()
    return 0
}
```

Coming from C

In C Amiga development, you'd manage this with Makefiles:

```
# Typical C Makefile
CC = vc
CFLAGS = -c99 -O2

all: app

app: app.o library.o
    $(CC) -o $@ $^

%.o: %.c
    $(CC) $(CFLAGS) -c $<
```

Novus workspaces provide similar functionality with less boilerplate. The `workspace.toml` and `project.toml` files replace Makefiles for most cases, though you can still use Make for complex build scenarios.

2.19 Debugging Techniques

Debugging on the Amiga requires different techniques than modern desktop development.

2.19.1 Print Debugging

The simplest technique—add print statements to trace execution:

```
fn process_data(data: &Data) -> Result<Output, Error> {
    println("DEBUG: entering process_data")
    print("DEBUG: data.len = ")
    print_u32(data.len)
    println("")

    let result = expensive_operation(data)?

    println("DEBUG: operation completed")
    Result::Ok(result)
}
```

For release builds, use conditional compilation:

```
const DEBUG: bool = true
```

```
fn debug_print(msg: *u8) {
    if DEBUG {
        println(msg)
    }
}
```

2.19.2 Assert for Invariants

Use assertions to catch bugs early:

```
fn get_element(arr: &[i32], index: u32) -> i32 {
    assert!(index < arr.len(), "index out of bounds")
    return arr[index].unwrap()
}

fn normalize(vec: &var Vec3) {
    let len = vec.length()
    assert!(len > 0.0, "cannot normalize zero vector")
    vec.x = vec.x / len
    vec.y = vec.y / len
    vec.z = vec.z / len
}
```

2.19.3 Memory Tracking

Track allocations to find leaks:

```
var g_alloc_count: u32 = 0
var g_free_count: u32 = 0

fn tracked_alloc(size: u32) -> *u8 {
    let ptr = unsafe { AllocMem(size, MEMF_PUBLIC) }
    if ptr != null {
        g_alloc_count += 1
    }
    return ptr
}

fn tracked_free(ptr: *u8, size: u32) {
    if ptr != null {
        g_free_count += 1
        unsafe { FreeMem(ptr, size) }
    }
}

fn report_leaks() {
    if g_alloc_count != g_free_count {
        print("WARNING: ")
        print_u32(g_alloc_count - g_free_count)
        println(" allocations not freed!")
    }
}
```

2.19.4 Using Amiga Debuggers

For complex issues, use an Amiga debugger:

- **Barfly** — Integrated assembler and debugger
- **SnoopDOS** — Monitor DOS and library calls
- **Enforcer** — Catch illegal memory accesses (requires MMU)
- **MuForce** — Similar to Enforcer, works on 68040/060

Compile with debug symbols for better stack traces:

```
novus compile myprogram.novus -o myprogram --debug-symbols
```

2.20 Build System Integration

While Novus has its own build system, you can integrate it with external tools.

2.20.1 Using Make with Novus

Create a **Makefile** for complex builds:

```
NOVUS = novus compile
OUTPUT_DIR = build
SRC = src/main.novus

.PHONY: all clean release test

all: $(OUTPUT_DIR)/myapp

$(OUTPUT_DIR)/myapp: $(SRC)
    @mkdir -p $(OUTPUT_DIR)
    $(NOVUS) $< -o $@

release: $(SRC)
    @mkdir -p $(OUTPUT_DIR)
    $(NOVUS) $< -o $(OUTPUT_DIR)/myapp --release

clean:
    rm -rf $(OUTPUT_DIR)

test:
    novus test tests/test_main.novus -o $(OUTPUT_DIR)/test_runner
```

2.20.2 CI/CD Integration

Example GitHub Actions workflow:

```
name: Build
on: [push, pull_request]

jobs:
  build:
```

```
runs-on: ubuntu-latest
steps:
  - uses: actions/checkout@v4

  - name: Setup .NET
    uses: actions/setup-dotnet@v4
    with:
      dotnet-version: '9.0.x'

  - name: Build Novus
    run: dotnet build

  - name: Compile Project
    run: |
      ./Novus/bin/Debug/net9.0/novus compile \
        src/main.novus -o build/myapp

  - name: Upload Artifact
    uses: actions/upload-artifact@v4
    with:
      name: myapp-amiga
      path: build/myapp
```

2.21 Summary

You now have everything needed to start developing with Novus:

- **Installation** — .NET SDK + clone the repository
- **First program** — `novus compile hello.novus`
- **Pipeline understanding** — `Source` → `IR` → `C` → `Object` → `HUNK`
- **Project structure** — `novus new` creates standard layout
- **Workflow** — Edit, compile, transfer, test, repeat
- **Emulator setup** — FS-UAE, WinUAE, or Amiberry
- **Debugging** — Print statements, assertions, and Amiga debuggers

The following chapters teach the language itself. Keep this chapter handy as a reference for the development workflow.

Chapter 3

Language Basics

“Simplicity is prerequisite for reliability.”
—Edsger W. Dijkstra

This chapter covers the fundamental building blocks of Novus: how to declare variables, what types are available, how literals work, and how expressions combine values. By the end, you’ll understand the core syntax and be ready to write meaningful programs.

Every language has its foundation—the atoms from which programs are built. In Novus, these foundations are designed with systems programming in mind: explicit types, predictable memory layout, and direct mapping to 68k machine operations. Understanding these basics thoroughly will make everything else fall into place.

3.1 Comments

Novus supports two comment styles:

3.1.1 Single-Line Comments

Use `//` for comments that extend to the end of the line:

```
// This is a single-line comment
let x = 42 // Comments can appear after code
```

3.1.2 Multi-Line Comments

Use `/* */` for comments spanning multiple lines:

```
/*
 * This is a multi-line comment.
 * Useful for longer explanations.
 */
let result = compute()
```

Multi-line comments do not nest. The first `*/` closes the comment, regardless of how many `/*` appeared inside.

3.2 Variables and Mutability

Novus provides three keywords for declaring bindings:

3.2.1 Mutable Variables with `var`

Variables declared with `var` can be reassigned:

```
var x = 10
x = 20    // OK - can reassign
x = x + 5 // OK
```

Use `var` for counters, accumulators, temporary values, and any data that changes during execution.

3.2.2 Immutable Bindings with `let`

Bindings declared with `let` cannot be reassigned:

```
let x = 10
x = 20    // ERROR - cannot reassign
```

Immutable bindings make code intent clearer and allow the compiler to optimize more aggressively. Use `let` by default; switch to `var` when you need mutability.

Note that immutability applies to the binding, not necessarily to the data it refers to. For example, a `let` binding to a mutable reference (`&var T`) still allows modifying the referenced data.

3.2.3 Compile-Time Constants with `const`

Constants declared with `const` are evaluated at compile time:

```
pub const SCREEN_WIDTH: u32 = 320
pub const SCREEN_HEIGHT: u32 = 200
pub const SCREEN_SIZE: u32 = SCREEN_WIDTH * SCREEN_HEIGHT
```

Constants must be literals or constant expressions—values that the compiler can evaluate without executing code. They cannot call functions or use runtime values.

Constants are typically declared at module scope with `pub` visibility.

3.2.4 Comparison

Keyword	Reassignable	Evaluation	Typical Scope
<code>var</code>	Yes	Runtime	Function/Block
<code>let</code>	No	Runtime	Function/Block
<code>const</code>	No	Compile-time	Module

Coming from C

In C, all variables are mutable by default. In Novus, prefer `let` (immutable) and only use `var` when you need to reassign.

C	Novus
<code>int x = 10;</code>	<code>var x = 10</code>
<code>const int X = 10;</code>	<code>let x = 10</code> or <code>const X = 10</code>

Novus's `const` is stricter than C's—it must be a compile-time constant expression, not just a runtime immutable value.

3.2.5 Type Annotations

You can explicitly annotate types:

```
let x: i32 = 42
var count: u32 = 0
```

When the type is obvious from context, annotations are optional:

```
let x = 42          // Inferred as i32
let y = 42u32       // Type suffix makes it u32
```

The compiler performs type inference within function bodies. Module-level constants typically require explicit types.

3.3 Primitive Types

Novus provides several categories of primitive types.

3.3.1 Integer Types

Novus supports signed and unsigned integers in multiple sizes:

- i8** Signed 8-bit integer (−128 to 127)
- i16** Signed 16-bit integer (−32,768 to 32,767)
- i32** Signed 32-bit integer (−2,147,483,648 to 2,147,483,647)
- i64** Signed 64-bit integer (wide range, requires two 32-bit registers on 68k)
- u8** Unsigned 8-bit integer (0 to 255)
- u16** Unsigned 16-bit integer (0 to 65,535)
- u32** Unsigned 32-bit integer (0 to 4,294,967,295)
- u64** Unsigned 64-bit integer (0 to 18,446,744,073,709,551,615)

The default integer type is **i32**. When you write a bare integer literal like 42, it defaults to **i32** unless context requires otherwise.

3.3.2 Floating-Point Types

- f32** 32-bit IEEE 754 single-precision float
- f64** 64-bit IEEE 754 double-precision float

On 68k systems without an FPU, floating-point operations use software emulation—slow but accurate. For performance-critical math on non-FPU systems, prefer fixed-point types.

3.3.3 Fixed-Point Types

Novus provides fixed-point arithmetic for efficient fractional math on integer hardware:

- fixed16** 16-bit fixed-point (8.8 format: 8 integer bits, 8 fractional bits)
- fixed32** 32-bit fixed-point (16.16 format: 16 integer bits, 16 fractional bits)

Fixed-point types use native 68k integer instructions, making them dramatically faster than soft-float on systems without an FPU.

Fixed-Point Representation

Fixed-point numbers store fractional values by scaling integers. A `fixed16` with 8.8 format stores the value multiplied by 256:

```
// fixed16 (8.8 format)
let half: fixed16 = 0.5      // Stored as 128 (0.5 * 256)
let quarter: fixed16 = 0.25  // Stored as 64 (0.25 * 256)
let pi: fixed16 = 3.14        // Stored as ~804 (3.14 * 256)

// fixed32 (16.16 format) - more precision
let precise_pi: fixed32 = 3.14159 // Stored as 205887 (3.14159 * 65536)
```

Fixed-Point Arithmetic

All standard arithmetic operations work on fixed-point types:

```
let a: fixed32 = 2.5
let b: fixed32 = 1.5

let sum = a + b      // 4.0
let diff = a - b     // 1.0
let product = a * b  // 3.75
let quotient = a / b // 1.666...
```

The compiler automatically handles the scaling during multiplication and division, using efficient shift operations where possible.

Converting Between Fixed and Integer

Convert between fixed-point and integers explicitly:

```
let f: fixed32 = 3.75
let i: i32 = f.to_int()      // 3 (truncates)
let rounded: i32 = f.round() // 4 (rounds to nearest)

let from_int: fixed32 = fixed32::from_int(5) // 5.0
```

When to Use Fixed-Point

Fixed-point excels in these scenarios:

- **Game physics** — Position, velocity, acceleration calculations
- **Audio processing** — Sample interpolation, volume scaling
- **Graphics** — Texture mapping, rotation, scaling
- **Animation** — Smooth movement, easing functions

Coming from C

C programmers often implement fixed-point manually with macros and shift operations:

C (manual)	Novus
<code>#define FP_SHIFT 16</code>	(built-in)
<code>int fp_mul(int a, int b)</code>	<code>a * b</code>
<code>{ return (a * b) » 16; }</code>	(automatic scaling)
<code>int x = (3 « 16) / 2;</code>	<code>let x: fixed32 = 3.0 / 2.0</code>

Novus handles the bit manipulation automatically while generating identical machine code.

3.3.4 Boolean Type

The `bool` type represents truth values:

```
let flag: bool = true
let done: bool = false
```

Booleans result from comparison and logical operations. Unlike C, integers and pointers do not implicitly convert to booleans—you must use explicit comparisons.

Coming from C

C allows `if (ptr)` and `if (count)`. Novus requires explicit comparisons:

C	Novus
<code>if (ptr)</code>	<code>if ptr != null</code>
<code>if (!ptr)</code>	<code>if ptr == null</code>
<code>if (count)</code>	<code>if count != 0</code>
<code>while (n-)</code>	<code>while n > 0 { n- ... }</code>

This catches bugs like `if (x = 10)` which in C silently assigns and evaluates to true.

3.4 Literals

3.4.1 Integer Literals

Integer literals support multiple bases and optional type suffixes.

Decimal

```
let x = 42
let y = 1_000_000 // Underscores improve readability
```

Underscores are ignored; they exist solely to make large numbers easier to read.

Hexadecimal

Novus supports both C-style and Amiga-style hexadecimal:

```
let color1 = 0xFF00AA // C-style
let color2 = $FF00AA // Amiga-style (same value)
```

Both forms are equivalent. Use whichever feels natural—Amiga developers often prefer `$`, while those from C backgrounds prefer `0x`.

Binary

Binary literals use the % prefix:

```
let mask = %1010    // Binary literal (value: 10)
let flags = %11110000
```

Binary literals are useful for bit patterns and masks when working with hardware registers.

Type Suffixes

Append a type suffix to specify the literal's type explicitly:

```
let a = 42u32    // u32
let b = 42i8     // i8
let c = 100u16   // u16
```

Type suffixes work with all bases:

```
let hex_byte = $FFu8
let binary_word = %1010u16
let decimal_long = 1000u32
```

Without a suffix, integer literals default to i32.

3.4.2 Floating-Point Literals

Floating-point literals require a decimal point:

```
let pi = 3.14
let e = 2.71828f32 // f32 with explicit suffix
```

Without a suffix, floating-point literals default to f64.

3.4.3 String Literals

String literals use double quotes and support escape sequences:

```
let greeting = "Hello, Amiga!"
let multiline = "Line 1\nLine 2\nLine 3"
let path = "SYS:Utilities\\" // Escaped backslash
```

Supported escape sequences include:

- `\n` Newline (line feed)
- `\r` Carriage return
- `\t` Horizontal tab
- `\0` Null byte
- `\\` Backslash
- `\"` Double quote
- `\'` Single quote
- `\xNN` Hexadecimal byte (e.g., `\x1B` for ESC)

3.4.4 F-Strings (Formatted Strings)

F-strings embed expressions directly in string literals:

```
let x = 42
let message = f"The answer is {x}"
// Results in: "The answer is 42"
```

Expressions inside `{...}` are evaluated and converted to strings. F-strings compile to efficient string concatenation or formatting calls.

3.4.5 Character Literals

Character literals use single quotes:

```
let ch: u8 = 'A'
let newline: u8 = '\n'
let escape: u8 = '\x1B' // ESC character
```

Character literals produce values of type `u8`. Novus does not have a separate `char` type—characters are simply 8-bit unsigned integers.

3.5 Operators

3.5.1 Arithmetic Operators

Standard arithmetic works on integer and floating-point types:

Operator	Meaning
+	Addition
-	Subtraction
*	Multiplication
/	Division
%	Modulus (remainder)

Division by zero is a runtime error in debug builds and undefined behavior in release builds. Always validate denominators when accepting untrusted input.

```
let sum = 10 + 20 // 30
let product = 5 * 6 // 30
let quotient = 20 / 5 // 4
let remainder = 17 % 5 // 2
```

3.5.2 Comparison Operators

Comparisons produce `bool` values:

Operator	Meaning
==	Equal to
!=	Not equal to
<	Less than
>	Greater than
<=	Less than or equal to
>=	Greater than or equal to

```
let a = 10
let b = 20
let equal = a == b    // false
let less = a < b      // true
let greater_eq = a >= b // false
```

3.5.3 Logical Operators

Logical operators combine boolean values:

Operator	Meaning
&&	Logical AND (short-circuiting)
	Logical OR (short-circuiting)
!	Logical NOT

```
let x = true
let y = false
let and_result = x && y    // false
let or_result = x || y    // true
let not_result = !x       // false
```

Both `&&` and `||` short-circuit: the right-hand side is not evaluated if the result is determined by the left-hand side alone.

3.5.4 Bitwise Operators

Bitwise operators manipulate individual bits:

Operator	Meaning
&	Bitwise AND
	Bitwise OR
^	Bitwise XOR
~	Bitwise NOT (one's complement)
<<	Left shift
>>	Right shift (arithmetic for signed, logical for unsigned)

```
let a = %1111 // 15
let b = %1010 // 10

let and_bits = a & b    // %1010 (10)
let or_bits = a | b     // %1111 (15)
let xor_bits = a ^ b    // %0101 (5)
let not_bits = ~a       // %0000 (-16 in two's complement)

let shifted_left = 1 << 3 // 8
let shifted_right = 16 >> 2 // 4
```

Bitwise operations compile to single 68k instructions and execute in one or two cycles—essential for custom chip programming.

Practical Bitwise Examples

Bitwise operations are fundamental for Amiga programming:

```
// Setting bits (turn ON)
var flags: u16 = 0
flags |= %00000001 // Set bit 0
flags |= %00000100 // Set bit 2
// flags is now %00000101

// Clearing bits (turn OFF)
flags &= ~%00000001 // Clear bit 0
// flags is now %00000100

// Toggling bits (flip)
flags ^= %00001111 // Toggle bottom 4 bits
// flags is now %00001011

// Testing bits
if (flags & %00001000) != 0 {
    // Bit 3 is set
}

// Extracting bit fields
let color: u16 = $0ABC
let red = (color >> 8) & $F // Extract bits 8-11: $A
let green = (color >> 4) & $F // Extract bits 4-7: $B
let blue = color & $F // Extract bits 0-3: $C
```

Hardware Register Manipulation

Custom chip programming relies heavily on bitwise operations:

```
// Enable DMA channels
const DMAF_SETCLR: u16 = $8000
const DMAF_COPPER: u16 = $0080
const DMAF_BITPLANE: u16 = $0100
const DMAF_BLITTER: u16 = $0040

// Turn on Copper, Bitplane, and Blitter DMA
let dmacon_value = DMAF_SETCLR | DMAF_COPPER | DMAF_BITPLANE |
    DMAF_BLITTER

// Mask and shift for sprite position
let x: u16 = 150
let y: u16 = 100
let sprite_pos = ((y & $FF) << 8) | ((x >> 1) & $FF)
```

3.5.5 Compound Assignment Operators

Compound assignment combines an operation with assignment:

```
var x = 10
x += 5 // x = x + 5
x -= 3 // x = x - 3
x *= 2 // x = x * 2
```

```
x /= 4    // x = x / 4
x %= 3    // x = x % 3

x &= $FF  // x = x & $FF
x |= $80  // x = x | $80
x ^= $AA  // x = x ^ $AA
x <<= 2   // x = x << 2
x >>= 1   // x = x >> 1
```

All arithmetic and bitwise operators have corresponding compound forms.

3.5.6 Increment and Decrement Operators

Novus supports both prefix and postfix increment/decrement:

```
var x = 5

++x // Pre-increment: x becomes 6, evaluates to 6
--x // Pre-decrement: x becomes 5, evaluates to 5

let a = x++ // Post-increment: a = 5, then x becomes 6
let b = x-- // Post-decrement: b = 6, then x becomes 5
```

Prefix operators modify the variable and return the new value. Postfix operators return the original value and then modify the variable.

3.5.7 Range Operators

Range operators create range values for iteration:

start..end Exclusive range (start to end - 1)

start..=end Inclusive range (start to end)

```
for i in 0..10 {
    // i goes from 0 to 9
}

for j in 0..=10 {
    // j goes from 0 to 10 (inclusive)
}
```

Ranges are covered in detail in Chapter 6.

3.5.8 Operator Precedence

When multiple operators appear in an expression, precedence determines evaluation order. Higher precedence operators bind more tightly:

Precedence	Operators	Associativity
1 (highest)	() [] . ->	Left-to-right
2	! ~ ++ - * & (unary)	Right-to-left
3	* / % (binary)	Left-to-right
4	+ - (binary)	Left-to-right
5	<< >>	Left-to-right
6	< <= > >=	Left-to-right
7	== !=	Left-to-right
8	& (bitwise AND)	Left-to-right
9	^ (bitwise XOR)	Left-to-right
10	(bitwise OR)	Left-to-right
11	&&	Left-to-right
12		Left-to-right
13	=	Left-to-right
14 (lowest)	= += -= etc.	Right-to-left

Precedence Examples

```
// Multiplication before addition
let x = 2 + 3 * 4      // 14, not 20

// Shifts before comparison
let y = 1 << 4 > 10    // true (16 > 10)

// Bitwise AND before OR
let z = 1 | 2 & 3      // 3 (2 & 3 = 2, then 1 | 2 = 3)

// Use parentheses for clarity
let clear = (flags & mask) == 0
let combined = (a | b) & (c | d)
```

Coming from C

Novus uses the same operator precedence as C, so your intuition transfers directly. The common gotcha in both languages is that bitwise operators have lower precedence than comparison:

```
// WRONG in both C and Novus:
if flags & MASK == 0 { ... } // Parsed as: flags & (MASK == 0)

// CORRECT:
if (flags & MASK) == 0 { ... }
```

Always parenthesize bitwise operations in conditions.

3.6 Type Conversions

3.6.1 C-Style Casts

Use C-style cast syntax to convert between numeric types:

```
let x: u32 = 42
let y: i32 = (i32)x
```

```
let z: *u8 = (*u8)y
```

Casts can chain:

```
let w: i64 = (i64)(i32)(u32)x
```

Casts are explicit and potentially unsafe—truncation, sign changes, and pointer conversions are your responsibility. The compiler trusts that you know what you’re doing.

3.6.2 The `as` Keyword

The `as` keyword provides an alternative syntax for type assertions:

```
let x = 42 as u32
```

Both cast styles are valid; choose whichever feels more natural.

3.7 References and Pointers

Novus distinguishes between references (non-null pointers checked at compile time) and raw pointers (nullable, unchecked).

3.7.1 Immutable References

Create a reference with the `&` operator:

```
var x = 10
let ref_x: &i32 = &x
```

The type is `&i32`—a reference to an `i32`. References are guaranteed non-null and automatically dereference in most contexts.

3.7.2 Mutable References

For references that allow modification, use `&var`:

```
var x = 10
let ref_x: &var i32 = &var x
*ref_x = 20 // OK - modifies x
```

The type is `&var i32`. The `&var` syntax appears both in the type annotation and when creating the reference.

Note that Novus uses `var` for mutable, not `mut` as in some other languages.

3.7.3 Raw Pointers

Raw pointers use the `*T` syntax:

```
let ptr: *u8 = (*u8)0xDFF000 // Custom chip register
```

Pointers can be null and require explicit null checks:

```
if ptr == null {
    return // Handle null case
}
```



```
}

```

Raw pointers are essential for FFI and hardware access, but prefer references when possible.

Coming from C

Novus distinguishes between references (guaranteed non-null) and pointers (nullable):

C	Novus	Meaning
<code>int *p</code>	<code>*i32</code>	Nullable pointer
(no equivalent)	<code>&i32</code>	Non-null reference
<code>int * const p</code>	<code>let p: *i32</code>	Immutable pointer
<code>const int *p</code>	<code>*i32 (read-only)</code>	Pointer to const

References eliminate an entire class of null pointer bugs by making non-null the default.

3.7.4 Dereferencing

Use `*` to dereference a pointer or reference:

```
let x = 10
let ref_x = &x
let value = *ref_x // 10

```

For pointers, dereferencing an invalid or null pointer is undefined behavior.

3.7.5 Pointer Arithmetic

Pointer arithmetic works like C—offsets are scaled by the pointed-to type’s size:

```
let arr = [10, 20, 30, 40, 50]
let ptr: *i32 = &arr[0u32]

let first = *ptr // 10
let second = *(ptr + 1) // 20 (advances by sizeof(i32) = 4 bytes)
let third = *(ptr + 2) // 30

```

Pointer subtraction returns the number of elements between two pointers:

```
let start: *i32 = &arr[0u32]
let end: *i32 = &arr[4u32]
let count = end - start // 4 elements

```

3.7.6 Auto-Dereferencing

Novus automatically dereferences pointers when accessing struct fields:

```
struct Point { x: i32, y: i32 }

let p = Point { x: 10, y: 20 }
let ptr: *Point = &p

// Both work identically:
let x1 = (*ptr).x // Explicit dereference
let x2 = ptr.x    // Auto-dereference (preferred)

```

This convenience applies to method calls as well—you don’t need `(*ptr).method()`.

3.7.7 Address-of Operator

The `&` operator takes the address of a value:

```
var x: i32 = 42
let ptr: *i32 = &x      // Address of x

// For mutable access:
let mut_ptr: *var i32 = &var x
*mut_ptr = 100          // Modifies x
```

You cannot take the address of a temporary value or literal—only named variables and array elements have stable addresses.

3.8 Arrays

3.8.1 Fixed-Size Arrays

Arrays have a fixed size known at compile time:

```
let numbers = [10, 20, 30]
let primes = [2, 3, 5, 7, 11]
```

The compiler infers both the element type and size from the literal. The type of `numbers` is `[i32; 3]`—an array of 3 `i32` values. The size is part of the type.

3.8.2 Repeat Syntax

Initialize arrays with a repeated value:

```
let zeros = [0; 100]      // 100 zeros
let buffer = [0u8; 1024]  // 1024 bytes
```

The syntax `[value; count]` creates an array of `count` elements, each initialized to `value`.

3.8.3 Indexing

Access array elements with bracket indexing:

```
let numbers = [10, 20, 30]
let first = numbers[0u32]  // 10
let second = numbers[1u32] // 20
```

Indices must be unsigned integers. In debug builds, out-of-bounds access triggers a panic. In release builds, bounds checking may be elided for performance.

3.9 Tuples

3.9.1 Tuple Types

Tuples group multiple values of potentially different types:

```
let point: (i32, i32) = (100, 200)
let rgb: (u8, u8, u8) = (255, 128, 64)
```

The type `(i32, i32)` means “a tuple of two `i32` values.”

3.9.2 Tuple Literals

Create tuples with parenthesized, comma-separated values:

```
let color = (192, 96, 32)
```

3.9.3 Destructuring

Extract tuple elements with destructuring:

```
let (r, g, b) = get_rgb()
```

This binds `r`, `g`, and `b` to the tuple’s three elements. Destructuring works in `let` and `var` bindings.

3.9.4 Returning Multiple Values

Tuples enable returning multiple values from functions:

```
fn get_rgb() -> (i32, i32, i32) {
    return (255, 128, 64)
}
```

3.10 Expressions and Statements

In Novus, most constructs are expressions—they produce values.

3.10.1 Expression vs Statement

An expression evaluates to a value:

```
let x = 2 + 2 // Expression: 4
```

A statement performs an action without producing a value:

```
x = 10 // Assignment statement
```

However, even assignments are expressions in Novus, allowing chained assignment:

```
var a = 0
var b = 0
a = b = 5 // Both a and b become 5
```

3.10.2 Expression Contexts

Many control flow constructs in Novus can be used as expressions. For example, `if-else` can produce a value:

```
let max = if a > b { a } else { b }
```

The `match` expression (covered in Chapter 6) similarly produces values based on pattern matching. Both branches must have the same type when used as expressions.

3.11 Ininsics

Novus provides compile-time intrinsics that look like function calls but are handled specially by the compiler:

```
// Get size of a type in bytes
let size = @sizeof(MyStruct)

// Create a zero-initialized value of a type
let data = @zeroed(MyStruct)
```

Additionally, `@drop_in_place(expr)` can be used to explicitly drop a value before it goes out of scope.

Note: The `alignof()` intrinsic is currently only available in inline assembly blocks. These intrinsics are resolved at compile time and have no runtime overhead.

3.11.1 Available Intrinsics

`@sizeof(T)` Returns the size of type `T` in bytes

`@zeroed(T)` Creates a zero-initialized value of type `T`

`@drop_in_place(expr)` Explicitly drops a value before scope end

`@unreachable()` Marks code as unreachable (optimization hint)

`@type_name(T)` Returns the name of type `T` as a string

3.11.2 Common Uses

The `@sizeof` intrinsic is essential for memory allocation:

```
let size = @sizeof(Player) * count
let mem = AllocMem(size, MEMF_PUBLIC)
```

The `@zeroed` intrinsic creates zero-initialized instances, commonly used in constructors:

```
impl Vec<T> {
    pub fn new() -> Vec<T> {
        return @zeroed(Vec<T>) // All fields zero/null
    }
}
```

3.12 Best Practices

Following these guidelines will help you write clear, efficient Novus code:

3.12.1 Variable Declarations

- **Prefer `let` over `var`** — Immutability is the default; only use `var` when you need to reassign.
- **Use meaningful names** — `player_count` is better than `pc` or `x`.
- **Initialize at declaration** — Avoid declaring variables without immediate initialization.
- **Use `const` for true constants** — Screen dimensions, hardware addresses, magic numbers.

```
// Good
const SCREEN_WIDTH: u16 = 320
let player_count = get_player_count()
var current_score = 0u32

// Avoid
var x = 0           // What is x?
let BUFFER_SIZE = 1024 // Should be const
```

3.12.2 Type Annotations

- **Omit when obvious** — `let x = 42` is fine; the type is clearly `i32`.
- **Annotate for clarity** — `let mask: u16 = $FFFF` makes the intended type explicit.
- **Always annotate function signatures** — Parameters and return types should be explicit.

3.12.3 Numeric Types

- **Match the hardware** — Use `u16` for hardware registers, `u32` for addresses.
- **Use unsigned for bit manipulation** — Bitwise operations on signed types can surprise you.
- **Prefer `fixed32` over `f32`** — On non-FPU systems, fixed-point is 10-100x faster.

3.13 Common Pitfalls

3.13.1 Integer Overflow

Unlike some languages, Novus does not check for integer overflow in release builds:

```
var x: u8 = 255
x += 1    // x becomes 0, not 256!

var y: i8 = 127
y += 1    // y becomes -128 (overflow wraps)
```

Use wider types or explicit overflow checks when necessary.

3.13.2 Signed vs Unsigned Comparisons

Comparing signed and unsigned integers can produce surprising results:

```
let signed: i32 = -1
let unsigned: u32 = 1

// This requires explicit casting
if (u32)signed > unsigned { // -1 as u32 is 4294967295
    // This branch executes!
}
```

The compiler warns about mixed signed/unsigned comparisons. Heed these warnings.

3.13.3 Precedence Surprises

Bitwise operators have lower precedence than comparison operators:

```
// WRONG: Parsed as flags & (MASK != 0)
if flags & MASK != 0 { }

// CORRECT: Explicit parentheses
if (flags & MASK) != 0 { }
```

When in doubt, add parentheses.

3.13.4 Reference vs Copy

Understanding when you have a reference vs a copy is crucial:

```
var x = 42
var y = x      // y is a COPY of x
y = 100        // x is still 42

var a = [1, 2, 3]
let b = &a     // b is a REFERENCE to a
// Modifying through b would modify a
```

Primitive types (`i32`, `bool`, etc.) are always copied. Arrays and structs may be copied or referenced depending on context.

3.14 Showcase: Temperature Converter

This example demonstrates many concepts from this chapter in a practical application:

```
// temperature_converter.novus
// Demonstrates: variables, types, operators, fixed-point,
// functions

from std::io import println

// Constants for conversion formulas
const FAHRENHEIT_OFFSET: fixed32 = 32.0
const CELSIUS_MULTIPLIER: fixed32 = 1.8 // 9/5

// Temperature scales as an enum
```

```

enum Scale {
    Celsius,
    Fahrenheit,
    Kelvin,
}

// A temperature value with its scale
struct Temperature {
    value: fixed32,
    scale: Scale,
}

impl Temperature {
    // Create a new temperature
    pub fn new(value: fixed32, scale: Scale) -> Temperature {
        return Temperature { value, scale }
    }

    // Convert to Celsius
    pub fn to_celsius(&self) -> fixed32 {
        match self.scale {
            Scale::Celsius => self.value,
            Scale::Fahrenheit => (self.value - FAHRENHEIT_OFFSET)
                / CELSIUS_MULTIPLIER,
            Scale::Kelvin => self.value - 273.15,
        }
    }

    // Convert to Fahrenheit
    pub fn to_fahrenheit(&self) -> fixed32 {
        let celsius = self.to_celsius()
        return celsius * CELSIUS_MULTIPLIER + FAHRENHEIT_OFFSET
    }

    // Convert to Kelvin
    pub fn to_kelvin(&self) -> fixed32 {
        let celsius = self.to_celsius()
        return celsius + 273.15
    }

    // Check if freezing
    pub fn is_freezing(&self) -> bool {
        return self.to_celsius() <= 0.0
    }

    // Check if boiling (water at sea level)
    pub fn is_boiling(&self) -> bool {
        return self.to_celsius() >= 100.0
    }
}

fn main() -> i32 {
    // Create some temperatures
    let room_temp = Temperature::new(72.0, Scale::Fahrenheit)
    let freezing = Temperature::new(0.0, Scale::Celsius)
    let absolute_zero = Temperature::new(0.0, Scale::Kelvin)

    // Convert and display

```

```
println(f"Room temperature:
    {room_temp.to_celsius().to_int()}C")
println(f"Freezing point:
    {freezing.to_fahrenheit().to_int()}F")
println(f"Absolute zero:
    {absolute_zero.to_celsius().to_int()}C")

// Boolean operations
if room_temp.is_freezing() {
    println("Brrr! It's freezing!")
} else {
    println("Temperature is above freezing.")
}

// Array of temperatures
var temps = [
    Temperature::new(32.0, Scale::Fahrenheit),
    Temperature::new(100.0, Scale::Celsius),
    Temperature::new(273.15, Scale::Kelvin),
]

// Find the hottest
var max_celsius: fixed32 = -1000.0
for i in 0u32..3 {
    let c = temps[i].to_celsius()
    if c > max_celsius {
        max_celsius = c
    }
}

println(f"Hottest: {max_celsius.to_int()}C")

// Bit manipulation example: encode temp + scale in u16
// Format: [scale:2][temp_int:14]
let temp_int = room_temp.to_celsius().to_int()
let scale_bits: u16 = 0u16 // Celsius = 0
let encoded: u16 = (scale_bits << 14) | ((u16)(temp_int) &
    $3FFF)

// Decode
let decoded_scale = encoded >> 14
let decoded_temp = (i16)(encoded & $3FFF)

println(f"Encoded: ${encoded:04X}, Scale: {decoded_scale},
    Temp: {decoded_temp}")

return 0
}
```

This showcase demonstrates:

- Constants with `const` and fixed-point values
- Enums for type-safe categories
- Structs with methods via `impl`
- Fixed-point arithmetic for precise calculations

- Boolean expressions and comparisons
- Arrays and iteration
- Bitwise operations for data packing
- F-strings for output formatting

3.15 Summary

This chapter covered Novus’s fundamental syntax:

- **Variables:** `var` for mutable, `let` for immutable, `const` for compile-time constants
- **Types:** Integers (`i8–i64`, `u8–u64`), floats (`f32`, `f64`), fixed-point (`fixed16`, `fixed32`), and `bool`
- **Literals:** Decimal, hex (`0x/$`), binary (`%`), strings, f-strings, and characters
- **Operators:** Arithmetic, comparison, logical, bitwise, and compound assignment
- **References:** `&T` (immutable), `&var T` (mutable), `*T` (raw pointers)
- **Arrays:** Fixed-size with type `[T; N]`, indexed with unsigned integers
- **Tuples:** Group values with `(T1, T2, ...)`, destructure with `let (a, b) = ...`

With these building blocks, you can write meaningful programs. The next chapter explores the type system in depth, covering structs, enums, and pattern matching.

Chapter 4

Types

“A type system is a syntactic method for enforcing levels of abstraction in programs.”
— John C. Reynolds

Novus features a static, strong type system designed to provide safety without sacrificing the control needed for systems programming on the Amiga. Every value has a known type at compile time, enabling the compiler to catch errors early and generate efficient 68k machine code.

This chapter explores Novus’s type system in depth: primitive types, composite types (structs and enums), traits for polymorphism, pattern matching for working with data, and generics for code reuse.

4.1 Type System Overview

Novus’s type system is built on these principles:

- **Static typing** — all types known at compile time
- **Strong typing** — no implicit conversions between unrelated types
- **Explicit conversions** — casts must be written explicitly
- **Zero-cost abstractions** — high-level types compile to efficient machine code
- **Null safety** — no null pointers; use `Option<T>` instead
- **Error handling** — errors are values via `Result<T, E>`

Every expression in Novus has a type, and the compiler verifies that types are used correctly throughout your program.

4.2 Primitive Types

Novus provides a set of built-in primitive types (covered in Chapter 2). Here’s a quick reference:

4.3 Structs

Structs are composite types that group related data together. They are the primary way to create custom types in Novus.

Type	Description
i8, i16, i32, i64	Signed integers
u8, u16, u32, u64	Unsigned integers
bool	Boolean (<code>true</code> or <code>false</code>)
f32, f64	Floating-point (soft-float on 68k without FPU)
fixed16, fixed32	Fixed-point (8.8 and 16.16 format)
*T	Raw pointer to type T (can be null)
&T, &var T	References (guaranteed non-null)

Table 4.1: Novus Primitive Types

4.3.1 Declaring Structs

A struct declaration defines a new type with named fields:

```
struct Point {  
    x: i32,  
    y: i32,  
}  
  
struct Color {  
    r: u8,  
    g: u8,  
    b: u8,  
    a: u8,  
}
```

Each field has a name and a type. Fields are separated by commas, and a trailing comma is allowed.

Coming from C

No `typedef` needed! In C you often write:

```
typedef struct { int x, y; } Point;
```

In Novus, the struct name is directly usable as a type:

```
struct Point { x: i32, y: i32 }
```

Also note: no semicolons after field definitions, and types come *after* names (`x: i32` not `int x`).

4.3.2 Struct Visibility

By default, structs are private to the module where they're defined. Use `pub` to make them public:

```
pub struct Point {  
    x: i32,  
    y: i32,  
}  
  
internal struct Config {  
    setting: bool,  
}  
  
struct Private {  
    data: i32,  
}
```

Important: Regardless of struct visibility, *all fields are always private*. Novus does not support field-level visibility modifiers. To expose field values, provide accessor methods in an `impl` block.

4.3.3 Creating Struct Instances

Create instances using struct literal syntax:

```
let origin = Point { x: 0, y: 0 }
let pixel = Point { x: 100, y: 200 }

let red = Color { r: 255, g: 0, b: 0, a: 255 }
```

All fields must be initialized. The order of fields in the literal doesn't matter.

4.3.4 Accessing Fields

Access struct fields using dot notation:

```
let p = Point { x: 10, y: 20 }
let x_coord = p.x // 10
let y_coord = p.y // 20
```

To modify fields, the binding must be mutable:

```
var p = Point { x: 10, y: 20 }
p.x = 15
p.y = 25
```

4.3.5 Struct Update Syntax

The spread operator `..` allows creating a new struct from an existing one, updating specific fields:

```
let p1 = Point { x: 10, y: 20 }
let p2 = Point { x: 100, ..p1 } // { x: 100, y: 20 }
let p3 = Point { y: 200, ..p1 } // { x: 10, y: 200 }
```

The `..expr` must appear last in the literal and copies all unspecified fields from the given expression. At least one field must be explicitly specified; you cannot use spread alone.

4.3.6 Generic Structs

Structs can be parameterized with type parameters:

```
struct Pair<T, U> {
    first: T,
    second: U,
}
```

Usage:

```
// ... (with Pair defined above)
let int_pair = Pair { first: 42, second: 100 }
let mixed = Pair { first: 10, second: "hello" }
```

The compiler infers type arguments from usage, or you can specify them explicitly:

```
let explicit: Pair<i32, i32> = Pair { first: 1, second: 2 }
```

4.3.7 Const Generic Parameters

Structs can also have compile-time constant parameters:

```
struct Buffer<const N: u32> {
    data: [u8; N],
    len: u32,
}
```

Usage:

```
// ... (with Buffer defined above)
let small_buf: Buffer<64> = Buffer {
    data: [0; 64],
    len: 0
}
let large_buf: Buffer<1024> = Buffer {
    data: [0; 1024],
    len: 0
}
```

Const parameters must be integer types and are known at compile time.

4.3.8 Where Clauses

Struct definitions can have **where** clauses to constrain generic parameters:

```
struct Container<T> where T: Clone {
    value: T,
    backup: T,
}
```

This requires that any type `T` used with `Container` must implement the `Clone` trait.

4.4 Implementation Blocks

Methods and associated functions are defined in `impl` blocks:

```
impl Point {
    fn new(x: i32, y: i32) -> Point {
        Point { x: x, y: y }
    }

    fn distance(&self) -> f32 {
        let dx = self.x as f32
        let dy = self.y as f32
    }
}
```

```

        sqrt(dx * dx + dy * dy)
    }

    fn translate(&var self, dx: i32, dy: i32) {
        self.x = self.x + dx
        self.y = self.y + dy
    }
}

```

4.4.1 Self Parameter Variants

Methods can take `self` in different forms:

Form	Meaning
<code>&self</code>	Immutable borrow; method reads but doesn't modify
<code>&var self</code>	Mutable borrow; method can modify the instance
<code>self</code>	Takes ownership; instance consumed by call
<code>consuming self</code>	Explicit ownership transfer; semantically identical to <code>self</code>

Table 4.2: Self Parameter Forms

```

impl Point {
    // Immutable borrow - reads only
    fn magnitude(&self) -> i32 {
        self.x * self.x + self.y * self.y
    }

    // Mutable borrow - modifies in place
    fn scale(&var self, factor: i32) {
        self.x = self.x * factor
        self.y = self.y * factor
    }

    // Consumes self - takes ownership
    fn into_tuple(self) -> (i32, i32) {
        (self.x, self.y)
    }

    // Explicit consuming - same as self
    fn consume(consuming self) {
        // self is consumed
    }
}

```

4.4.2 Associated Functions

Functions without a `self` parameter are associated functions (like static methods):

```

impl Point {
    fn origin() -> Point {
        Point { x: 0, y: 0 }
    }
}

```

```
fn from_polar(r: f32, theta: f32) -> Point {
    Point {
        x: (r * cos(theta)) as i32,
        y: (r * sin(theta)) as i32,
    }
}
```

Usage:

```
// ... (with Point defined above)
let origin = Point::origin()
let p = Point::from_polar(10.0, 3.14159 / 4.0)
```

4.4.3 Generic Implementations

Implementations can be generic:

```
// ... (illustrative - with Pair defined above)
impl<T, U> Pair<T, U> {
    fn new(first: T, second: U) -> Pair<T, U> {
        Pair { first: first, second: second }
    }

    fn get_first(&self) -> &T {
        &self.first
    }
}
```

Specialized implementations can use where clauses:

```
// ... (with Pair defined above)
impl<T> Pair<T, T> where T: Clone {
    fn duplicate_first(self) -> Option<Pair<T, T>> {
        match self.first.clone() {
            Option::Some(copy) => {
                return Option::Some(Pair { first: self.first,
                    second: copy })
            },
            Option::None => {
                return Option::None
            },
        }
    }
}
```

4.4.4 Multiple Impl Blocks

A type can have multiple `impl` blocks:

```
impl Point {
    fn new(x: i32, y: i32) -> Point {
        Point { x: x, y: y }
    }
}
```



```

}

impl Point {
    fn distance(&self) -> f32 {
        sqrt((self.x * self.x + self.y * self.y) as f32)
    }
}

```

This is useful for organizing code or when implementing different traits.

4.5 Enums

Enumerations define a type by enumerating its possible variants. Each variant can carry associated data.

4.5.1 Unit Variants

Simple enums list named constants:

```

enum Direction {
    North,
    South,
    East,
    West,
}

```

Usage:

```

// ... (with Direction defined above)
let heading = Direction::North

```

4.5.2 Tuple Variants

Variants can carry data as tuples:

```

enum Command {
    Quit,
    Move(i32, i32),
    Write(i32),
    ChangeColor(u8, u8, u8),
}

```

Usage:

```

// ... (with Command defined above)
let cmd1 = Command::Quit
let cmd2 = Command::Move(10, 20)
let cmd3 = Command::ChangeColor(255, 0, 0)

```

Each variant can have a different number and type of fields.

4.5.3 Generic Enums

Enums can be parameterized with type parameters:

```
enum Option<T> {
    Some(T),
    None,
}

enum Result<T, E> {
    Ok(T),
    Err(E),
}
```

These are the foundation of Novus's error handling and null safety (covered in detail later).

4.5.4 Enum Limitations

Important: Novus enums support only *unit variants* and *tuple variants*. Struct-style variants are **not supported**:

```
// NOT SUPPORTED - struct variants (example only - invalid syntax)
enum Message {
    Move { x: i32, y: i32 }, // ... (compile error!)
}
```

Instead, use tuple variants:

```
enum Message {
    Move(i32, i32),
    Quit,
}
```

If you need named fields, define a separate struct:

```
struct MoveData {
    x: i32,
    y: i32,
}

enum Message {
    Move(MoveData),
    Quit,
}
```

4.6 Traits

Traits define shared behavior across types. They're similar to interfaces in other languages but more powerful.

4.6.1 Defining Traits

A trait declares a set of methods that implementing types must provide:

```
trait Clone {
    fn clone(&self) -> Option<Self> // Returns Option because
                                    // allocation can fail
}
```

```

trait Eq {
    fn eq(&self, other: &Self) -> bool
}

trait Hash {
    fn hash(&self) -> u32
}

trait Display {
    fn fmt(&self, f: &var Formatter) -> bool
}

```

4.6.2 Default Implementations

Traits can provide default implementations:

```

trait Printable {
    fn print(&self) {
        // Default implementation
        println("Printable object")
    }

    fn debug_print(&self) // Must be implemented
}

```

Types implementing the trait can use the default or override it.

4.6.3 Implementing Traits

Use `impl TraitName for TypeName` to implement a trait:

```

impl Clone for Point {
    fn clone(&self) -> Option<Point> {
        return Option::Some(Point { x: self.x, y: self.y })
    }
}

impl Eq for Point {
    fn eq(&self, other: &Point) -> bool {
        return self.x == other.x && self.y == other.y
    }
}

```

Once implemented, you can call trait methods:

```

let p1 = Point { x: 10, y: 20 }
let p2 = Point { x: 30, y: 40 }

// Clone returns Option<T> because allocation can fail
match p1.clone() {
    Option::Some(copy) => {
        // use copy
    },
    Option::None => {
        // handle allocation failure
    }
}

```

```

    }
}

// Eq provides equality comparison
if p1.eq(&p2) {
    // points are equal
}

```

4.6.4 Built-in Traits

Novus provides several important built-in traits:

Trait	Purpose
Drop	Custom cleanup when value goes out of scope
Clone	Explicit duplication (returns <code>Option<Self></code>)
Copy	Marker trait for implicitly copyable types
Eq	Equality comparison
Hash	Hashing for use in collections
From<T>	Type conversion from T (method: <code>convert</code>)
Default	Provides a default value
Display	User-facing formatted output
Debug	Programmer-facing debug output

Table 4.3: Core Built-in Traits

The Drop Trait

Drop provides custom cleanup logic:

```

struct FileHandle {
    fd: i32,
}

impl Drop for FileHandle {
    fn drop(&var self) {
        close_file(self.fd)
    }
}

```

When a `FileHandle` goes out of scope, `drop` is called automatically.

The From Trait

`From<T>` enables type conversions. The method is named `convert`:

```

impl From<i32> for Point {
    fn convert(value: i32) -> Point {
        return Point { x: value, y: value }
    }
}

```

Usage:

```

// ... (with Point and From impl defined above)

```

```
let p = Point::convert(42)
```

4.6.5 Trait Bounds

Traits can be used as bounds on generic parameters:

```
fn are_equal<T>(a: &T, b: &T) -> bool where T: Eq {
    return a.eq(b)
}

fn clone_if_equal<T>(a: &T, b: &T) -> Option<T> where T: Clone +
Eq {
    if a.eq(b) {
        return a.clone() // clone() already returns Option<T>
    }
    return Option::None
}
```

The first function requires `T` to implement `Eq`. The second requires both `Clone` and `Eq`, using `+` to combine bounds.

4.7 Pattern Matching

Pattern matching is a powerful feature for destructuring and inspecting data. It's the primary way to work with enums and extract data from composite types.

4.7.1 Pattern Types

Novus supports several pattern forms:

Literal Patterns

Match exact values:

```
42           // Matches the number 42
true         // Matches boolean true
"hello"      // Matches the string "hello"
null         // Matches null pointer
```

Identifier Patterns

Bind a value to a name:

```
x           // Binds value to variable x
count       // Binds value to variable count
```

Wildcard Pattern

Ignore a value:

```
-           // Matches anything, discards value
```

Variant Patterns

Destructure enum variants:

```
// Pattern syntax examples (illustrative)
Option::Some(x)           // Matches Some, binds inner value to x
Option::None              // Matches None
Command::Move(x, y)       // Matches Move, binds coordinates
Result::Err(e)            // Matches Err, binds error value to e
```

Or Patterns

Match multiple patterns:

```
1 | 2 | 3                 // Matches 1, 2, or 3
Direction::North | Direction::South // Matches either variant
```

Reference Patterns

Match through references:

```
&x           // Matches a reference, binds referent to x
&42          // Matches reference to 42
```

Guarded Patterns

Add conditional guards:

```
x if x > 0           // Matches any value > 0
Some(x) if x < 100    // Matches Some with value < 100
```

4.7.2 Match Expressions

The `match` expression compares a value against patterns:

```
match value {
  Pattern1 => expression1,
  Pattern2 => expression2,
  Pattern3 if guard => expression3,
  _ => default_expression,
}
```

Each arm consists of a pattern, optional guard, and an expression. The first matching pattern's expression is evaluated.

Basic Matching

```
let number = 42
let description = match number {
  0 => "zero",
  1 => "one",
  2 | 3 | 5 | 7 => "small prime",
```

```

42 => "the answer",
_   => "something else",
}

```

Matching Enums

```

let cmd = Command::Move(10, 20)

match cmd {
  Command::Quit => {
    println("Exiting")
    exit()
  },
  Command::Move(x, y) => {
    println("Moving to ({}, {})", x, y)
    move_cursor(x, y)
  },
  Command::Write(ch) => {
    println("Writing character {}", ch)
    write_char(ch)
  },
  Command::ChangeColor(r, g, b) => {
    set_color(r, g, b)
  },
}

```

Pattern Guards

Guards add conditional logic to patterns:

```

let number = 15

match number {
  x if x < 0 => println("negative"),
  x if x == 0 => println("zero"),
  x if x < 10 => println("small positive"),
  x if x < 100 => println("medium positive"),
  _ => println("large positive"),
}

```

Exhaustiveness

Match expressions must be *exhaustive* — they must cover all possible values:

```

// ERROR: non-exhaustive
match direction {
  Direction::North => "north",
  Direction::South => "south",
  // Missing East and West!
}

// OK: exhaustive with wildcard
match direction {
  Direction::North => "north",
}

```

```
    Direction::South => "south",
    _ => "other",
}

// OK: exhaustive with all variants
match direction {
    Direction::North => "north",
    Direction::South => "south",
    Direction::East  => "east",
    Direction::West  => "west",
}
```

The compiler verifies exhaustiveness and reports errors if any case is unhandled.

4.7.3 If Let Expressions

Novus provides `if let` for binding values that pass a truthiness check. This is useful for checking pointers for null or integers for zero:

```
let ptr: *u8 = get_pointer()

if let p = ptr {
    // p is bound to ptr (non-null)
    use_pointer(p)
} else {
    // ptr was null
}
```

The condition succeeds if the value is “truthy” (non-null for pointers, non-zero for integers). The bound variable is available in the `if` block.

For `Option` and `Result` types, use `match` for pattern matching:

```
match opt {
    Option::Some(x) => {
        println("Got value: {}", x)
    },
    Option::None => {
        println("No value")
    },
}
```

4.8 Option and Result

Novus eliminates null pointer errors and exception-based error handling by using two core types: `Option<T>` and `Result<T, E>`.

4.8.1 Option<T>

`Option<T>` represents an optional value — either `Some(T)` or `None`:

```
enum Option<T> {
    Some(T),
    None,
}
```


Use `Option` when a value might be absent:

```
fn find_user(id: u32) -> Option<User> {
    if user_exists(id) {
        Option::Some(get_user(id))
    } else {
        Option::None
    }
}
```

Working with Option

The preferred way to work with `Option` is pattern matching:

```
let opt = Option::Some(42)

// Pattern matching (preferred - safe and explicit)
match opt {
    Option::Some(x) => println("Value: %ld", x),
    Option::None => println("No value"),
}
```

For quick checks, use the `is_some()` and `is_none()` methods. Note that these methods **consume** the `Option`:

```
let opt1 = Option::Some(42)
if opt1.is_some() {
    println("Has value")
}
// opt1 is consumed - cannot use it again

let opt2 = Option::None
if opt2.is_none() {
    println("No value")
}
```

To extract the value directly:

```
let opt = Option::Some(42)

// unwrap_or (safe - provides default)
let x = opt.unwrap_or(0) // Returns 42, or 0 if None

// unwrap (panics on None - use sparingly)
let opt2 = Option::Some(100)
let y = opt2.unwrap() // Returns 100
```

Option Methods

Common `Option` methods. All methods consume `self` unless otherwise noted:

Note: Because these methods consume `self`, you cannot call multiple methods on the same `Option` value. Use pattern matching when you need to inspect and then use the value.

Method	Description
<code>is_some(self) -> bool</code>	Returns <code>true</code> if <code>Some</code> (consumes)
<code>is_none(self) -> bool</code>	Returns <code>true</code> if <code>None</code> (consumes)
<code>unwrap(self) -> T</code>	Returns value, panics on <code>None</code>
<code>unwrap_or(self, T) -> T</code>	Returns value or default
<code>ok_or(self, E) -> Result<T,E></code>	Converts to <code>Result</code>

Table 4.4: Common Option Methods

4.8.2 Result<T, E>

`Result<T, E>` represents either success (`Ok(T)`) or failure (`Err(E)`):

```
enum Result<T, E> {
    Ok(T),
    Err(E),
}
```

Use `Result` for operations that can fail:

```
fn parse_number(s: &str) -> Result<i32, ParseError> {
    if is_valid_number(s) {
        Result::Ok(parse(s))
    } else {
        Result::Err(ParseError::InvalidFormat)
    }
}
```

Working with Result

The preferred way to work with `Result` is pattern matching:

```
let result = parse_number("42")

// Pattern matching (preferred - handles both cases)
match result {
    Result::Ok(n) => println("Number: %ld", n),
    Result::Err(e) => println("Parse failed"),
}

// Another match if you only care about Ok
match result {
    Result::Ok(n) => println("Got number: %ld", n),
    Result::Err(_) => {},
}
```

For quick checks, `is_ok()` and `is_err()` consume the result:

```
let r1 = parse_number("42")
if r1.is_ok() {
    println("Parse succeeded")
}
// r1 is consumed

let r2 = parse_number("bad")
if r2.is_err() {
```

```
println("Parse failed")
}
```

To extract values directly:

```
let result = parse_number("42")

// unwrap_or (safe - provides default)
let n = result.unwrap_or(0) // Returns parsed value or 0

// unwrap (panics on Err - use sparingly)
let result2 = parse_number("100")
let m = result2.unwrap()
```

Result Methods

Common Result methods. All methods consume `self`:

Method	Description
<code>is_ok(self) -> bool</code>	Returns <code>true</code> if <code>Ok</code> (consumes)
<code>is_err(self) -> bool</code>	Returns <code>true</code> if <code>Err</code> (consumes)
<code>unwrap(self) -> T</code>	Returns value, panics on <code>Err</code>
<code>unwrap_or(self, T) -> T</code>	Returns value or default
<code>ok(self) -> Option<T></code>	Converts to <code>Option</code> , discards error
<code>err(self) -> Option<E></code>	Gets error as <code>Option</code>

Table 4.5: Common Result Methods

4.8.3 The ? Operator

The `?` operator provides early return on error:

```
fn read_number_from_file(path: &str) -> Result<i32, Error> {
    let file = open_file(path)? // Returns Err early
    let contents = read_file(&file)? // Returns Err early
    let number = parse_number(&contents)? // Returns Err early
    Result::Ok(number)
}
```

The `?` operator:

1. Checks if the `Result` is `Err`
2. If `Err`, returns the error immediately
3. If `Ok`, unwraps the value and continues

This is equivalent to:

```
fn read_number_from_file(path: &str) -> Result<i32, Error> {
    let file = match open_file(path) {
        Result::Ok(f) => f,
        Result::Err(e) => return Result::Err(e),
    }

    let contents = match read_file(&file) {
```

```
        Result::Ok(c) => c,
        Result::Err(e) => return Result::Err(e),
    }

    let number = match parse_number(&contents) {
        Result::Ok(n) => n,
        Result::Err(e) => return Result::Err(e),
    }

    Result::Ok(number)
}
```

The `?` operator works with both `Result` and `Option`:

```
fn get_first_element(vec: &Vec<i32>) -> Option<i32> {
    let first = vec.get(0)? // Returns None early if empty
    Option::Some(*first * 2)
}
```

4.9 Generics

Generics enable code reuse by parameterizing types and functions with type variables.

4.9.1 Generic Functions

Functions can be generic over types:

```
fn identity<T>(value: T) -> T {
    value
}
```

Usage:

```
// ... (with identity defined above)
let x = identity(42)           // T = i32
let s = identity("hello")     // T = &str
```

Multiple type parameters:

```
fn pair<T, U>(first: T, second: U) -> (T, U) {
    (first, second)
}
```

Usage:

```
// ... (with pair defined above)
let p = pair(42, "hello")     // T = i32, U = &str
```

4.9.2 Generic Structs

Structs can be generic (as shown earlier):

```
struct Vec<T> {
    ptr: *T,
```

```

    len: u32,
    capacity: u32,
}

```

Implementation:

```

// ... (with Vec defined above)
impl<T> Vec<T> {
    fn new() -> Vec<T> {
        return @zeroed(Vec<T>)
    }

    fn push(&var self, value: T) {
        // ... Implementation
    }

    fn get(&self, index: u32) -> Option<&T> {
        if index < self.len {
            Option::Some(&self.ptr[index])
        } else {
            Option::None
        }
    }
}

```

4.9.3 Const Generic Parameters

Functions and types can be parameterized with compile-time constants:

```

fn create_array<const N: u32>() -> [i32; N] {
    [0; N]
}

```

Usage:

```

// ... (with create_array defined above)
let arr1 = create_array::<10>() // [i32; 10]
let arr2 = create_array::<100>() // [i32; 100]

```

Structs with const generics:

```

struct Matrix<const ROWS: u32, const COLS: u32> {
    data: [f32; ROWS * COLS],
}

```

Const parameters must be integer types and known at compile time.

4.9.4 Where Clauses

Generic parameters can be constrained with `where` clauses:

```

fn hash_clone<T>(value: &T) -> u32 where T: Clone + Hash {
    let copy = value.clone()
    return copy.hash()
}

```

Where clauses can appear on:

- Function definitions
- Struct definitions
- Impl blocks
- Trait definitions

Multiple constraints with +:

```
fn process<T>(value: T) where T: Clone + Eq + Hash {  
    // T must implement Clone, Eq, and Hash  
}
```

4.9.5 Turbofish Syntax

Type arguments can be specified explicitly using turbofish syntax `::<>`:

```
let vec = Vec::<i32>::new()  
let result = parse::<u32>("42")  
let pair = Pair::<i32, bool>::new(10, true)
```

This is necessary when the compiler cannot infer type arguments:

```
// Ambiguous - compiler doesn't know T  
let vec = Vec::new() // ERROR  
  
// Explicit - compiler knows T = i32  
let vec = Vec::<i32>::new() // OK
```

4.9.6 Generic Constraints in Practice

Here's a complete example showing generic constraints:

```
trait Comparable {  
    fn compare(&self, other: &Self) -> i32  
}
```

Function with trait constraint:

```
fn find_max<T>(items: &[T]) -> Option<T> where T: Comparable {  
    if items.len() == 0 {  
        return Option::None  
    }  
  
    var max = &items[0]  
    var i = 1  
  
    while i < items.len() {  
        if items[i].compare(max) > 0 {  
            max = &items[i]  
        }  
        i = i + 1  
    }  
}
```

```

    Option::Some(max)
}

```

Implementing the trait:

```

// ... (with Comparable defined above)
impl Comparable for i32 {
    fn compare(&self, other: &i32) -> i32 {
        if *self < *other {
            return -1
        } else if *self > *other {
            return 1
        } else {
            return 0
        }
    }
}

```

Usage:

```

// ... (with Comparable, find_max, and impl above)
let numbers = [10, 5, 20, 15]
let max = find_max(&numbers) // Some(20)

```

4.10 Type Coercions and Casts

Novus has strict type checking with limited implicit conversions.

4.10.1 Integer Widening

Smaller integer types can be implicitly widened to larger ones in some contexts:

```

let x: i8 = 42
let y: i32 = x // Implicit widening OK (i8 -> i32)

```

4.10.2 Integer Narrowing

Narrowing requires explicit casts:

```

let x: i32 = 1000
let y: i8 = x as i8 // Explicit cast required (truncates)

```

4.10.3 Reference to Pointer

References can be converted to raw pointers:

```

let x = 42
let r: &i32 = &x
let p: *i32 = r as *i32 // Reference to pointer

```

4.10.4 Array to Pointer

Arrays decay to pointers in some contexts:

```
let arr = [1, 2, 3, 4, 5]
let ptr: *i32 = arr as *i32 // Array to pointer
```

4.10.5 Explicit Casts

Use `as` for explicit type conversions:

```
let x: i32 = 42
let y: f32 = x as f32 // Int to float
let z: u32 = x as u32 // Signed to unsigned
let b: u8 = (x & 0xFF) as u8 // Truncation
```

Always prefer explicit casts to make your intent clear.

4.11 Type Inference

Novus uses type inference to reduce verbosity while maintaining static typing.

4.11.1 Local Variable Inference

The `let` keyword allows the compiler to infer types:

```
let x = 42 // Inferred as i32
let s = "hello" // Inferred as &str
let v = Vec::new() // ERROR: cannot infer T
let v = Vec::<i32>::new() // OK: explicit type argument
```

You can always specify types explicitly:

```
let x: i32 = 42
let s: &str = "hello"
let v: Vec<i32> = Vec::new()
```

4.11.2 Function Return Type Inference

Function return types must be explicitly declared:

```
// ERROR: missing return type
fn add(a: i32, b: i32) {
    a + b
}

// OK: explicit return type
fn add(a: i32, b: i32) -> i32 {
    a + b
}
```


4.11.3 Generic Type Inference

Generic type arguments are often inferred from usage:

```
var vec = Vec::<i32>::new()
vec.push(42) // T inferred as i32

let opt = Option::Some(42) // T inferred as i32
```

When inference fails, use turbofish syntax:

```
let result = parse::<i32>("42")
```

4.12 Summary

Novus's type system provides:

- **Structs** for grouping related data with private fields
- **Enums** for sum types with unit and tuple variants
- **Traits** for shared behavior and polymorphism
- **Pattern matching** for safe data inspection and destructuring
- **Option and Result** for null safety and error handling
- **Generics** for code reuse with type and const parameters
- **Type inference** to reduce verbosity
- **Explicit casts** for type conversions

These features combine to create a type system that catches errors at compile time while remaining expressive and efficient. The emphasis on explicitness (no hidden allocations, no implicit conversions) and safety (no null, errors as values) makes Novus code predictable and robust.

In the next chapter, we'll explore Novus's memory management model, including ownership, borrowing, and the interaction with AmigaOS memory pools.

Chapter 5

Memory Management

“Memory management is the heart of every systems program.” — Brian Kernighan

The Amiga’s cooperative multitasking environment demands careful memory management. Unlike modern systems with virtual memory and garbage collection, the Amiga uses flat physical memory shared between all tasks. A memory leak or dangling pointer affects the entire system, not just one process.

Novus provides a memory management model that combines the safety of modern languages with the control needed for 68k systems programming. This chapter covers ownership and move semantics, references and borrowing, the Drop trait for automatic cleanup, defer blocks for explicit resource management, and the standard library’s memory abstractions.

5.1 The Amiga Memory Model

Before diving into Novus specifics, it’s essential to understand AmigaOS memory management.

5.1.1 Memory Types

The Amiga has distinct memory types with different characteristics:

Chip RAM Accessible by both CPU and custom chips (Agnus, Denise, Paula). Required for graphics, audio, and DMA operations. Limited to 2MB on most systems.

Fast RAM Accessible only by CPU. Faster access since no DMA contention. Can be much larger (up to 2GB on 68040/060 systems).

Public Memory Safe for inter-task communication. Other tasks may access it.

When allocating memory, you specify flags to request the appropriate type:

```
// ... (with memory flags imported from std::ffi::amiga_consts)
// Chip RAM for graphics buffers
let bitmap = MemoryBlock::alloc(size, MEMF_CHIP | MEMF_CLEAR)

// Fast RAM for general data (falls back to Chip if unavailable)
let buffer = MemoryBlock::alloc(size, MEMF_FAST | MEMF_PUBLIC)

// Any memory (system chooses)
let data = MemoryBlock::alloc(size, MEMF_PUBLIC | MEMF_CLEAR)
```

5.1.2 Memory Flag Reference

5.1.3 No Size Parameter on Free

A critical feature of Novus memory types: you don’t pass the size when freeing memory. Unlike raw `FreeMem(ptr, size)` calls, Novus types track size automatically:

Flag	Description
MEMF_ANY	Any memory type (value: 0)
MEMF_PUBLIC	Memory safe for inter-task access
MEMF_CHIP	Chip RAM (required for DMA)
MEMF_FAST	Fast RAM (CPU-only, faster)
MEMF_CLEAR	Zero-fill the allocated memory
MEMF_LARGEST	Allocate largest available block
MEMF_LOCAL	Memory local to the current task
MEMF_24BITDMA	Memory in 24-bit address space
MEMF_KICK	Memory for Kickstart/resident modules
MEMF_REVERSE	Allocate from high addresses down
MEMF_NO_EXPUNGE	Prevent memory from being expunged
MEMF_TOTAL	Query total memory (not for allocation)

Table 5.1: AmigaOS Memory Flags

```
// AmigaOS C style - error-prone
void* ptr = AllocMem(1024, MEMF_PUBLIC);
// ... must remember size is 1024 ...
FreeMem(ptr, 1024); // Wrong size = memory corruption!

// Novus style - size tracked automatically
var block = MemoryBlock::alloc(1024, MEMF_PUBLIC)?
// ... no need to track size ...
block.free() // Size remembered internally!
```

This eliminates an entire class of bugs common in AmigaOS programming.

5.2 Ownership and Move Semantics

Every value in Novus has a single *owner*—the variable that holds it. When the owner goes out of scope, the value is dropped (cleaned up). This ownership model prevents double-free and use-after-free bugs.

5.2.1 Move by Default

When you pass a value to a function that takes ownership, the value is *moved*:

```
fn consume(consuming s: String) {
    // s is owned by this function
    // it will be dropped when consume() returns
}

fn main() -> i32 {
    let x = String::new()

    consume(x) // x is moved into consume()

    // ERROR: x has been moved, cannot use it
    // let len = x.len() // Compile error!

    return 0
}
```

The `consuming` keyword explicitly marks parameters that take ownership. After a move, the original variable is invalid—the compiler prevents you from using it.

5.2.2 Borrowing to Avoid Moves

To use a value without transferring ownership, *borrow* it with a reference:

```
fn print_length(s: &String) {
    // s is borrowed immutably
    let len = s.len()
    println("Length: {}", len)
}

fn main() -> i32 {
    let x = String::new_from_str("Hello")?

    print_length(&x) // Borrow x

    // x is still valid here
    let len = x.len() // OK!

    return 0
}
```

The `&` operator creates a reference—a pointer that borrows the value without taking ownership.

5.2.3 The Borrow Checker

Novus includes a borrow checker that tracks moves and prevents use-after-move errors:

```
fn test_move_tracking() {
    let s = String::new()

    if some_condition {
        consume(s) // s moved in this branch
    }

    // ERROR: s may have been moved
    // let len = s.len() // Compile error!
}
```

The borrow checker is control-flow sensitive—it tracks moves across `if`, `match`, and `while` branches.

Important: Novus’s borrow checker tracks *moves* and *reference lifetimes*. References cannot outlive their source, and converting references to raw pointers requires `unsafe`. Novus does not enforce reference *exclusivity* (you can have multiple mutable references), but it does prevent dangling references. See Section 5.3.5 for details.

5.3 References

References are non-null pointers that borrow values without taking ownership.

5.3.1 Immutable References

An immutable reference `&T` provides read-only access:

```
fn read_point(p: &Point) -> i32 {
    return p.x + p.y // Can read fields
}
```

Usage:

```
// ... (with Point and read_point defined above)
let pt = Point { x: 10, y: 20 }
let sum = read_point(&pt) // Borrow pt immutably
```

5.3.2 Mutable References

A mutable reference `&var T` allows modification:

```
fn scale_point(p: &var Point, factor: i32) {
    p.x = p.x * factor
    p.y = p.y * factor
}
```

Usage:

```
// ... (with Point and scale_point defined above)
var pt = Point { x: 10, y: 20 }
scale_point(&var pt, 2) // Borrow pt mutably
// pt is now { x: 20, y: 40 }
```

Note that the variable being borrowed must be declared with `var` (mutable).

5.3.3 Self Parameter Forms

Methods use special self parameter syntax:

```
impl Point {
    // Immutable borrow - reads only
    fn magnitude(&self) -> i32 {
        return self.x * self.x + self.y * self.y
    }

    // Mutable borrow - can modify
    fn translate(&var self, dx: i32, dy: i32) {
        self.x = self.x + dx
        self.y = self.y + dy
    }

    // Consuming - takes ownership
    fn into_tuple(consuming self) -> (i32, i32) {
        return (self.x, self.y)
    }
}
```

5.3.4 References vs Raw Pointers

References (`&T`, `&var T`) are guaranteed non-null. Raw pointers (`*T`) can be null and require `unsafe` for dereferencing:

```

let x = 42
let ref: &i32 = &x           // Reference - guaranteed valid
let ptr: *i32 = &x as *i32   // Raw pointer - could be null

// References can be used directly
let y = *ref // OK

// Raw pointers require unsafe
unsafe {
    let z = *ptr // Must be in unsafe block
}

```

Use references when possible. Reserve raw pointers for FFI and low-level operations.

5.3.5 Reference Lifetimes

References in Novus have *lifetimes*—they are tied to the scope of the value they borrow. The compiler tracks these lifetimes and prevents you from using a reference after its source has been dropped. This eliminates an entire class of bugs: dangling pointers.

Why Lifetimes Matter

Consider this common pattern when working with AmigaOS screens:

```

fn dangerous() {
    let rp: &RastPort
    {
        let screen = ScreenHandle::lores("Demo", 5)?
        rp = screen.rastport() // rp borrows from screen
    } // screen dropped here!
    SetAPen(rp, 2) // BUG: rp is dangling!
}

```

Without lifetime checking, this code would compile and crash at runtime—a Guru Meditation. With Novus’s lifetime tracking, the compiler catches this at compile time:

```

error[E0597]: ‘screen’ does not live long enough
--> example.novus:5:14
|
4 |         let screen = ScreenHandle::lores("Demo", 5)?
|           ^^^^^^^ borrowed value
5 |         rp = screen.rastport()
|           ----- borrow occurs here
6 |     }
|     ~ ‘screen’ dropped here while still borrowed

```

Lifetime Rules

Novus enforces these lifetime rules at compile time:

1. **References cannot outlive their source.** A reference must not be used after the value it borrows from goes out of scope.
2. **References cannot be stored in struct fields.** This is a v1 limitation—lifetime parameters on structs are planned for a future version. Use raw pointers if you need to store references in structs.

3. **Method return lifetimes are inferred.** When a method returns a reference, it's tied to `&self` (or the single reference parameter if no `self`). Multiple reference parameters without `&self` is ambiguous and rejected.
4. **Converting reference to pointer requires `unsafe`.** To escape lifetime tracking, you must explicitly use an `unsafe` block.

Error Reference

E0597: Does not live long enough Triggered when a reference outlives the value it borrows from.

```
fn bad() {
    let r: &i32
    {
        let x: i32 = 42
        r = &x           // x is in inner scope
    }                   // x dropped here
    let y = *r           // ERROR: x doesn't live long enough
}
```

Fix: Move the reference into the same scope as its source, or restructure to avoid storing the reference.

E0106: Cannot contain reference / Cannot infer lifetime Triggered when storing a reference in a struct field, or when lifetime inference is ambiguous.

```
// ERROR: struct cannot contain reference
struct BadCache {
    rp: &RastPort // Not allowed in v1
}

// ERROR: cannot infer lifetime
fn pick(a: &i32, b: &i32) -> &i32 { // Which input?
    return a
}
```

Fix for structs: Use a raw pointer (`*RastPort`) and manage safety manually.

Fix for functions: Add `&self` parameter, or reduce to single reference parameter.

E0133: Requires `unsafe` Triggered when converting a reference to a raw pointer outside an `unsafe` block.

```
let x: i32 = 42
let r: &i32 = &x
let p: *i32 = (*i32)r // ERROR: requires unsafe
```

Fix: Wrap in `unsafe` if you can guarantee pointer validity:

```
let p: *i32 = unsafe { (*i32)r } // OK
```

E0515: Cannot return reference to local Triggered when returning a reference to a local variable.


```
fn bad() -> &i32 {
    let x: i32 = 42
    return &x // ERROR: x will be dropped when function returns
}
```

Fix: Return an owned value, or take the value as a reference parameter and return a reference tied to it.

Escape Hatches

Sometimes you need to break the rules—especially when interfacing with AmigaOS libraries that manage their own memory. Novus provides explicit escape hatches:

Raw pointers (**T*) have no lifetime tracking. Use them for FFI and cases where you manually guarantee validity:

```
// FFI function returns pointer with unknown lifetime
extern fn GetSomePointer() -> *Thing

fn use_ffi() {
    let ptr: *Thing = GetSomePointer()
    unsafe {
        // You guarantee ptr is valid
        do_something(*ptr)
    }
}
```

Unsafe blocks let you convert references to pointers when needed:

```
fn pass_to_legacy_api(data: &MyData) {
    let ptr: *MyData = unsafe { (*MyData)data }
    LegacyFunction(ptr) // You guarantee data outlives this call
}
```

The philosophy is *power with guardrails*: safe by default, explicit opt-out when you need control. The `unsafe` keyword marks exactly where you’re taking responsibility for memory safety.

5.4 The Drop Trait

The Drop trait enables automatic cleanup when values go out of scope—Novus’s form of RAI (Resource Acquisition Is Initialization).

5.4.1 Implementing Drop

```
from std::core import Drop

struct FileHandle {
    fd: i32,
}

impl Drop for FileHandle {
    fn drop(&var self) {
        // Called automatically when FileHandle goes out of scope
        close_file(self.fd)
    }
}
```

```
    }  
}  
  
fn main() -> i32 {  
    {  
        let file = FileHandle { fd: open_file("data.txt") }  
        // ... use file ...  
    } // file.drop() called automatically here  
  
    return 0  
}
```

5.4.2 Drop Execution Order

Drop is called in reverse declaration order (LIFO):

```
fn main() -> i32 {  
    let a = Resource::new(1) // Created first  
    let b = Resource::new(2) // Created second  
    let c = Resource::new(3) // Created third  
  
    return 0  
    // Drop order: c, b, a (reverse)  
}
```

5.4.3 Move Prevents Double Drop

When a value is moved, the compiler deactivates its drop:

```
fn take_ownership(consuming r: Resource) {  
    // r is dropped here when function returns  
}  
  
fn main() -> i32 {  
    let r = Resource::new(42)  
    take_ownership(r) // r is moved  
  
    return 0  
    // r's drop is NOT called - it was moved  
}
```

This prevents double-free bugs.

5.4.4 Drop for Enums

Enums with data-carrying variants have Drop called on their payloads:

```
fn main() -> i32 {  
    {  
        let opt: Option<String> = Option::Some(String::new())  
        // opt goes out of scope  
    } // String inside is automatically dropped  
  
    return 0  
}
```

5.5 Defer Blocks

While Drop provides automatic cleanup, sometimes you need explicit control over cleanup order or want to clean up resources that don't implement Drop. The `defer` statement handles this.

5.5.1 Basic Defer

A defer block executes when the enclosing function returns:

```
fn process_file() -> i32 {
    let fd = open_file("data.txt")

    defer {
        close_file(fd) // Guaranteed to run
    }

    // ... process file ...
    // Even if we return early, defer runs

    if error_occurred {
        return -1 // defer still runs
    }

    return 0 // defer runs here too
}
```

5.5.2 LIFO Execution Order

Multiple defers execute in Last-In-First-Out order:

```
fn main() -> i32 {
    var resource1 = open_resource(1)
    defer {
        cleanup_resource(1) // Runs SECOND
    }

    var resource2 = open_resource(2)
    defer {
        cleanup_resource(2) // Runs FIRST
    }

    return 0
    // Execution: cleanup(2), then cleanup(1)
}
```

This matches the natural pattern of “last opened, first closed.”

5.5.3 Return Value Computed Before Defer

The return value is computed *before* defer blocks execute:

```
fn compute() -> i32 {
    var result = 42

    defer {
```

```
        result = 0 // This does NOT change the return value!
    }

    return result // Returns 42, not 0
}
```

This is important: defer blocks cannot modify the return value.

5.5.4 Defer vs Drop

Use Drop when:

- Cleanup is tied to a type's lifetime
- You want automatic cleanup without explicit code
- The resource is always cleaned up the same way

Use defer when:

- Cleanup depends on runtime conditions
- You need cleanup for values that don't implement Drop
- You want explicit control over cleanup order
- The cleanup logic is function-specific

5.6 The Using Statement

The `using` statement combines resource acquisition with guaranteed cleanup:

```
struct Resource {
    value: i32
}

impl Drop for Resource {
    fn drop(&var self) {
        // Cleanup
    }
}

fn create_resource(v: i32) -> Resource {
    return Resource { value: v }
}

fn main() -> i32 {
    using create_resource(42) {
        // Resource exists in this scope
        let x = 10
    } // Resource is dropped here

    return 0
}
```

This is syntactic sugar for creating a resource and ensuring it's dropped at block end.

5.7 RAI: Resource Acquisition Is Initialization

RAI is a programming pattern where resource lifetime is tied to object lifetime. When you acquire a resource (memory, file handle, hardware access), you wrap it in an object. When that object is destroyed, the resource is automatically released. Novus implements RAI through the `Drop` trait.

5.7.1 Why RAI Matters on the Amiga

The Amiga's cooperative multitasking makes RAI especially valuable:

- **No memory protection** — a leaked resource affects the entire system
- **Limited resources** — only 8 sprites, 4 audio channels, finite Chip RAM
- **Complex cleanup sequences** — closing a window requires freeing menus, gadgets, and screen locks in the right order
- **Early return paths** — error handling often needs cleanup at multiple exit points

RAI ensures resources are released even when functions return early due to errors. The standard library uses RAI extensively—over 70 types implement `Drop`.

5.7.2 The Three RAI Patterns

Novus uses three complementary RAI patterns throughout the standard library:

Handle Types

Handle types wrap OS resources and manage their complete lifecycle:

```
from std::ui::window import WindowHandle, WindowBuilder
from std::ui::screen import ScreenHandle

fn show_window() -> Result<(), IntuitionError> {
    // WindowHandle manages the Intuition window
    var window = WindowBuilder::new()
        .title("My App")
        .size(640, 480)
        .build()?

    // Use the window...
    window.draw_text(10, 20, "Hello, Amiga!")

    return Result::Ok(())
} // window.drop() closes the window automatically
```

Common handle types in the standard library:

Guard Types

Guard types provide temporary exclusive access to a shared resource. When the guard goes out of scope, the resource is released for others to use:

```
from std::graphics::blitter import BlitterGuard, own_blitter
from std::os::exec import SemaphoreHandle, SemaphoreGuard
```

Module	Handle Types
<code>std::ui</code>	<code>WindowHandle</code> , <code>ScreenHandle</code> , <code>BufferedWindow</code>
<code>std::graphics</code>	<code>BitMapHandle</code> , <code>FontHandle</code> , <code>SpriteHandle</code>
<code>std::audio</code>	<code>AudioChannel</code> , <code>SampleHandle</code> , <code>ModPlayer</code>
<code>std::os</code>	<code>TimerHandle</code> , <code>MsgPortHandle</code> , <code>SignalHandle</code>
<code>std::net</code>	<code>TcpStream</code> , <code>TcpListener</code> , <code>UdpSocket</code>
<code>std::thread</code>	<code>ThreadHandle</code> , <code>ProcessHandle</code>

Table 5.2: Common Handle Types

```
fn blit_safely() {
    // BlitterGuard gives exclusive blitter access
    match own_blitter() {
        Option::Some(guard) => {
            // Blitter is ours exclusively
            blit_copy(src, dst, width, height)

            // guard.drop() calls DisownBlitter()
        },
        Option::None => {
            // Couldn't get blitter
        }
    }
}

fn critical_update(sem: &SemaphoreHandle, data: &var SharedData) {
    // SemaphoreGuard holds the lock
    var lock = sem.obtain()

    // Safe to modify shared data
    data.counter = data.counter + 1
} // lock.drop() calls ReleaseSemaphore()
```

Guard types follow a common pattern:

- Created by calling an “obtain” or “lock” function
- Hold exclusive (or shared) access while they exist
- Release access in their `Drop` implementation
- Cannot be copied or cloned

Common guard types:

Builder Types

Builder types accumulate configuration, then create a resource. They clean up intermediate allocations if building fails:

```
from std::ui::mui import MUIAppBuilder, MUIApp

fn create_app() -> Option<MUIApp> {
    // Builder accumulates configuration
    var builder = MUIAppBuilder::new()
```

Guard Type	Purpose
BlitterGuard	Exclusive blitter access (OwnBlitter/DisownBlitter)
SemaphoreGuard	Exclusive semaphore lock (ObtainSemaphore/ReleaseSemaphore)
SharedSemaphoreGuard	Shared semaphore lock (read access)
CriticalSection	Task switching disabled (Forbid/Permit)
InterruptSection	Interrupts disabled (Disable/Enable)
RefreshGuard	Window refresh lock (BeginRefresh/EndRefresh)

Table 5.3: Guard Types in the Standard Library

```
builder.title("My MUI Application")
builder.version("$VER: MyApp 1.0")
builder.author("Developer")

// Build creates the MUI application
// Returns None if creation fails
let app = builder.build()?

return Option::Some(app)
} // If build() fails, builder.drop() cleans up

// Alternative: method chaining with @chain attribute
fn create_app_chained() -> Option<MUIApp> {
    return MUIAppBuilder::new()
        .title("My App")
        .version("1.0")
        .build()
}
```

Builder types are especially useful for AmigaOS APIs that require many configuration options (MUI, GadTools, Intuition).

Common builder types:

Builder Type	Creates
MUIAppBuilder	MUI application objects
MUIGroupBuilder	MUI group objects
GadToolsMenuBuilder	GadTools menu strips
CopperListBuilder	Hardware copper lists
WindowBuilder	Intuition windows
ReActionBuilder	ReAction GUI elements

Table 5.4: Builder Types in the Standard Library

5.7.3 Implementing Your Own RAI Types

When wrapping AmigaOS resources, follow these patterns:

```
from std::core import Drop, Option
from std::ffi::exec import FreeMem
from std::ffi::amiga_consts import MEMF_ANY

// Handle pattern: wrap a resource pointer
```

```

struct MyResourceHandle {
    ptr: *MyResource,
    owned: bool,
}

impl MyResourceHandle {
    // Constructor acquires the resource
    pub fn new() -> Option<MyResourceHandle> {
        let ptr = allocate_my_resource()
        if ptr == null {
            return Option::None
        }
        return Option::Some(MyResourceHandle {
            ptr: ptr,
            owned: true,
        })
    }

    // into_raw() transfers ownership out
    pub fn into_raw(consuming self) -> *MyResource {
        // Caller now owns the pointer
        self.owned = false
        return self.ptr
    }
}

impl Drop for MyResourceHandle {
    fn drop(&var self) {
        // Only free if we still own it
        if self.owned && self.ptr != null {
            free_my_resource(self.ptr)
        }
    }
}

```

Key implementation points:

- Track ownership with a boolean flag
- Provide `into_raw()` to transfer ownership out
- Check for null before freeing
- Make constructors return `Option` or `Result`

5.7.4 RAII with Error Handling

RAII and the `?` operator work together beautifully:

```

fn complex_operation() -> Result<(), Error> {
    // Each resource is cleaned up if later steps fail
    var window = open_window()? // Auto-closed on error
    var screen = open_screen()? // Auto-closed on error
    var bitmap = alloc_bitmap()? // Auto-freed on error

    // If this fails, all three are cleaned up
    do_work(&window, &screen, &bitmap)?
}

```



```

    return Result::Ok(())
}
// Normal exit: all three dropped in reverse order
// Error exit: all acquired resources dropped

```

Without RAI, this would require deeply nested error handling or goto-based cleanup.

5.7.5 When NOT to Use RAI

RAI isn't always appropriate:

- **Borrowed resources** — if you're given a pointer you don't own, don't wrap it in a handle that frees it
- **Shared ownership** — if multiple owners need to keep a resource alive, use explicit reference counting
- **Async cleanup** — if cleanup must happen in a different context (like a callback), use explicit management
- **Performance-critical paths** — very hot loops might benefit from manual pooling

Coming from C

In C, you typically use `goto cleanup` patterns or deeply nested `if` statements to handle resource cleanup. Novus's RAI eliminates these patterns:

```

// C: Manual cleanup with goto
int do_work(void) {
    void *a = NULL, *b = NULL;
    int result = -1;

    a = AllocMem(100, MEMF_ANY);
    if (!a) goto cleanup;

    b = AllocMem(200, MEMF_ANY);
    if (!b) goto cleanup;

    // ... actual work ...
    result = 0;

cleanup:
    if (b) FreeMem(b, 200);
    if (a) FreeMem(a, 100);
    return result;
}

```

In Novus, the same code is much simpler:

```

// Novus: RAI handles cleanup
fn do_work() -> Result<(), ExecError> {
    var a = MemoryBlock::try_alloc(100, MEMF_ANY)?
    var b = MemoryBlock::try_alloc(200, MEMF_ANY)?

    // ... actual work ...

    return Result::Ok(())
}

```

```
} // Both freed automatically, in correct order
```

The `?` operator propagates errors, and `Drop` ensures cleanup—no goto, no nested ifs, no size tracking.

5.8 Standard Library Memory Types

Novus provides several memory types in `std::memory` that wrap AmigaOS allocation with safety and convenience.

5.8.1 MemoryBlock

`MemoryBlock` is the fundamental building block—raw bytes with automatic size tracking:

```
from std::memory::block import MemoryBlock
from std::ffi::amiga_consts import MEMF_FAST, MEMF_CLEAR

fn main() -> i32 {
    // Allocate 1024 bytes
    match MemoryBlock::alloc(1024, MEMF_FAST | MEMF_CLEAR) {
        Option::Some(var block) => {
            defer {
                block.free()
            }

            let ptr: *u8 = block.ptr()
            let size: u32 = block.size()

            // Use the memory...
            ptr[0] = 42
        },
        Option::None => {
            // Allocation failed
            return 1
        }
    }

    return 0
}
```

MemoryBlock Methods

`MemoryBlock` implements `Drop`, so it's freed automatically when it goes out of scope.

5.8.2 MemHandle

For common allocation patterns, `MemHandle` provides convenient factory methods with preset flags:

```
from std::memory::amiga import MemHandle

fn main() -> i32 {
    // Chip RAM for graphics/audio (MEMF_CHIP | MEMF_CLEAR |
    MEMF_PUBLIC)
```

Method	Description
<code>alloc(size, flags)</code>	Allocate raw memory, returns <code>Option<MemoryBlock></code>
<code>try_alloc(size, flags)</code>	Allocate with <code>Result<MemoryBlock, ExecError></code>
<code>ptr()</code>	Get pointer to memory
<code>size()</code>	Get size in bytes
<code>is_empty()</code>	Check if zero-size block
<code>into_raw()</code>	Take ownership of pointer, returns <code>(*u8, u32)</code>
<code>free()</code>	Free the memory
<code>resize(new_size, flags)</code>	Resize block, returns success <code>bool</code>

Table 5.5: MemoryBlock Methods

```

match MemHandle::new_chip(1024) {
  Option::Some(var chip_mem) => {
    defer { chip_mem.drop() }
    // Use chip memory for DMA operations
  },
  Option::None => return 1
}

// Fast RAM for general data (MEMF_FAST | MEMF_CLEAR |
// MEMF_PUBLIC)
match MemHandle::new_fast(2048) {
  Option::Some(var fast_mem) => {
    defer { fast_mem.drop() }
    // Use fast memory for CPU operations
  },
  Option::None => return 1
}

// Any available memory
match MemHandle::new_any(512) {
  Option::Some(var any_mem) => {
    // System chooses chip or fast
  },
  Option::None => return 1
}

return 0
}

```

`MemHandle` uses `AllocVec/FreeVec` internally, which stores the size in the allocation header. This is slightly less memory-efficient than `MemoryBlock` (4 bytes overhead) but provides the same convenience.

Use `MemHandle` for common cases; use `MemoryBlock` when you need full control over memory flags.

5.8.3 Allocation<T>

`Allocation<T>` provides type-safe bulk allocation with element count tracking:

```

from std::memory::block import Allocation
from std::ffi::amiga_consts import MEMF_CHIP, MEMF_CLEAR

```

```
fn main() -> i32 {  
    // Allocate space for 100 u32 values  
    match Allocation::<u32>::new(100, MEMF_CHIP | MEMF_CLEAR) {  
        Option::Some(var alloc) => {  
            defer {  
                alloc.free()  
            }  
  
            let count: u32 = alloc.count()           // 100  
            let bytes: u32 = alloc.size_bytes()      // 400  
            let ptr: *u32 = alloc.as_mut_ptr()  
  
            // Initialize elements  
            var i: u32 = 0  
            while i < count {  
                ptr[i] = i * 2  
                i = i + 1  
            }  
        },  
        Option::None => {  
            return 1  
        }  
    }  
  
    return 0  
}
```

`Allocation<T>` checks for overflow when computing `count * sizeof(T)`, preventing a common security vulnerability.

5.8.4 `Box<T>`

`Box<T>` allocates a single value on the heap:

```
from std::memory::block import Box  
  
struct BigData {  
    buffer: [u8; 4096],  
}  
  
fn main() -> i32 {  
    // Allocate BigData on heap instead of stack  
    match Box::new(BigData { buffer: [0; 4096] }) {  
        Option::Some(var boxed) => {  
            defer {  
                boxed.free()  
            }  
  
            let ptr: *BigData = boxed.get_mut()  
            ptr.buffer[0] = 42  
        },  
        Option::None => {  
            return 1  
        }  
    }  
  
    return 0  
}
```

```
}

```

Use `Box<T>` for:

- Large structs that don't fit on the stack
- Recursive data structures
- Values that need heap allocation with automatic cleanup

5.8.5 `Slice<T>`

`Slice<T>` is a fat pointer—a pointer plus length—for passing arrays without losing size information:

```
from std::memory::slice import Slice

fn sum(numbers: &Slice<i32>) -> i32 {
    var total: i32 = 0
    var i: u32 = 0

    while i < numbers.len() {
        match numbers.get(i) {
            Option::Some(n) => {
                total = total + n
            },
            Option::None => {}
        }
        i = i + 1
    }

    return total
}

fn main() -> i32 {
    let arr = [1, 2, 3, 4, 5]
    let slice = Slice::new(&arr as *i32, 5)

    let result = sum(&slice) // No need to pass count separately!

    return result
}
```

Slice Methods

Method	Description
<code>new(ptr, len)</code>	Create slice from pointer and length
<code>len()</code>	Get length
<code>is_empty()</code>	Check if empty
<code>get(index)</code>	Get element with bounds check, returns <code>Option<T></code>
<code>get_unchecked(index)</code>	Get element without bounds check (unsafe)

Table 5.6: Slice Methods

5.9 Copy and Clone Traits

Some types can be duplicated safely. Novus provides two traits for this.

5.9.1 The Copy Trait

Copy is a marker trait for types that can be implicitly copied:

Primitive types are implicitly Copy:

```
// Primitive types are Copy
let x: i32 = 42
let y = x    // x is copied, both x and y are valid
```

Structs can be marked as Copy if all fields are Copy:

```
struct Point {
    x: i32,
    y: i32,
}

impl Copy for Point {}
```

Usage:

```
// ... (with Point defined as Copy above)
let p1 = Point { x: 1, y: 2 }
let p2 = p1    // p1 is copied, both valid
```

Types that manage resources (like `String`, `Vec`, `MemoryBlock`) should NOT implement `Copy`.

5.9.2 The Clone Trait

`Clone` provides explicit deep copying. It returns `Option<Self>` because cloning types that allocate memory (like `Vec` or `String`) can fail:

```
// ... (with Point defined above)
impl Clone for Point {
    fn clone(&self) -> Option<Point> {
        return Option::Some(Point { x: self.x, y: self.y })
    }
}
```

Usage:

```
// ... (with Point and Clone impl above)
let p1 = Point { x: 1, y: 2 }

match p1.clone() {
    Option::Some(p2) => {
        // p2 is a deep copy of p1
    },
    Option::None => {
        // Clone failed - only happens for types that allocate
    },
}
```

For types that manage resources, `clone()` must allocate new memory, which can fail on a memory-constrained Amiga:

```
// Vec::clone() allocates new memory for the copy
let original = Vec::<i32>::new()?
// ... populate original ...

match original.clone() {
    Option::Some(copy) => {
        // copy is independent - modifying it doesn't affect
        // original
    },
    Option::None => {
        // Allocation failed - handle gracefully
    }
}
```

5.10 Common Patterns

5.10.1 Resource Acquisition with Error Handling

```
fn process_with_resources() -> Result<i32, Error> {
    // Acquire first resource
    var mem = MemoryBlock::try_alloc(1024, MEMF_PUBLIC)?
    defer {
        mem.free()
    }

    // Acquire second resource
    var file = open_file("data.txt")?
    defer {
        close_file(file)
    }

    // Use resources...

    return Result::Ok(0)
}
```

The `?` operator propagates errors, and `defer` ensures cleanup regardless of how the function exits.

5.10.2 Passing Arrays to Functions

```
fn process_buffer(data: &Slice<u8>) {
    var i: u32 = 0
    while i < data.len() {
        let byte = data.get_unchecked(i)
        // Process byte...
        i = i + 1
    }
}

fn main() -> i32 {
```

```
let buffer = [1u8, 2, 3, 4, 5]
let slice = Slice::new(&buffer as *u8, 5)
process_buffer(&slice)
return 0
}
```

5.10.3 Moving Out of a Container

```
fn take_from_block() -> (*u8, u32) {
    var block = MemoryBlock::alloc(1024, MEMF_PUBLIC)?

    // Transfer ownership to caller
    let (ptr, size) = block.into_raw()

    // block is now empty - Drop won't free anything
    return (ptr, size)

    // Caller must manually: FreeMem(ptr, size)
}
```

Use `into_raw()` when you need to transfer ownership to code that manages memory manually.

5.11 Memory Safety Guarantees

Novus provides several safety guarantees:

No Double-Free The borrow checker deactivates drop for moved values.

No Use-After-Move The compiler prevents using variables after they've been moved.

Automatic Size Tracking `MemoryBlock` and `Allocation<T>` eliminate size-tracking bugs.

Overflow Protection `Allocation<T>` checks for integer overflow before allocation.

Bounds Checking `Slice::get()` returns `Option<T>`, forcing explicit handling of out-of-bounds access.

Non-Null References References (`&T`, `&var T`) are guaranteed non-null.

5.11.1 What Is NOT Guaranteed

For full transparency, Novus does NOT prevent:

- Multiple simultaneous mutable references (no exclusivity enforcement)
- Use-after-free with raw pointers (`*T`)
- Data races (no thread-safety guarantees beyond `Exec` primitives)
- Memory leaks from forgotten cleanup (though `Drop` helps)

These trade-offs keep the language simple while providing significant safety improvements over C.

5.12 Summary

Novus memory management provides:

- **Ownership** — every value has a single owner
- **Move semantics** — `consuming` parameters take ownership
- **References** — `&T` (immutable) and `&var T` (mutable) borrow without ownership transfer
- **Borrow checker** — prevents use-after-move at compile time
- **Drop trait** — automatic cleanup when values go out of scope
- **RAII patterns** — `Handle`, `Guard`, and `Builder` types for safe resource management
- **Defer blocks** — explicit cleanup with LIFO execution order
- **Using statement** — scoped resource management
- **Standard library types** — `MemoryBlock`, `MemHandle`, `Allocation<T>`, `Box<T>`, `Slice<T>`
- **AmigaOS integration** — proper memory flags for Chip/Fast/Public RAM

The combination of RAII, ownership tracking, and size-safe abstractions makes Novus significantly safer than C for AmigaOS programming. The standard library's 70+ types with `Drop` implementations ensure that resources are properly cleaned up even in the presence of errors.

In the next chapter, we'll explore control flow in depth: conditionals, loops, pattern matching, and error handling with the `?` operator.

Chapter 6

Control Flow

“The art of programming is the art of organizing complexity.”
—Edsger W. Dijkstra

Control flow constructs determine the order in which statements are executed. Novus provides a comprehensive set of control flow mechanisms, from basic conditionals to sophisticated pattern matching. This chapter covers all the ways you can direct program execution in Novus.

Programs aren’t just sequences of statements—they branch, loop, and make decisions. Novus provides control flow constructs that are both familiar to C programmers and enhanced with modern safety features. Pattern matching ensures you handle all cases; labeled loops prevent off-by-one errors in nested structures; and the try operator makes error handling concise without hiding the control flow.

6.1 Conditional Statements

6.1.1 if and else

The if statement executes code based on a boolean condition:

```
fn classify_temperature(temp: i32) -> i32 {  
    if temp < 0 {  
        return -1 // Freezing  
    } else if temp < 20 {  
        return 0 // Cold  
    } else if temp < 30 {  
        return 1 // Comfortable  
    } else {  
        return 2 // Hot  
    }  
}
```

The condition must be of type `bool`. Unlike C, Novus does not implicitly convert integers to booleans—you must use explicit comparisons.

Coming from C

C allows any integer or pointer as a condition. Novus requires explicit boolean expressions:

C	Novus
if (ptr)	if ptr != null
if (count)	if count != 0
while (n-)	while n > 0 { n--; ... }
if (flags & MASK)	if (flags & MASK) != 0

This catches common bugs like `if (x = 10)` (assignment instead of comparison), which

would silently succeed in C but is a type error in Novus.

Blocks are required. Unlike C, the braces around the body are mandatory, even for single statements:

```
// Correct
if x > 0 {
    return x
}

// Syntax error: missing braces
// if x > 0 return x
```

6.1.2 if as an Expression

In Novus, `if` can be used as an expression when both branches return a value:

```
let max = if a > b { a } else { b }

// Also valid with blocks
let result = if condition {
    compute_something()
} else {
    compute_fallback()
}
```

When used as an expression, both the `if` and `else` branches must be present and must return the same type:

```
// Error: branches return different types
let bad = if condition { 42 } else { "hello" }
```

6.1.3 Short-Circuit Evaluation

The logical operators `&&` (and) and `||` (or) use short-circuit evaluation:

- `a && b`: If `a` is false, `b` is not evaluated
- `a || b`: If `a` is true, `b` is not evaluated

This allows safe idioms like:

```
// ptr is only dereferenced if non-null
if ptr != 0 && (*ptr).value > 0 {
    // Safe to use ptr here
}

// check_signal is only called if running is true
if running && check_signal() {
    handle_signal()
}
```

6.1.4 if let

The `if let` construct checks whether a value is non-null (for pointers) or non-zero (for integers) and binds it to a variable. This is primarily used when working with raw pointers from AmigaOS FFI calls:

```
fn process_pointer(ptr: *u8) -> i32 {
    if let p = ptr {
        // p is bound to ptr and is guaranteed non-null here
        return (i32)p
    } else {
        // ptr was null
        return -1
    }
}
```

Note: For pattern matching on `Option<T>` and `Result<T, E>` types, use `let-else` (Section 6.1.5) or `match` (Section 6.4) instead.

The `if let` construct is particularly useful with pointers returned from AmigaOS functions that may return null on failure:

```
fn open_library(name: *u8) -> i32 {
    let lib = OpenLibrary(name, 0)

    if let base = lib {
        // Library opened successfully, base is non-null
        CloseLibrary(base)
        return 0
    } else {
        // Library not found
        return 1
    }
}
```

6.1.5 let-else

The `let-else` statement handles pattern matching where the `else` block must diverge (i.e., return, break, continue, or panic). This is ideal for early-exit patterns:

```
enum MaybeValue {
    Some(i32),
    None
}

fn process(input: MaybeValue) -> i32 {
    let MaybeValue::Some(value) = input else {
        // This block runs if pattern doesn't match
        // Must diverge - cannot fall through
        return -1
    }

    // value is now bound and available here
    return value * 2
}
```

The `else` block *must* diverge. The compiler will reject code where the `else` block could fall through to the subsequent statements.

This pattern works well with `Result` and `Option` types:

```
from std::core import Result
from std::error::errors import DosError

fn read_file(path: *u8) -> Result<i32, DosError> {
    let Result::Ok(handle) = open_file(path) else {
        return Result::Err(DosError::NotFound)
    }

    // handle is now available for use
    let bytes = read_data(handle)
    close_file(handle)
    return Result::Ok(bytes)
}
```

6.2 Loops

Novus provides several loop constructs, each suited to different use cases.

6.2.1 while Loops

The `while` loop executes its body as long as the condition is true:

```
var count = 0
while count < 10 {
    count++
}
// count is now 10
```

The condition is evaluated before each iteration. If the condition is initially false, the body never executes.

Coming from C

Novus `while` loops work exactly like C's:

C	Novus
<code>while (x > 0) { ... }</code>	<code>while x > 0 { ... }</code>
<code>do { ... } while (x);</code>	(not supported; use <code>forever + break</code>)

Novus doesn't have `do-while` loops. Use `forever` with a conditional `break` at the end instead.

Practical While Loop Patterns

Common patterns using `while` loops:

```
// Wait for hardware ready with timeout
fn wait_blitter_with_timeout(max_frames: i32) -> bool {
    var frames = 0
    while (custom.dmaconr & DMAF_BLTDONE) != 0 {
        frames++
        if frames > max_frames {
            return false // Timeout
        }
    }
}
```

```

    }
    WaitTOF() // Wait for frame
}
return true // Blitter finished
}

// Process a linked list
fn count_nodes(mut head: *Node) -> i32 {
    var count = 0
    while head != null {
        count++
        head = head.next
    }
    return count
}

// Read until delimiter
fn read_until_newline(buffer: *var u8, max_len: i32) -> i32 {
    var pos = 0
    while pos < max_len {
        let ch = read_char()
        if ch == '\n' || ch == 0 {
            break
        }
        buffer[pos] = ch
        pos++
    }
    return pos
}

```

```

fn find_first_even(start: i32, max: i32) -> i32 {
    var n = start
    while n < max {
        if n % 2 == 0 {
            return n // Found an even number
        }
        n++
    }
    return -1 // None found
}

```

6.2.2 forever Loops

The forever loop creates an infinite loop that must be exited with **break** or **return**:

```

fn wait_for_signal() -> i32 {
    var iterations = 0
    forever {
        iterations++
        if check_signal() {
            break // Exit the loop
        }
        if iterations > 1000 {
            return -1 // Timeout
        }
    }
}

```

```

    return iterations
}

```

The `forever` keyword is clearer than `while true` and signals the programmer's intent that the loop is designed to be exited via control flow rather than a condition becoming false.

Common Forever Loop Patterns

```

// Game main loop
fn game_loop() {
    forever {
        handle_input()
        update_game_state()
        render_frame()

        if game_over {
            break
        }

        WaitTOF() // Sync to vertical blank
    }
}

// State machine
fn run_state_machine() -> i32 {
    var state = State::Init

    forever {
        state = match state {
            State::Init => {
                initialize()
                State::Running
            },
            State::Running => {
                if process_tick() {
                    State::Running
                } else {
                    State::Finished
                }
            },
            State::Finished => break,
        }
    }

    return 0
}

```

This is particularly useful for event loops in Amiga applications:

```

fn main_event_loop(window: *Window) {
    forever {
        let msg = GetMsg(window.UserPort)
        if let m = msg {
            let class = m.Class
            ReplyMsg(m)

            if class == IDCMP_CLOSEWINDOW {

```



```

        break // User closed the window
    }

    handle_message(class)
}

WaitPort(window.UserPort)
}
}

```

6.2.3 C-Style for Loops

Novus supports traditional C-style **for** loops with initialization, condition, and update expressions:

```

// Sum numbers 0 through 9
var sum = 0
for var i = 0; i < 10; i++ {
    sum = sum + i
}
// sum is now 45

```

The initialization can declare a new variable (with **var** or **let**) or assign to an existing one. The loop variable declared in the initialization is scoped to the loop.

```

// Skip even numbers with continue
var odd_sum = 0
for var j = 0; j < 10; j++ {
    if j % 2 == 0 {
        continue // Skip even numbers
    }
    odd_sum = odd_sum + j
}
// odd_sum is 1 + 3 + 5 + 7 + 9 = 25

```

6.2.4 Range-Based for Loops

Novus provides syntactic sugar for iterating over numeric ranges:

```

// Exclusive range: 0, 1, 2, 3, 4 (excludes 5)
var sum = 0
for i in 0..5 {
    sum = sum + i
}
// sum is 10

```

```

// Inclusive range: 1, 2, 3, 4, 5 (includes 5)
var sum = 0
for i in 1..=5 {
    sum = sum + i
}
// sum is 15

```

Ranges can use variables:

```
let start = 2
let end = 7
var sum = 0
for i in start..end {
    sum = sum + i
}
// sum is 2 + 3 + 4 + 5 + 6 = 20
```

An empty range (where start equals end) executes zero iterations:

```
var count = 0
for i in 5..5 {
    count++ // Never executes
}
// count is 0
```

6.2.5 for-in Loops with Collections

The `for-in` loop iterates over collections that provide `get()` and `len()` methods:

```
// ... (with Vec imported from std::collections::vec)
let vec: Vec<i32> = Vec { [10, 20, 30, 40, 50] }

var sum: i32 = 0
for value in vec {
    sum += value
}
// sum is 150
```

The loop variable receives a copy of each element. The collection itself remains usable after the loop completes. For collections containing non-copyable types, the iteration behavior depends on the collection's implementation.

Any type that provides `get(index: i32)` and `len()` methods can be used with `for-in`. The compiler generates code equivalent to a C-style for loop that calls these methods.

You can declare the loop variable as mutable if you need to modify the local copy:

```
for var item in collection {
    item = transform(item) // Modifies local copy
    process(item)
}
```

6.3 Loop Control Statements

6.3.1 break and continue

The `break` statement immediately exits the innermost loop:

```
var i = 0
while i < 100 {
    if i == 42 {
        break // Exit immediately
    }
    i++
}
```

```
}
// i is 42
```

The `continue` statement skips the rest of the current iteration and proceeds to the next:

```
var sum = 0
var i = 0
while i < 10 {
  i++
  if i == 5 {
    continue // Skip adding 5
  }
  sum = sum + i
}
// sum is 1+2+3+4+6+7+8+9+10 = 50 (5 is skipped)
```

6.3.2 Labeled Loops

When working with nested loops, you can label loops and use those labels with `break` and `continue` to control outer loops:

```
var count = 0
outer: while true {
  var inner = 0
  while inner < 5 {
    count++
    if count == 7 {
      break outer // Exit the outer loop
    }
    inner++
  }
}
// count is 7
```

Labels work with all loop types:

```
outer: forever {
  for i in 0..10 {
    for j in 0..10 {
      if i * j > 50 {
        break outer // Exit all three loops
      }
    }
  }
}
```

The `continue` statement also works with labels:

```
var sum = 0
var i = 0
outer: while i < 5 {
  i++
  var j = 0
  while j < 3 {
    j++
    if i == 3 && j == 2 {
      continue outer
    }
    sum = sum + i
  }
}
```

```

        continue outer // Skip to next i iteration
    }
    sum++
}
}

```

Note: Unlike C labels, Novus labels are only valid for loop control. Novus does not support `goto` statements—labels can only be used with `break` and `continue`.

6.3.3 Postfix Conditions

Novus supports postfix `if` and `unless` for concise conditional statements:

```

fn classify(x: i32) -> i32 {
    return -1 if x < 0 // Return -1 if negative
    return 1 if x > 0 // Return 1 if positive
    return 0 // Zero
}

```

The `unless` keyword is the logical inverse of `if`:

```

fn check(valid: bool) -> i32 {
    return -1 unless valid // Return -1 if NOT valid
    return 0
}

```

Postfix conditions work with `break` and `continue`:

```

var sum = 0
var m = 0
while m < 10 {
    m++
    continue if m % 3 == 0 // Skip multiples of 3
    sum = sum + m
}
// sum is 1+2+4+5+7+8+10 = 37

outer: while true {
    var k = 0
    while k < 100 {
        k++
        break outer if k == 5 // Exit outer when k reaches 5
    }
}

```

6.4 Pattern Matching with `match`

The `match` expression provides exhaustive pattern matching over values. It's more powerful than a simple `switch` statement and is the idiomatic way to handle enums in Novus.

6.4.1 Basic match Expressions

```

fn http_status_category(code: i32) -> i32 {
    match code {

```

```

    200 => return 1, // Success
    201 => return 1,
    204 => return 1,
    400 => return 2, // Client Error
    401 => return 2,
    404 => return 2,
    500 => return 3, // Server Error
    503 => return 3,
    -   => return 0 // Unknown (wildcard)
  }
}

```

The underscore (`_`) is a wildcard pattern that matches any value not handled by previous arms.

Coming from C

`match` is similar to C's `switch`, but safer and more powerful:

- **No fallthrough:** Each arm is independent; no `break` needed
- **Exhaustive:** The compiler ensures all cases are handled
- **Expressions:** `match` returns a value, not just executes code
- **Patterns:** Can destructure enums, tuples, and structs

C switch

```

case 1: ...; break;
default: ...

```

Novus match

```

1 => ...,
_ => ...

```

Forget `break`—the absence of fallthrough eliminates a whole class of bugs.

6.4.2 Matching on Enums

Pattern matching is essential for working with enums, especially those with associated data:

```

enum Status {
    Success,
    Pending,
    Error(i32)
}

fn handle_status(status: Status) -> i32 {
    match status {
        Status::Success => return 0,
        Status::Pending => return 1,
        Status::Error(code) => return code // Bind the error code
    }
}

```

When an enum variant carries data, the pattern can bind that data to a variable for use in the match arm.

6.4.3 Match with Generic Enums

Pattern matching works seamlessly with generic enums like `Option<T>` and `Result<T, E>`:

```

from std::core import Option

fn get_value_or_default(opt: Option<i32>) -> i32 {
    match opt {
        Option::Some(val) => return val,
        Option::None => return 0
    }
}

from std::core import Result
from std::error::errors import DosError

fn process_result(res: Result<i32, DosError>) -> i32 {
    match res {
        Result::Ok(value) => return value,
        Result::Err(_) => return -1 // Ignore error details
    }
}

```

6.4.4 Match with Blocks

When a match arm requires multiple statements, use a block:

```

fn sign(n: i32) -> i32 {
    match n {
        0 => return 0,
        - => {
            if n > 0 {
                return 1 // Positive
            }
            return -1 // Negative
        }
    }
}

```

6.4.5 Match with Guards

Work in Progress

Match guards are parsed but code generation has known issues. Use `if/else` chains as an alternative until this is resolved.

Match arms can include guard clauses with `if` to add additional conditions:

```

fn classify_number(n: i32) -> i32 {
    match n {
        x if x < 0 => return -1, // Negative
        0 => return 0, // Zero
        x if x <= 10 => return 1, // Small positive
        - => return 2 // Large positive
    }
}

```

The guard is evaluated after the pattern matches but before the arm is executed. If the guard evaluates to false, matching continues with the next arm.

6.4.6 Matching Integer Literals

Novus supports matching on decimal, hexadecimal, and binary integer literals:

```
fn parse_permission(bits: i32) -> i32 {
  match bits {
    %111 => return 7, // rwx (binary)
    %110 => return 6, // rw-
    %100 => return 4, // r--
    -    => return 0
  }
}

fn hex_component(hex: i32) -> i32 {
  match hex {
    $00 => return 0, // Hexadecimal
    $FF => return 255,
    $80 => return 128,
    -    => return hex
  }
}
```

6.4.7 Exhaustiveness Checking

The Novus compiler verifies that match expressions are exhaustive—every possible value must be handled. For enums, this means every variant must have a corresponding arm (or a wildcard):

```
enum Direction {
  North,
  South,
  East,
  West
}

fn direction_to_int(dir: Direction) -> i32 {
  match dir {
    Direction::North => return 0,
    Direction::South => return 1,
    Direction::East  => return 2
    // Error: non-exhaustive pattern, Direction::West not
    // covered
  }
}
```

Either add a case for `Direction::West` or use a wildcard pattern.

6.5 Error Propagation with the Try Operator

The try operator (?) provides concise error propagation for functions returning `Result<T, E>`.

6.5.1 Basic Usage

Without the try operator, error handling requires verbose match expressions:

```
fn process_file_manual(path: *u8) -> Result<i32, DosError> {
    let handle_result = open_file(path)
    let handle = match handle_result {
        Result::Ok(h) => h,
        Result::Err(e) => return Result::Err(e)
    }

    let bytes_result = read_file(handle)
    let bytes = match bytes_result {
        Result::Ok(b) => b,
        Result::Err(e) => return Result::Err(e)
    }

    close_file(handle)
    return Result::Ok(bytes)
}
```

With the try operator, this becomes:

```
fn process_file(path: *u8) -> Result<i32, DosError> {
    let handle = open_file(path)?
    let bytes = read_file(handle)?
    close_file(handle)?

    return Result::Ok(bytes)
}
```

The ? operator:

1. Unwraps `Result::Ok(value)` and yields `value`
2. On `Result::Err(e)`, immediately returns `Result::Err(e)` from the current function

6.5.2 Try Operator in Expressions

The try operator can be used within expressions:

```
fn get_file_size(path: *u8) -> Result<i32, DosError> {
    let handle = open_file(path)?
    let bytes = read_file(handle)?
    return Result::Ok(bytes + 100) // Add header overhead
}
```

6.5.3 Chaining Operations

The try operator excels at chaining multiple fallible operations:

```
fn complex_operation() -> Result<i32, DosError> {
    let file = open_file("data.txt")?
    let header = read_header(file)?
    let body = read_body(file, header.length)?
    let result = process_data(body)?
    close_file(file)?
}
```



```
    return Result::Ok(result)
}
```

Each `?` acts as an early-exit point—if any operation fails, the error propagates immediately to the caller.

6.5.4 Combining with defer

The `try` operator pairs well with `defer` for resource cleanup:

```
fn safe_file_operation(path: *u8) -> Result<i32, DosError> {
    let handle = open_file(path)?
    defer {
        close_file(handle)
    }

    let header = read_header(handle)?    // If this fails...
    let body = read_body(handle)?       // ...or this...

    // The defer block still runs, closing the file
    return Result::Ok(process(body))
}
```

6.6 The defer Statement

The `defer` statement schedules code to run when the current scope exits, regardless of how it exits (normal completion, early return, or error propagation).

6.6.1 Basic Usage

```
fn example() -> i32 {
    let buffer = allocate_buffer(1024)
    defer {
        free_buffer(buffer)
    }

    // Use buffer...
    if some_condition() {
        return -1 // defer block still runs
    }

    return 0 // defer block runs before returning
}
```

6.6.2 Multiple Defers

Multiple `defer` statements execute in reverse order (last-in, first-out):

```
fn multi_resource() {
    let resource1 = acquire_first()
    defer { release_first(resource1) }
```

```

let resource2 = acquire_second()
defer { release_second(resource2) }

let resource3 = acquire_third()
defer { release_third(resource3) }

// Use resources...

// On exit: release_third, then release_second, then
//          release_first
}

```

This LIFO order ensures resources are released in the reverse order of acquisition, which is typically the correct order for dependent resources.

6.6.3 Defer with AmigaOS Resources

The `defer` statement is particularly useful for managing AmigaOS resources:

```

fn open_window_safe() -> i32 {
    let screen = LockPubScreen(0)
    if screen == 0 {
        return -1
    }
    defer { UnlockPubScreen(0, screen) }

    let visual_info = GetVisualInfo(screen, 0)
    if visual_info == 0 {
        return -2
    }
    defer { FreeVisualInfo(visual_info) }

    let window = OpenWindowTags(screen)
    if window == 0 {
        return -3
    }
    defer { CloseWindow(window) }

    // Use window...
    process_events(window)

    return 0
    // Resources released in order: window, visual_info, screen
}

```

6.7 The using Statement

The `using` statement provides scoped resource management for types that implement the `Drop` trait. When the block exits, the resource's `drop` method is called automatically.

```

pub struct Resource {
    value: i32
}

impl Drop for Resource {
    fn drop(&var self) {

```

```

        // Cleanup happens here automatically
    }
}

fn create_resource(v: i32) -> Resource {
    return Resource { value: v }
}

fn process_data() -> i32 {
    using create_resource(42) {
        // The resource exists for this block
        let x: i32 = 10
        // ... use x ...
    }
    // Resource's drop() called automatically here

    return 0
}

```

The `using` statement takes an expression that returns a type implementing `Drop`. The resource is created, the block executes, and then `drop()` is called—regardless of how the block exits.

6.7.1 using vs defer

Both `using` and `defer` provide deterministic cleanup, but they serve different purposes:

Use `defer` when:

- You need to bind the resource to a variable for use in the scope
- The resource must live for the entire function scope
- You need cleanup after early returns from anywhere in the function
- You're working with raw FFI resources that don't implement `Drop`

Use `using` when:

- The resource is only needed to exist during a specific block
- You don't need to access the resource directly (e.g., a lock or guard)
- The resource type implements the `Drop` trait

Example showing the difference:

```

fn process_files() -> i32 {
    // defer: resource is bound to a variable, cleanup at
    // function exit
    let handle = open_file("log.txt")
    defer { close_file(handle) }

    // using: resource exists for block scope only
    // Useful for locks, guards, or temporary resources
    using acquire_lock() {
        // Lock is held for this block
        do_protected_work()
    }
    // Lock released here
}

```

```
// We can still use handle
write_to_file(handle, "Done")
return 0
// handle closed here by defer
}
```

6.8 Common Control Flow Patterns

This section presents idiomatic patterns that combine multiple control flow constructs effectively.

6.8.1 Early Exit Pattern

Use early returns to handle error cases first, keeping the happy path unindented:

```
fn process_data(input: *Data) -> Result<i32, Error> {
    // Guard clauses first
    if input == null {
        return Result::Err(Error::NullInput)
    }

    if input.size == 0 {
        return Result::Err(Error::EmptyData)
    }

    if input.size > MAX_SIZE {
        return Result::Err(Error::TooLarge)
    }

    // Happy path - not nested inside if/else
    let result = compute(input)?
    return Result::Ok(result)
}
```

Compare this to the nested alternative:

```
// Less readable - deep nesting
fn process_data_nested(input: *Data) -> Result<i32, Error> {
    if input != null {
        if input.size > 0 {
            if input.size <= MAX_SIZE {
                let result = compute(input)?
                return Result::Ok(result)
            } else {
                return Result::Err(Error::TooLarge)
            }
        } else {
            return Result::Err(Error::EmptyData)
        }
    } else {
        return Result::Err(Error::NullInput)
    }
}
```

6.8.2 Resource Acquisition Pattern

Acquire resources with early exit, using defer for cleanup:

```
fn safe_window_operation() -> Result<i32, Error> {
    // Acquire screen
    let screen = LockPubScreen(null)
    if screen == null {
        return Result::Err(Error::NoScreen)
    }
    defer { UnlockPubScreen(null, screen) }

    // Acquire visual info
    let vi = GetVisualInfo(screen, null)
    if vi == null {
        return Result::Err(Error::NoVisualInfo)
    }
    defer { FreeVisualInfo(vi) }

    // Acquire window
    let window = OpenWindow(screen, vi)
    if window == null {
        return Result::Err(Error::NoWindow)
    }
    defer { CloseWindow(window) }

    // All resources acquired - do work
    // Any return path will properly clean up all resources
    let result = process_window_events(window)?
    return Result::Ok(result)
}
```

6.8.3 Retry Pattern

Retry operations with exponential backoff:

```
fn retry_connect(addr: *u8, max_retries: i32) ->
Result<*Connection, Error> {
    var delay_ms = 100
    var attempt = 0

    while attempt < max_retries {
        attempt++

        match try_connect(addr) {
            Result::Ok(conn) => return Result::Ok(conn),
            Result::Err(e) => {
                if attempt >= max_retries {
                    return Result::Err(e)
                }
                // Wait before retry
                Delay(delay_ms / 20) // Convert to ticks
                delay_ms *= 2 // Exponential backoff
            }
        }
    }
}
```

```

    return Result::Err(Error::MaxRetriesExceeded)
}

```

6.8.4 State Machine Pattern

Implement state machines using match:

```

enum ParserState {
    Start,
    InTag,
    InAttribute,
    InValue,
    Done,
}

fn parse_tag(input: *u8, len: i32) -> Result<Tag, ParseError> {
    var state = ParserState::Start
    var pos = 0

    while pos < len {
        let ch = input[pos]

        state = match state {
            ParserState::Start => {
                if ch == '<' {
                    ParserState::InTag
                } else {
                    return
                        Result::Err(ParseError::ExpectedOpenTag)
                }
            },
            ParserState::InTag => {
                if ch == ' ' {
                    ParserState::InAttribute
                } else if ch == '>' {
                    ParserState::Done
                } else {
                    ParserState::InTag // Keep reading tag name
                }
            },
            ParserState::InAttribute => {
                if ch == '=' {
                    ParserState::InValue
                } else if ch == '>' {
                    ParserState::Done
                } else {
                    ParserState::InAttribute
                }
            },
            ParserState::InValue => {
                if ch == ' ' {
                    ParserState::InAttribute
                } else if ch == '>' {
                    ParserState::Done
                } else {
                    ParserState::InValue
                }
            }
        }
    }
}

```

```

        },
        ParserState::Done => break,
    }

    pos++
}

return build_tag(input, len)
}

```

6.8.5 Search Pattern with Multiple Collections

Search through multiple sources with early exit:

```

fn find_resource(name: *u8) -> Option<*Resource> {
    // Check cache first
    if let res = cache_lookup(name) {
        return Option::Some(res)
    }

    // Check local directory
    if let res = local_lookup(name) {
        cache_store(name, res)
        return Option::Some(res)
    }

    // Check system paths
    for i in 0u32..system_paths.len() {
        let path = system_paths.get(i)
        if let res = path_lookup(path, name) {
            cache_store(name, res)
            return Option::Some(res)
        }
    }

    return Option::None
}

```

6.9 Showcase: Menu System

This showcase demonstrates control flow constructs in a practical menu system:

```

// menu_system_showcase.novus
// Demonstrates: conditionals, loops, match, defer, error handling

from std::io import println
from std::core import Result, Option

// Menu item types
enum MenuItem {
    Action { id: i32, label: *u8 },
    Submenu { id: i32, label: *u8, items: *MenuList },
    Separator,
}

```

```

// Menu result from user interaction
enum MenuResult {
    Selected(i32),           // Item ID selected
    Cancelled,              // User cancelled
    Timeout,                // No selection in time
}

struct MenuState {
    current_item: i32,
    is_open: bool,
    parent: Option<*MenuState>,
}

impl MenuState {
    pub fn new() -> MenuState {
        return MenuState {
            current_item: 0,
            is_open: false,
            parent: Option::None,
        }
    }
}

// Process a single menu item
fn process_menu_item(item: &MenuItem, state: &var MenuState) ->
MenuResult {
    match *item {
        MenuItem::Action { id, label } => {
            // Action item - return selection
            return MenuResult::Selected(id)
        },
        MenuItem::Submenu { id, label, items } => {
            // Open submenu - recursive handling
            var sub_state = MenuState::new()
            sub_state.parent = Option::Some(state as *MenuState)

            forever {
                let result = run_menu_loop(items, &var sub_state)
                match result {
                    MenuResult::Selected(sub_id) => {
                        return MenuResult::Selected(sub_id)
                    },
                    MenuResult::Cancelled => {
                        // Return to parent menu
                        break
                    },
                    MenuResult::Timeout => {
                        return MenuResult::Timeout
                    },
                }
            }

            // Continue in parent menu
            return MenuResult::Cancelled
        },
        MenuItem::Separator => {
            // Skip separators
            return MenuResult::Cancelled
        }
    }
}

```



```

    },
  }
}

// Navigate menu with keyboard
fn handle_menu_key(key: u8, state: &var MenuState, item_count:
i32) -> Option<MenuResult> {
  // Use match for key handling
  match key {
    'j' => {
      // Move down, skip separators
      var next = state.current_item + 1
      while next < item_count {
        if !is_separator(next) {
          state.current_item = next
          break
        }
        next++
      }
      return Option::None
    },
    'k' => {
      // Move up, skip separators
      var prev = state.current_item - 1
      while prev >= 0 {
        if !is_separator(prev) {
          state.current_item = prev
          break
        }
        prev--
      }
      return Option::None
    },
    '\r' => {
      // Enter - select current
      return
        Option::Some(MenuResult::Selected(state.current_item))
    },
    27 => {
      // Escape - cancel
      return Option::Some(MenuResult::Cancelled)
    },
    _ => {
      // Unknown key
      return Option::None
    },
  }
}

// Main menu loop
fn run_menu_loop(items: *MenuList, state: &var MenuState) ->
MenuResult {
  var timeout_counter = 0
  const MAX_TIMEOUT: i32 = 3000 // 30 seconds at 100 ticks/sec

  state.is_open = true
  defer {
    state.is_open = false
  }
}

```

```

}

forever {
    // Draw current state
    draw_menu(items, state)

    // Wait for input with timeout
    let key_result = wait_for_key_with_timeout(10)

    match key_result {
        Option::Some(key) => {
            timeout_counter = 0 // Reset timeout

            // Handle navigation
            let nav_result = handle_menu_key(key, state,
                items.len())

            if let result = nav_result {
                match result {
                    MenuResult::Selected(idx) => {
                        // Process the selected item
                        let item = items.get(idx)
                        return process_menu_item(item, state)
                    },
                    MenuResult::Cancelled => {
                        return MenuResult::Cancelled
                    },
                    _ => { }
                }
            }
        },
        Option::None => {
            timeout_counter++
            if timeout_counter > MAX_TIMEOUT {
                return MenuResult::Timeout
            }
        },
    }
}

}

// Example usage
fn main() -> i32 {
    println("=== Menu System Showcase ===")

    // Build a sample menu structure
    var state = MenuState::new()

    // Simulate menu interaction
    let result = simulate_menu_interaction(&var state)

    match result {
        MenuResult::Selected(id) => {
            println(f"Selected item: {id}")
            return 0
        },
        MenuResult::Cancelled => {
            println("Menu cancelled")
        }
    }
}

```

```

        return 1
    },
    MenuResult::Timeout => {
        println("Menu timed out")
        return 2
    },
}
}
}

```

This showcase demonstrates:

- **Match expressions** for handling enums with associated data
- **Forever loops** for event-driven menu processing
- **While loops** for skipping separators during navigation
- **If let** for handling optional values
- **Defer** for cleanup when menu closes
- **Nested match** for complex input handling
- **Early exit** with `break` and `return`

6.10 Summary

Novus provides a complete set of control flow constructs:

- **Conditionals:** `if/else`, `if let`, `let-else` for pattern matching with early exit
- **Short-circuit evaluation:** `&&` and `||` operators for efficient boolean logic
- **Loops:** `while`, `forever`, C-style `for`, range-based `for`, and `for-in`
- **Loop control:** `break`, `continue`, labeled loops for nested control
- **Postfix conditions:** `if` and `unless` for concise conditional statements
- **Pattern matching:** `match` with exhaustiveness checking, guards, and literal patterns
- **Error propagation:** The `?` operator for concise `Result` handling
- **Resource management:** `defer` and `using` for deterministic cleanup

These constructs work together to enable clear, safe, and expressive code. The emphasis on explicit control flow and exhaustive pattern matching helps prevent bugs while keeping the code readable.

In the next chapter, we'll explore functions and methods, including function definitions, parameter passing, generic functions, and `impl` blocks.

Chapter 7

Functions and Methods

“Functions should do one thing. They should do it well. They should do it only.”
—Robert C. Martin

Functions are the foundation of structured programming. In Novus, they’re designed with systems programming in mind—explicit parameter passing, clear ownership semantics, and direct integration with the Amiga’s calling conventions. Whether you’re writing a simple utility function or implementing complex data structures, understanding Novus’s function system is essential.

Functions are the fundamental building blocks of Novus programs. This chapter covers everything from basic function definitions to generic methods, closures, and compile-time evaluation. By the end, you’ll understand how to write reusable, type-safe code that integrates seamlessly with the Novus type system.

7.1 Function Basics

7.1.1 Defining Functions

Functions are declared with the `fn` keyword:

```
fn add(a: i32, b: i32) -> i32 {  
    return a + b  
}  
  
fn greet(name: *u8) {  
    // Prints the name to console  
    println(name)  
}
```

A function declaration consists of:

- The `fn` keyword
- The function name (must be a valid identifier)
- Parameters in parentheses, each with name and type
- An optional return type after `->`
- The function body in braces

If no return type is specified, the function returns the unit type `()`.

7.1.2 Implicit Returns

When a function body ends with an expression (not a statement), that expression becomes the function's return value:

```
fn square(x: i32) -> i32 {
    x * x // Implicit return - no return keyword needed
}

fn add(a: i32, b: i32) -> i32 {
    a + b
}
```

You can also use explicit `return` statements, which is required for early returns:

```
fn safe_sqrt(x: i32) -> i32 {
    if x < 0 {
        return 0 // Early return requires 'return'
    }
    // ... compute sqrt ...
    result // Implicit return at end
}
```

Both styles are valid. Use implicit returns for simple functions and explicit returns when you need early exit or want to be more explicit about control flow.

7.1.3 Parameters

Parameters are passed by value by default. Each parameter requires an explicit type:

```
fn calculate(x: i32, y: i32, z: i32) -> i32 {
    return x + y * z
}
```

Important: Parameters are always immutable within the function body. To work with a mutable copy, assign to a local variable:

```
fn process(x: i32) -> i32 {
    // x = x + 1 // Error: cannot mutate parameter
    var local = x // Create mutable copy
    local = local + 1
    return local * 2
}
```

7.1.4 Reference Parameters

Pass data by reference to avoid copying large structures:

```
struct Player {
    x: i32,
    y: i32,
    health: i32,
    score: u32,
}

// Pass by immutable reference - cannot modify
```

```

fn get_position(player: &Player) -> (i32, i32) {
    return (player.x, player.y)
}

// Pass by mutable reference - can modify
fn damage(player: &var Player, amount: i32) {
    player.health = player.health - amount
    if player.health < 0 {
        player.health = 0
    }
}

fn main() -> i32 {
    var p = Player { x: 100, y: 200, health: 100, score: 0 }

    let (x, y) = get_position(&p)
    damage(&var p, 25)

    return 0
}

```

Coming from C

In C, you use pointers for pass-by-reference and must manually track whether a function modifies data:

```

// C: Is player modified? Must check implementation!
void process(Player *player);
void print_info(const Player *player);

```

Novus makes mutability explicit at the call site:

```

// Novus: Call site shows intent clearly
process(&var player)    // Will modify player
print_info(&player)     // Won't modify player

```

The key differences:

- `&T` (immutable reference) $\approx$ `const T*`
- `&var T` (mutable reference) $\approx$ `T*`
- References are never null; use `*T` for nullable pointers
- The call site `&var` makes modifications visible

The reference type syntax mirrors the declaration:

- `&T` — immutable reference, pass with `&value`
- `&var T` — mutable reference, pass with `&var value`

7.1.5 The `consuming` Keyword

The `consuming` keyword indicates that a function takes ownership of a value, consuming it so it cannot be used afterward:

```
fn take_ownership(consuming s: String) {  
    // s is now owned by this function  
    // ... (The original variable is no longer usable)  
}  
  
fn use_string() {  
    var s = String::new()  
    take_ownership(s)  
    // s is no longer valid here - compilation error if used  
}
```

This is particularly useful for builder patterns and type-state transformations:

```
impl StringBuilder {  
    // Consumes the builder and returns the final String  
    pub fn build(consuming self) -> String {  
        // Convert builder to String, consuming self  
    }  
}
```

Consuming Parameters in AmigaOS FFI

Many AmigaOS system calls use `consuming` to indicate ownership transfer. For example, `FreeMem` consumes its pointer parameter:

```
from std::ffi::exec import AllocMem, FreeMem  
from std::ffi::amiga_consts import MEMF_FAST  
  
fn example() {  
    let ptr = unsafe { AllocMem(1024, MEMF_FAST) }  
    if ptr != null {  
        // ... use ptr ...  
        // ptr is consumed by FreeMem - cannot use it afterward  
        unsafe { FreeMem(ptr, 1024) }  
        // ptr is now invalid  
    }  
}
```

7.1.6 Early Returns

Use explicit `return` for early exit from functions:

```
fn safe_divide(a: i32, b: i32) -> i32 {  
    if b == 0 {  
        return 0 // Early return for error case  
    }  
    a / b // Implicit return for normal case  
}
```

7.1.7 Returning Multiple Values

Use tuples to return multiple values:

```
fn divmod(a: i32, b: i32) -> (i32, i32) {
```



```

    let quotient = a / b
    let remainder = a % b
    (quotient, remainder) // Return tuple
}

fn main() -> i32 {
    let (q, r) = divmod(17, 5)
    // q = 3, r = 2
    return 0
}

```

7.2 Visibility

7.2.1 Public Functions

By default, functions are private to their module. Use `pub` to make them accessible from other modules:

```

// Private - only accessible within this module
fn helper() -> i32 {
    return 42
}

// Public - accessible from other modules
pub fn api_function() -> i32 {
    return helper()
}

```

7.2.2 Internal Visibility

The `internal` modifier makes a function visible within the same package but not to external code:

```

// ... (illustrative - internal visibility modifier)
internal fn package_utility() -> i32 {
    return 100
}

```

7.3 Traits

Traits define shared behavior that types can implement. Understanding traits is essential before discussing generic functions, as generics often use trait bounds.

7.3.1 Defining Traits

A trait declares method signatures that implementing types must provide:

```

pub trait Eq {
    fn eq(&self, other: &Self) -> bool
}

```

The `Self` type (capitalized) refers to the implementing type. In the example above, when implementing `Eq` for `Point`, `Self` becomes `Point`.

7.3.2 Implementing Traits

Use `impl TraitName for Type` to implement a trait:

```
struct Point {
    x: i32,
    y: i32,
}

impl Eq for Point {
    fn eq(&self, other: &Point) -> bool {
        return self.x == other.x && self.y == other.y
    }
}
```

7.3.3 Common Standard Library Traits

Novus defines several important traits in `std::core`:

Drop Destructor called when value goes out of scope

Clone Deep copy with potential allocation (returns `Option<Self>`)

Copy Marker trait for bitwise-copyable types

Eq Equality comparison

Ord Total ordering for comparison

Hash Hashing for use in `HashMap`

Display Format value to a `Formatter`; returns `bool` for success

From<T> Type conversion from `T`

Into<T> Type conversion to `T`

Default Provide a default value for a type

Example implementing `Drop`:

```
from std::ffi::exec import FreeMem

impl Drop for Buffer {
    fn drop(&var self) {
        if self.size > 0 {
            unsafe { FreeMem(self.ptr, self.size) }
            self.ptr = null
            self.size = 0
        }
    }
}
```

7.4 Methods and Impl Blocks

Methods are functions associated with a type, defined within `impl` blocks.

7.4.1 Defining Methods

```
struct Rectangle {
    width: u32,
    height: u32,
}

impl Rectangle {
    // Associated function (no self) - called as Rectangle::new()
    pub fn new(width: u32, height: u32) -> Rectangle {
        return Rectangle { width: width, height: height }
    }

    // Method with immutable self reference
    pub fn area(&self) -> u32 {
        return self.width * self.height
    }

    // Method with mutable self reference
    pub fn scale(&var self, factor: u32) {
        self.width = self.width * factor
        self.height = self.height * factor
    }

    // Method that consumes self
    pub fn into_tuple(self) -> (u32, u32) {
        return (self.width, self.height)
    }
}
```

7.4.2 Self Parameter Types

Methods can receive `self` in different ways:

Receiver	Meaning
<code>&self</code>	Immutable reference; can read but not modify
<code>&var self</code>	Mutable reference; can read and modify
<code>self</code>	Takes ownership by value; consumes the instance
<code>consuming self</code>	Explicit ownership transfer (clearer intent)

Both `self` and `consuming self` have the same semantics—they consume the value. The `consuming` keyword makes the intent more explicit and is recommended for clarity.

```
impl Resource {
    // Read-only access
    fn info(&self) -> i32 {
        return self.value
    }

    // Modify in place
    fn update(&var self, new_value: i32) {
```

```

        self.value = new_value
    }

    // Consume and transform - explicit consuming
    fn into_raw(consuming self) -> *u8 {
        let ptr = self.ptr
        // Don't run destructor - ownership transferred
        return ptr
    }
}

```

7.4.3 Associated Functions

Functions without a `self` parameter are called *associated functions*. They're called using the type name:

```

// ... (with Vec defined elsewhere)
impl<T> Vec<T> {
    // Associated function - no self
    // @zeroed creates a zero-initialized instance
    pub fn new() -> Vec<T> {
        return @zeroed(Vec<T>)
    }

    // Associated function with parameters
    pub fn with_capacity(cap: u32) -> Result<Vec<T>, ExecError> {
        // ... Allocate memory and return Vec
    }
}

```

Calling associated functions:

```

// ... (with Vec and impl defined above)
var v = Vec::i32::new()
var v2 = Vec::i32::with_capacity(100)?

```

Associated functions are commonly used for constructors (`new`, `with_capacity`) and factory methods.

7.4.4 Calling Methods

Methods are called using dot notation:

```

var rect = Rectangle::new(10, 20)
let a = rect.area()           // Call method with &self
rect.scale(2)                 // Call method with &var self
let (w, h) = rect.into_tuple() // Consumes rect

```

The compiler automatically borrows `self` as needed. You don't need to write `(&rect).area()`—the compiler inserts the borrow for you.

Coming from C

In C, you pass structs to functions explicitly and use `->` for pointer access:

```
// C: Explicit struct passing
struct Rectangle *r = create_rect(10, 20);
int area = rect_area(r);           // Pass pointer explicitly
rect_scale(r, 2);                  // No indication it modifies r
int w = r->width;                  // Arrow for pointer access
```

Novus uses method syntax with automatic borrowing:

```
// Novus: Method syntax, uniform dot access
var rect = Rectangle::new(10, 20)
let area = rect.area()             // Auto-borrow as &self
rect.scale(2)                      // Auto-borrow as &var self
let w = rect.width                 // Dot works for values too
```

Key differences:

- No `->` operator—dot (`.`) works for both values and pointers
- Methods are called `value.method()`, not `method(&value)`
- Auto-deref: `ptr.field` works like C's `ptr->field`
- Constructors are `Type::new()`, not separate `create_type()` functions

7.5 Generic Functions

Generic functions work with multiple types using type parameters.

7.5.1 Type Parameters

```
fn identity<T>(value: T) -> T {
    return value
}

fn swap<T>(a: &var T, b: &var T) {
    let temp = *a
    *a = *b
    *b = temp
}
```

Type parameters are declared in angle brackets after the function name.

7.5.2 Trait Bounds

Use `where` clauses to constrain type parameters to types that implement specific traits:

```
// ... (with Eq and Display imported)
pub fn expect_eq<T>(actual: T, expected: T, message: Str) where
    T: Eq + Display {
    if actual != expected {
        // ... Report failure - can use Eq and Display methods on
        T
    }
}
```

Multiple bounds are combined with +:

```
fn sort<T>(items: &var [T], count: u32) where T: Ord + Clone {  
    // Can use Ord for comparison, Clone for copying  
}
```

7.5.3 Turbofish Syntax

When the compiler cannot infer generic types, use the turbofish syntax (`::<T>`) to specify them explicitly:

```
// Type inferred from variable declaration  
let v: Vec<i32> = Vec::new()  
  
// Explicit type with turbofish  
let v = Vec::<i32>::new()  
  
// Turbofish on function calls  
let chan = bounded_channel::<i32>()? 
```

The turbofish is especially useful when calling generic functions where the type cannot be inferred from context.

7.5.4 Generic Trait Implementations

Implement traits for generic types:

```
// ... (with HashMap defined elsewhere)  
impl<K, V> Drop for HashMap<K, V> where K: Hash + Eq {  
    fn drop(&var self) {  
        // ... Free all buckets and entries  
    }  
}
```

7.6 Function Overloading

Novus supports function overloading—multiple functions with the same name but different parameter signatures.

7.6.1 Overloading Rules

Functions can be overloaded when they differ by:

- Parameter types (i32 vs i64)
- Parameter count
- Reference vs value (T vs &T)
- Mutability (&T vs &var T)
- Pointer vs value (T vs *T)

```

fn abs(x: i32) -> i32 {
    if x < 0 { return -x }
    return x
}

fn abs(x: i16) -> i16 {
    if x < 0 { return -x }
    return x
}

fn print_value(x: i32) {
    // Print i32
}

fn print_value(x: i32, y: i32) {
    // Print two values
}

fn print_value(x: i32, y: i32, z: i32) {
    // Print three values
}

```

7.6.2 Overloading Restrictions

Overloading is **not** allowed when functions differ only by:

- Return type (with identical parameters)
- Parameter names (with identical types)

```

// ERROR: Same parameter types, different return types
fn get() -> i32 { return 42 }
fn get() -> i64 { return 42 } // Error: duplicate signature

// ERROR: Same types, different names
fn test(x: i32) -> i32 { return x }
fn test(y: i32) -> i32 { return y } // Error: duplicate signature

```

7.7 Const Functions

Const functions can be evaluated at compile time when called with constant arguments.

7.7.1 Defining Const Functions

```

const fn times_two(x: i32) -> i32 {
    return x * 2
}

const fn max(a: i32, b: i32) -> i32 {
    if a > b {
        return a
    }
    return b
}

```

```
}  
  
// Const fn calling another const fn  
const fn clamp(value: i32, low: i32, high: i32) -> i32 {  
    if value < low { return low }  
    if value > high { return high }  
    return value  
}
```

7.7.2 Using Const Functions

Const functions enable compile-time computation for constants:

```
const WIDTH: i32 = 320  
const HEIGHT: i32 = 200  
const DOUBLED_WIDTH: i32 = times_two(WIDTH)    // 640 at compile  
        time  
const SCREEN_SIZE: i32 = WIDTH * HEIGHT        // 64000 at  
        compile time  
const CLAMPED: i32 = clamp(999, 0, 100)        // 100 at compile  
        time
```

Const functions can also be called at runtime with non-constant arguments:

```
fn process(x: i32) -> i32 {  
    return times_two(x)    // Runtime call  
}
```

7.7.3 Const Function Restrictions

Const functions have limitations:

- No heap allocation
- No I/O operations
- No external function calls
- Only call other const functions
- Limited recursion depth

7.8 Closures

Closures are anonymous functions that can capture variables from their environment.

7.8.1 Closure Syntax

```
// Basic closure with type annotations  
let add = |x: i32, y: i32| -> i32 { x + y }  
  
// Closure capturing environment  
let multiplier = 10  
let scale = |x: i32| -> i32 { x * multiplier }
```



```
let result = scale(5) // 50

// No parameters (note: space between pipes required)
let get_value = | | -> i32 { 42 }
```

Closures use pipes (|) to delimit parameters:

```
|parameters| -> ReturnType { body }
```

7.8.2 Capturing Variables

Closures capture variables from their enclosing scope. By default, captures are by value (copied):

```
fn example() -> i32 {
    let offset = 10
    let add_offset = |x: i32| -> i32 {
        x + offset // Captures 'offset' by value
    }
    add_offset(5) // 15
}
```

7.9 External Functions

External functions declare interfaces to C code or AmigaOS system libraries.

7.9.1 Extern Declarations

External functions declare interfaces to AmigaOS system libraries. Here are common examples from `exec.library`:

```
// Memory allocation
extern pub fn AllocMem(byteSize: u32, requirements: u32) -> *u8
extern pub fn FreeMem(consuming memoryBlock: *u8, byteSize: u32)

// Task synchronization
extern pub fn Wait(signalSet: u32) -> u32
extern pub fn Signal(task: *Task, signals: u32)

// Message passing
extern pub fn PutMsg(port: *MsgPort, message: *Message)
extern pub fn GetMsg(port: *MsgPort) -> *Message
```

From `dos.library` (note BCPL pointer conventions—file handles are `i32`):

```
// File operations (using BCPL pointers)
extern pub fn Open(name: *u8, accessMode: i32) -> i32
extern pub fn Close(consuming file: i32) -> i32
extern pub fn Read(file: i32, buffer: *u8, length: i32) -> i32
extern pub fn Write(file: i32, buffer: *u8, length: i32) -> i32
```

External functions have no body—they're implemented in ROM or disk-based libraries and linked at runtime.

7.9.2 Calling External Functions

AmigaOS system calls require `unsafe` blocks because they bypass Novus’s safety guarantees:

```
from std::ffi::exec import AllocMem, FreeMem
from std::ffi::amiga_consts import MEMF_FAST, MEMF_CLEAR

fn allocate_buffer(size: u32) -> *u8 {
    var ptr: *u8 = null
    unsafe {
        ptr = AllocMem(size, MEMF_FAST | MEMF_CLEAR)
    }
    return ptr
}

fn free_buffer(ptr: *u8, size: u32) {
    unsafe {
        FreeMem(ptr, size)
    }
}
```

7.9.3 AmigaOS Library Bindings

Novus provides pre-generated FFI bindings for all AmigaOS libraries. These are auto-generated from SFD (System Function Definition) files:

```
// Import from pre-generated FFI modules
from std::ffi::exec import AllocMem, FreeMem, Wait, Signal
from std::ffi::dos import Open, Close, Read, Write
from std::ffi::graphics import BltBitMap, BltTemplate
from std::ffi::intuition import OpenWindow, CloseWindow

// Import structure definitions
from std::ffi::amiga_structs import Task, MsgPort, BitMap,
    RastPort

// Import constants
from std::ffi::amiga_consts import MEMF_CHIP, MEMF_FAST,
    MODE_OLDFILE
```

Developers should import from these modules rather than writing `extern fn` declarations manually.

7.9.4 Variadic Functions

External functions can accept variable arguments using `...args`:

```
extern fn VPrintf(format: *u8, ...args) -> i32
```

Note that variadic functions are only supported for external declarations; Novus code cannot define variadic functions.

7.10 Best Practices

7.10.1 Parameter Passing Guidelines

Situation	Recommendation
Small values (integers, bools)	Pass by value
Large structs, read-only	Pass by reference (<code>&T</code>)
Large structs, modify	Pass by mutable reference (<code>&var T</code>)
Transfer ownership	Pass by value or use <code>consuming</code>

7.10.2 Method Receiver Guidelines

Receiver	Use When
<code>&self</code>	Reading data, computing derived values
<code>&var self</code>	Modifying internal state
<code>self / consuming self</code>	Transforming into another type, builder finalization

7.10.3 AmigaOS FFI Best Practices

When calling AmigaOS system libraries:

- Always wrap library calls in `unsafe` blocks
- Check return values—many functions return `NULL/0` on failure
- Track allocation sizes—`FreeMem` requires the original size
- Match memory flags—free with the same size you allocated
- Respect `consuming` parameters—don't use pointers/handles after consumption
- Use `stdlib` wrappers (`MemoryBlock`, `Vec`, etc.) when possible for safety
- BCPL pointers (`DOS.library`) are `i32` values, not `*u8`

Example of safe memory management:

```
from std::memory::block import MemoryBlock
from std::ffi::amiga_consts import MEMF_FAST

// Prefer MemoryBlock over raw AllocMem/FreeMem
match MemoryBlock::alloc(1024, MEMF_FAST) {
  Option::Some(var block) => {
    // Use block.ptr() to access raw pointer
    let ptr = block.ptr()
    // ... use ptr ...
    // Automatic cleanup when block goes out of scope
  },
  Option::None => {
    // Handle allocation failure
  }
}
```

7.10.4 Naming Conventions

- Use `snake_case` for function names: `get_value`, `set_position`
- Use `new` for constructors: `Vec::new()`, `String::new()`
- Use `with_*` for constructors with parameters: `with_capacity`
- Use `is_*` for boolean queries: `is_empty`, `is_valid`
- Use `into_*` for consuming conversions: `into_raw`, `into_string`
- Use `as_*` for borrowing conversions: `as_str`, `as_slice`
- Use `try_*` for fallible operations returning `Result`: `try_alloc`

7.11 Common Function Patterns

This section presents idiomatic patterns for function design in Novus.

7.11.1 The Builder Pattern

Builders accumulate configuration before producing a final value:

```
struct WindowBuilder {
    title: *u8,
    width: u16,
    height: u16,
    left: u16,
    top: u16,
    flags: u32,
}

impl WindowBuilder {
    pub fn new(title: *u8) -> WindowBuilder {
        return WindowBuilder {
            title: title,
            width: 640,
            height: 400,
            left: 0,
            top: 0,
            flags: 0,
        }
    }

    // Each setter returns self for chaining
    pub fn size(&var self, width: u16, height: u16) -> &var
        WindowBuilder {
        self.width = width
        self.height = height
        return self
    }

    pub fn position(&var self, left: u16, top: u16) -> &var
        WindowBuilder {
        self.left = left
        self.top = top
        return self
    }
}
```

```

}

pub fn borderless(&var self) -> &var WindowBuilder {
    self.flags |= WFLG_BORDERLESS
    return self
}

pub fn backdrop(&var self) -> &var WindowBuilder {
    self.flags |= WFLG_BACKDROP
    return self
}

// Final build consumes the builder
pub fn build(consuming self) -> Result<*Window,
IntuitionError> {
    // Create the window from accumulated settings
    let window = unsafe {
        OpenWindowTags(
            null,
            WA_Title, self.title,
            WA_Width, (u32)self.width,
            WA_Height, (u32)self.height,
            WA_Left, (u32)self.left,
            WA_Top, (u32)self.top,
            WA_Flags, self.flags,
            TAG_DONE
        )
    }

    if window == null {
        return Result::Err(IntuitionError::OpenFailed)
    }
    return Result::Ok(window)
}

}

// Usage
fn open_game_window() -> Result<*Window, IntuitionError> {
    var builder = WindowBuilder::new("My Game")
    let window = builder
        .size(320, 256)
        .position(0, 0)
        .borderless()
        .backdrop()
        .build()?

    return Result::Ok(window)
}

```

7.11.2 The Factory Pattern

Factories create objects based on runtime conditions:

```

enum AudioDriver {
    Paula { channels: u8 },
    Ahi { device_id: i32 },
    Null,
}

```

```

}

impl AudioDriver {
    // Factory method chooses implementation at runtime
    pub fn create(prefer_ahi: bool) -> Result<AudioDriver,
        AudioError> {
        if prefer_ahi {
            // Try AHI first
            match try_open_ahi() {
                Result::Ok(id) => {
                    return Result::Ok(AudioDriver::Ahi {
                        device_id: id })
                },
                Result::Err(_) => {
                    // Fall through to Paula
                }
            }
        }

        // Try Paula hardware
        if is_paula_available() {
            return Result::Ok(AudioDriver::Paula { channels: 4 })
        }

        // No audio available - return null driver
        return Result::Ok(AudioDriver::Null)
    }

    pub fn play(&self, sample: &Sample) -> Result<(), AudioError>
    {
        match *self {
            AudioDriver::Paula { channels } => play_paula(sample,
                channels),
            AudioDriver::Ahi { device_id } => play_ahi(sample,
                device_id),
            AudioDriver::Null => Result::Ok(()), // Silent
                success
        }
    }
}

```

7.11.3 The Callback Pattern

Pass functions as parameters for customizable behavior:

```

// Type alias for callback signature
type Comparator<T> = fn(&T, &T) -> i32

// Sort with custom comparator
pub fn sort_by<T>(items: &var [T], count: u32, compare:
    Comparator<T>) {
    // Simple bubble sort for demonstration
    for i in 0u32..count {
        for j in 0u32..(count - 1 - i) {
            if compare(&items[j], &items[j + 1]) > 0 {
                // Swap
                let temp = items[j]

```

```

        items[j] = items[j + 1]
        items[j + 1] = temp
    }
}
}

// Usage
fn example() {
    var scores = [100, 50, 75, 25, 90]

    // Sort ascending
    sort_by(&var scores, 5, |a: &i32, b: &i32| -> i32 { *a - *b })

    // Sort descending
    sort_by(&var scores, 5, |a: &i32, b: &i32| -> i32 { *b - *a })
}

```

7.11.4 The RAII Guard Pattern

Functions that return guards for automatic cleanup:

```

struct ForbidGuard {
    _dummy: u8,
}

impl ForbidGuard {
    pub fn new() -> ForbidGuard {
        unsafe { Forbid() }
        return ForbidGuard { _dummy: 0 }
    }
}

impl Drop for ForbidGuard {
    fn drop(&var self) {
        unsafe { Permit() }
    }
}

// Usage - interrupts disabled for the scope
fn critical_section() {
    let _guard = ForbidGuard::new()
    // Forbid() is active here
    modify_shared_data()
    // Permit() called automatically when _guard goes out of scope
}

```

7.11.5 The Visitor Pattern

Process structures without knowing their internals:

```

trait FileVisitor {
    fn visit_file(&var self, path: &Str, size: u32)
    fn visit_directory(&var self, path: &Str)
    fn should_descend(&self, path: &Str) -> bool
}

```

```
fn walk_directory(path: &Str, visitor: &var dyn FileVisitor) ->
Result<(), DosError> {
    let lock = Lock(path.as_ptr(), ACCESS_READ)
    if lock == 0 {
        return Result::Err(DosError::NotFound)
    }
    defer { UnLock(lock) }

    var fib = @zeroed(FileInfoBlock)
    if !Examine(lock, &var fib) {
        return Result::Err(DosError::ExamineFailed)
    }

    while ExNext(lock, &var fib) {
        let name = fib_name(&fib)

        if fib.fib_DirEntryType > 0 {
            // Directory
            visitor.visit_directory(&name)
            if visitor.should_descend(&name) {
                let sub_path = path.concat(&name)
                walk_directory(&sub_path, visitor)?
            }
        } else {
            // File
            visitor.visit_file(&name, fib.fib_Size as u32)
        }
    }

    return Result::Ok(())
}

// Example visitor
struct FileSizeCounter {
    total_bytes: u64,
    file_count: u32,
}

impl FileVisitor for FileSizeCounter {
    fn visit_file(&var self, path: &Str, size: u32) {
        self.total_bytes += (u64)size
        self.file_count += 1
    }

    fn visit_directory(&var self, path: &Str) {
        // Ignore directories
    }

    fn should_descend(&self, path: &Str) -> bool {
        return true // Descend into all directories
    }
}
```


7.12 Showcase: Animation System

This showcase demonstrates functions and methods in a practical animation system:

```
// animation_system_showcase.novus
// Demonstrates: functions, methods, traits, generics, closures

from std::io import println
from std::core import Option, Clone
from std::collections::vec import Vec

//
// =====
// Easing Functions
//
// =====

type EasingFn = fn(fixed32) -> fixed32

// Linear interpolation (no easing)
pub fn ease_linear(t: fixed32) -> fixed32 {
    return t
}

// Quadratic ease-in
pub fn ease_in_quad(t: fixed32) -> fixed32 {
    return t * t
}

// Quadratic ease-out
pub fn ease_out_quad(t: fixed32) -> fixed32 {
    return t * (2.0 - t)
}

// Quadratic ease-in-out
pub fn ease_in_out_quad(t: fixed32) -> fixed32 {
    if t < 0.5 {
        return 2.0 * t * t
    }
    return -1.0 + (4.0 - 2.0 * t) * t
}

//
// =====
// Animatable Trait
//
// =====

trait Animatable {
    fn lerp(&self, target: &Self, t: fixed32) -> Self
}

impl Animatable for i32 {
    fn lerp(&self, target: &i32, t: fixed32) -> i32 {
        let diff = *target - *self
        return *self + (fixed32::from_int(diff) * t).to_int()
    }
}
```

```

impl Animatable for fixed32 {
    fn lerp(&self, target: &fixed32, t: fixed32) -> fixed32 {
        return *self + (*target - *self) * t
    }
}

//
// =====
// Animation State
//
// =====

enum AnimationState {
    Playing,
    Paused,
    Finished,
}

//
// =====
// Tween Structure
//
// =====

struct Tween<T> where T: Animatable + Clone {
    start_value: T,
    end_value: T,
    duration_frames: u32,
    current_frame: u32,
    easing: EasingFn,
    state: AnimationState,
}

impl<T> Tween<T> where T: Animatable + Clone {
    // Constructor with default linear easing
    pub fn new(start: T, end: T, duration: u32) -> Tween<T> {
        return Tween {
            start_value: start,
            end_value: end,
            duration_frames: duration,
            current_frame: 0,
            easing: ease_linear,
            state: AnimationState::Playing,
        }
    }

    // Builder method to set easing function
    pub fn with_easing(&var self, easing: EasingFn) -> &var
        Tween<T> {
        self.easing = easing
        return self
    }

    // Get current interpolated value
    pub fn value(&self) -> T {
        if self.duration_frames == 0 {
            return self.end_value.clone().unwrap()
        }
    }
}

```

```

    }

    // Calculate normalized time (0.0 to 1.0)
    let t_raw = fixed32::from_int((i32)self.current_frame)
        / fixed32::from_int((i32)self.duration_frames)

    // Clamp to [0, 1]
    let t_clamped = if t_raw > 1.0 { 1.0 } else { t_raw }

    // Apply easing function
    let t_eased = (self.easing)(t_clamped)

    // Interpolate between start and end
    return self.start_value.lerp(&self.end_value, t_eased)
}

// Advance animation by one frame
pub fn tick(&var self) {
    match self.state {
        AnimationState::Playing => {
            self.current_frame += 1
            if self.current_frame >= self.duration_frames {
                self.state = AnimationState::Finished
            }
        },
        _ => { }
    }
}

// Check if animation is complete
pub fn is_finished(&self) -> bool {
    match self.state {
        AnimationState::Finished => true,
        _ => false,
    }
}

// Pause the animation
pub fn pause(&var self) {
    match self.state {
        AnimationState::Playing => {
            self.state = AnimationState::Paused
        },
        _ => { }
    }
}

// Resume the animation
pub fn resume(&var self) {
    match self.state {
        AnimationState::Paused => {
            self.state = AnimationState::Playing
        },
        _ => { }
    }
}

// Reset to beginning

```

```

    pub fn reset(&var self) {
        self.current_frame = 0
        self.state = AnimationState::Playing
    }
}

//
// =====
// Animation Sequence
// =====

struct AnimationSequence<T> where T: Animatable + Clone {
    tweens: Vec<Tween<T>>,
    current_index: u32,
    loop_count: i32, // -1 for infinite
    loops_remaining: i32,
}

impl<T> AnimationSequence<T> where T: Animatable + Clone {
    pub fn new() -> Option<AnimationSequence<T>> {
        let tweens = Vec::<Tween<T>>::new()
        return Option::Some(AnimationSequence {
            tweens: tweens,
            current_index: 0,
            loop_count: 1,
            loops_remaining: 1,
        })
    }

    pub fn add(&var self, tween: Tween<T>) {
        self.tweens.push(tween)
    }

    pub fn set_loop_count(&var self, count: i32) {
        self.loop_count = count
        self.loops_remaining = count
    }

    pub fn value(&self) -> Option<T> {
        if self.current_index >= self.tweens.len() {
            return Option::None
        }
        return
            Option::Some(self.tweens.get(self.current_index).unwrap().value())
    }

    pub fn tick(&var self) -> bool {
        if self.current_index >= self.tweens.len() {
            return false // Sequence complete
        }

        let current =
            self.tweens.get_mut(self.current_index).unwrap()
            current.tick()

        if current.is_finished() {
            self.current_index += 1
        }
    }
}

```

```

        // Check if we should loop
        if self.current_index >= self.tweens.len() {
            if self.loop_count < 0 || self.loops_remaining >
            1 {
                // Reset for next loop
                self.current_index = 0
                for i in 0u32..self.tweens.len() {
                    self.tweens.get_mut(i).unwrap().reset()
                }
                if self.loops_remaining > 0 {
                    self.loops_remaining -= 1
                }
                return true // Still playing
            }
            return false // Done
        }
    }

    return true // Still playing
}

//
=====
// Example Usage
//
=====

fn main() -> i32 {
    println("=== Animation System Showcase ===\n")

    // Simple tween with linear easing
    println("--- Linear Tween (0 to 100, 10 frames) ---")
    var linear_tween = Tween::new(0, 100, 10)

    while !linear_tween.is_finished() {
        println(f"Frame {linear_tween.current_frame}: value =
        {linear_tween.value()}")
        linear_tween.tick()
    }

    // Tween with custom easing
    println("\n--- Ease-Out Quad Tween ---")
    var eased_tween = Tween::new(0, 100, 10)
    eased_tween.with_easing(ease_out_quad)

    while !eased_tween.is_finished() {
        println(f"Frame {eased_tween.current_frame}: value =
        {eased_tween.value()}")
        eased_tween.tick()
    }

    // Animation sequence
    println("\n--- Animation Sequence (3 tweens, 2 loops) ---")
    match AnimationSequence::<i32>::new() {
        Option::Some(var seq) => {
            seq.add(Tween::new(0, 50, 5))

```

```
seq.add(Tween::new(50, 25, 5))
seq.add(Tween::new(25, 100, 5))
seq.set_loop_count(2)

var frame = 0
while seq.tick() {
  match seq.value() {
    Option::Some(val) => {
      println(f"Frame {frame}: value = {val}")
    },
    Option::None => { }
  }
  frame += 1
},
Option::None => {
  println("Failed to create sequence")
}

println("\n=== Showcase Complete ===")
return 0
}
```

This showcase demonstrates:

- **Function types** as easing function callbacks
- **Traits** for defining animatable behavior
- **Generic structs and methods** with trait bounds
- **Builder pattern** with `with_easing`
- **Consuming and reference self** in methods
- **Match expressions** for state handling
- **Closures** (implicitly in function type usage)

7.13 Summary

This chapter covered the complete function system in Novus:

- **Function basics:** Declaration, parameters, return values
- **Reference parameters:** `&T` for reading, `&var T` for modifying
- **Traits:** Defining and implementing shared behavior
- **Methods:** Impl blocks, self receivers, associated functions
- **Generics:** Type parameters, trait bounds, turbofish syntax
- **Overloading:** Multiple functions with the same name
- **Const functions:** Compile-time evaluation
- **Closures:** Anonymous functions capturing environment

- **External functions:** AmigaOS FFI and system calls

Functions in Novus are designed for clarity and safety. Explicit parameter passing, trait bounds, and ownership semantics help prevent bugs while keeping the code readable.

The next chapter explores modules and project organization, showing how to structure larger Novus programs.

Chapter 8

Modules and Project Organization

“The purpose of abstraction is not to be vague, but to create a new semantic level in which one can be absolutely precise.”

—Edsger W. Dijkstra

As programs grow beyond a single file, organization becomes critical. Novus provides a module system that enforces explicit dependencies and visibility, preventing the tangled “include spaghetti” that plagues large C codebases. This chapter covers how to structure your code into modules, manage visibility, and organize multi-file projects.

Novus provides a module system for organizing code into logical units with explicit visibility control. This chapter covers imports, visibility modifiers, project structure, and the standard library organization.

8.1 Module Basics

In Novus, each source file (`.novus`) is a module. The module’s namespace is determined by its location in the directory structure.

8.1.1 Module Naming

Module paths use `::` as the separator:

File Path	Module Path
<code>src/main.novus</code>	<code>main</code>
<code>src/renderer.novus</code>	<code>renderer</code>
<code>src/graphics/bitmap.novus</code>	<code>graphics::bitmap</code>
<code>src/sound/player.novus</code>	<code>sound::player</code>

The standard library uses the `std` prefix:

File Path	Module Path
<code>std/core.novus</code>	<code>std::core</code>
<code>std/collections/vec.novus</code>	<code>std::collections::vec</code>
<code>std/ffi/exec.novus</code>	<code>std::ffi::exec</code>

8.2 Imports

Use the `from ... import ...` syntax to bring items from other modules into scope.

8.2.1 Basic Imports

Import specific items by name:

```
// Import specific types
from std::core import Option, Result

// Import functions and constants
from std::ffi::exec import AllocMem, FreeMem
from std::ffi::amiga_consts import MEMF_FAST, MEMF_CLEAR
```

Coming from C

C's `#include` literally copies header files into your source, leading to:

- Include order dependencies
- Include guards (`#ifndef/#define/#endif`)
- Slow compilation from re-parsing headers
- Polluted namespaces with everything in headers

```
// C: Headers are textually included
#include <exec/memory.h>
#include <proto/exec.h>      // Order matters!
```

Novus imports are semantic, not textual:

```
// Novus: Import specific symbols
from std::ffi::exec import AllocMem, FreeMem
from std::ffi::amiga_consts import MEMF_FAST
```

Benefits:

- No include guards needed
- Order doesn't matter
- Only imported names are in scope
- Faster compilation (no re-parsing)
- No “header file” vs “source file” distinction

8.2.2 Wildcard Imports

Import all public items from a module:

```
// Import everything from amiga_structs
from std::ffi::amiga_structs import *

// Now Task, MsgPort, Library, etc. are all in scope
```

Use wildcards sparingly—they can make it unclear where symbols come from.

8.2.3 Pattern Imports

Import multiple items matching a prefix or suffix pattern:

```
// Import all IDCMP_* constants (prefix pattern)
from std::ffi::amiga_consts import IDCMP_*

// This imports IDCMP_CLOSEWINDOW, IDCMP_MOUSEBUTTONS, etc.

// Import all *Error types (suffix pattern)
from std::error::errors import *Error

// This imports DosError, ExecError, GraphicsError, etc.
```

8.2.4 Import Aliases

Rename imports to avoid conflicts or for convenience:

```
from std::collections::vec import Vec as DynamicArray
from std::collections::hashmap import HashMap as Map
```

Using the aliases:

```
// ... (with DynamicArray and Map imported above)
var items = DynamicArray::<i32>::new()
var lookup = Map::<Str, i32>::new()
```

8.2.5 Multiple Imports

Import from multiple modules at the top of your file:

```
// Recommended import order for stdlib files:
// 1. std::core - fundamental types
from std::core import Option, Result, Drop, Clone

// 2. std::collections - data structures
from std::collections::vec import Vec

// 3. std::memory - memory management
from std::memory::block import MemoryBlock

// 4. std::strings - string types
from std::strings::core import Str

// 5. std::error - error types
from std::error::errors import ExecError

// 6. std::ffi - AmigaOS bindings
from std::ffi::amiga_structs import Task, MsgPort
from std::ffi::amiga_consts import MEMF_PUBLIC
from std::ffi::exec import AllocMem, FreeMem
```

8.3 Visibility Modifiers

Novus has four visibility levels that control where items can be accessed.

8.3.1 Private (Default)

Items without a visibility modifier are private to their module:

```
// Only visible in this file
const BUFFER_SIZE: u32 = 1024

fn helper() -> i32 {
    42
}

struct InternalData {
    value: i32,
}
```

Private items cannot be imported by other modules.

8.3.2 Public (pub)

The `pub` keyword exports items, making them visible to other modules:

```
// Exported - can be imported by other modules
pub const VERSION: u32 = 1

pub fn create_bitmap(w: u32, h: u32) -> Result<Bitmap,
    GraphicsError> {
    // ...
}

pub struct Bitmap {
    width: u32,
    height: u32,
    data: *u8,
}
```

Use `pub` for your module’s public API—functions, types, and constants that other code should use.

Coming from C

C uses `static` for file-local symbols and relies on headers for “public” APIs:

```
// C: static = private to this file
static int helper(void) { return 42; }

// C: non-static = visible (if declared in header)
int public_api(void) { return helper(); }
```

Novus inverts this: private by default, explicit `pub` for public:

```
// Novus: private by default
fn helper() -> i32 { return 42 }
```

```
// Novus: pub = exported
pub fn public_api() -> i32 { return helper() }
```

The mapping:

- C `static` $\approx$ Novus (default, no modifier)
- C non-static + header $\approx$ Novus `pub`
- C `extern` $\approx$ Novus `extern` (for FFI only)

8.3.3 Internal (`internal`)

The `internal` modifier makes items visible within the same package (workspace) but not exported to external code:

```
// ... (internal visibility - illustrative)
internal const SHARED_BUFFER_SIZE: u32 = 4096

internal fn shared_helper() -> i32 {
    return 100
}
```

Use `internal` for implementation details shared between modules in a library.

8.3.4 Extern (`extern`)

The `extern` keyword declares items defined elsewhere, resolved at link time:

```
// ... (extern declarations - illustrative)
extern var SysBase: *ExecBase

// Hardware register at specific address
extern var CUSTOM: *Custom at 0xdff000

// Function from another object file or assembly
extern fn asm_memcpy(dest: *u8, src: *u8, len: u32)
```

The `at ADDRESS` syntax is used for memory-mapped hardware registers. Common examples include `CUSTOM at 0xdff000` (custom chips), `CIA_A at 0xbfe001`, and `CIA_B at 0xbfd000`.

Note: AmigaOS library functions are imported from the standard library FFI modules, not declared as `extern` in user code:

```
// Import AmigaOS functions from std::ffi
from std::ffi::exec import AllocMem, FreeMem, Wait, Signal
from std::ffi::dos import Open, Close, Read, Write
```

External declarations have no implementation in Novus—they’re provided by AmigaOS, assembly files, or C libraries.

8.4 Re-exports

Re-exports allow modules to expose items from other modules as part of their own API. This is useful for creating convenient public interfaces.

8.4.1 The `pub use` Syntax

```
// ... (example from std/prelude.novus - illustrative)

// Re-export core types
pub use std::core::Option, Result, Drop, From

// Re-export string types
pub use std::strings::core::Str

// Re-export error types
pub use std::error::errors::DosError, ExecError, IntuitionError
```

Now users can import from `std::prelude` instead of multiple modules:

```
// Instead of:
from std::core import Option, Result
from std::strings::core import Str
from std::error::errors import DosError

// Users can write:
from std::prelude import Option, Result, Str, DosError
```

8.4.2 Multiple Re-exports

Re-export multiple items to create a convenient API:

```
// ... (example from std/async/mod.novus - illustrative)

// Re-export future types
pub use std::async::future::Future, Poll, Context, Waker
pub use std::async::future::Ready, Pending, ready, pending

// Re-export executor functions
pub use std::async::executor::block_on_sleep, poll_sleep_once

// Re-export sleep utilities
pub use std::async::sleep::Sleep, sleep, sleep_secs, sleep_millis
```

8.5 Constants and Static Variables

8.5.1 Constants (`const`)

Constants are compile-time values that are inlined at every use site:

```
const MAX_SPRITES: u32 = 8
const PI: i32 = 3_14159 // Fixed-point representation
pub const VERSION: u32 = 1
```

```
// Constants can use const functions
const fn double(x: i32) -> i32 { x * 2 }
const DOUBLED: i32 = double(21) // 42 at compile time
```

Properties of constants:

- Must be computable at compile time
- No memory address (inlined at use sites)
- Always immutable
- Preferred for configuration values

8.5.2 Static Variables (`static`)

Static variables have a fixed memory location and live for the entire program:

```
// Immutable static - initialized data (size inferred from
// elements)
static SINE_TABLE: [i16] = [
    0, 804, 1608, 2410, 3212, 4011, 4808, // ...
]

// Mutable static - use with caution
static var frame_count: u32 = 0
static var initialized: bool = false
```

Properties of static variables:

- Have a memory address (stored in `.data` or `.bss`)
- Live for the entire program lifetime
- Immutable by default (`static var` for mutable)
- Must be initialized with a constant expression

8.5.3 Mutable Statics

Mutable static variables are declared with `static var`:

```
static var g_chip_pool: ChipPool = ChipPool {
    blocks: null,
    count: 0,
}

pub fn get_pool() -> *ChipPool {
    return &var g_chip_pool
}
```

Warning: Mutable statics are dangerous in multitasking environments. For thread-safe global state, consider using:

- Atomic types for simple counters
- Semaphores for complex data
- Immutable statics with interior mutability

8.6 Project Structure

Novus uses TOML configuration files to define projects and workspaces.

8.6.1 Single Project

A simple project structure:

```
myapp/  
  project.toml      # Project configuration  
  src/  
    main.novus      # Entry point  
    renderer.novus  # Additional module  
    sound/  
      player.novus  # Submodule  
      mixer.novus
```

The `project.toml` file:

```
[package]  
name = "myapp"  
version = "0.1.0"  
type = "cli"           # cli, workbench, dual, library, device  
entry = "src/main.novus" # Entry point (optional, defaults to  
    main.novus)  
description = "My Amiga application"  
authors = ["Your Name"]  
  
[build]  
target_cpu = "68020"    # 68000, 68020, 68040, 68060  
fpu = "auto"           # soft, 68881, 68040, auto  
output = "build"  
optimization_level = 2  
  
[paths]  
src = "src"
```

8.6.2 Workspaces

Workspaces contain multiple related projects:

```
myworkspace/  
  workspace.toml      # Workspace configuration  
  mylib/  
    project.toml  
    src/  
      lib.novus  
  myapp/  
    project.toml  
    src/  
      main.novus  
  examples/  
    project.toml  
    src/  
      demo.novus
```


The `workspace.toml` file:

```
[workspace]
name = "myworkspace"
version = "0.1.0"
members = ["mylib", "myapp", "examples"] # Directory names of
    member projects

[workspace.build]
target_cpu = "68020"
fpu = "auto"
optimization_level = 2
```

The `members` array lists directory names of projects within the workspace.

8.6.3 Project Types

Novus supports several project types, each with different startup behavior:

cli Command-line application launched from Shell with `argc/argv` arguments

workbench Application launched by double-clicking icons; receives `WBStartup` message with file arguments

dual Application that handles both CLI and Workbench launches

library AmigaOS shared library (`.library`) with `ROMTag` and vector table

device AmigaOS device driver (`.device`) with `BeginIO/AbortIO` interface

Important: Workbench applications *must* reply to the `WBStartup` message using `ReplyMsg()` when exiting, or Workbench will hang. Use `Input() == 0` to detect Workbench launches.

8.6.4 Creating Projects

Use the `novus new` command to create projects:

```
# Create a new workspace
novus new myworkspace

# Create a CLI project in the workspace
cd myworkspace
novus new mytool --type cli

# Create a Workbench application
novus new mygui --type workbench

# Create a library
novus new mylib --type library
```

8.6.5 Building Projects

Build commands:

```
# Build all projects in workspace
novus build

# Build a specific project
novus build mytool

# Build with specific CPU target
novus build --cpu 68040
```

8.7 Standard Library Organization

The Novus standard library is organized into focused modules:

8.7.1 Core Modules

std::core Fundamental types: `Option`, `Result`, `Drop`, `Clone`, `Eq`, etc.

std::prelude Common re-exports for convenience

8.7.2 Data Structures

std::collections::vec Dynamic array (`Vec<T>`)

std::collections::hashmap Hash map (`HashMap<K, V>`)

std::collections::hashset Hash set (`HashSet<T>`)

std::collections::vecdeque Double-ended queue

std::collections::ringbuffer Fixed-size ring buffer

std::collections::bitset Bit set

std::collections::slotmap Slot map for entity storage

std::collections::freelist Free list allocator

std::collections::smallvec Small-buffer-optimized vector

std::collections::intrusive_list Intrusive doubly-linked list

8.7.3 Memory Management

std::memory::block Raw memory blocks (`MemoryBlock`)

std::memory::chip_pool Chip memory pool

std::memory::chip_cache Chip memory cache

8.7.4 Strings

std::strings::core String slice type (`Str`)

std::strings::format Formatting and Display trait

8.7.5 Error Handling

std::error::errors Error types: `DosError`, `ExecError`, `GraphicsError`, etc.

8.7.6 I/O and Networking

std::io Basic I/O functions

std::net::socket BSD socket interface

std::net::tcp TCP client/server

std::net::udp UDP sockets

std::net::dns DNS resolution

8.7.7 Concurrency

std::async Async/await runtime

std::sync Synchronization primitives

std::thread Task management

8.7.8 Graphics and UI

std::graphics::bitmap Bitmap handling

std::graphics::sprite Hardware sprites

std::graphics::copper Copper list generation

std::graphics::blitter Blitter operations

std::graphics::intuition Intuition window/screen wrappers

std::ui::window Window management

std::ui::screen Screen management

std::ui::menu Menu creation

std::ui::events Event handling

8.7.9 AmigaOS FFI

The FFI modules provide direct bindings to AmigaOS libraries:

std::ffi::exec Exec library (memory, tasks, signals)

std::ffi::dos DOS library (files, paths, CLI)

std::ffi::graphics Graphics library (rendering, text)

std::ffi::intuition Intuition library (windows, screens)

std::ffi::layers Layers library (layer management)

std::ffi::gadtools GadTools library (GUI toolkit)

std::ffi::utility Utility library (tags, hooks)

std::ffi::workbench Workbench library (icons, AppIcons)

std::ffi::icon Icon library (icon handling)

std::ffi::asl ASL library (file/font requesters)

std::ffi::diskfont Diskfont library (font loading)

std::ffi::commodities Commodities library (hotkeys)

std::ffi::iffparse IFFParse library (IFF files)

std::ffi::timer_device Timer device (delays, timing)

std::ffi::audio_device Audio device (Paula access)

std::ffi::bsdsocket BSD sockets (networking)

std::ffi::amiga_structs AmigaOS structure definitions

std::ffi::amiga_consts AmigaOS constants

8.7.10 Other Modules

std::math Fixed-point math, trigonometry

std::audio Paula audio, ProTracker player

std::test Test framework

std::os Operating system utilities

8.8 The Prelude

The prelude (**std::prelude**) re-exports commonly used types:

```
// ... (contents of std/prelude.novus - illustrative)
pub use std::core::Option, Result, Drop, From
pub use std::strings::core::Str
pub use std::error::errors::DosError, ExecError, IntuitionError,
    GraphicsError
```

Import the prelude for quick access to common types:

```
from std::prelude import *

pub fn main() -> i32 {
    let result: Result<i32, DosError> = do_something()
    match result {
        Ok(value) => return value,
        Err(_) => return 1
    }
}
```

8.9 Best Practices

8.9.1 Visibility Guidelines

1. **Default to private** — only make items `pub` when needed for external API
2. **Use internal for shared implementation** — better than making everything `pub`
3. **Document public APIs** — anything `pub` should have doc comments
4. **Keep modules focused** — one module, one responsibility

8.9.2 Import Guidelines

1. **Prefer specific imports** over wildcards for clarity
2. **Use wildcards sparingly** — only for well-known modules like `amiga_consts`
3. **Group imports logically** — core types, then collections, then FFI
4. **Use aliases** to avoid name conflicts

8.9.3 Project Organization

1. Use **workspaces** for multi-project development
2. **Separate library and application code** into different projects
3. **Keep tests in a separate project** that depends on the library
4. Use **meaningful directory structure** that matches module paths

8.9.4 Constants vs Statics

Use	When
<code>const</code>	Configuration values, magic numbers, version strings
<code>static</code>	Lookup tables, large data that shouldn't be inlined
<code>static var</code>	Global state (use sparingly, prefer safer alternatives)

8.10 Common Module Patterns

This section presents idiomatic patterns for organizing Novus projects.

8.10.1 The Facade Pattern

Create a simple public interface that hides complex internal modules:

```
graphics/
  mod.novus          # Facade - re-exports public API
  bitmap.novus       # Internal: bitmap operations
  sprite.novus       # Internal: sprite handling
  copper.novus       # Internal: copper list generation
  blitter.novus      # Internal: blitter operations
```

The facade module (`mod.novus`):

```
// graphics/mod.novus - Public facade

// Re-export only what users need
pub use crate::graphics::bitmap::Bitmap, BitmapFlags
pub use crate::graphics::sprite::Sprite, SpriteSheet
pub use crate::graphics::copper::CopperList
pub use crate::graphics::blitter::Blit, BlitMode

// Convenience functions that combine internal modules
pub fn create_display(width: u16, height: u16, depth: u8)
    -> Result<Display, GraphicsError>
{
    let bitmap = bitmap::create_bitmap(width, height, depth)?
    let copper = copper::create_copper_list(&bitmap)?
    return Result::Ok(Display { bitmap, copper })
}
```

Users import from the facade:

```
// Clean import from facade
from mylib::graphics import Bitmap, Sprite, create_display

// Instead of multiple internal imports:
// from mylib::graphics::bitmap import Bitmap
// from mylib::graphics::sprite import Sprite
// etc.
```

8.10.2 The Error Module Pattern

Centralize error types in a single module:

```
// errors.novus
pub enum AppError {
    Io(DosError),
    Graphics(GraphicsError),
    Audio(AudioError),
    InvalidInput(*u8),
    NotFound,
    OutOfMemory,
}

impl AppError {
    pub fn from_dos(e: DosError) -> AppError {
        return AppError::Io(e)
    }

    pub fn from_graphics(e: GraphicsError) -> AppError {
        return AppError::Graphics(e)
    }
}

impl Display for AppError {
    fn display(&self, f: &var Formatter) -> bool {
        match *self {
            AppError::Io(ref e) => f.write_str("I/O error"),
            AppError::Graphics(ref e) => f.write_str("Graphics")
        }
    }
}
```

```

        error"),
    AppError::Audio(ref e) => f.write_str("Audio error"),
    AppError::InvalidInput(msg) => {
        f.write_str("Invalid input: ")
        f.write_str(Str::from_ptr(msg))
    },
    AppError::NotFound => f.write_str("Not found"),
    AppError::OutOfMemory => f.write_str("Out of memory"),
}
}
}

```

All modules use the centralized error type:

```

from crate::errors import AppError

pub fn load_asset(path: *u8) -> Result<Asset, AppError> {
    let file = open_file(path).map_err(AppError::from_dos)?
    // ...
}

```

8.10.3 The Configuration Module Pattern

Keep configuration constants in one place:

```

// config.novus
pub const APP_NAME: *u8 = "MyGame"
pub const VERSION: u32 = 0x0001_0000 // 1.0

// Display settings
pub const SCREEN_WIDTH: u16 = 320
pub const SCREEN_HEIGHT: u16 = 256
pub const SCREEN_DEPTH: u8 = 5

// Memory budgets
pub const CHIP_BUDGET: u32 = 512 * 1024 // 512 KB
pub const FAST_BUDGET: u32 = 1024 * 1024 // 1 MB

// Timing
pub const FRAME_RATE: u32 = 50 // PAL
pub const TICK_RATE: u32 = 100

// Build-time options
pub const DEBUG_ENABLED: bool = true
pub const PROFILE_ENABLED: bool = false

```

Import configuration where needed:

```

from crate::config import SCREEN_WIDTH, SCREEN_HEIGHT, CHIP_BUDGET

fn initialize_display() -> Result<Display, Error> {
    let display = create_display(SCREEN_WIDTH, SCREEN_HEIGHT)?
    // ...
}

```

8.10.4 The Plugin Architecture Pattern

Define a trait interface for plugins:

```
// plugin.novus - Plugin interface
pub trait Plugin {
    fn name(&self) -> *u8
    fn version(&self) -> u32
    fn init(&var self) -> Result<(), PluginError>
    fn shutdown(&var self)
    fn process(&var self, data: &var PluginData) -> Result<(),
        PluginError>
}

pub struct PluginManager {
    plugins: Vec<Box<dyn Plugin>>,
}

impl PluginManager {
    pub fn new() -> PluginManager {
        return PluginManager {
            plugins: Vec::new(),
        }
    }

    pub fn register(&var self, plugin: Box<dyn Plugin>) ->
        Result<(), PluginError> {
        plugin.init()?
        self.plugins.push(plugin)
        return Result::Ok(())
    }

    pub fn process_all(&var self, data: &var PluginData) ->
        Result<(), PluginError> {
        for i in 0u32..self.plugins.len() {
            self.plugins.get_mut(i).unwrap().process(data)?
        }
        return Result::Ok(())
    }
}
```

Implement plugins in separate modules:

```
// plugins/audio_visualizer.novus
from crate::plugin import Plugin, PluginData, PluginError

pub struct AudioVisualizer {
    active: bool,
    intensity: fixed32,
}

impl Plugin for AudioVisualizer {
    fn name(&self) -> *u8 { "AudioVisualizer" }
    fn version(&self) -> u32 { 0x0001_0000 }

    fn init(&var self) -> Result<(), PluginError> {
        self.active = true
        self.intensity = 1.0
    }
}
```



```

        return Result::Ok(())
    }

    fn shutdown(&var self) {
        self.active = false
    }

    fn process(&var self, data: &var PluginData) -> Result<(),
        PluginError> {
        if !self.active {
            return Result::Ok(())
        }
        // Process audio data and update visuals
        return Result::Ok(())
    }
}

```

8.10.5 The Module Index Pattern

For large modules with many files, create an index module:

```

collections/
  mod.novus           # Index - re-exports all collections
  vec.novus
  hashmap.novus
  hashset.novus
  vecdeque.novus
  ringbuffer.novus

```

The index module:

```

// collections/mod.novus
pub use crate::collections::vec::Vec
pub use crate::collections::hashmap::HashMap
pub use crate::collections::hashset::HashSet
pub use crate::collections::vecdeque::VecDeque
pub use crate::collections::ringbuffer::RingBuffer

```

Users can import from the index or specific modules:

```

// Import from index (common items)
from std::collections import Vec, HashMap

// Import from specific module (less common items)
from std::collections::ringbuffer import RingBuffer,
    RingBufferOptions

```

8.11 Showcase: Modular Game Framework

This showcase demonstrates module organization in a game framework:

```

amigame/
  workspace.toml
  engine/

```

```

project.toml
src/
  lib.novus          # Engine facade
  config.novus       # Configuration constants
  errors.novus       # Error types
  display/
    mod.novus        # Display subsystem facade
    screen.novus
    bitmap.novus
    palette.novus
  input/
    mod.novus        # Input subsystem facade
    keyboard.novus
    joystick.novus
    mouse.novus
  audio/
    mod.novus        # Audio subsystem facade
    paula.novus
    sample.novus
    module_player.novus
  resources/
    mod.novus        # Resource management facade
    loader.novus
    cache.novus
    asset.novus
game/
  project.toml
  src/
    main.novus       # Game entry point
    game_state.novus # Game state management
    entities/
      mod.novus
      player.novus
      enemy.novus
      projectile.novus

```

Engine Facade (engine/src/lib.novus):

```

// engine/src/lib.novus
// Public API facade for the game engine

// Re-export configuration
pub use crate::config::*

// Re-export error types
pub use crate::errors::{EngineError, DisplayError, InputError,
  AudioError}

// Re-export subsystem facades
pub use crate::display::{Display, Screen, Bitmap, Palette}
pub use crate::input::{Input, KeyCode, JoyButton, MouseButton}
pub use crate::audio::{Audio, Sample, ModulePlayer}
pub use crate::resources::{ResourceManager, Asset, AssetId}

// Engine core
pub struct Engine {
  display: Display,
  input: Input,

```

```

    audio: Audio,
    resources: ResourceManager,
    running: bool,
}

impl Engine {
    pub fn new() -> Result<Engine, EngineError> {
        let display = Display::init()?
        let input = Input::init()?
        let audio = Audio::init()?
        let resources = ResourceManager::new()?

        return Result::Ok(Engine {
            display,
            input,
            audio,
            resources,
            running: false,
        })
    }

    pub fn run(&var self, game: &var dyn Game) -> Result<(),
    EngineError> {
        self.running = true
        game.init(&var self.resources)?

        while self.running {
            self.input.poll()

            if self.input.quit_requested() {
                self.running = false
                break
            }

            game.update(&self.input)?
            game.render(&var self.display)?

            self.display.flip()
        }

        game.shutdown()
        return Result::Ok(())
    }

    pub fn quit(&var self) {
        self.running = false
    }
}

// Game trait - implemented by user's game
pub trait Game {
    fn init(&var self, resources: &var ResourceManager) ->
    Result<(), EngineError>
    fn update(&var self, input: &Input) -> Result<(), EngineError>
    fn render(&self, display: &var Display) -> Result<(),
    EngineError>
    fn shutdown(&var self)
}

```

Display Subsystem (engine/src/display/mod.novus):

```
// engine/src/display/mod.novus
// Display subsystem facade

pub use crate::display::screen::Screen, ScreenMode
pub use crate::display::bitmap::Bitmap, BitmapDepth
pub use crate::display::palette::Palette, Color

use crate::errors::DisplayError
use crate::config::{SCREEN_WIDTH, SCREEN_HEIGHT, SCREEN_DEPTH}

pub struct Display {
    screen: Screen,
    front_buffer: Bitmap,
    back_buffer: Bitmap,
    palette: Palette,
}

impl Display {
    pub fn init() -> Result<Display, DisplayError> {
        let screen = Screen::open(SCREEN_WIDTH, SCREEN_HEIGHT,
                                   SCREEN_DEPTH)?
        let front = Bitmap::create(SCREEN_WIDTH, SCREEN_HEIGHT,
                                   SCREEN_DEPTH)?
        let back = Bitmap::create(SCREEN_WIDTH, SCREEN_HEIGHT,
                                   SCREEN_DEPTH)?
        let palette = Palette::default()

        return Result::Ok(Display {
            screen,
            front_buffer: front,
            back_buffer: back,
            palette,
        })
    }

    pub fn clear(&var self) {
        self.back_buffer.clear()
    }

    pub fn blit(&var self, src: &Bitmap, x: i16, y: i16) {
        self.back_buffer.blit(src, x, y)
    }

    pub fn flip(&var self) {
        // Swap buffers
        let temp = self.front_buffer
        self.front_buffer = self.back_buffer
        self.back_buffer = temp

        // Update screen
        self.screen.show(&self.front_buffer)
    }

    pub fn set_palette(&var self, palette: &Palette) {
        self.palette = palette.clone().unwrap()
        self.screen.load_palette(&self.palette)
    }
}
```

```

    }
}

impl Drop for Display {
    fn drop(&var self) {
        // Cleanup handled by member drops
    }
}

```

Game Implementation (game/src/main.novus):

```

// game/src/main.novus
// Main game using the engine

from engine import Engine, Game, ResourceManager, Display, Input
from engine import EngineError, KeyCode
from crate::game_state import GameState
from crate::entities import Player, Enemy

struct MyGame {
    state: GameState,
    player: Player,
    enemies: Vec<Enemy>,
}

impl MyGame {
    pub fn new() -> MyGame {
        return MyGame {
            state: GameState::Title,
            player: Player::new(),
            enemies: Vec::new(),
        }
    }
}

impl Game for MyGame {
    fn init(&var self, resources: &var ResourceManager) ->
        Result<(), EngineError> {
        // Load game assets
        let player_gfx = resources.load_bitmap("player.iff"?
        self.player.set_graphics(player_gfx)

        // Initialize enemies
        for i in 0..5 {
            var enemy = Enemy::new()
            enemy.set_position(50 + i * 40, 50)
            self.enemies.push(enemy)
        }

        return Result::Ok(())
    }

    fn update(&var self, input: &Input) -> Result<(),
        EngineError> {
        match self.state {
            GameState::Title => {
                if input.key_pressed(KeyCode::Space) {
                    self.state = GameState::Playing
                }
            }
        }
    }
}

```

```

        },
        GameState::Playing => {
            // Update player based on input
            if input.key_down(KeyCode::Left) {
                self.player.move_left()
            }
            if input.key_down(KeyCode::Right) {
                self.player.move_right()
            }
            if input.key_pressed(KeyCode::Space) {
                self.player.fire()
            }

            // Update enemies
            for i in 0u32..self.enemies.len() {
                self.enemies.get_mut(i).unwrap().update()
            }

            // Check collisions
            self.check_collisions()
        },
        GameState::GameOver => {
            if input.key_pressed(KeyCode::Space) {
                self.reset()
                self.state = GameState::Title
            }
        },
    },
}

return Result::Ok(())
}

fn render(&self, display: &var Display) -> Result<(),
EngineError> {
    display.clear()

    match self.state {
        GameState::Title => {
            // Draw title screen
        },
        GameState::Playing => {
            self.player.draw(display)
            for i in 0u32..self.enemies.len() {
                self.enemies.get(i).unwrap().draw(display)
            }
        },
        GameState::GameOver => {
            // Draw game over screen
        },
    }
}

return Result::Ok(())
}

fn shutdown(&var self) {
    // Cleanup
}

```

```

}

fn main() -> i32 {
    // Initialize engine
    let engine_result = Engine::new()
    match engine_result {
        Result::Ok(var engine) => {
            var game = MyGame::new()

            match engine.run(&var game) {
                Result::Ok(()) => return 0,
                Result::Err(e) => {
                    println(f"Game error: {e}")
                    return 1
                }
            }
        },
        Result::Err(e) => {
            println(f"Failed to initialize engine: {e}")
            return 1
        }
    }
}

```

This showcase demonstrates:

- **Workspace structure** with engine library and game application
- **Facade pattern** hiding subsystem complexity
- **Re-exports** creating clean public APIs
- **Trait-based plugin architecture** with the `Game` trait
- **Module hierarchy** organizing related functionality
- **Visibility control** keeping internals private
- **Configuration module** for centralized constants

8.12 Summary

This chapter covered the Novus module system:

- **Imports:** `from module::path` import items
- **Visibility:** `Private` (default), `pub`, `internal`, `extern`
- **Re-exports:** `pub use` for API convenience
- **Constants:** `const` for compile-time values
- **Statics:** `static` for fixed-location data
- **Projects:** `project.toml` and `workspace.toml`
- **Standard library:** Well-organized modules for all needs

Good module organization makes code easier to understand, maintain, and reuse. Start with private visibility and only expose what's needed for your public API.

The next chapter covers error handling with `Result` and `Option`.

Chapter 9

Attributes

“Metadata is a love note to the future.” — Jason Scott

Attributes provide a way to attach metadata to code elements—functions, structs, enums, and variables. The compiler uses this metadata to alter behavior, generate code, or produce warnings. This chapter covers all implemented attributes with practical examples.

9.1 Understanding Attributes

In systems programming, we often need to tell the compiler things that can’t be expressed in the code itself. Should this function be inlined? Is this struct packed for hardware access? Is this function safe to call from an interrupt handler?

Attributes provide these hints and directives. Unlike comments, attributes are part of the language—the compiler parses them, validates them, and acts on them.

Coming from C

GCC/Clang Attributes: C has `__attribute__((packed))`, `__attribute__((always_inline))`, etc. Novus uses cleaner `@packed` and `@inline` syntax. Same ideas, better ergonomics.

9.2 Attribute Syntax

Novus supports two equivalent attribute syntaxes. Choose whichever feels more natural for your codebase.

9.2.1 At-Sign Syntax (Recommended)

The primary syntax uses `@` followed by the attribute name:

```
@test
fn test_addition() {
    expect_eq(2 + 2, 4, "basic addition")
}

@inline
fn fast_add(a: i32, b: i32) -> i32 {
    return a + b
}
```

This syntax is concise and visually distinct from code.

9.2.2 Bracket Syntax

An alternative bracket syntax is also supported:

```
#[test]
fn test_multiplication() {
    expect_eq(3 * 7, 21, "multiplication")
}

#[derive(Eq, Hash)]
struct Point {
    x: i32,
    y: i32,
}
```

The bracket syntax is useful when attributes have complex arguments.

9.2.3 Attribute Arguments

Attributes can accept arguments—both named and positional:

```
// Named arguments
@library(name = "example.library", version = 1, revision = 0)
pub struct ExampleLibrary {
    counter: u32,
}

// Positional arguments
@test("Verify sorting performance")
fn test_sort_perf() {
    // ...
}

// Flag arguments
@test(skip = "Not yet implemented")
fn test_future() {
    // Skipped during test runs
}

// Mixed arguments
@libfunc(offset = -30) // LVO offset -30 bytes
pub fn my_library_function() -> i32 {
    return 42
}
```

9.2.4 Multiple Attributes

Multiple attributes can be stacked on a single item:

```
@test
@inline
pub fn test_fast_path() {
    // This function is both a test case and should be inlined
}

@interrupt
@export
```

```
pub fn vblank_handler() {
    // Interrupt handler that must be preserved in the binary
}
```

The order of attributes generally doesn't matter, though some combinations are invalid (the compiler will warn you).

9.3 Testing Attributes

Novus has first-class support for testing through attributes.

9.3.1 @test

Marks a function as a test case. The `novus test` command discovers and compiles these into a test runner.

```
@test
fn test_basic() {
    expect_eq(1 + 1, 2, "one plus one")
}

@test("Descriptive test name")
fn test_with_description() {
    expect_eq(true, true, "truth is true")
}
```

Arguments:

positional string Optional test description

skip = "reason" Skip this test with an explanation

```
@test(skip = "Hardware not available in CI")
fn test_blitter_operations() {
    // This test requires real Amiga hardware
}

@test(skip)
fn test_known_broken() {
    // Skipped without explanation (not recommended)
}
```

9.3.2 @bench / @benchmark

Marks a function for performance testing. Both names are equivalent:

```
@bench
fn bench_sort() {
    var data = generate_random_data(1000)
    sort(&var data)
}

@benchmark
fn bench_hash() {
    var map = HashMap::new()
```

```
for i in 0..1000 {  
    map.insert(i, i * 2)  
}  
}
```

Benchmark functions are included when you run `novus test -b`. The runner executes each benchmark multiple times and reports timing statistics.

9.4 Optimization Attributes

These attributes guide the compiler's optimization decisions.

9.4.1 @inline

Hints that the compiler should inline this function at call sites:

```
@inline  
fn min(a: i32, b: i32) -> i32 {  
    if a < b { return a } else { return b }  
}  
  
@inline  
fn abs(x: i32) -> i32 {  
    if x < 0 { return -x } else { return x }  
}
```

Small, frequently-called functions benefit most from inlining. Benefits include:

- Elimination of function call overhead
- Better register allocation across the inlined code
- Opportunities for constant folding when arguments are known

The compiler may ignore this hint for:

- Recursive functions
- Very large functions
- Functions with complex control flow

9.4.2 @noinline

Prevents inlining even when the compiler might otherwise choose to inline:

```
@noinline  
fn error_handler() {  
    // Keep error handling out of hot paths  
    // to improve instruction cache locality  
    log_error("Something went wrong")  
    cleanup_resources()  
}  
  
@noinline  
fn cold_path_check(value: i32) -> bool {  
    // Rarely-taken branch - don't pollute hot code
```

```
    return validate_complex_condition(value)
}
```

Use `@noinline` for:

- Error handling code (rarely executed)
- Large functions called from tight loops
- Functions where you need predictable stack usage

9.5 Layout Attributes

9.5.1 @packed

Removes padding between struct fields:

```
// Without @packed: 8 bytes (a=1, padding=3, b=4)
struct Normal {
    a: u8,
    b: u32,
}

// With @packed: 5 bytes (a=1, b=4, no padding)
@packed
struct Packed {
    a: u8,
    b: u32,
}
```

Coming from C

C Equivalent: `#pragma pack(1)` or `__attribute__((packed))`. Same effect, but Novus makes it explicit per-struct.

Warning: Packed structs may have unaligned fields, which can cause:

- Slower memory access (multiple bus cycles)
- Address errors on 68000/68010 (crashes!)
- Correct but slow access on 68020+

Use when:

- Matching external data formats (file formats, network protocols)
- Mapping hardware registers
- Memory-constrained situations where every byte matters

9.6 Trait Derivation

9.6.1 @derive / #[derive]

Automatically generates trait implementations for structs:

```
#[derive(Eq)]
struct Point {
    x: i32,
    y: i32,
}

#[derive(Eq, Hash)]
struct UserId {
    id: u32,
}
```

Supported traits:

Eq Generates `eq(&self, &Self) -> bool` comparing all fields

Hash Generates `hash(&self) -> u32` combining field hashes

You can derive multiple traits at once:

```
#[derive(Eq, Hash)]
struct CacheKey {
    module: u32,
    offset: u32,
}

// Now CacheKey can be used as HashMap key
var cache = HashMap::<CacheKey, Data>::new()
```

9.7 Method Chaining

9.7.1 @chain

Enables fluent method chaining by making a method return `&var Self`:

```
struct WindowBuilder {
    width: i32,
    height: i32,
    title: *u8,
    flags: u32,
}

impl WindowBuilder {
    pub fn new() -> WindowBuilder {
        return WindowBuilder {
            width: 640,
            height: 480,
            title: "Window",
            flags: 0,
        }
    }

    @chain
    pub fn width(&var self, w: i32) {
        self.width = w
    }
}
```

```

@chain
pub fn height(&var self, h: i32) {
    self.height = h
}

@chain
pub fn title(&var self, t: *u8) {
    self.title = t
}

pub fn build(&self) -> Window {
    // Create actual window...
}
}

```

Usage:

```

var window = WindowBuilder::new()
    .width(800)
    .height(600)
    .title("My Application")
    .build()

```

The `@chain` attribute automatically adds an implicit `return self` at the end of the method.

9.8 Interoperability Attributes

9.8.1 @export

Marks a function for export with an unmangled name, preventing dead code elimination:

```

@export
pub fn my_callback(data: *u8) -> i32 {
    // Can be called from external C code or AmigaOS
    return process(data)
}

@export
pub fn _UserLibInit() -> i32 {
    // Entry point called by AmigaOS
    return initialize_library()
}

```

Use this when:

- Creating callback functions for AmigaOS (hooks, interrupt handlers)
- Functions that need to be called from assembly or C
- Entry points that might appear unused to the optimizer

9.8.2 @extern_type

Marks a struct as an external C type, preventing Novus from emitting its definition:

```
@extern_type
struct timeval {
    tv_sec: i32,
    tv_usec: i32,
}

@extern_type
struct Library {
    node: Node,
    flags: u8,
    // ... matches C header definition
}
```

This tells the compiler: “This type is defined elsewhere (in C headers or runtime); just use it, don’t emit a definition.”

9.9 Synchronization Attributes

These attributes provide automatic synchronization wrappers.

9.9.1 @atomic

Wraps the function body in a Forbid/Permit critical section:

```
@atomic
pub fn update_shared_counter(counter: *u32) {
    // Forbid() called on entry
    *counter = *counter + 1
    // Permit() called on exit (even if returning early)
}

@atomic
pub fn swap_values(a: *i32, b: *i32) {
    let temp = *a
    *a = *b
    *b = temp
    // No task can switch during this operation
}
```

This prevents task switching while the function executes. The compiler ensures Permit() is called on all exit paths.

Use for:

- Short, fast operations on shared data
- Lock-free data structure updates
- Counter increments visible to multiple tasks

Avoid for:

- Functions that might block (I/O, waiting)
- Long-running computations
- Anything that allocates memory

9.9.2 @no_interrupts

Wraps the function body in Disable/Enable:

```
@no_interrupts
pub fn poke hardware_register(addr: *u16, value: u16) {
    // Disable() called on entry - ALL interrupts blocked
    *addr = value
    // Enable() called on exit
}

@no_interrupts
pub fn read_modify_write(reg: *u16, mask: u16, value: u16) {
    let current = *reg
    *reg = (current & ~mask) | (value & mask)
}
```

Warning: Use sparingly! This blocks *all* interrupts including:

- Audio DMA (causes audio glitches)
- Disk I/O (can corrupt data)
- Vertical blank (freezes display updates)
- Timer interrupts (breaks timing)

Only use for timing-critical hardware access that must be atomic.

Coming from C

AmigaOS: Disable() disables all hardware interrupts by setting the processor interrupt mask. Enable() restores it. These are Exec functions, not direct hardware access.

9.10 Interrupt Handling Attributes

These attributes are essential for writing interrupt-driven code.

9.10.1 @interrupt

Marks a function as an interrupt handler. The compiler generates appropriate prologue/epilogue code:

```
@interrupt
pub fn vblank_handler() {
    // Called during vertical blank interrupt
    // All registers properly saved/restored
    update_frame_counter()
}

@interrupt
pub fn audio_handler() {
    // Called when audio channel needs more data
    feed_next_sample()
}
```

The compiler generates:

- Register save (MOVEM.L to stack)
- Your function body
- Register restore (MOVEM.L from stack)
- RTE (Return from Exception)

9.10.2 @interrupt_safe

Documents that a function is safe to call from interrupt context:

```
@interrupt_safe
pub fn increment_counter(counter: *u32) {
    // Safe: no allocations, no blocking calls
    *counter = *counter + 1
}

@interrupt_safe
fn calculate_checksum(data: *u8, len: u32) -> u32 {
    // Safe: pure computation, no system calls
    var sum: u32 = 0
    for i in 0..len {
        sum = sum + (data[i] as u32)
    }
    return sum
}
```

Functions marked @interrupt_safe must:

- Never allocate memory
- Never call blocking functions (Wait, WaitIO, etc.)
- Never call Permit() without matching Forbid()
- Complete quickly (microseconds, not milliseconds)

This is currently documentation-only but helps communicate intent and catch errors during code review.

9.11 Diagnostic Attributes

9.11.1 @suppress

Suppresses specific compiler warnings:

```
@suppress("W0002", "Known safe: buffer size validated externally")
pub fn write_unchecked(buffer: *u8, data: *u8, len: i32) {
    // Warning W0002 (unchecked buffer access) suppressed
    // We accept responsibility for bounds checking
}

@suppress("W0001") // Minimum suppression - just the code
pub fn use_deprecated_api() {
    call_old_function()
}
```

Arguments:

warning code The warning ID to suppress (e.g., "W0001")

justification Explanation of why the warning is safe to ignore (recommended)

Always prefer fixing warnings over suppressing them. When suppression is necessary, document why.

9.12 Program Configuration

9.12.1 `#[stack_size]`

Sets the stack size for the compiled program:

```
#[stack_size(131072)] // 128KB stack

fn main() -> i32 {
    // Programs with deep recursion need more stack
    return 0
}
```

Stack Size Guidelines:

Use Case	Recommended Size
Simple programs	8KB – 16KB
Typical applications	32KB – 64KB (default)
Deep recursion (parsers, tree traversal)	128KB – 256KB
Large stack arrays	Array size + 32KB buffer

Table 9.1: Stack Size Recommendations

Note: AmigaOS CLI default stack is only 4KB, far too small for most Novus programs. The Novus runtime automatically checks and increases stack size if needed.

9.13 Library Attributes

These attributes are used when creating AmigaOS shared libraries.

9.13.1 `@library`

Marks a struct as an Amiga shared library:

```
@library(name = "example.library", version = 1, revision = 0)
pub struct ExampleLibrary {
    base: Library, // Must include Library base
    counter: u32,
    initialized: bool,
}
```

Arguments:

name Library name (must end with `.library`)

version Major version number

revision Minor revision number

The compiler generates:

- ROMTag structure
- Library vector table
- Open/Close/Expunge handlers
- Proper HUNK structure for library loading

9.13.2 @libfunc

Marks a method as a library vector entry point:

```
impl ExampleLibrary {
    @libfunc
    pub fn add_numbers(&self, a: i32, b: i32) -> i32 {
        return a + b
    }

    @libfunc
    pub fn get_counter(&self) -> u32 {
        return self.counter
    }
}
```

Library functions are called through the library base pointer, following AmigaOS calling conventions.

9.13.3 Library Lifecycle Hooks

```
impl ExampleLibrary {
    @libinit
    pub fn init() -> bool {
        // One-time initialization when library loads
        // Allocate resources, open dependencies
        return true
    }

    @libopen
    pub fn open(version: u32) -> bool {
        // Called each time library is opened
        if version > 1 { return false } // Version check
        self.counter = self.counter + 1
        return true
    }

    @libclose
    pub fn close() {
        // Called each time library is closed
        self.counter = self.counter - 1
    }

    @libexpunge
    pub fn expunge() -> bool {
        // Return true if library can be unloaded
    }
}
```

```

        // Return false to keep library resident
        return self.counter == 0
    }
}

```

9.14 Device Attributes

These attributes are used when creating AmigaOS devices.

9.14.1 @device

Marks a struct as an Amiga device:

```

@device(name = "example.device", version = 1, revision = 0)
pub struct ExampleDevice {
    base: Device,
    unit: u32,
    active: bool,
}

```

9.14.2 @devicecmd

Marks a method as a device command handler:

```

impl ExampleDevice {
    @devicecmd(CMD_READ)
    pub fn handle_read(&self, io: *IORequest) {
        // Handle read command
        io.io_Actual = read_data(io.io_Data, io.io_Length)
    }

    @devicecmd(CMD_WRITE)
    pub fn handle_write(&self, io: *IORequest) {
        // Handle write command
        write_data(io.io_Data, io.io_Length)
        io.io_Actual = io.io_Length
    }
}

```

9.14.3 @beginio / @abortio

Standard device I/O handlers:

```

impl ExampleDevice {
    @beginio
    pub fn begin_io(&self, io: *IORequest) {
        // Start I/O operation
        let cmd = io.io_Command
        match cmd {
            CMD_READ => self.handle_read(io),
            CMD_WRITE => self.handle_write(io),
            _ => io.io_Error = IOERR_NOCMD,
        }
    }
}

```

```
    }

    @abortio
    pub fn abort_io(&self, io: *IORequest) -> i32 {
        // Abort pending I/O
        // Return 0 if aborted, error otherwise
        return 0
    }
}
```

9.14.4 Device Lifecycle Hooks

Similar to libraries: @deviceinit, @deviceopen, @deviceclose, @deviceexpunge.

9.15 Common Attribute Patterns

9.15.1 Interrupt-Safe State Updates

Combining @atomic with @interrupt_safe:

```
static var frame_count: u32 = 0
static var frames_per_second: u32 = 0

@interrupt_safe
@inline
fn increment_frame_count() {
    frame_count = frame_count + 1
}

@atomic
pub fn get_and_reset_fps() -> u32 {
    let fps = frames_per_second
    frames_per_second = 0
    return fps
}

@interrupt
@export
pub fn vblank_interrupt() {
    increment_frame_count()
    if frame_count >= 50 { // PAL: 50 frames = 1 second
        frames_per_second = frame_count
        frame_count = 0
    }
}
```

9.15.2 Builder Pattern with Validation

```
struct ConfigBuilder {
    width: Option<i32>,
    height: Option<i32>,
    depth: i32,
}
```

```

impl ConfigBuilder {
    pub fn new() -> ConfigBuilder {
        return ConfigBuilder {
            width: Option::None,
            height: Option::None,
            depth: 8,
        }
    }

    @chain
    pub fn width(&var self, w: i32) {
        self.width = Option::Some(w)
    }

    @chain
    pub fn height(&var self, h: i32) {
        self.height = Option::Some(h)
    }

    @chain
    pub fn depth(&var self, d: i32) {
        self.depth = d
    }

    pub fn build(&self) -> Result<Config, ConfigError> {
        // Validate all required fields are set
        let w = match self.width {
            Option::Some(v) => v,
            Option::None => return
                Result::Err(ConfigError::MissingWidth),
        }
        let h = match self.height {
            Option::Some(v) => v,
            Option::None => return
                Result::Err(ConfigError::MissingHeight),
        }
        return Result::Ok(Config { width: w, height: h, depth:
            self.depth })
    }
}

```

9.16 Showcase: Interrupt-Driven Music Player

This showcase demonstrates interrupt and synchronization attributes for audio playback.

```

// music_player_showcase.novus - Demonstrates interrupt attributes

from std::core import Option
from std::io::file import print, println
from std::ffi::exec import Forbid, Permit, Disable, Enable

// Player state stored as globals for interrupt access
static var g_is_playing: bool = false
static var g_current_position: u32 = 0
static var g_sample_length: u32 = 0
static var g_tempo_counter: u32 = 0

```

```
static var g_tempo_divider: u32 = 6 // ~8 ticks/sec at PAL
static var g_volume: u16 = 64

// @interrupt_safe: Documents this is safe to call from interrupt
// @inline: Hint to inline for performance
@interrupt_safe
@inline
fn advance_position() -> bool {
    if g_current_position >= g_sample_length {
        return false
    }
    g_current_position = g_current_position + 1u32
    return true
}

// @no_interrupts: Wraps in Disable/Enable for hardware access
@no_interrupts
fn write_audio_volume(volume: u16) {
    g_volume = volume
    // In real code: write to Paula register here
}

// @atomic: Wraps in Forbid/Permit for task-safe state access
@atomic
fn start_playback(length: u32, tempo: u32) -> bool {
    if g_is_playing { return false }
    g_sample_length = length
    g_current_position = 0u32
    g_tempo_divider = tempo
    g_is_playing = true
    return true
}

@atomic
fn stop_playback() {
    if g_is_playing {
        g_is_playing = false
        write_audio_volume(0u16)
    }
}

// @interrupt: Generates proper save/restore and RTE
// @export: Prevents dead code elimination
@interrupt
@export
pub fn vblank_music_handler() {
    if !g_is_playing { return }

    g_tempo_counter = g_tempo_counter + 1u32
    if g_tempo_counter < g_tempo_divider { return }
    g_tempo_counter = 0u32

    if !advance_position() {
        g_is_playing = false
    }
}

pub fn main() -> i32 {
```



```

println("Key Attributes Demonstrated:")
println("  @interrupt      - VBlank handler")
println("  @interrupt_safe - Safe helper functions")
println("  @atomic          - Critical sections")
println("  @no_interrupts   - Hardware access")
println("  @inline          - Performance hints")
println("  @export          - Callback preservation")
return 0
}

```

This pattern is typical for:

- MOD/XM music players
- Sample playback engines
- Sound effect systems
- Any audio requiring precise timing

9.17 Attribute Reference

Attribute	Applies To	Purpose
@test	functions	Mark as test case
@bench	functions	Mark for benchmarking
@inline	functions	Hint to inline
@noinline	functions	Prevent inlining
@packed	structs	Remove field padding
@derive	structs	Auto-implement traits
@chain	methods	Enable method chaining
@export	functions	Keep and don't mangle name
@extern_type	structs	External C type
@atomic	functions	Wrap in Forbid/Permit
@no_interrupts	functions	Wrap in Disable/Enable
@interrupt	functions	Mark as interrupt handler
@interrupt_safe	functions	Safe for interrupt context
@suppress	functions	Suppress warning
#[stack_size]	program	Set stack size
@library	structs	Define Amiga library
@libfunc	methods	Library vector entry
@device	structs	Define Amiga device
@devicecmd	methods	Device command handler

Table 9.2: Complete Attribute Reference

9.18 Unknown Attributes

If you use an unknown attribute, the compiler emits a warning:

```

warning[W2001]: unknown attribute 'foo'
--> test.novus:3:1
   |
3 | @foo

```

```
| ^^^  
|  
help: This attribute is not recognized and will be ignored  
help: Known attributes: library, libfunc, inline, test, ...
```

This helps catch typos while allowing forward compatibility with future attributes.

9.19 Future Attributes

The following attributes are planned for future versions:

@deprecated Mark functions as deprecated with migration notes

@must_use Require return value to be used

@cold Mark function as unlikely to be called (affects code layout)

@hot Mark function as frequently called (affects optimization)

@pure Function has no side effects (enables memoization)

Chapter 10

Error Handling

“Errors should never pass silently. Unless explicitly silenced.” — The Zen of Python

Novus treats errors as values, not exceptions. Every operation that can fail returns a `Result` type, making error handling explicit and impossible to ignore. This chapter covers Novus’s error handling philosophy, the standard error types, and patterns for writing robust code.

10.1 Philosophy: Errors as Values

Unlike languages with exceptions, Novus requires explicit handling of errors:

- **No hidden control flow** — You always know when a function can fail
- **No uncaught exceptions** — Errors can’t silently propagate up the stack
- **Compile-time enforcement** — Ignoring a `Result` produces a warning
- **Zero runtime cost** — No exception tables or unwinding machinery

10.1.1 Why Not Exceptions?

Many languages use exceptions for error handling. While convenient in some ways, exceptions have significant drawbacks for systems programming:

1. Stack Unwinding Cost: Exception handling requires runtime machinery to unwind the stack, search for handlers, and call destructors. On a 68020 running at 14 MHz, this overhead is unacceptable.

2. Invisible Control Flow: Any function call might throw, making control flow analysis difficult. In Novus, functions that can fail are marked by their return type.

3. Unpredictable Cleanup: With exceptions, you need `try/finally` blocks to ensure resources are released. In Novus, `Result` combinators and pattern matching make cleanup explicit.

4. Binary Size: Exception tables and unwinding code add significant size to executables. Every byte counts on classic Amiga systems.

10.1.2 Why Not `errno`?

The traditional C approach of global error codes has its own problems:

```
// C with errno
FILE *f = fopen("data.txt", "r");
if (f == NULL) {
    if (errno == ENOENT) { /* not found */ }
    else if (errno == EACCES) { /* permission denied */ }
```

```
// Easy to forget, hard to get right
}
```

Problems with `errno`:

- Global mutable state is inherently error-prone
- Easy to forget to check it
- Can be overwritten by intermediate calls
- Not type-safe — any integer could be assigned
- Thread-unsafe without special handling

AmigaOS uses a similar pattern with `IoErr()`:

```
// AmigaOS with IoErr()
BPTR lock = Lock("myfile", ACCESS_READ);
if (lock == 0) {
    LONG err = IoErr(); // Must call immediately!
    // But what if you call something else first?
    PrintFault(err, "Error");
}
```

The danger: if you call *any* other DOS function before `IoErr()`, the error code may be overwritten.

Coming from C

In C, errors are typically indicated by return values (`-1`, `NULL`) or global `errno`/`IoErr()`. This leads to bugs when callers forget to check.

In Novus, `Result<T, E>` forces you to handle both success and failure cases. You cannot accidentally use an error value as if it were successful.

C (AmigaOS)	Novus
if (ptr == NULL)	match result { ... }
if (fd < 0)	let file = open(path)?
IoErr()	Result::Err(DosError::NotFound)
Manual conversion	From trait auto-conversion

10.2 The Result Type

The core error handling type is `Result<T, E>`:

```
enum Result<T, E> {
    Ok(T), // Success with value
    Err(E), // Failure with error
}
```

Every function that can fail returns a `Result`:

```
fn parse_number(s: *u8) -> Result<i32, ParseError> {
    // Returns Ok(value) or Err(error)
}
```

```
fn open_file(path: *u8) -> Result<FileHandle, DosError> {
    // Returns Ok(handle) or Err(error)
}
```

10.2.1 Result Methods Reference

The `Result` type provides several methods for working with values:

Method	Description
<code>is_ok()</code>	Returns <code>true</code> if the result is <code>Ok</code>
<code>is_err()</code>	Returns <code>true</code> if the result is <code>Err</code>
<code>unwrap()</code>	Extracts the <code>Ok</code> value, panics on <code>Err</code>
<code>unwrap_or(default)</code>	Returns the <code>Ok</code> value or the provided default
<code>expect(msg)</code>	Like <code>unwrap()</code> but with a custom panic message
<code>ok()</code>	Converts to <code>Option<T></code> , discarding the error
<code>err()</code>	Converts to <code>Option<E></code> , discarding the success value

```
let result: Result<i32, Error> = some_operation()

// Query the state
if result.is_ok() {
    println("Operation succeeded")
}

// Get value with default
let value = result.unwrap_or(0)

// Convert to Option
let maybe_value: Option<i32> = result.ok()
```

10.3 The Option Type

While `Result` represents success or failure, `Option` represents presence or absence:

```
enum Option<T> {
    Some(T), // Value present
    None,    // No value
}
```

10.3.1 When to Use Option vs Result

- Use **Option** when the only failure mode is “not found” or “not available”
- Use **Result** when there are multiple failure modes to distinguish
- Use **Result** for operations that interact with the OS or hardware

```
// Option: "not found" is the only failure
fn find_user(id: u32) -> Option<User> { ... }
```

```
// Result: multiple failure modes possible
fn load_user(id: u32) -> Result<User, DbError> { ... }
```

10.3.2 Option Methods Reference

Method	Description
<code>is_some()</code>	Returns <code>true</code> if value is present
<code>is_none()</code>	Returns <code>true</code> if no value
<code>unwrap()</code>	Extracts the value, panics on <code>None</code>
<code>unwrap_or(default)</code>	Returns value or provided default
<code>expect(msg)</code>	Like <code>unwrap()</code> with custom panic message

10.3.3 Converting Between Option and Result

Sometimes you need to convert between the two types:

```
// Option -> Result: provide an error for None
fn get_config_value(key: *u8) -> Result<*u8, ConfigError> {
    match config.get(key) {
        Option::Some(val) => Result::Ok(val),
        Option::None => Result::Err(ConfigError::MissingKey)
    }
}
```

Coming from C

In C, null pointers serve double duty: representing “no value” (`Option`) and “error occurred” (`Result`). This conflation causes bugs.

```
// C: Is NULL "not found" or "error"?
Node *node = find_node(list, key);
if (node == NULL) {
    // Did find fail, or is the key just missing?
    // Have to check errno to know...
}
```

In Novus, the type system distinguishes these cases. `Option<T>` means “might not exist” while `Result<T, E>` means “might have failed”.

10.4 The ? Operator In Depth

The `?` operator is syntactic sugar for early return on error. It transforms verbose error checking into concise, readable code.

10.4.1 What ? Actually Does

When you write:

```
let file = open_file(path)?
```

The compiler generates equivalent code to:

```
let file = match open_file(path) {
  Result::Ok(val) => val,
  Result::Err(e) => return Result::Err(e)
}
```

The ? operator:

1. Checks if the result is `Err`
2. If `Err`, returns immediately with that error
3. If `Ok`, unwraps the value and continues

10.4.2 Chaining Multiple ? Calls

The real power of ? shows when chaining fallible operations:

```
fn process_file(path: *u8) -> Result<Data, DosError> {
  let file = open_file(path)?           // Returns Err early if
    failed
  let header = read_header(&file)?       // Returns Err early if
    failed
  let body = read_body(&file)?           // Returns Err early if
    failed
  let data = parse_data(&body)?          // Returns Err early if
    failed
  Result::Ok(data)
}
```

Without ?, this would require deeply nested match statements:

```
fn process_file(path: *u8) -> Result<Data, DosError> {
  match open_file(path) {
    Result::Ok(file) => {
      match read_header(&file) {
        Result::Ok(header) => {
          match read_body(&file) {
            Result::Ok(body) => {
              match parse_data(&body) {
                Result::Ok(data) =>
                  Result::Ok(data),
                Result::Err(e) => Result::Err(e)
              }
            },
            Result::Err(e) => Result::Err(e)
          }
        },
        Result::Err(e) => Result::Err(e)
      }
    },
    Result::Err(e) => Result::Err(e)
  }
}
```

10.4.3 Using ? in main()

In Novus, `main()` returns `i32` (the exit code). You can still use ? by wrapping your logic:

```

fn run() -> Result<(), AppError> {
    let config = load_config()?
    let window = open_window()?
    main_loop(&window)?
    Result::Ok(())
}

pub fn main() -> i32 {
    match run() {
        Result::Ok(_) => return 0,
        Result::Err(err) => {
            println(err.message())
            return 1
        }
    }
}

```

10.4.4 Limitations of ?

The ? operator can only be used in functions that return `Result`:

```

fn valid() -> Result<(), Error> {
    let x = fallible_op()? // OK: function returns Result
    Result::Ok(())
}

fn invalid() {
    let x = fallible_op()? // ERROR: function doesn't return
                           Result
}

```

10.5 Standard Error Types

Novus provides error types for each AmigaOS subsystem.

10.5.1 DosError

Errors from the DOS library (file operations, paths):

```

pub enum DosError {
    NotFound,           // File/directory not found (IoErr 205)
    AlreadyExists,      // Object already exists (IoErr 203)
    DirNotFound,        // Directory not found (IoErr 212)
    DirectoryNotEmpty,  // Can't delete non-empty dir (IoErr 216)
    PermissionDenied,   // Access denied (IoErr 214)
    WriteProtected,     // Disk is write-protected (IoErr 223)
    ReadProtected,      // File is read-protected (IoErr 224)
    DiskFull,           // No space on disk (IoErr 221)
    NoDisk,             // No disk in drive (IoErr 218)
    DeviceNotMounted,   // Device not available (IoErr 225)
    InvalidLock,        // Bad file lock (IoErr 211)
    Unknown,            // Unmapped error code
}

```


DosError Variant	IoErr() Code	Description
NotFound	205	Object not found
AlreadyExists	203	Object already exists
DirNotFound	212	Directory not found
DirectoryNotEmpty	216	Directory not empty
PermissionDenied	214	Object is protected
WriteProtected	223	Disk is write-protected
ReadProtected	224	File is read-protected
DiskFull	221	Disk is full
NoDisk	218	No disk in drive
DeviceNotMounted	225	Device not mounted
InvalidLock	211	Invalid lock

10.5.2 ExecError

Errors from the Exec library (memory, tasks, signals):

```
pub enum ExecError {
    NoMem,           // Memory allocation failed
    BadSignal,       // Invalid signal number
    NullReturn,      // Unexpected null return
    BadTask,         // Invalid task pointer
    BadPort,         // Invalid message port
    SignalInUse,     // Signal already allocated
    TaskNotFound,    // FindTask failed
    LibraryNotFound, // OpenLibrary failed
    IndexOutOfBounds, // Array bounds violation
    CapacityExceeded, // Fixed buffer full
}
```

10.5.3 IntuitionError

Errors from the Intuition library (windows, screens, gadgets):

```
pub enum IntuitionError {
    WindowOpenFailed, // OpenWindow failed
    ScreenOpenFailed, // OpenScreen failed
    GadgetCreateFailed, // Gadget creation failed
    MenuCreateFailed, // Menu creation failed
    InvalidWindow, // Bad window pointer
    InvalidScreen, // Bad screen pointer
    NoIDCMP, // IDCMP setup failed
    BadDisplayMode, // Invalid display mode
}
```

10.5.4 GraphicsError

Errors from the Graphics library (bitmaps, sprites, fonts):

```
pub enum GraphicsError {
    BitMapAllocFailed, // BitMap allocation failed
    FontOpenFailed, // OpenFont failed
    InvalidBitMap, // Bad BitMap pointer
    InvalidRastPort, // Bad RastPort pointer
}
```

```
InvalidSpriteChannel, // Bad sprite channel (0-7)
SpriteChannelInUse,   // Channel already allocated
NoSpritesAvailable,   // All sprites in use
}
```

10.6 The NovusError Wrapper

When a function may fail with errors from multiple subsystems, use `NovusError`:

```
pub enum NovusError {
    Dos(DosError),
    Exec(ExecError),
    Intuition(IntuitionError),
    Graphics(GraphicsError),
}
```

The `From` trait implementations enable automatic conversion with the `?` operator:

```
fn load_and_display(path: *u8) -> Result<(), NovusError> {
    let data = read_file(path)?           // DosError -> NovusError
    let screen = open_screen()?           // IntuitionError ->
        NovusError
    let bitmap = load_bitmap(&data)?      // GraphicsError ->
        NovusError
    display(screen, bitmap)
    Result::Ok(())
}
```

10.7 Error Conversion with From

The `From` trait enables automatic error type conversion when using `?`:

```
trait From<T> {
    fn convert(value: T) -> Self
}
```

When you use `?` in a function returning `Result<T, NovusError>`, and the inner call returns `Result<T, DosError>`, the compiler automatically calls `NovusError::From<DosError>::convert()` to wrap the error.

10.7.1 Manual Conversion

You can also convert errors manually:

```
from std::error::errors import novus_error_from_dos

fn wrapper() -> Result<(), NovusError> {
    let file = match open("file.txt", MODE_OLDFILE) {
        Result::Ok(f) => f,
        Result::Err(e) => return
            Result::Err(novus_error_from_dos(e))
    }
    Result::Ok(())
}
```

```
}
```

10.7.2 Getting Raw Error Codes

For debugging or interop with C code, get raw AmigaOS error codes:

```
from std::error::errors import dos_error_to_code

let code = dos_error_to_code(DosError::NotFound) // Returns 205
```

10.8 Error Messages

All standard error types implement a `message()` method:

```
match result {
  Result::Err(err) => {
    print("Error: ")
    println(err.message())
  },
  _ => {}
}
```

Error messages are human-readable strings:

- `DosError::NotFound` → "Object not found"
- `ExecError::NoMem` → "Memory allocation failed"
- `IntuitionError::WindowOpenFailed` → "Failed to open window"

10.9 Error Handling Patterns

This section covers common patterns for handling errors effectively.

10.9.1 Early Return Pattern

Validate inputs at the start of a function and return errors immediately. This keeps the main logic clean:

```
fn process(data: *u8, len: u32) -> Result<(), Error> {
  // Validate all inputs first
  if data == null {
    return Result::Err(Error::InvalidInput)
  }
  if len == 0 {
    return Result::Err(Error::EmptyData)
  }
  if len > MAX_SIZE {
    return Result::Err(Error::TooLarge)
  }

  // All validation passed - main logic here
  do_actual_work(data, len)
}
```

Coming from C

In C, you often see scattered validation throughout functions:

```
int process(char *data, size_t len) {
    int result = init_something();
    if (result < 0) return result;

    if (data == NULL) return -EINVAL; // Mixed with logic!

    result = do_more_work();
    if (result < 0) goto cleanup;

    if (len == 0) goto cleanup; // Validation too late!
    // ...
}
```

The early return pattern in Novus collects all validation at the top, making the “happy path” obvious.

10.9.2 Fallback Chains

Try multiple approaches, falling back on failure. This is common for finding configuration files or resources:

```
fn find_config() -> Result<Config, ConfigError> {
    // Try user config first
    match load_from("ENV:app.cfg") {
        Result::Ok(cfg) => return Result::Ok(cfg),
        Result::Err(_) => {}
    }

    // Fall back to archive
    match load_from("ENVARC:app.cfg") {
        Result::Ok(cfg) => return Result::Ok(cfg),
        Result::Err(_) => {}
    }

    // Fall back to system config
    match load_from("S:app.cfg") {
        Result::Ok(cfg) => return Result::Ok(cfg),
        Result::Err(_) => {}
    }

    // Use defaults as last resort
    Result::Ok(default_config())
}
```

10.9.3 Cleanup on Error with defer

Use `defer` to ensure cleanup even when returning early on error:

```
fn process_file(path: *u8) -> Result<Data, DosError> {
    let file = open(path, MODE_OLDFILE)?
    defer { close(file) } // Always runs, even on error return
}
```

```

let header = read_header(&file)? // May return early
let body = read_body(&file)?     // May return early

Result::Ok(combine(header, body))
// defer runs here on success too
}

```

10.9.4 Error Context

Add context when propagating errors by matching and re-wrapping with domain-specific information:

```

fn load_user_config() -> Result<Config, ConfigError> {
    let path = get_config_path()?

    let file = match open(path, MODE_OLDFILE) {
        Result::Ok(f) => f,
        Result::Err(_) => return
            Result::Err(ConfigError::FileNotFound)
    }

    let data = match read_all(&file) {
        Result::Ok(d) => d,
        Result::Err(_) => return
            Result::Err(ConfigError::ReadFailed)
    }

    parse_config(&data)
}

```

10.9.5 Retry Pattern

For transient failures (disk not ready, network timeout), retry a few times:

```

fn open_with_retry(path: *u8, max_retries: u32) ->
    Result<FileHandle, DosError> {
    var attempts: u32 = 0

    while attempts < max_retries {
        match open(path, MODE_OLDFILE) {
            Result::Ok(file) => return Result::Ok(file),
            Result::Err(DosError::NoDisk) => {
                // Disk not ready - wait and retry
                delay(Duration::from_secs(1))
                attempts = attempts + 1u32
            },
            Result::Err(e) => return Result::Err(e) //
                Non-retriable
        }
    }

    Result::Err(DosError::NoDisk)
}

```

10.10 Panic: When Errors Are Bugs

A **panic** is for situations where continuing is impossible—invariant violations, impossible states, or programmer errors.

```
// Panic on invariant violation
fn get_first(vec: &Vec<i32>) -> i32 {
    if vec.len() == 0 {
        panic!("get_first called on empty vector")
    }
    return vec[0].unwrap()
}
```

10.10.1 Panic vs Error

Use	When	Example
Result::Err	Failure is expected	File not found
panic!()	Failure is a bug	Index out of bounds

10.10.2 assert!() for Debug Checks

Use `assert!()` to catch programming errors during development:

```
fn process_buffer(buf: *u8, len: u32) {
    assert!(buf != null, "buffer must not be null")
    assert!(len > 0, "length must be positive")
    // ...
}
```

Assertions are checked in debug builds. They verify invariants that should always hold if the code is correct.

10.11 Creating Custom Errors

Define domain-specific error types for your application:

```
pub enum ConfigError {
    /// Config file not found in any search location
    FileNotFound,

    /// File exists but cannot be read
    ReadFailed(DosError),

    /// File is too large to load
    FileTooLarge(u32),

    /// Syntax error at specific line
    ParseError(u32),

    /// Required key is missing
    MissingKey,

    /// Value is invalid for expected type
    InvalidValue,
```

```

    /// Memory allocation failed
    OutOfMemory,
}

```

10.11.1 Implementing message()

Provide human-readable messages:

```

impl ConfigError {
    pub fn message(&self) -> *u8 {
        match *self {
            ConfigError::FileNotFound =>
                return "Config file not found in any search
                    location",
            ConfigError::ReadFailed(_) =>
                return "Failed to read config file",
            ConfigError::FileTooLarge(_) =>
                return "Config file exceeds maximum size",
            ConfigError::ParseError(_) =>
                return "Syntax error in config file",
            ConfigError::MissingKey =>
                return "Required configuration key is missing",
            ConfigError::InvalidValue =>
                return "Configuration value is invalid",
            ConfigError::OutOfMemory =>
                return "Out of memory loading config",
        }
    }
}

```

10.11.2 Extracting Error Details

For errors with associated data, provide accessor methods:

```

impl ConfigError {
    /// Get the line number for parse errors
    pub fn line_number(&self) -> Option<u32> {
        match *self {
            ConfigError::ParseError(line) => Option::Some(line),
            _ => Option::None
        }
    }

    /// Get the file size for too-large errors
    pub fn file_size(&self) -> Option<u32> {
        match *self {
            ConfigError::FileTooLarge(size) => Option::Some(size),
            _ => Option::None
        }
    }
}

```

10.12 Showcase: Config Loader

This complete program demonstrates comprehensive error handling in Novus. It loads configuration files with fallback locations, custom error types, and proper error reporting.

```
// config_loader.novus - Robust Config Loader

from std::core import Option, Result
from std::error::errors import DosError
from std::ffi::dos import Open, Close, Read, Lock, UnLock, Examine
from std::ffi::amiga_consts import MODE_OLDFILE, ACCESS_READ
from std::io::file import print, println

// -----
// Custom Error Type with Rich Information
// -----

pub enum ConfigError {
    FileNotFound,
    ReadFailed(DosError),
    FileTooLarge(u32),
    ParseError(u32), // Line number
    MissingKey,
    InvalidValue,
    OutOfMemory,
}

impl ConfigError {
    pub fn message(&self) -> *u8 {
        match *self {
            ConfigError::FileNotFound =>
                return "Config file not found in any search
                    location",
            ConfigError::ReadFailed(_) =>
                return "Failed to read config file",
            ConfigError::FileTooLarge(_) =>
                return "Config file exceeds maximum size (64KB)",
            ConfigError::ParseError(_) =>
                return "Syntax error in config file",
            ConfigError::MissingKey =>
                return "Required configuration key is missing",
            ConfigError::InvalidValue =>
                return "Configuration value is invalid",
            ConfigError::OutOfMemory =>
                return "Out of memory loading config",
        }
    }

    pub fn line_number(&self) -> Option<u32> {
        match *self {
            ConfigError::ParseError(line) => Option::Some(line),
            _ => Option::None
        }
    }
}
```



```
// -----
// Fallback Chain: Try Multiple Locations
// -----

pub fn load_config(name: *u8) -> Result<Config, ConfigError> {
    var path_buffer = [0u8; 256]

    // Try ENV: first (volatile)
    copy_path(&path_buffer[0], "ENV:", name)
    match try_load_from(&path_buffer[0]) {
        Result::Ok(cfg) => return Result::Ok(cfg),
        Result::Err(_) => {}
    }

    // Try ENVARC: (persistent)
    copy_path(&path_buffer[0], "ENVARC:", name)
    match try_load_from(&path_buffer[0]) {
        Result::Ok(cfg) => return Result::Ok(cfg),
        Result::Err(_) => {}
    }

    // Try S: (startup scripts)
    copy_path(&path_buffer[0], "S:", name)
    match try_load_from(&path_buffer[0]) {
        Result::Ok(cfg) => return Result::Ok(cfg),
        Result::Err(_) => {}
    }

    // Try current directory
    match try_load_from(name) {
        Result::Ok(cfg) => return Result::Ok(cfg),
        Result::Err(e) => return Result::Err(e)
    }
}

// -----
// Main: Demonstrating Error Handling
// -----

pub fn main() -> i32 {
    println("Config Loader - Error Handling Showcase")
    println("")

    match load_config("test.config") {
        Result::Ok(config) => {
            println("Config loaded successfully!")
            print("  Source: ")
            println(&config.source_path[0])
            return 0
        },
        Result::Err(err) => {
            print("Error: ")
            println(err.message())
            return 1
        }
    }
}
```

```
    }  
  }  
}
```

This program demonstrates:

- **Custom error types** with variant-specific data (line numbers, file sizes)
- **Fallback chains** trying multiple locations (ENV:, ENVARC:, S:, current dir)
- **Error messages** via the `message()` method
- **Error context** extraction via `line_number()`

The full source code is available in `Novus.Tests/Examples/config_loader_showcase.novus`.

10.13 Summary

- Use `Result<T, E>` for all fallible operations
- Use `Option<T>` when “not found” is the only failure mode
- **Pattern match** for explicit handling of all cases
- Use `?` to propagate errors concisely
- Use `unwrap_or` for safe defaults
- **Avoid `unwrap()`** except when failure is impossible
- Use **standard errors** (`DosError`, `ExecError`, etc.) for AmigaOS operations
- Use `NovusError` when crossing subsystem boundaries
- Use `defer` to ensure cleanup on error paths
- Use `panic` only for invariant violations and bugs
- **Create custom errors** for domain-specific failure modes

Error handling in Novus is explicit but ergonomic. The `?` operator and `Result` combinators make it easy to write code that correctly handles failures without obscuring the happy path.

10.14 Advanced Error Patterns

This section covers advanced patterns for complex error handling scenarios.

10.14.1 Collecting Multiple Errors

Sometimes you want to validate everything before reporting, collecting all errors rather than stopping at the first:

```
pub struct ValidationErrors {  
    errors: Vec<*u8>,  
}  
  
impl ValidationErrors {  
    pub fn new() -> ValidationErrors {
```

```

        return ValidationErrors { errors: Vec::new() }
    }

    pub fn add(&var self, message: *u8) {
        self.errors.push(message)
    }

    pub fn has_errors(&self) -> bool {
        return self.errors.len() > 0
    }

    pub fn to_result(&self) -> Result<(), ValidationError> {
        if self.has_errors() {
            return
                Result::Err(ValidationError::Multiple(self.errors.len()))
        }
        Result::Ok(())
    }
}

fn validate_user(user: &User) -> Result<(), ValidationError> {
    var errors = ValidationErrors::new()

    if user.name.len() == 0 {
        errors.add("Name is required")
    }
    if user.age < 0 {
        errors.add("Age must be non-negative")
    }
    if user.email.len() == 0 {
        errors.add("Email is required")
    }
    if !user.email.contains('@') {
        errors.add("Email must contain @")
    }

    // Print all errors
    for i in 0..errors.errors.len() {
        print("  - ")
        println(errors.errors[i])
    }

    errors.to_result()
}

```

10.14.2 Error Recovery Strategies

Different errors require different recovery strategies. Here's a pattern that encapsulates recovery logic:

```

pub enum RecoveryAction {
    Retry,           // Try the operation again
    Skip,            // Skip this item, continue with others
    Fallback,        // Use a fallback value or method
    Abort,           // Stop completely
}

```

```

fn determine_recovery(err: &DosError) -> RecoveryAction {
    match *err {
        // Transient errors - retry makes sense
        DosError::NoDisk => RecoveryAction::Retry,

        // Missing files - skip or use fallback
        DosError::NotFound => RecoveryAction::Fallback,

        // Permission errors - cannot recover
        DosError::PermissionDenied => RecoveryAction::Abort,

        // Other errors - skip by default
        _ => RecoveryAction::Skip
    }
}

fn process_files(paths: &[*u8]) -> Result<u32, DosError> {
    var processed: u32 = 0

    for i in 0..paths.len() {
        let path = paths[i]

        loop {
            match open(path, MODE_OLDFILE) {
                Result::Ok(file) => {
                    defer { close(file) }
                    process_file(&file)?
                    processed += 1
                    break
                },
                Result::Err(err) => {
                    match determine_recovery(&err) {
                        RecoveryAction::Retry => {
                            delay(Duration::from_millis(500))
                            continue
                        },
                        RecoveryAction::Skip => break,
                        RecoveryAction::Fallback => {
                            use_fallback(path)
                            processed += 1
                            break
                        },
                        RecoveryAction::Abort => return
                            Result::Err(err)
                    }
                }
            }
        }
    }

    Result::Ok(processed)
}

```

10.14.3 Error Logging Pattern

For production code, logging errors before propagating them helps debugging:

```

fn log_error(location: *u8, err: &NovusError) {
    print("[ERROR] ")
    print(location)
    print(": ")
    println(err.message())
}

fn load_with_logging(path: *u8) -> Result<Data, NovusError> {
    let file = match open(path, MODE_OLDFILE) {
        Result::Ok(f) => f,
        Result::Err(e) => {
            let err = novus_error_from_dos(e)
            log_error("load_with_logging:open", &err)
            return Result::Err(err)
        }
    }
    defer { close(file) }

    let data = match read_all(&file) {
        Result::Ok(d) => d,
        Result::Err(e) => {
            let err = novus_error_from_dos(e)
            log_error("load_with_logging:read", &err)
            return Result::Err(err)
        }
    }

    Result::Ok(data)
}

```

Coming from C

C programs often use preprocessor macros for error logging:

```

#define CHECK_ERROR(expr, msg) \
    do { \
        if (!(expr)) { \
            fprintf(stderr, "%s:%d: %s\n", __FILE__, __LINE__, \
                msg); \
            return -1; \
        } \
    } while(0)

```

In Novus, logging is explicit and type-safe. You can create helper functions that log and return errors in a single call, but the pattern remains clear and debuggable.

10.14.4 Error Wrapping for Context

Add high-level context to low-level errors:

```

pub enum AppError {
    ConfigLoad(ConfigError),
    DatabaseConnect(DbError),
    NetworkTimeout(NetworkError),
    Initialization(*u8), // Free-form context message
}

```

```
impl AppError {
    pub fn message(&self) -> *u8 {
        match *self {
            AppError::ConfigLoad(_) => "Failed to load
                configuration",
            AppError::DatabaseConnect(_) => "Failed to connect to
                database",
            AppError::NetworkTimeout(_) => "Network operation
                timed out",
            AppError::Initialization(msg) => msg,
        }
    }

    pub fn inner(&self) -> Option<*u8> {
        match *self {
            AppError::ConfigLoad(ref e) =>
                Option::Some(e.message()),
            AppError::DatabaseConnect(ref e) =>
                Option::Some(e.message()),
            AppError::NetworkTimeout(ref e) =>
                Option::Some(e.message()),
            AppError::Initialization(_) => Option::None,
        }
    }
}

fn init_app() -> Result<App, AppError> {
    // Wrap errors with context
    let config = load_config("app.cfg")
        .map_err(|e| AppError::ConfigLoad(e))?

    let db = connect_database(&config.db_path)
        .map_err(|e| AppError::DatabaseConnect(e))?

    Result::Ok(App { config, db })
}

// Display both layers
fn report_error(err: &AppError) {
    print("Error: ")
    println(err.message())

    match err.inner() {
        Option::Some(inner) => {
            print("  Cause: ")
            println(inner)
        },
        Option::None => {}
    }
}
```

10.15 Testing Error Conditions

Testing that your code handles errors correctly is just as important as testing the happy path.

10.15.1 Triggering Specific Errors

Create test helpers that return specific errors:

```
#[test]
fn test_handles_not_found() {
    // Use a path that definitely doesn't exist
    let result = load_config("NONEXISTENT:fake.cfg")

    match result {
        Result::Err(ConfigError::FileNotFound) => {
            // Expected - test passes
        },
        Result::Ok(_) => {
            panic!("Expected FileNotFound, got Ok")
        },
        Result::Err(e) => {
            panic!("Expected FileNotFound, got different error")
        }
    }
}

#[test]
fn test_handles_parse_error() {
    // Create a file with invalid content
    let file = create_test_file("test_bad.cfg", "invalid {
content")
    defer { delete_test_file("test_bad.cfg") }

    let result = parse_config(&file)

    match result {
        Result::Err(ConfigError::ParseError(line)) => {
            assert!(line == 1, "Error should be on line 1")
        },
        _ => panic!("Expected ParseError")
    }
}
```

10.15.2 Error Injection with Function Parameters

Design functions to accept fallible operations as parameters:

```
// Production signature
type FileOpener = fn(*u8) -> Result<FileHandle, DosError>

// Use dependency injection for testing
fn process_with_opener(
    path: *u8,
    opener: FileOpener
) -> Result<Data, DosError> {
    let file = opener(path)?
    defer { close(file) }
    read_and_parse(&file)
}

// In tests, provide a mock opener
```

```
#[test]
fn test_handles_open_failure() {
    fn failing_opener(path: *u8) -> Result<FileHandle, DosError> {
        Result::Err(DosError::PermissionDenied)
    }

    let result = process_with_opener("test.txt", failing_opener)

    match result {
        Result::Err(DosError::PermissionDenied) => { /* pass */ },
        _ => panic!("Expected PermissionDenied")
    }
}
```

10.15.3 Testing Error Messages

Verify that error messages are correct:

```
#[test]
fn test_error_messages() {
    assert_eq!(
        DosError::NotFound.message(),
        "Object not found"
    )
    assert_eq!(
        DosError::DiskFull.message(),
        "Disk is full"
    )
    assert_eq!(
        ConfigError::MissingKey.message(),
        "Required configuration key is missing"
    )
}
```

10.16 Performance Considerations

Error handling in Novus is designed to be zero-cost when errors don't occur.

10.16.1 Result is Zero-Cost

The `Result<T, E>` type is a tagged union. When `T` and `E` have the same size, `Result` is just `sizeof(T) + 1` byte for the discriminant (often padded to alignment). There's no heap allocation, no vtable, no exception tables.

```
// Result<i32, i32> is exactly 8 bytes:
// - 4 bytes for the value (i32)
// - 4 bytes for alignment (includes discriminant)

// Result<*u8, DosError> is 8 bytes:
// - 4 bytes for the pointer or error code
// - 4 bytes for alignment

// Compare to C++ exceptions:
// - Exception tables in every function
// - RTTI for type matching
```



```
// - Stack unwinding machinery
```

10.16.2 The Cost of the ? Operator

The ? operator compiles to a simple conditional branch:

```
// This Novus code:
let value = operation()?

// Compiles to roughly this 68k:
//      move.l    d0,d2          ; save result
//      tst.l     d0             ; check discriminant
//      bne.s     .error         ; branch if error
//      move.l    4(sp),d0       ; extract value
//      bra.s     .continue
// .error:
//      move.l    d2,d0          ; return error
//      rts
// .continue:
```

This is essentially the same cost as checking a return value in C.

10.16.3 Avoid Unwrap in Hot Paths

While `unwrap()` is convenient, it has overhead from the panic path:

```
// Slower - has panic machinery
fn sum_values(results: &[Result<i32, Error>]) -> i32 {
    var total: i32 = 0
    for i in 0..results.len() {
        total += results[i].unwrap() // Panic check each
                                   iteration
    }
    return total
}

// Faster - single error check
fn sum_values_checked(results: &[Result<i32, Error>]) ->
Result<i32, Error> {
    var total: i32 = 0
    for i in 0..results.len() {
        total += results[i]? // Early return, no panic code
    }
    Result::Ok(total)
}
```

10.16.4 When to Use `unwrap_or`

For performance-critical code with known-safe defaults, `unwrap_or` avoids branch misprediction:

```
// For optional lookups with a sensible default
fn get_color(palette: &Palette, index: u32) -> u32 {
    // Returns black (0) if index is out of range
    return palette.get(index).unwrap_or(0)
}
```

```
}  
  
// This compiles to efficient branchless code on 68020+
```

10.17 Real-World Error Handling

This section shows complete error handling for common Amiga scenarios.

10.17.1 File Processing Pipeline

A complete file processing example with proper cleanup:

```
pub fn convert_file(  
    input_path: *u8,  
    output_path: *u8  
) -> Result<u32, ConvertError> {  
    // Open input file  
    let input = match open(input_path, MODE_OLDFILE) {  
        Result::Ok(f) => f,  
        Result::Err(e) => return  
            Result::Err(ConvertError::InputOpen(e))  
    }  
    defer { close(input) }  
  
    // Get file size for progress reporting  
    let size = match examine(&input) {  
        Result::Ok(info) => info.size,  
        Result::Err(e) => return  
            Result::Err(ConvertError::InputRead(e))  
    }  
  
    // Open output file  
    let output = match open(output_path, MODE_NEWFILE) {  
        Result::Ok(f) => f,  
        Result::Err(e) => return  
            Result::Err(ConvertError::OutputOpen(e))  
    }  
    defer {  
        close(output)  
        // Delete partial output on error (defers run in reverse  
        order)  
        if error_occurred {  
            delete_file(output_path)  
        }  
    }  
  
    var error_occurred = false  
    var bytes_written: u32 = 0  
  
    // Process in chunks  
    var buffer = [0u8; 4096]  
    loop {  
        let bytes_read = match read(&input, &buffer) {  
            Result::Ok(0) => break,    // EOF  
            Result::Ok(n) => n,  
            Result::Err(e) => {
```

```

        error_occurred = true
        return Result::Err(ConvertError::InputRead(e))
    }
}

let converted = convert_chunk(&buffer, bytes_read)?

match write(&output, &converted, converted.len()) {
    Result::Ok(n) => bytes_written += n,
    Result::Err(e) => {
        error_occurred = true
        return Result::Err(ConvertError::OutputWrite(e))
    }
}

Result::Ok(bytes_written)
}

```

10.17.2 GUI Error Feedback

Reporting errors through Intuition requesters:

```

pub fn show_error(window: *Window, err: &NovusError) {
    var body = EasyStruct {
        struct_size: size_of::<EasyStruct>(),
        flags: 0,
        title: "Error",
        text_format: err.message(),
        gadget_format: "OK",
    }

    EasyRequest(window, &body, null, null)
}

pub fn show_error_with_retry(
    window: *Window,
    err: &NovusError
) -> bool {
    var body = EasyStruct {
        struct_size: size_of::<EasyStruct>(),
        flags: 0,
        title: "Error",
        text_format: err.message(),
        gadget_format: "Retry|Cancel",
    }

    // Returns 1 for Retry, 0 for Cancel
    return EasyRequest(window, &body, null, null) == 1
}

// Usage with retry loop
fn load_with_ui_retry(
    window: *Window,
    path: *u8
) -> Result<Data, NovusError> {
    loop {

```

```

        match load_data(path) {
            Result::Ok(data) => return Result::Ok(data),
            Result::Err(err) => {
                if !show_error_with_retry(window, &err) {
                    return Result::Err(err)
                }
                // User chose Retry - loop continues
            }
        }
    }
}

```

10.17.3 Multi-Step Initialization

Initializing multiple subsystems with proper cleanup on partial failure:

```

pub struct Application {
    screen: *Screen,
    window: *Window,
    audio: AudioContext,
    font: *TextFont,
}

impl Application {
    pub fn init() -> Result<Application, AppError> {
        // Step 1: Open screen
        let screen = match open_screen(&SCREEN_SPEC) {
            Result::Ok(s) => s,
            Result::Err(e) => return
                Result::Err(AppError::Screen(e))
        }

        // Step 2: Open window (cleanup screen on failure)
        let window = match open_window(screen, &WINDOW_SPEC) {
            Result::Ok(w) => w,
            Result::Err(e) => {
                close_screen(screen)
                return Result::Err(AppError::Window(e))
            }
        }

        // Step 3: Initialize audio (cleanup window and screen on failure)
        let audio = match AudioContext::init() {
            Result::Ok(a) => a,
            Result::Err(e) => {
                close_window(window)
                close_screen(screen)
                return Result::Err(AppError::Audio(e))
            }
        }

        // Step 4: Open font (cleanup all previous on failure)
        let font = match open_font("topaz.font", 8) {
            Result::Ok(f) => f,
            Result::Err(e) => {
                audio.shutdown()
            }
        }
    }
}

```

```

        close_window(window)
        close_screen(screen)
        return Result::Err(AppError::Font(e))
    }
}

Result::Ok(Application { screen, window, audio, font })
}

impl Drop for Application {
    fn drop(&var self) {
        close_font(self.font)
        self.audio.shutdown()
        close_window(self.window)
        close_screen(self.screen)
    }
}

```

Coming from C

In C, multi-step initialization often uses the “goto cleanup” pattern:

```

int init_app(App *app) {
    app->screen = OpenScreen(&spec);
    if (!app->screen) goto fail_screen;

    app->window = OpenWindow(&spec);
    if (!app->window) goto fail_window;

    app->audio = init_audio();
    if (!app->audio) goto fail_audio;

    return 0; // Success

fail_audio:
    CloseWindow(app->window);
fail_window:
    CloseScreen(app->screen);
fail_screen:
    return -1;
}

```

While functional, this pattern is error-prone. Novus’s RAII with `Drop` ensures cleanup happens automatically and in the correct order.

10.18 Error Handling Best Practices

10.18.1 Do’s

- **Do** use `?` for concise error propagation
- **Do** pattern match to handle specific error variants
- **Do** use `defer` for cleanup that must happen
- **Do** create custom error types for your domain

- **Do** implement `message()` for human-readable errors
- **Do** use `unwrap_or` when a default makes sense
- **Do** test error paths, not just the happy path
- **Do** log errors before propagating in complex systems

10.18.2 Don'ts

- **Don't** use `unwrap()` unless failure is impossible
- **Don't** ignore `Result` return values
- **Don't** use `panic!()` for recoverable errors
- **Don't** catch and re-throw without adding context
- **Don't** use sentinel values (0, -1, null) instead of `Result`
- **Don't** forget to handle all error variants in `match`
- **Don't** mix error handling strategies in the same codebase

10.18.3 Error Design Guidelines

When designing error types for a library or module:

1. **Be specific:** Create variants for each distinct failure mode
2. **Include context:** Add data that helps diagnose the problem
3. **Be recoverable:** If an error can be retried or worked around, make that obvious
4. **Be consistent:** Use similar patterns across your codebase
5. **Document errors:** Each variant should have a doc comment

```
/// Errors that can occur during network operations
pub enum NetworkError {
    /// Connection to server failed
    /// Contains: host address, port
    ConnectionFailed { host: *u8, port: u16 },

    /// Connection timed out after max_wait milliseconds
    /// Recoverable: retry with longer timeout
    Timeout { operation: *u8, max_wait: u32 },

    /// Server sent invalid response
    /// Not recoverable: protocol error
    InvalidResponse { code: u32 },

    /// Network stack not initialized
    /// Recoverable: call network_init() first
    NotInitialized,
}
```

Chapter 11

Advanced Traits and Generics

“Abstraction is not about vagueness, it is about being precise on a new semantic level.” — Edsger W. Dijkstra

Chapter 4 introduced traits as a way to define shared behavior across types. This chapter goes deeper, covering operator overloading, the complete trait ecosystem, advanced generic patterns, and real-world idioms that make Novus code both flexible and efficient.

By the end of this chapter, you’ll understand how to:

- Make your types work with Novus operators (+, -, *, [], etc.)
- Implement comparison and ordering for custom types
- Create fallible and infallible type conversions
- Build custom iterators for any data structure
- Write generic numeric algorithms
- Design flexible APIs with reference traits
- Apply common patterns like newtype, extension traits, and blanket implementations

11.1 Operator Overloading

Novus allows you to define how operators work with your custom types by implementing operator traits. When you implement the `Add` trait for a type, the `+` operator works with values of that type.

11.1.1 Arithmetic Operator Traits

The arithmetic operator traits are defined in `std::core`:

```
pub trait Add {
    fn add(&self, other: &Self) -> Self
}

pub trait Sub {
    fn sub(&self, other: &Self) -> Self
}

pub trait Mul {
    fn mul(&self, other: &Self) -> Self
}

pub trait Div {
    fn div(&self, other: &Self) -> Self
}
```

```

}

pub trait Rem {
    fn rem(&self, other: &Self) -> Self
}

pub trait Neg {
    fn neg(&self) -> Self
}

```

Each trait method takes references and returns an owned value. This design avoids unnecessary copies while keeping the interface simple.

Complete Operator-to-Trait Mapping

Operator	Trait	Method Signature
<code>a + b</code>	Add	<code>fn add(&self, other: &Self) -> Self</code>
<code>a - b</code>	Sub	<code>fn sub(&self, other: &Self) -> Self</code>
<code>a * b</code>	Mul	<code>fn mul(&self, other: &Self) -> Self</code>
<code>a / b</code>	Div	<code>fn div(&self, other: &Self) -> Self</code>
<code>a % b</code>	Rem	<code>fn rem(&self, other: &Self) -> Self</code>
<code>-a</code>	Neg	<code>fn neg(&self) -> Self</code>

Table 11.1: Arithmetic Operator Traits

11.1.2 Implementing Arithmetic Operators

Let's implement arithmetic for a 2D vector type:

```

from std::core import Add, Sub, Neg

struct Vec2 {
    x: i32,
    y: i32,
}

impl Vec2 {
    pub fn new(x: i32, y: i32) -> Vec2 {
        return Vec2 { x: x, y: y }
    }
}

impl Add for Vec2 {
    fn add(&self, other: &Vec2) -> Vec2 {
        return Vec2 {
            x: self.x + other.x,
            y: self.y + other.y
        }
    }
}

impl Sub for Vec2 {
    fn sub(&self, other: &Vec2) -> Vec2 {
        return Vec2 {
            x: self.x - other.x,

```



```

        y: self.y - other.y
    }
}

impl Neg for Vec2 {
    fn neg(&self) -> Vec2 {
        return Vec2 { x: -self.x, y: -self.y }
    }
}

```

Now you can use operators naturally:

```

// ... (with Vec2 and operator impls defined above)
let a = Vec2::new(10, 20)
let b = Vec2::new(5, 15)

let sum = a + b           // Vec2 { x: 15, y: 35 }
let diff = a - b          // Vec2 { x: 5, y: 5 }
let negated = -a          // Vec2 { x: -10, y: -20 }

```

11.1.3 Fixed-Point Arithmetic Example

Here's how `fixed32` (16.16 fixed-point) implements multiplication:

```

// ... (illustrative - actual impl is in compiler)
impl Mul for fixed32 {
    fn mul(&self, other: &fixed32) -> fixed32 {
        // Fixed-point multiplication requires shifting
        // a * b = (a_raw * b_raw) >> 16
        // This is handled by the compiler intrinsic
        return fixed32_mul(*self, *other)
    }
}

```

Performance note: On 68000, 32-bit multiplication is done in software. The 68020+ have hardware `MULU.L` and `MULS.L` instructions. The compiler generates appropriate code based on your `-cpu` target.

11.1.4 Common Mistakes

Mistake 1: Forgetting that methods take references

```

// WRONG - takes self by value
impl Add for Vec2 {
    fn add(self, other: Vec2) -> Vec2 { // Wrong signature!
        // ...
    }
}

// CORRECT - takes references
impl Add for Vec2 {
    fn add(&self, other: &Vec2) -> Vec2 {
        // ...
    }
}

```

Mistake 2: Division by zero

Operator traits don't have built-in error handling. If your type can encounter division by zero, you have options:

```
// Option 1: Panic (simple but crashes)
impl Div for Fraction {
    fn div(&self, other: &Fraction) -> Fraction {
        if other.numerator == 0 {
            panic!("Division by zero")
        }
        // ... normal division
    }
}

// Option 2: Saturate or return a sentinel value
impl Div for Fixed32 {
    fn div(&self, other: &Fixed32) -> Fixed32 {
        if other.is_zero() {
            return Fixed32::max_value() // Saturate to max
        }
        // ... normal division
    }
}
```

For fallible operations, consider providing a separate `try_div` method that returns `Result`.

11.1.5 Index Operators

The `Index` and `IndexMut` traits enable the `[]` operator for your types.

```
pub trait Index<I, T> {
    fn index(&self, idx: I) -> T
}

pub trait IndexMut<I, T> {
    fn index_set(&var self, idx: I, value: T)
}
```

Note the asymmetry: `Index` returns a value, while `IndexMut` takes a value to set. This differs from some other languages where both return references.

Implementing Index for a Grid

```
from std::core import Index, IndexMut

struct Grid {
    data: Vec<i32>,
    width: u32,
    height: u32,
}

impl Grid {
    pub fn new(width: u32, height: u32) -> Option<Grid> {
        let size = width * height
        match Vec::with_capacity(size) {
            Option::Some(var v) => {
```

```

        // Initialize with zeros
        for i in 0..size {
            v.push(0)
        }
        return Option::Some(Grid {
            data: v,
            width: width,
            height: height
        })
    },
    Option::None => return Option::None
}

fn calc_index(&self, x: u32, y: u32) -> u32 {
    return y * self.width + x
}

// Read access: let value = grid[(x, y)]
impl Index<(u32, u32), i32> for Grid {
    fn index(&self, pos: (u32, u32)) -> i32 {
        let (x, y) = pos
        let idx = self.calc_index(x, y)
        return self.data.get(idx).unwrap()
    }
}

// Write access: grid[(x, y)] = value
impl IndexMut<(u32, u32), i32> for Grid {
    fn index_set(&var self, pos: (u32, u32), value: i32) {
        let (x, y) = pos
        let idx = self.calc_index(x, y)
        self.data.set(idx, value)
    }
}

```

Usage:

```

// ... (with Grid defined above)
var grid = Grid::new(10, 10).unwrap()

grid[(5, 3)] = 42 // Write using IndexMut
let value = grid[(5, 3)] // Read using Index

```

Bounds Checking

In debug builds, you should add bounds checking to catch programming errors:

```

// ... (with Grid defined above)
impl Index<(u32, u32), i32> for Grid {
    fn index(&self, pos: (u32, u32)) -> i32 {
        let (x, y) = pos
        if x >= self.width || y >= self.height {
            panic!("Grid index out of bounds")
        }
        let idx = self.calc_index(x, y)
    }
}

```

```

        return self.data.get(idx).unwrap()
    }
}

```

For production code where bounds are guaranteed, you can use unchecked access for performance.

11.2 Comparison and Equality Traits

Beyond the basic `Eq` trait (covered in Chapter 4), Novus provides `PartialOrd` and `Ord` for ordering comparisons.

11.2.1 The Equality Contract

The `Eq` trait requires these mathematical properties:

- **Reflexive:** `a.eq(&a)` must return `true`
- **Symmetric:** `a.eq(&b)` must equal `b.eq(&a)`
- **Transitive:** if `a.eq(&b)` and `b.eq(&c)`, then `a.eq(&c)`

Violating these properties will cause collections like `HashMap` to behave incorrectly.

11.2.2 PartialOrd — Comparison Operators

`PartialOrd` enables the `<`, `<=`, `>`, and `>=` operators:

```

pub trait PartialOrd {
    fn lt(&self, other: &Self) -> bool    // <
    fn le(&self, other: &Self) -> bool    // <=
    fn gt(&self, other: &Self) -> bool    // >
    fn ge(&self, other: &Self) -> bool    // >=
}

```

Implementing PartialOrd

```

from std::core import Eq, PartialOrd

struct Version {
    major: u32,
    minor: u32,
    patch: u32,
}

impl Eq for Version {
    fn eq(&self, other: &Version) -> bool {
        if self.major != other.major { return false }
        if self.minor != other.minor { return false }
        return self.patch == other.patch
    }
}

impl PartialOrd for Version {

```

```

fn lt(&self, other: &Version) -> bool {
    if self.major != other.major {
        return self.major < other.major
    }
    if self.minor != other.minor {
        return self.minor < other.minor
    }
    return self.patch < other.patch
}

fn le(&self, other: &Version) -> bool {
    return self.lt(other) || self.eq(other)
}

fn gt(&self, other: &Version) -> bool {
    return other.lt(self)
}

fn ge(&self, other: &Version) -> bool {
    return other.le(self)
}
}

```

Usage:

```

// ... (with Version defined above)
let v1 = Version { major: 1, minor: 2, patch: 0 }
let v2 = Version { major: 1, minor: 3, patch: 0 }

if v1 < v2 {
    println("v1 is older than v2")
}

if v1 >= v2 {
    println("v1 is at least as new as v2")
}

```

11.2.3 Ord — Total Ordering

`Ord` is for types where every value is comparable. It's required for sorting and binary search.

```

pub trait Ord {
    /// Returns: negative if self < other, zero if equal,
    ///           positive if self > other
    fn cmp(&self, other: &Self) -> i32
}

```

The return value convention:

- Negative (< 0): `self` is less than `other`
- Zero (== 0): `self` equals `other`
- Positive (> 0): `self` is greater than `other`

```

from std::core import Ord

```

```
// ... (with Version and Eq defined above)
impl Ord for Version {
    fn cmp(&self, other: &Version) -> i32 {
        if self.major != other.major {
            return (i32)self.major - (i32)other.major
        }
        if self.minor != other.minor {
            return (i32)self.minor - (i32)other.minor
        }
        return (i32)self.patch - (i32)other.patch
    }
}
```

With Ord implemented, your type works with sorting:

```
// ... (with Version and Ord defined above)
var versions: Vec<Version> = load_versions()
versions.sort() // Sorts using Ord::cmp
```

11.2.4 When to Use PartialOrd vs Ord

- Use PartialOrd when comparison might not always be possible (e.g., floating-point with NaN)
- Use Ord when all values are comparable (integers, strings, versions)

Floating-point types implement PartialOrd but not Ord because NaN is not comparable to anything, including itself.

11.2.5 Deriving Comparison Traits

For simple structs, use @derive:

```
#[derive(Eq)]
struct Point {
    x: i32,
    y: i32,
}
```

The derived Eq compares all fields in declaration order. Note that @derive currently supports Eq and Hash, but not Ord.

11.2.6 Common Mistakes

Mistake 1: Inconsistent Eq and Hash

If two values are equal via Eq, they *must* produce the same hash:

```
// WRONG - case-insensitive eq but case-sensitive hash
impl Eq for CaseInsensitiveString {
    fn eq(&self, other: &Self) -> bool {
        return self.to_lowercase() == other.to_lowercase()
    }
}

impl Hash for CaseInsensitiveString {
```

```

    fn hash(&self) -> u32 {
        return self.data.hash() // BUG: "Hello" and "hello" hash
                                differently!
    }
}

// CORRECT - both use lowercase
impl Hash for CaseInsensitiveString {
    fn hash(&self) -> u32 {
        return self.to_lowercase().hash()
    }
}

```

Mistake 2: Ord that disagrees with Eq

If `a.cmp(&b) == 0`, then `a.eq(&b)` should return `true`. Violating this causes subtle bugs in sorted collections.

11.3 Conversion Traits

Novus provides a family of traits for type conversion, covered in Chapter 4's **From** section. Here we cover the complete picture.

11.3.1 The Conversion Trait Family

Trait	Method	Use Case
<code>From<T></code>	<code>convert(T) -> Self</code>	Infallible conversion from T
<code>Into<T></code>	<code>into(self) -> T</code>	Infallible conversion to T
<code>TryFrom<T, E></code>	<code>try_from(T) -> Result<Self, E></code>	Fallible conversion from T
<code>TryInto<T, E></code>	<code>try_into(self) -> Result<T, E></code>	Fallible conversion to T

Table 11.2: Conversion Traits

11.3.2 From and Into

`From<T>` was covered in Chapter 4. The key insight is that `Into` is the reverse direction:

```

pub trait From<T> {
    fn convert(value: T) -> Self
}

pub trait Into<T> {
    fn into(consuming self) -> T
}

```

If you implement `From<A>` for B, Novus can automatically provide `Into<B>` for A. Always prefer implementing `From` over `Into`.

11.3.3 TryFrom and TryInto

When conversion can fail, use the fallible variants:

```

pub trait TryFrom<T, E> {
    fn try_from(value: T) -> Result<Self, E>
}

```

```
pub trait TryInto<T, E> {
    fn try_into(consuming self) -> Result<T, E>
}
```

Example: Parsing with TryFrom

```
from std::core import TryFrom, Result

pub enum ParseError {
    InvalidFormat,
    OutOfRange,
    Empty,
}

impl TryFrom<*u8, ParseError> for u8 {
    fn try_from(s: *u8) -> Result<u8, ParseError> {
        if s == null {
            return Result::Err(ParseError::Empty)
        }

        // Simple ASCII digit parsing
        var value: u32 = 0
        var ptr = s
        while *ptr != 0 {
            let c = *ptr
            if c < '0' || c > '9' {
                return Result::Err(ParseError::InvalidFormat)
            }
            value = value * 10 + (u32)(c - '0')
            if value > 255 {
                return Result::Err(ParseError::OutOfRange)
            }
            ptr = ptr + 1
        }
        return Result::Ok((u8)value)
    }
}
```

Usage:

```
// ... (with u8 TryFrom impl defined above)
match u8::try_from("255") {
    Result::Ok(n) => println(f"Parsed: {n}"),
    Result::Err(ParseError::OutOfRange) => println("Number too
        large"),
    Result::Err(ParseError::InvalidFormat) => println("Not a
        number"),
    Result::Err(ParseError::Empty) => println("Empty string"),
}
```

11.3.4 Designing Error Types for Conversions

Good conversion error types are:

- **Specific:** Different failure modes get different variants
- **Informative:** Include context when useful
- **Lightweight:** Don't allocate in the error path if possible

```
// Good - specific variants, no allocation
pub enum NumericParseError {
    Empty,
    InvalidCharacter,
    Overflow,
    Underflow,
}

// Less good - loses information
pub enum GenericError {
    Failed,
}
```

11.3.5 Choosing the Right Conversion Trait

- Use `From/Into` when conversion always succeeds
- Use `TryFrom/TryInto` when conversion can fail
- Implement `From`, not `Into` (you get `Into` for free)
- Never panic in `From` — if it can fail, use `TryFrom`

11.4 The Iterator System

Iterators provide a uniform way to traverse sequences. Implementing `Iterator` lets your type work with loops and iterator-consuming functions.

11.4.1 The Iterator Trait

```
pub trait Iterator<T> {
    fn next(&var self) -> Option<T>
}
```

The contract:

- `next()` returns `Some(value)` for each element
- When exhausted, returns `None`
- After returning `None`, should continue returning `None`

11.4.2 Creating a Custom Iterator

Let's implement an iterator for a simple range:

```

from std::core import Iterator, Option

struct RangeIter {
    current: i32,
    end: i32,
}

impl RangeIter {
    pub fn new(start: i32, end: i32) -> RangeIter {
        return RangeIter { current: start, end: end }
    }
}

impl Iterator<i32> for RangeIter {
    fn next(&var self) -> Option<i32> {
        if self.current >= self.end {
            return Option::None
        }
        let value = self.current
        self.current = self.current + 1
        return Option::Some(value)
    }
}

```

Usage:

```

// ... (with RangeIter defined above)
var iter = RangeIter::new(0, 5)
loop {
    match iter.next() {
        Option::Some(n) => println(f"{n}"),
        Option::None => break,
    }
}
// Output: 0, 1, 2, 3, 4

```

11.4.3 Iterating Collections

The `Iterable` trait connects collections to for-loops:

```

pub trait Iterable<T> {
    fn get(&self, index: u32) -> Option<T>
    fn len(&self) -> u32
}

```

When you write `for item in collection`, the compiler uses `len()` to determine bounds and `get()` to access each element by index.

Vec Iterator Example

```

// ... (illustrative - actual impl in stdlib)
struct VecIter<T> {
    vec: *Vec<T>,
    index: u32,
}

```

```
impl<T> Iterator<T> for VecIter<T> {
    fn next(&var self) -> Option<T> {
        if self.index >= (*self.vec).len() {
            return Option::None
        }
        let value = (*self.vec).get(self.index)
        self.index = self.index + 1
        return value
    }
}
```

11.4.4 Iterator for a Linked List

Here's a more complex example — iterating a singly-linked list:

```
from std::core import Iterator, Option

struct Node<T> {
    value: T,
    next: *Node<T>,
}

struct LinkedList<T> {
    head: *Node<T>,
    len: u32,
}

struct LinkedListIter<T> {
    current: *Node<T>,
}

impl<T> LinkedList<T> {
    pub fn iter(&self) -> LinkedListIter<T> {
        return LinkedListIter { current: self.head }
    }
}

impl<T> Iterator<T> for LinkedListIter<T> {
    fn next(&var self) -> Option<T> {
        if self.current == null {
            return Option::None
        }
        let value = (*self.current).value
        self.current = (*self.current).next
        return Option::Some(value)
    }
}
```

11.4.5 Performance Considerations

Iterators in Novus are designed for zero overhead:

- Iterator structs are typically small (a pointer and an index)
- `next()` calls are inlined by the optimizer

- No heap allocation for iteration itself
- Performance matches hand-written loops

When iterators might be slower:

- Very tight loops where function call overhead matters
- When the optimizer can't inline (rare)
- Complex iterator chains (consider breaking into loops)

11.4.6 Common Mistakes

Mistake 1: Modifying the collection while iterating

```
// ... (illustrative - shows incorrect pattern)
// WRONG - undefined behavior
var vec = Vec::new()
vec.push(1)
vec.push(2)
vec.push(3)

var iter = vec.iter()
loop {
    match iter.next() {
        Option::Some(n) => {
            if n == 2 {
                vec.push(4) // BUG: modifying while iterating!
            }
        },
        Option::None => break,
    }
}
```

Mistake 2: Using an exhausted iterator

```
// ... (illustrative - shows incorrect pattern)
var iter = vec.iter()
// Consume all elements
loop {
    match iter.next() {
        Option::Some(_) => {},
        Option::None => break,
    }
}

// This won't iterate anything - iterator is exhausted
loop {
    match iter.next() {
        Option::Some(n) => println(f"{n}"), // Never reached
        Option::None => break,
    }
}
```

Create a new iterator if you need to iterate again.

11.5 Numeric Traits

These traits enable generic programming over numeric types — essential for math libraries and algorithms.

11.5.1 Zero and One

```
pub trait Zero {
    fn zero() -> Self
    fn is_zero(&self) -> bool
}

pub trait One {
    fn one() -> Self
}
```

These represent the additive and multiplicative identity elements.

Implementations for Primitive Types

All numeric primitives implement these traits:

```
// ... (implemented in std::core)
impl Zero for i32 {
    fn zero() -> i32 { return 0i32 }
    fn is_zero(&self) -> bool { return *self == 0i32 }
}

impl One for i32 {
    fn one() -> i32 { return 1i32 }
}
```

11.5.2 Writing Generic Numeric Functions

```
from std::core import Zero, One, Add, Mul

fn sum<T>(items: *T, count: u32) -> T where T: Zero + Add {
    var total = T::zero()
    for i in 0..count {
        total = total + *(items + i)
    }
    return total
}

fn product<T>(items: *T, count: u32) -> T where T: One + Mul {
    var result = T::one()
    for i in 0..count {
        result = result * *(items + i)
    }
    return result
}
```

Usage:

```
// ... (with sum and product defined above)
let numbers = [1, 2, 3, 4, 5]
let total = sum(&numbers[0], 5)    // 15
let prod = product(&numbers[0], 5) // 120
```

11.5.3 Bounded — Min and Max Values

```
pub trait Bounded {
    fn min_value() -> Self
    fn max_value() -> Self
}
```

Bounds for Novus Numeric Types

Type	min_value()	max_value()
i8	-128	127
i16	-32,768	32,767
i32	-2,147,483,648	2,147,483,647
u8	0	255
u16	0	65,535
u32	0	4,294,967,295

Table 11.3: Numeric Type Bounds

11.5.4 Practical Use: Saturating Arithmetic

Saturating arithmetic clamps at bounds instead of overflowing:

```
// ... (illustrative - generic numeric functions)
from std::core import Add, Bounded, PartialOrd

fn saturating_add<T>(a: T, b: T) -> T where T: Add + Bounded +
    PartialOrd {
    let result = a + b
    // Check for overflow: if result wrapped around
    if result.lt(&a) {
        return T::max_value()
    }
    return result
}

fn clamp<T>(value: T, min: T, max: T) -> T where T: PartialOrd {
    if value.lt(&min) { return min }
    if value.gt(&max) { return max }
    return value
}
```

11.5.5 Fixed-Point and Bounded

Fixed-point types also implement `Bounded`:

```
// ... (illustrative - actual values depend on representation)
impl Bounded for fixed32 {
    fn min_value() -> fixed32 {
        // -32768.0 in 16.16 format
        return fixed32_from_raw(-2147483648i32)
    }
    fn max_value() -> fixed32 {
        // 32767.99998... in 16.16 format
        return fixed32_from_raw(2147483647i32)
    }
}
```

11.6 Reference Traits

`AsRef` and `AsMut` enable functions to accept multiple types that can be “viewed as” a common type.

11.6.1 `AsRef` — Cheap Reference Conversion

```
pub trait AsRef<T> {
    fn as_ref(&self) -> T
}
```

Note: In Novus, `AsRef` returns `T` directly (typically a slice or reference type), not `&T`.

11.6.2 Designing Flexible APIs

Without `AsRef`:

```
// Inflexible - only accepts Str
fn count_lines(s: Str) -> u32 {
    var count: u32 = 0
    for c in s.bytes() {
        if c == '\n' {
            count = count + 1
        }
    }
    return count
}
```

With `AsRef`:

```
// Flexible - accepts String, Str, or anything AsRef<Str>
fn count_lines<S>(s: &S) -> u32 where S: AsRef<Str> {
    let str_view = s.as_ref()
    var count: u32 = 0
    // ... count newlines
    return count
}
```

11.6.3 `AsMut` — Mutable View Conversion

```
pub trait AsMut<T> {  
    fn as_mut(&var self) -> T  
}
```

Use when your function needs to modify the underlying data.

11.6.4 Standard Library Usage

Many stdlib functions use `AsRef` for flexibility:

```
// ... (illustrative - file opening API)  
pub fn open<P>(path: &P, mode: u32) -> Result<FileHandle,  
    DosError>  
    where P: AsRef<Str>  
{  
    let path_str = path.as_ref()  
    // ... open the file using path_str  
}
```

Callers benefit:

```
// Works with string literal (Str)  
let f1 = open(&"config.txt", MODE_OLDFILE)  
  
// Works with owned String  
var path = String::new()  
path.push_str("data/file.dat")  
let f2 = open(&path, MODE_OLDFILE)
```

11.6.5 When Not to Use `AsRef`

- When you need ownership (use `Into` instead)
- When the conversion is expensive (`AsRef` should be cheap)
- When it makes the API confusing (simpler is sometimes better)

11.7 Advanced Generics

This section covers more sophisticated generic patterns.

11.7.1 Multiple Trait Bounds

Use `+` to require multiple traits:

```
fn debug_and_hash<T>(value: &T) -> u32 where T: Debug + Hash {  
    println(value.debug_string())  
    return value.hash()  
}
```

11.7.2 Multiple Type Parameters


```
// ... (illustrative - multiple generic parameters)
fn convert_and_store<T, U>(input: T, storage: &var Vec<U>) where
    T: Into<U> {
    let converted: U = input.into()
    storage.push(converted)
}
```

11.7.3 Where Clauses on Structs

Constrain which types can instantiate a generic struct:

```
// ... (illustrative - where clause on struct)
struct SortedVec<T> where T: Ord {
    items: Vec<T>,
}

impl<T> SortedVec<T> where T: Ord {
    pub fn new() -> Option<SortedVec<T>> {
        match Vec::new() {
            Option::Some(v) => {
                return Option::Some(SortedVec { items: v })
            },
            Option::None => return Option::None
        }
    }

    pub fn insert(&var self, value: T) {
        // Find insertion point to maintain sorted order
        var idx: u32 = 0
        for i in 0..self.items.len() {
            match self.items.get(i) {
                Option::Some(item) => {
                    if item.cmp(&value) > 0 {
                        break
                    }
                },
                Option::None => break,
            }
            idx = i + 1
        }
        self.items.insert(idx, value)
    }
}
```

11.7.4 Generic Implementations

Implement traits generically for families of types:

```
// ... (illustrative - Display for all Option<T> where T: Display)
impl<T> Display for Option<T> where T: Display {
    fn fmt(&self, f: &var Formatter) -> bool {
        match *self {
            Option::Some(ref v) => {
                f.write_str("Some(")
                v.fmt(f)
            }
        }
    }
}
```

```
        f.write_str(""))
    },
    Option::None => {
        f.write_str("None")
    },
}
return true
}
}
```

11.7.5 Bounds on Individual Methods

Add bounds to specific methods rather than the entire impl:

```
// ... (illustrative - method-level where clause)
impl<T> Vec<T> {
    // Available for all Vec<T>
    pub fn len(&self) -> u32 {
        return self.length
    }

    // Only available when T: Eq
    pub fn contains(&self, value: &T) -> bool where T: Eq {
        for i in 0..self.len() {
            match self.get(i) {
                Option::Some(item) => {
                    if item.eq(value) {
                        return true
                    }
                },
                Option::None => {},
            }
        }
        return false
    }
}
```

11.7.6 Const Generic Parameters

For compile-time size parameters:

```
struct FixedBuffer<const N: u32> {
    data: [u8; N],
    len: u32,
}

impl<const N: u32> FixedBuffer<N> {
    pub fn new() -> FixedBuffer<N> {
        return FixedBuffer {
            data: [0u8; N],
            len: 0,
        }
    }

    pub fn capacity(&self) -> u32 {
        return N
    }
}
```

```
    }
}
```

Usage:

```
// ... (with FixedBuffer defined above)
let small = FixedBuffer::<64>::new() // 64-byte buffer
let large = FixedBuffer::<1024>::new() // 1KB buffer
```

11.7.7 Monomorphization

Generics in Novus are monomorphized — the compiler generates specialized code for each concrete type used:

```
fn max<T>(a: T, b: T) -> T where T: Ord {
    if a.cmp(&b) > 0 { return a }
    return b
}

// When you call:
let m1 = max(10i32, 20i32)
let m2 = max(5u8, 3u8)

// Compiler generates:
// fn max_i32(a: i32, b: i32) -> i32 { ... }
// fn max_u8(a: u8, b: u8) -> u8 { ... }
```

Benefits: Zero runtime overhead, full optimization.

Cost: Code size increases with each instantiation. On memory-constrained 68k systems, be mindful of instantiating generics with many different types.

11.8 Common Patterns and Idioms

These patterns appear frequently in well-designed Novus code.

11.8.1 The Newtype Pattern

Wrap a type to give it distinct identity:

```
struct Pixels {
    value: i32,
}

struct Bytes {
    value: u32,
}

struct ChipAddr {
    ptr: *u8,
}

impl Pixels {
    pub fn new(value: i32) -> Pixels {
        return Pixels { value: value }
    }
}
```

```

    pub fn get(&self) -> i32 {
        return self.value
    }
}

```

Now the compiler prevents mixing them up:

```

// ... (with Pixels and Bytes defined above)
fn allocate_pixels(width: Pixels, height: Pixels) -> *u8 {
    // ...
}

fn allocate_bytes(size: Bytes) -> *u8 {
    // ...
}

let w = Pixels::new(320)
let h = Pixels::new(256)
let size = Bytes { value: 1024 }

allocate_pixels(w, h)      // OK
allocate_pixels(w, size)  // Compile error! Bytes != Pixels

```

Newtype for Chip RAM Addresses

On the Amiga, chip RAM and fast RAM serve different purposes. Use newtypes to prevent bugs:

```

struct ChipPtr {
    ptr: *u8, // Pointer to chip RAM
}

struct FastPtr {
    ptr: *u8, // Pointer to fast RAM
}

fn setup_copper_list(gfx: ChipPtr) {
    // Copper can only access chip RAM
    // Type system ensures we can't accidentally pass fast RAM
}

fn cache_data(buf: FastPtr) {
    // Data caching works best in fast RAM
}

```

11.8.2 Extension Traits

Add methods to types you don't own:

```

trait I32Ext {
    fn is_even(&self) -> bool
    fn is_positive(&self) -> bool
    fn abs(&self) -> i32
}

```

```
impl I32Ext for i32 {
    fn is_even(&self) -> bool {
        return *self % 2 == 0
    }

    fn is_positive(&self) -> bool {
        return *self > 0
    }

    fn abs(&self) -> i32 {
        if *self < 0 { return -*self }
        return *self
    }
}
```

Usage:

```
// ... (with I32Ext defined above)
let x: i32 = -42
if x.is_even() {
    println("Even!")
}
let positive = x.abs() // 42
```

11.8.3 Blanket Implementations

Implement a trait for all types matching a bound:

```
// ... (illustrative - every Display type gets to_string)
impl<T> ToString for T where T: Display {
    fn to_string(&self) -> Option<String> {
        match String::new() {
            Option::Some(var s) => {
                // Use Display to write to string
                self.fmt(&var s.as_formatter())
                return Option::Some(s)
            },
            Option::None => return Option::None,
        }
    }
}
```

Now any type implementing `Display` automatically gets `to_string()`.

11.8.4 Marker Traits

Traits with no methods, used to mark capabilities:

```
pub trait Copy {}
```

The `Copy` marker tells the compiler a type can be implicitly copied (bitwise) rather than moved. Primitives like `i32`, `u8`, and `bool` are `Copy`.

11.8.5 The Builder Pattern

For complex object construction:

```
struct WindowConfig {
    title: *u8,
    width: u32,
    height: u32,
    borderless: bool,
}

struct WindowBuilder {
    config: WindowConfig,
}

impl WindowBuilder {
    pub fn new() -> WindowBuilder {
        return WindowBuilder {
            config: WindowConfig {
                title: "Untitled",
                width: 320,
                height: 256,
                borderless: false,
            }
        }
    }

    @chain
    pub fn title(&var self, t: *u8) {
        self.config.title = t
    }

    @chain
    pub fn size(&var self, w: u32, h: u32) {
        self.config.width = w
        self.config.height = h
    }

    @chain
    pub fn borderless(&var self) {
        self.config.borderless = true
    }

    pub fn build(&self) -> Result<WindowHandle, IntuitionError> {
        // ... create the window
    }
}
```

Usage:

```
// ... (with WindowBuilder defined above)
var builder = WindowBuilder::new()
let window = builder
    .title("My App")
    .size(640, 480)
    .borderless()
    .build()?
```

11.8.6 The Handle Pattern

Wrap OS resources with RAII cleanup:

```
struct FileHandle {
    lock: *FileLock,
}

impl FileHandle {
    pub fn open(path: *u8, mode: i32) -> Result<FileHandle,
        DosError> {
        let lock = Lock(path, mode)
        if lock == null {
            return Result::Err(dos_last_error())
        }
        return Result::Ok(FileHandle { lock: lock })
    }
}

impl Drop for FileHandle {
    fn drop(&var self) {
        if self.lock != null {
            UnLock(self.lock)
            self.lock = null
        }
    }
}
```

The file is automatically closed when the `FileHandle` goes out of scope.

11.9 Summary

This chapter covered advanced trait and generic features:

- **Operator traits** — Add, Sub, Mul, Div, Rem, Neg, Index, IndexMut
- **Comparison traits** — Eq, PartialOrd, Ord
- **Conversion traits** — From, Into, TryFrom, TryInto
- **Iterator traits** — Iterator, Iterable
- **Numeric traits** — Zero, One, Bounded
- **Reference traits** — AsRef, AsMut
- **Advanced generics** — multiple bounds, where clauses, const generics
- **Common patterns** — newtype, extension traits, blanket implementations, builder, handle

Traits are the foundation of generic programming in Novus. Master them, and you'll write code that is both flexible and efficient — code worthy of the Amiga.

Chapter 12


Standard Library Overview

“A language is only as good as its libraries.” — Bjarne Stroustrup

The Novus standard library provides everything you need for Amiga development: collections, strings, memory management, I/O, and comprehensive AmigaOS bindings. This chapter introduces the library’s organization, conventions, and design philosophy. Subsequent chapters cover each module in depth.

12.1 Library Organization

The standard library is organized into modules under the `std` namespace:



12.1.1 Module Hierarchy

std::core Fundamental types: `Option`, `Result`, core traits (`Drop`, `Clone`, `From`)

std::collections Data structures: `Vec`, `HashMap`, `HashSet`, `VecDeque`, `BitSet`, `SlotMap`

std::strings String handling: `Str` (borrowed), `String` (owned), `StringBuilder`, formatting

std::memory Memory management: `MemoryBlock`, `Slice`, chip memory pools, allocators

std::io Input/output: file operations, console I/O, ANSI terminal support

std::ffi AmigaOS bindings: `Exec`, `DOS`, `Graphics`, `Intuition`, and 70+ other libraries

std::math Mathematics: fixed-point operations, vectors, angles, square roots

std::async Asynchronous programming: futures, executors, sleep operations

std::hardware Hardware access: CPU detection, chipset identification, custom registers

std::ui User interface: windows, screens, menus, gadgets, events

std::graphics Graphics operations: drawing, sprites, copper lists, blitter

std::audio Audio playback: Paula channels, MOD players, sample management

std::net Networking: TCP/UDP sockets, DNS, TLS (requires `bsdsocket.library`)

std::test Testing framework: assertions, test attributes, benchmarking

std::error Error types: `DosError`, `ExecError`, `NovusError`

12.1.2 Physical Layout

The library lives in the `std/` directory of the Novus installation:

```
std/
  core.novus           # Core types (Option, Result, traits)
  prelude.novus        # Common re-exports
  collections/
    vec.novus          # Dynamic array
    hashmap.novus      # Hash map
    hashset.novus      # Hash set
    vecdeque.novus     # Double-ended queue
    bitset.novus       # Bit array
    slotmap.novus      # Generational arena
    ...
  strings/
    core.novus         # Str and String types
    format.novus       # Formatting utilities
    display.novus      # Display trait
  memory/
    block.novus        # MemoryBlock type
    slice.novus        # Slice operations
    chip.novus         # Chip memory utilities
    chip_pool.novus    # Chip memory pooling
    ...
  ffi/
    exec.novus         # Exec library
    dos.novus          # DOS library
    graphics.novus     # Graphics library
    intuition.novus    # Intuition library
    amiga_consts.novus # All AmigaOS constants
    amiga_structs.novus # All AmigaOS structures
    ...
  ...
```

12.2 The Prelude

The prelude provides commonly-used types that you'll need in almost every program. Import it to get started quickly:

```
from std::prelude import *
```

This imports:

- `Option<T>` and `Result<T, E>` — for safe error handling
- `Drop`, `From` — core traits
- `Str` — borrowed string type
- `DosError`, `ExecError`, `IntuitionError`, `GraphicsError`, `NovusError` — error types

12.2.1 Selective Imports

For explicit control, import only what you need:

```
from std::prelude import Option, Result
```

Or import directly from the source module:

```
from std::core import Option, Result, Drop
from std::strings::core import Str, String
from std::error::errors import DosError
```

Coming from C

In C, you might include multiple headers to get the types you need:

```
#include <exec/types.h>
#include <dos/dos.h>
#include <string.h>
```

In Novus, the prelude gives you essentials in one import. For AmigaOS types, the `std::ffi` modules are organized by library, not by NDK header structure.

12.3 Import Conventions

Novus uses a consistent import ordering convention throughout the standard library and in idiomatic user code.

12.3.1 Recommended Import Order

```
// 1. Core types first
from std::core import Option, Result, Drop, Clone

// 2. Collections
from std::collections::vec import Vec
from std::collections::hashmap import HashMap

// 3. Memory management
from std::memory::block import MemoryBlock
from std::memory::slice import Slice

// 4. Strings
from std::strings::core import Str, String
from std::strings::format import format

// 5. Error types
from std::error::errors import DosError, ExecError, NovusError

// 6. AmigaOS structures
from std::ffi::amiga_structs import Window, Screen, RastPort

// 7. AmigaOS constants
from std::ffi::amiga_consts import MEMF_PUBLIC, MEMF_CHIP

// 8. AmigaOS library functions
from std::ffi::exec import AllocMem, FreeMem, Wait
from std::ffi::dos import Open, Close, Read, Write
from std::ffi::intuition import OpenWindow, CloseWindow
```

```
// 9. Other stdlib modules
from std::io::file import println
from std::math::vec2 import Vec2
```

12.3.2 Wildcard Imports

Use wildcard imports sparingly—they can cause name collisions:

```
// OK: prelude is designed for wildcard import
from std::prelude import *

// OK: constants are namespaced and rarely collide
from std::ffi::amiga_consts import *

// Avoid: too many names, potential collisions
from std::ffi::exec import *
from std::ffi::dos import *

// Better: import specific functions
from std::ffi::exec import AllocMem, FreeMem, Wait
from std::ffi::dos import Open, Close, Lock, UnLock
```

12.3.3 Module Aliases

For frequently-used modules with long paths, create local aliases:

```
// Not yet supported - planned feature
// use std::ffi::amiga_consts as consts
// use std::collections::hashmap as map

// Current workaround: import specific items
from std::ffi::amiga_consts import MEMF_PUBLIC, MEMF_CHIP,
    MEMF_FAST
from std::ffi::amiga_consts import MODE_OLDFILE, MODE_NEWFILE
```

12.4 Core Types

The `std::core` module defines the fundamental types that permeate all Novus code.

12.4.1 Option<T>

Represents a value that may or may not be present:

```
pub enum Option<T> {
    Some(T),
    None,
}
```

Key methods:

Method	Description
<code>is_some()</code>	Returns <code>true</code> if value is present
<code>is_none()</code>	Returns <code>true</code> if no value
<code>unwrap()</code>	Returns value, panics if <code>None</code>
<code>unwrap_or(default)</code>	Returns value or default
<code>or(other)</code>	Returns self if <code>Some</code> , otherwise other
<code>FromPointer(ptr)</code>	Creates <code>Some</code> if non-null, <code>None</code> if null

```
from std::core import Option

fn find_user(id: u32) -> Option<User> {
    if id == 0 {
        return Option::None
    }
    // ... lookup user ...
    return Option::Some(user)
}

// Usage
match find_user(42) {
    Option::Some(user) => println(user.name),
    Option::None => println("User not found")
}

// Or with unwrap_or
let user = find_user(42).unwrap_or(default_user)
```

12.4.2 Result<T, E>

Represents either success with a value or failure with an error:

```
pub enum Result<T, E> {
    Ok(T),
    Err(E),
}
```

Key methods:

Method	Description
<code>is_ok()</code>	Returns <code>true</code> if success
<code>is_err()</code>	Returns <code>true</code> if error
<code>unwrap()</code>	Returns value, panics if error
<code>unwrap_or(default)</code>	Returns value or default
<code>ok()</code>	Converts to <code>Option<T></code> , discarding error
<code>err()</code>	Converts to <code>Option<E></code> , discarding value

```
from std::core import Result
from std::error::errors import DosError

fn load_config(path: *u8) -> Result<Config, DosError> {
    let file = open(path, MODE_OLDFILE)?
    defer { close(file) }

    let data = read_all(&file)?
```

```

    parse_config(&data)
}

// Usage with ? operator
fn init() -> Result<(), DosError> {
    let config = load_config("app.cfg"?
    // ... use config ...
    Result::Ok(())
}

// Or with match
match load_config("app.cfg") {
    Result::Ok(config) => use_config(&config),
    Result::Err(e) => println(e.message())
}

```

12.4.3 Core Traits

Drop — Automatic cleanup when a value goes out of scope:

```

pub trait Drop {
    fn drop(&var self)
}

// Example: auto-closing file handle
impl Drop for FileHandle {
    fn drop(&var self) {
        if self.fh != null {
            unsafe { Close(self.fh) }
            self.fh = null
        }
    }
}

```

Clone — Explicit value copying:

```

pub trait Clone {
    fn clone(&self) -> Option<Self>
}

// Returns Option because cloning may allocate and fail
let original = Vec::new()
match original.clone() {
    Option::Some(copy) => { /* use copy */ },
    Option::None => { /* allocation failed */ }
}

```

From — Type conversion:

```

pub trait From<T> {
    fn convert(value: T) -> Self
}

// Example: DosError to NovusError
impl From<DosError> for NovusError {
    fn convert(e: DosError) -> NovusError {

```

```

        return NovusError::Dos(e)
    }
}


// Used automatically by ? operator
fn example() -> Result<(), NovusError> {
    let file = open("test.txt", MODE_OLDFILE)? // DosError ->
        NovusError
    Result::Ok(())
}

```

12.5 Error Types

The standard library provides error types for each AmigaOS subsystem.

12.5.1 Error Type Hierarchy



12.5.2 DosError

Errors from the DOS library (most common variants shown):

```

pub enum DosError {
    NotFound,           // Object not found (IoErr 205)
    AlreadyExists,      // Object already exists (IoErr 203)
    DirNotFound,        // Directory not found (IoErr 204)
    DirectoryNotEmpty,  // Can't delete non-empty dir (IoErr 216)
    PermissionDenied,   // Access denied
    WriteProtected,     // File write protected (IoErr 223)
    DiskFull,           // No space on disk (IoErr 221)
    NoDisk,             // No disk in drive (IoErr 226)
    InvalidLock,        // Bad file lock (IoErr 211)
    NoFreeStore,        // Out of memory (IoErr 103)
    ObjectInUse,        // Object is in use (IoErr 202)
    // ... plus 15+ more variants
    Unknown,            // Unmapped error code
}

```

12.5.3 ExecError

Errors from the Exec library:

```

pub enum ExecError {
    NoMem,              // Memory allocation failed
    BadSignal,          // Invalid signal number
    NullReturn,         // Function returned NULL unexpectedly
    BadTask,            // Invalid task pointer
    BadNode,            // Invalid list node
    BadPort,            // Invalid message port
}

```

```
SignalInUse,      // Signal already allocated
TaskNotFound,     // FindTask failed
LibraryNotFound,  // OpenLibrary failed
InvalidAddress,   // Invalid memory address
AlignmentError,   // Memory alignment violation
IndexOutOfBounds, // Array bounds violation
CapacityExceeded, // Fixed buffer full
}
```

12.5.4 NovusError

Wrapper for cross-subsystem error handling:

```
pub enum NovusError {
    Dos(DosError),
    Exec(ExecError),
    Intuition(IntuitionError),
    Graphics(GraphicsError),
}

// All error types implement message()
impl DosError {
    pub fn message(&self) -> *u8 {
        match *self {
            DosError::NotFound => "Object not found",
            DosError::DiskFull => "Disk is full",
            // ...
        }
    }
}
```

12.5.5 Converting Between Error Types

The `?` operator automatically converts errors using the `From` trait:

```
fn load_and_display(path: *u8) -> Result<(), NovusError> {
    // DosError automatically converts to NovusError
    let file = open(path, MODE_OLDFILE)?

    // ExecError also converts
    let buffer = MemoryBlock::alloc(1024, MEMF_PUBLIC)?

    Result::Ok(())
}
```

12.6 Memory Allocation Conventions

The standard library follows consistent patterns for memory allocation.

12.6.1 Fallible Allocation

All allocating functions return `Option` or `Result`:


```
// Vec allocation can fail - returns Result with error details
let vec = Vec::with_capacity(1000)?
// Or with explicit match:
match Vec::<i32>::with_capacity(1000) {
    Result::Ok(v) => { /* use v */ },
    Result::Err(e) => { /* out of memory: e.message() */ }
}

// MemoryBlock returns Option (simpler: worked or didn't)
match MemoryBlock::alloc(4096, MEMF_CHIP) {
    Option::Some(block) => { /* use block */ },
    Option::None => { /* allocation failed */ }
}
```

12.6.2 Memory Type Flags

AmigaOS memory flags are in `std::ffi::amiga_consts`:

```
from std::ffi::amiga_consts import MEMF_PUBLIC, MEMF_CHIP,
    MEMF_FAST
from std::ffi::amiga_consts import MEMF_CLEAR, MEMF_LARGEST,
    MEMF_TOTAL

// Memory flags:
// MEMF_PUBLIC - Accessible by other tasks
// MEMF_CHIP - Chip RAM (DMA accessible)
// MEMF_FAST - Fast RAM (CPU only)
// MEMF_CLEAR - Zero-fill memory
// MEMF_LARGEST - Return largest available block
// MEMF_TOTAL - Return total memory of type

// Typical patterns
let chip_buffer = MemoryBlock::alloc(size, MEMF_CHIP |
    MEMF_CLEAR)?
let fast_buffer = MemoryBlock::alloc(size, MEMF_FAST)?
let any_buffer = MemoryBlock::alloc(size, MEMF_PUBLIC)?
```

12.6.3 RAII Memory Management

All stdlib types that allocate implement `Drop`:

```
fn process_data() -> Result<(), Error> {
    let vec = Vec::with_capacity(100).unwrap()
    let buffer = MemoryBlock::alloc(1024, MEMF_PUBLIC)?

    // ... use vec and buffer ...

    Result::Ok(())
} // vec and buffer automatically freed here

// Even on early return:
fn early_exit() -> Result<(), Error> {
    let buffer = MemoryBlock::alloc(1024, MEMF_PUBLIC)?

    if some_condition {
```

```
        return Result::Err(Error::SomeError) // buffer freed here
    }

    Result::Ok(()) // buffer freed here too
}
```

12.7 AmigaOS FFI Organization

The `std::ffi` module provides bindings to AmigaOS libraries.

12.7.1 Structure

```
std/ffi/
  amiga_consts.novus      # ALL constants (combined)
  amiga_structs.novus     # ALL structures (combined)
  exec.novus              # Exec library functions
  dos.novus               # DOS library functions
  graphics.novus          # Graphics library functions
  intuition.novus         # Intuition library functions
  gadtools.novus          # GadTools library
  layers.novus            # Layers library
  utility.novus           # Utility library
  ... 70+ more libraries ...
```

12.7.2 Using FFI Bindings

```
// Import structures
from std::ffi::amiga_structs import Window, Screen, NewWindow

// Import constants
from std::ffi::amiga_consts import WFLG_CLOSEGADGET,
    WFLG_DRAGBAR, WFLG_ACTIVATE
from std::ffi::amiga_consts import IDCMP_CLOSEWINDOW,
    IDCMP_VANILLAKEY

// Import functions
from std::ffi::intuition import OpenWindow, CloseWindow, WaitPort
from std::ffi::exec import Wait

// Use in unsafe blocks
unsafe {
    let window = OpenWindow(&new_window_spec)
    if window != null {
        // ... use window ...
        CloseWindow(window)
    }
}
```

12.7.3 Safe Wrappers

Many FFI functions have safe wrappers in higher-level modules:

```
// Low-level FFI (requires unsafe)
from std::ffi::intuition import OpenWindowTags, CloseWindow

// High-level safe wrapper
from std::ui::window import Window, WindowBuilder

// Safe usage (no unsafe needed)
let window = WindowBuilder::new()
    .title("My Window")
    .size(320, 200)
    .close_gadget(true)
    .build()?

// Window automatically closed when dropped
```

Coming from C

In C, you call AmigaOS functions directly and manage lifetimes manually:

```
struct Window *win = OpenWindowTags(NULL,
    WA_Title, "My Window",
    WA_Width, 320,
    WA_Height, 200,
    TAG_DONE);
// ... must remember to call CloseWindow(win) ...
```

In Novus, the safe wrappers handle cleanup automatically. You can still use the raw FFI when needed, but the wrappers are usually sufficient.

12.8 Collections Overview

The `std::collections` module provides commonly-needed data structures.

12.8.1 Available Collections

Type	Module	Description
<code>Vec<T></code>	<code>vec</code>	Dynamic array, contiguous storage
<code>HashMap<K, V></code>	<code>hashmap</code>	Key-value map with $O(1)$ lookup
<code>HashSet<T></code>	<code>hashset</code>	Unique values with $O(1)$ contains
<code>VecDeque<T></code>	<code>vecdeque</code>	Double-ended queue
<code>BitSet</code>	<code>bitset</code>	Compact bit array
<code>SlotMap<T></code>	<code>slotmap</code>	Generational arena with stable keys
<code>RingBuffer<T></code>	<code>ringbuffer</code>	Fixed-size circular buffer
<code>FreeList<T></code>	<code>freelist</code>	Object pool with recycling
<code>SmallVec<T, N></code>	<code>smallvec</code>	Stack-allocated up to N elements
<code>IntrusiveList<T></code>	<code>intrusive_list</code>	Linked list with embedded nodes

12.8.2 Common Patterns

```
from std::collections::vec import Vec
from std::collections::hashmap import HashMap
```

```
// Dynamic array
var scores: Vec<i32> = Vec::new()
scores.push(100)?
scores.push(85)?
scores.push(92)?

for i in 0..scores.len() {
    let score = scores.get(i).unwrap()
    println(f"Score: {score}")
}

// Hash map
var names = HashMap::<u32, *u8>::new()?
names.insert(1, "Alice")?
names.insert(2, "Bob")?

match names.get(1) {
    Option::Some(name) => println(name),
    Option::None => println("Not found")
}
```

12.9 Strings Overview

Novus provides two main string types for different use cases.

12.9.1 String Types

Type	Ownership	Use Case
*u8	Borrowed	Literal strings, C interop
Str	Borrowed	String views with length
String	Owned	Mutable, growable strings
StringBuilder	Owned	Efficient string building

12.9.2 Usage Examples

```
from std::strings::core import Str, String, StringBuilder

// Literal (null-terminated, borrowed)
let greeting: *u8 = "Hello, Amiga!"

// Str - borrowed view with known length
let s = Str::from_cstr("Hello").unwrap()
println(f"Length: {s.len()}")

// String - owned, mutable
var name = String::from("Hello")
name.push_str(", World!")
println(name.as_cstr())

// StringBuilder - efficient concatenation
var builder = StringBuilder::new()
builder.push_str("Line 1\n")
builder.push_str("Line 2\n")
builder.push_str("Line 3\n")
```

```
let result = builder.to_string()
```

12.10 I/O Overview

The `std::io` module handles input and output operations.

12.10.1 Console Output

```
from std::io::file import print, println

// Basic output
println("Hello, World!")

// With formatting (f-strings)
let name = "Amiga"
let year = 1985
println(f"The {name} was released in {year}")

// Without newline
print("Enter name: ")
```

12.10.2 File Operations

```
from std::os::dos import open, close, read, write
from std::ffi::amiga_consts import MODE_OLDFILE, MODE_NEWFILE

// Reading a file
fn read_file(path: *u8) -> Result<Vec<u8>, DosError> {
    let file = open(path, MODE_OLDFILE)?
    defer { close(file) }

    var buffer: Vec<u8> = Vec::with_capacity(4096).unwrap()
    // ... read into buffer ...
    Result::Ok(buffer)
}

// Writing a file
fn write_file(path: *u8, data: &[u8]) -> Result<(), DosError> {
    let file = open(path, MODE_NEWFILE)?
    defer { close(file) }

    write(&file, data)?
    Result::Ok(())
}
```

12.11 Design Principles

The standard library follows several guiding principles.

12.11.1 Amiga First

The library is designed specifically for AmigaOS, not as a portable abstraction layer:

- Memory types (chip/fast) are first-class concepts
- AmigaOS error codes map directly to error types
- Custom chip access is supported, not hidden
- RAII works with AmigaOS resource management patterns

12.11.2 Zero-Cost Abstractions

High-level types compile to efficient code:

```
// Vec<i32> with bounds checking
let value = vec.get(index)?

// In release mode, compiles to roughly:
//  move.l (a0,d0.l*4),d0
// The Option wrapper is optimized away
```

12.11.3 Explicit Allocation

Every allocation is visible and can fail:

```
// Allocations return Option or Result
let vec = Vec::with_capacity(1000)? // Returns Result<Vec,
    ExecError>
let map = HashMap::new()? // Returns Result<HashMap,
    ExecError>
map.insert(key, value)? // Allocation happens here
```

12.11.4 Safe Defaults, Unsafe Escape Hatch

Most code uses safe wrappers. Raw FFI is available when needed:

```
// Safe (preferred)
let window = Window::open(&spec)?

// Unsafe (when needed)
unsafe {
    let raw_window = OpenWindow(&raw_spec)
}
```

12.12 What's Covered in Part II

The following chapters cover each module area in depth:

Chapter 13: Core Types Complete coverage of `Option`, `Result`, and core traits

Chapter 14: Collections All collection types with performance characteristics and usage patterns

Chapter 15: Strings String types, formatting, parsing, and Unicode handling

Chapter 16: Memory Memory management, allocators, pools, and chip memory handling

Chapter 17: I/O File operations, console I/O, formatting, and buffering

Chapter 18: Math Fixed-point arithmetic, vectors, matrices, and trigonometry

Chapter 19: AmigaOS Bindings Using the FFI layer, calling conventions, and safe wrappers

12.13 Quick Reference

12.13.1 Common Imports

```
// Everything you need for most programs
from std::prelude import *
from std::collections::vec import Vec
from std::io::file import println
from std::ffi::amiga_consts import MEMF_PUBLIC, MEMF_CHIP
```

12.13.2 Error Handling Pattern

```
fn main() -> i32 {
    match run() {
        Result::Ok(_) => return 0,
        Result::Err(e) => {
            println(e.message())
            return 1
        }
    }
}

fn run() -> Result<(), NovusError> {
    let config = load_config()?
    let window = open_window(&config)?
    main_loop(&window)?
    Result::Ok(())
}
```

12.13.3 Memory Allocation Pattern

```
fn process() -> Result<(), ExecError> {
    // Allocation returns Result
    let buffer = MemoryBlock::alloc(4096, MEMF_PUBLIC)?
    defer { /* buffer automatically freed */ }

    // Use buffer...
    do_work(buffer.as_slice())

    Result::Ok(())
}
```

12.14 Summary

The Novus standard library provides:

- **Core types** — `Option`, `Result`, and traits for safe error handling
- **Collections** — `Vec`, `HashMap`, and specialized containers
- **Strings** — Borrowed and owned strings with formatting
- **Memory** — Safe allocation with RAII cleanup
- **I/O** — File and console operations
- **FFI** — Complete AmigaOS library bindings
- **Math** — Fixed-point arithmetic and vectors
- **Async** — Cooperative multitasking via Exec signals

The library is designed for the Amiga: every feature respects the platform's constraints while providing modern ergonomics. Use the prelude for quick starts, import specific modules for explicit control.

The following chapters explore each module in depth, with complete API references and practical examples.

Chapter 13

Core Types and Traits

The `std::core` module is the foundation of Novus’s type system. It defines the fundamental types and traits that enable safe error handling, resource management, and generic programming. This chapter provides complete documentation for every type and trait in `std::core`.

13.1 `Option<T>`

`Option<T>` represents a value that may or may not be present. Use it instead of null pointers to make absence explicit and force handling.

```
pub enum Option<T> {  
    Some(T),  
    None,  
}
```

13.1.1 When to Use Option

- Function might not have a value to return (e.g., `Vec::pop()`)
- A field is optional
- Lookup operations that may fail to find an element
- Simple failure modes where “not found” is the only error

Note

Use `Option` when there’s only one failure mode. Use `Result` when you need to distinguish between different error conditions.

13.1.2 Creating Options

```
// Explicit construction  
let some_value: Option<i32> = Option::Some(42)  
let no_value: Option<i32> = Option::None  
  
// From a pointer (non-null check)  
let ptr: *u8 = get_some_pointer()  
let opt = Option::FromPointer(ptr) // Some(ptr) if non-null,  
    None otherwise
```

13.1.3 Option Methods

Method	Description
<code>is_some(self) -> bool</code>	Returns <code>true</code> if the option contains a value
<code>is_none(self) -> bool</code>	Returns <code>true</code> if the option is <code>None</code>
<code>unwrap(self) -> T</code>	Returns the value, panics if <code>None</code>
<code>unwrap_or(self, default: T) -> T</code>	Returns the value or a default
<code>unwrap_or_else(self, default: T) -> T</code>	Same as <code>unwrap_or</code> (closures planned)
<code>or(self, optb: Option<T>) -> Option<T></code>	Returns <code>self</code> if <code>Some</code> , otherwise <code>optb</code>
<code>contains(self, val: T) -> bool</code>	Returns <code>true</code> if contains the value (requires <code>Eq</code>)
<code>ok_or<E>(self, err: E) -> Result<T, E></code>	Converts to <code>Result</code> , mapping <code>None</code> to <code>Err(err)</code>
<code>FromPointer(ptr: *T) -> Option<*T></code>	Creates <code>Some(ptr)</code> if non-null, <code>None</code> otherwise

13.1.4 Pattern Matching with Option

The idiomatic way to handle `Option` is with pattern matching:

```
fn find_user(id: u32) -> Option<User> {
    // ... lookup logic ...
}

// Exhaustive match
match find_user(42) {
    Option::Some(user) => {
        println(user.name)
    },
    Option::None => {
        println("User not found")
    }
}

// If-let for single-arm matches
if let Option::Some(user) = find_user(42) {
    println(user.name)
}

// While-let for iteration
var stack = create_stack()
while let Option::Some(item) = stack.pop() {
    process(item)
}
```

13.1.5 Combining Options

```
// Chain with or()
let primary = get_primary_config()
let fallback = get_fallback_config()
let config = primary.or(fallback) // Use fallback if primary is
    None

// Convert to Result for richer error info
let user = find_user(id).ok_or(UserError::NotFound)?

// Check contents
```

```
let x = Option::Some(42)
if x.contains(42) {
    println("Found 42!")
}
```

Warning

Avoid unwrap() in library code. It panics on None, crashing the program. Use unwrap_or(), pattern matching, or the ? operator with ok_or() instead.

13.2 Result<T, E>

Result<T, E> represents an operation that can succeed with value T or fail with error E. It's the primary mechanism for error handling in Novus.

```
pub enum Result<T, E> {
    Ok(T),
    Err(E),
}
```

13.2.1 When to Use Result

- Operations that can fail in multiple ways
- I/O operations (file, network, hardware)
- Memory allocation
- Parsing and validation
- Any operation where callers need error details

13.2.2 Result Methods

Method	Description
is_ok(self) -> bool	Returns true if the result is Ok
is_err(self) -> bool	Returns true if the result is Err
unwrap(self) -> T	Returns the value, panics if Err
unwrap_err(self) -> E	Returns the error, panics if Ok
unwrap_or(self, default: T) -> T	Returns the value or a default
or(self, res: Result<T, E>) -> Result<T, E>	Returns self if Ok, otherwise res
ok(self) -> Option<T>	Converts Ok(v) to Some(v), Err to None
err(self) -> Option<E>	Converts Err(e) to Some(e), Ok to None

13.2.3 The ? Operator

The ? operator is syntactic sugar for early return on error:

```
fn read_config(path: *u8) -> Result<Config, DosError> {
    let file = open(path, MODE_OLDFILE)? // Returns Err early if
    failed
    defer { close(file) }
```

```

    let data = read_all(file)?
    let config = parse_config(data)?

    return Result::Ok(config)
}

// Equivalent without ?:
fn read_config_verbose(path: *u8) -> Result<Config, DosError> {
    match open(path, MODE_OLDFILE) {
        Result::Ok(file) => {
            defer { close(file) }
            match read_all(file) {
                Result::Ok(data) => {
                    match parse_config(data) {
                        Result::Ok(config) => return
                            Result::Ok(config),
                        Result::Err(e) => return Result::Err(e)
                    }
                },
                Result::Err(e) => return Result::Err(e)
            }
        },
        Result::Err(e) => return Result::Err(e)
    }
}

```

13.2.4 Error Type Conversion with ?

The ? operator automatically converts error types using the From trait:

```

// DosError -> NovusError conversion is automatic with ?
fn load_and_display(path: *u8) -> Result<(), NovusError> {
    let data = read_file(path)?           // DosError -> NovusError
    let window = open_window()?           // IntuitionError ->
        NovusError
    draw_data(window, data)?              // GraphicsError ->
        NovusError
    return Result::Ok(())
}

```

This works because `impl From<DosError> for NovusError` exists in the error module.

13.2.5 Pattern Matching with Result

```

match allocate_buffer(size) {
    Result::Ok(buffer) => {
        // Use buffer
        process(buffer)
    },
    Result::Err(ExecError::NoMem) => {
        println("Out of memory!")
    },
    Result::Err(e) => {
        // Handle other errors
        println(e.message())
    }
}

```

```
}
```

13.2.6 Converting Between Option and Result

```
// Option -> Result: provide an error for None case
let user = find_user(id).ok_or(UserError::NotFound)?

// Result -> Option: discard error information
let maybe_value = risky_operation().ok()
```

13.3 LibraryVersion

LibraryVersion represents semantic version information for AmigaOS libraries:

```
pub struct LibraryVersion {
    major: u16,
    minor: u16,
    patch: u16,
}
```

13.3.1 LibraryVersion Methods

```
// Creation
let version = LibraryVersion::new(39, 5, 0)

// Access components
let major = version.major() // 39
let minor = version.minor() // 5
let patch = version.patch() // 0

// Comparison (returns: negative if <, zero if ==, positive if >)
let cmp_result = version.cmp(&other_version)
if cmp_result < 0 {
    println("version is older")
}
```

13.4 Resource Management Traits

These traits enable automatic resource management (RAII) and explicit copying.

13.4.1 Drop — Automatic Cleanup

The Drop trait enables automatic resource cleanup when values go out of scope:

```
pub trait Drop {
    fn drop(&var self)
}
```

When `drop()` is called:

- At the end of a block scope
- Before reassignment to a variable
- Before early returns
- When a variable is shadowed

When `drop()` is NOT called:

- If the value was moved to another location
- If a field was partially moved (only non-moved fields are dropped)

```
impl Drop for MemoryBlock {
    fn drop(&var self) {
        if self.size > 0 {
            unsafe { FreeMem(self.ptr, self.size) }
            self.ptr = (*u8)0
            self.size = 0
        }
    }
}

// Usage - cleanup is automatic
fn example() -> Result<(), ExecError> {
    let mem = MemoryBlock::alloc(1024, MEMF_FAST).unwrap()
    // ... use mem ...
    return Result::Ok(())
} // mem.drop() called automatically here
```

Note

Drop order: Variables are dropped in reverse order of creation. Struct fields are dropped in declaration order. `defer` blocks run before automatic drop.

13.4.2 Copy — Implicit Bitwise Copy

The `Copy` trait is a marker trait indicating a type can be copied bitwise:

```
pub trait Copy {}
```

Types implementing `Copy` are implicitly copied when passed by value, rather than moved.

When to implement `Copy`:

- Simple value types (integers, bools, pointers)
- Small structs containing only `Copy` types
- Wrapper types around primitives

When NOT to implement `Copy`:

- Types that own heap memory (`Vec`, `String`, `Box`)
- Types with `Drop` implementations
- Types containing non-`Copy` fields

```

struct Point { x: i32, y: i32 }
impl Copy for Point {}

let p1 = Point { x: 10, y: 20 }
let p2 = p1 // Copies p1 instead of moving it
// p1 is still valid here!

```

Note

Primitive types (`i32`, `u32`, `bool`, etc.) and pointer types are automatically `Copy` without requiring explicit `impl`.

13.4.3 Clone — Explicit Deep Copy

The `Clone` trait provides explicit deep copying for types where copying may allocate:

```

pub trait Clone {
    fn clone(&self) -> Option<Self>
}

```

Key difference from `Copy`:

- `Copy` is implicit and infallible (bitwise)
- `Clone` is explicit and can fail (returns `Option`)

```

impl Clone for Vec<T> {
    fn clone(&self) -> Option<Self> {
        // with_capacity returns Result, convert to Option
        var new_vec = match Vec::with_capacity(self.len()) {
            Result::Ok(v) => v,
            Result::Err(_) => return Option::None
        }
        for i in 0..self.len {
            match new_vec.push(self.get(i).unwrap()) {
                Result::Ok(_) => {},
                Result::Err(_) => return Option::None
            }
        }
        return Option::Some(new_vec)
    }
}

// Usage
let original = Vec::new()
original.push(1).ok() // Ignore result for simple example
original.push(2)?

match original.clone() {
    Option::Some(copy) => {
        // copy is independent of original
    },
    Option::None => {
        // Allocation failed during clone
    }
}

```

13.5 Type Conversion Traits

These traits enable type conversions between different types.

13.5.1 From<T> — Infallible Conversion

The `From` trait enables automatic type conversions:

```
pub trait From<T> {
    fn convert(value: T) -> Self
}
```

Note

The method is named `convert` instead of `from` because `from` is a keyword in Novus (used in import statements).

Primary use case: automatic error conversions with the `?` operator.

```
impl From<DosError> for NovusError {
    fn convert(err: DosError) -> NovusError {
        return NovusError::Dos(err)
    }
}

// Now ? automatically converts DosError to NovusError
fn foo() -> Result<i32, NovusError> {
    let file = open_file("test.txt")? // DosError -> NovusError
    return Result::Ok(42)
}
```

13.5.2 Into<T> — Consuming Conversion

The complement to `From`:

```
pub trait Into<T> {
    fn into(consuming self) -> T
}
```

If you implement `From<A>` for `B`, you conceptually get `Into<B>` for `A`.

Tip

Prefer implementing `From` over `Into`. `From` is more flexible and provides better type inference.

13.5.3 TryFrom and TryInto — Fallible Conversion

For conversions that can fail:

```
pub trait TryFrom<T, E> {
    fn try_from(value: T) -> Result<Self, E>
}

pub trait TryInto<T, E> {
```



```
fn try_into(consuming self) -> Result<T, E>
}
```

```
// Example: parsing strings to numbers
impl TryFrom<Str, ParseError> for i32 {
    fn try_from(s: Str) -> Result<i32, ParseError> {
        return s.parse_i32()
    }
}

let num = i32::try_from("42")? // Result<i32, ParseError>
```

13.6 Comparison Traits

13.6.1 Eq — Equality

The Eq trait enables equality comparison with == and !=:

```
pub trait Eq {
    fn eq(&self, other: &Self) -> bool
}
```

Contract:

- Reflexive: `a.eq(&a)` must return `true`
- Symmetric: `a.eq(&b)` must equal `b.eq(&a)`
- Transitive: if `a.eq(&b)` and `b.eq(&c)`, then `a.eq(&c)`

```
impl Eq for Point {
    fn eq(&self, other: &Point) -> bool {
        return self.x == other.x && self.y == other.y
    }
}
```

13.6.2 Ord — Total Ordering

The Ord trait provides total ordering:

```
pub trait Ord {
    fn cmp(&self, other: &Self) -> i32
}
```

Return value semantics:

- Negative: `self < other`
- Zero: `self == other`
- Positive: `self > other`

```
impl Ord for Version {
    fn cmp(&self, other: &Version) -> i32 {
        if self.major != other.major {
            return (i32)self.major - (i32)other.major
        }
        if self.minor != other.minor {
            return (i32)self.minor - (i32)other.minor
        }
        return (i32)self.patch - (i32)other.patch
    }
}
```

13.6.3 PartialOrd — Comparison Operators

For types that support <, <=, >, >=:

```
pub trait PartialOrd {
    fn lt(&self, other: &Self) -> bool
    fn le(&self, other: &Self) -> bool
    fn gt(&self, other: &Self) -> bool
    fn ge(&self, other: &Self) -> bool
}
```

13.7 Hash Trait

The Hash trait enables types to be used as keys in hash-based collections:

```
pub trait Hash {
    fn hash(&self) -> u32
}
```

Contract:

- If two values are equal (via Eq), they **must** produce the same hash
- Different values *may* produce the same hash (collision)
- Hash should distribute values uniformly across the u32 range

```
impl Hash for Point {
    fn hash(&self) -> u32 {
        // Combine hashes using XOR and FNV-1a mixing constant
        var h: u32 = self.x.hash()
        h = h ^ (self.y.hash() * 2246822507u32)
        return h
    }
}
```

Warning

If you implement both Hash and Eq, maintain the invariant: `a.eq(&b)` implies `a.hash() == b.hash()`.

All primitive types (`u8`, `u16`, `u32`, `i8`, `i16`, `i32`, `bool`) have built-in `Hash` implementations using FNV-1a inspired mixing.

13.8 Operator Traits

These traits enable operator overloading for custom types.

13.8.1 Arithmetic Operators

```
pub trait Add { fn add(&self, other: &Self) -> Self }
pub trait Sub { fn sub(&self, other: &Self) -> Self }
pub trait Mul { fn mul(&self, other: &Self) -> Self }
pub trait Div { fn div(&self, other: &Self) -> Self }
pub trait Rem { fn rem(&self, other: &Self) -> Self }
pub trait Neg { fn neg(&self) -> Self }
```

```
impl Add for Point {
    fn add(&self, other: &Point) -> Point {
        return Point { x: self.x + other.x, y: self.y + other.y }
    }
}

let p1 = Point { x: 10, y: 20 }
let p2 = Point { x: 5, y: 15 }
let p3 = p1 + p2 // Point { x: 15, y: 35 }
```

13.8.2 Index Operators

```
pub trait Index<I, T> {
    fn index(&self, idx: I) -> T
}

pub trait IndexMut<I, T> {
    fn index_set(&var self, idx: I, value: T)
}
```

```
impl Index<u32, i32> for MyArray {
    fn index(&self, idx: u32) -> i32 {
        return self.data[idx]
    }
}

impl IndexMut<u32, i32> for MyArray {
    fn index_set(&var self, idx: u32, value: i32) {
        self.data[idx] = value
    }
}

let arr = MyArray::new()
let val = arr[0] // Calls index()
arr[0] = 42 // Calls index_set()
```

13.9 Iterator Traits

13.9.1 Iterable — For-In Loop Support

The `Iterable` trait enables types to be used in for-in loops:

```
pub trait Iterable<T> {
    fn get(&self, index: u32) -> Option<T>
    fn len(&self) -> u32
}

impl<T> Iterable<T> for Vec<T> {
    fn get(&self, index: u32) -> Option<T> { ... }
    fn len(&self) -> u32 { ... }
}

// For-in loop desugars to index-based iteration
for item in vec {
    process(item)
}
```

13.9.2 Iterator — Stateful Iteration

For more advanced iteration patterns:

```
pub trait Iterator<T> {
    fn next(&var self) -> Option<T>
}
```

Unlike `Iterable`, `Iterator` is stateful and consumed as it advances:

```
struct RangeIter {
    current: u32,
    end: u32,
}

impl Iterator<u32> for RangeIter {
    fn next(&var self) -> Option<u32> {
        if self.current >= self.end {
            return Option::None
        }
        let val = self.current
        self.current = self.current + 1
        return Option::Some(val)
    }
}

// Usage with while-let
var iter = string.split(",")
while let Option::Some(part) = iter.next() {
    process(part)
}
```

13.9.3 Collection Building Traits

```
pub trait Extend<T> {
    fn extend_from_ptr(&var self, source: *T, count: u32) -> bool
}

pub trait FromIterator<T> {
    fn from_slice(ptr: *T, count: u32) -> Option<Self>
}
```

13.10 Numeric Traits

13.10.1 Default — Default Values

```
pub trait Default {
    fn default() -> Self
}
```

Built-in defaults:

- Integers: 0
- bool: false
- Option<T>: None

```
impl Default for MyConfig {
    fn default() -> MyConfig {
        return MyConfig { timeout: 1000, retries: 3 }
    }
}

let config = MyConfig::default()
```

13.10.2 Zero and One — Identity Elements

```
pub trait Zero {
    fn zero() -> Self
    fn is_zero(&self) -> bool
}

pub trait One {
    fn one() -> Self
}
```

Used in generic numeric algorithms:

```
fn sum<T: Zero + Add>(values: &Vec<T>) -> T {
    var total = T::zero()
    for v in values {
        total = total.add(&v)
    }
}
```

```
    return total
}
```

13.10.3 Bounded — Min/Max Values

```
pub trait Bounded {
    fn min_value() -> Self
    fn max_value() -> Self
}
```

Type	min_value()	max_value()
u8	0	255
u16	0	65,535
u32	0	4,294,967,295
i8	-128	127
i16	-32,768	32,767
i32	-2,147,483,648	2,147,483,647

13.10.4 Primitive Trait Implementations Summary

The following table shows which traits are implemented for primitive types:

Type	Hash	Eq	Ord	Default	Zero	One	Bounded
u8, u16, u32	•	•	•	•	•	•	•
i8, i16, i32	•	•	•	•	•	•	•
bool	•	•	•	•	—	—	—
Option<T>	—	—	—	•	—	—	—
LibraryVersion	—	—	•	—	—	—	—

All primitive integer and boolean types have complete trait coverage. The Copy trait is implicit for all primitives and pointer types.

13.11 View Conversion Traits

13.11.1 AsRef and AsMut

Zero-cost reference-to-reference conversions:

```
pub trait AsRef<T> {
    fn as_ref(&self) -> T
}

pub trait AsMut<T> {
    fn as_mut(&var self) -> T
}
```

```
// String can be viewed as Str
impl AsRef<Str> for String {
    fn as_ref(&self) -> Str {
        return self.as_str()
    }
}
```

```

}

// Generic function accepting any string-like type
fn process<S: AsRef<Str>>(s: &S) {
    let str_view: Str = s.as_ref()
    // Works with String, PathString, etc.
}

```

13.11.2 Write — Unified Output

```

pub trait Write {
    fn write_bytes(&var self, ptr: *u8, len: u32) -> bool
    fn write_byte(&var self, b: u8) -> bool
}

```

Implemented by `StringBuilder`, `String`, formatters, and (future) file handles.

13.12 Error Trait

```

pub trait Error {
    fn message(&self) -> *u8
}

```

All stdlib error types implement this trait:

```

impl Error for DosError {
    fn message(&self) -> *u8 {
        match *self {
            DosError::NotFound => return "File not found",
            DosError::DiskFull => return "Disk full",
            _ => return "Unknown error"
        }
    }
}

// Usage
fn handle_error(err: &dyn Error) {
    println(err.message())
}

```

13.13 Function Traits

These traits represent callable types (for future closure support):

```

pub trait FnOnce<Args, Output> {
    fn call_once(consuming self, args: Args) -> Output
}

pub trait FnMut<Args, Output> {
    fn call_mut(&var self, args: Args) -> Output
}

```

```
pub trait Fn<Args, Output> {
    fn call(&self, args: Args) -> Output
}
```

Hierarchy:

- **Fn** — can be called repeatedly without mutation
- **FnMut** — can be called repeatedly with mutation
- **FnOnce** — can be called at least once (consumes self)

Note

Closures are not yet implemented in Novus. These traits are defined for future compatibility.

13.14 Summary

The `std::core` module provides:

Fundamental Types `Option<T>`, `Result<T, E>`, `LibraryVersion`

Resource Management `Drop` (automatic cleanup), `Copy` (bitwise copy), `Clone` (explicit deep copy)

Type Conversion `From/Into` (infallible), `TryFrom/TryInto` (fallible)

Comparison `Eq` (equality), `Ord` (total ordering), `PartialOrd` (operators), `Hash`

Operators `Add`, `Sub`, `Mul`, `Div`, `Rem`, `Neg`, `Index`, `IndexMut`

Iteration `Iterable` (for-in), `Iterator` (stateful), `Extend`, `FromIterator`

Numeric `Default`, `Zero`, `One`, `Bounded`

Views `AsRef`, `AsMut`, `Write`

Error Handling `Error`

Functions `Fn`, `FnMut`, `FnOnce` (future)

Tip

Import commonly used types from the prelude:

```
from std::prelude import *
// Imports: Option, Result, Drop, From, Str, and error types
```


Chapter 14

Collections

The `std::collections` module provides a comprehensive set of data structures optimized for Amiga systems. Each collection is designed with the constraints and opportunities of 68k processors in mind: deterministic memory management, cache-friendly layouts, and zero hidden allocations.

14.1 Overview

Novus collections fall into several categories based on their use cases:

Category	Collections	Use Case
Sequential	<code>Vec</code> , <code>VecDeque</code> , <code>SmallVec</code>	Ordered elements
Associative	<code>HashMap</code> , <code>HashSet</code>	Key-value lookup
Bit-level	<code>BitSet</code>	Compact boolean storage
Streaming	<code>RingBuffer</code> , <code>RingBufferPow2</code>	Fixed-size queues
Intrusive	<code>IntrusiveList</code> , <code>PriorityList</code>	AmigaOS interop
Arena	<code>SlotMap</code> , <code>ObjectPool</code>	Stable handles

Table 14.1: Collection categories and their primary use cases

14.1.1 Choosing the Right Collection

- **Need a growable array?** Use `Vec<T>`
- **Need a queue or deque?** Use `VecDeque<T>`
- **Know the typical size?** Use `SmallVec<T, N>`
- **Need key-value storage?** Use `HashMap<K, V>`
- **Need unique values?** Use `HashSet<T>`
- **Need compact flags?** Use `BitSet`
- **Need a fixed-size buffer?** Use `RingBuffer<T>`
- **Interfacing with AmigaOS lists?** Use `IntrusiveList<T>`
- **Need stable handles to objects?** Use `SlotMap<T>`
- **Need a pool allocator?** Use `ObjectPool<T>`

14.2 Vec<T> — Dynamic Array

`Vec<T>` is a contiguous, growable array type with heap-allocated contents. It's the most commonly used collection in Novus, similar to Rust's `Vec` or C++'s `std::vector`.

```
from std::collections::vec import Vec

fn example() -> Result<(), ExecError> {
    // Create an empty vector
    var numbers = Vec::<u32>::new()

    // Push elements (may allocate)
    numbers.push(1)?
    numbers.push(2)?
    numbers.push(3)?

    // Access by index
    match numbers.get(0) {
        Option::Some(val) => print_u32(val),
        Option::None => {}
    }

    // Pop from the end
    match numbers.pop() {
        Option::Some(val) => print_u32(val), // 3
        Option::None => {}
    }

    return Result::Ok(())
}
// Memory automatically freed when numbers goes out of scope
```

14.2.1 Creating Vectors

```
// Empty vector (zero capacity, no allocation)
var v1 = Vec::<u32>::new()

// Pre-allocated capacity (reduces reallocations)
var v2 = Vec::<u32>::with_capacity(100)?
```

Note

`Vec::new()` performs no allocation. Memory is only allocated on the first `push()` or `reserve()` call.

14.2.2 Core Operations

14.2.3 Growth Strategy

When a `Vec` needs more capacity, it doubles its current capacity (minimum 4 elements). This amortized doubling strategy provides $O(1)$ push operations on average while minimizing reallocations.

```
var v = Vec::<u32>::new()
```

Method	Returns	Complexity	Description
push(val)	Result<(), ExecError>	O(1)*	Append to end
pop()	Option<T>	O(1)	Remove from end
get(idx)	Option<T>	O(1)	Read element
set(idx, val)	Result<(), ExecError>	O(1)	Write element
insert(idx, val)	Result<(), ExecError>	O(n)	Insert at position
remove(idx)	Result<T, ExecError>	O(n)	Remove at position
len()	u32	O(1)	Element count
capacity()	u32	O(1)	Current capacity
is_empty()	bool	O(1)	Check if empty
clear()	()	O(1)	Remove all elements
reserve(n)	Result<(), ExecError>	O(n)	Ensure capacity
shrink_to_fit()	Result<(), ExecError>	O(n)	Reduce capacity

Table 14.2: Vec<T> method summary. *Amortized O(1), O(n) when growing.

```
// capacity: 0

v.push(1)? // allocates 4 elements
// capacity: 4

v.push(2)?
v.push(3)?
v.push(4)?
v.push(5)? // reallocates to 8 elements
// capacity: 8
```

Coming from C

Unlike C's `realloc()`, Novus always allocates new memory and copies data. This is because AmigaOS's `AllocMem/FreeMem` don't support in-place growth. Pre-allocate with `with_capacity()` when you know the expected size.

14.2.4 Iteration

Vec<T> implements the `Iterable<T>` trait, enabling use in `for-in` loops:

```
var numbers = Vec::<u32>::new()
numbers.push(10)?
numbers.push(20)?
numbers.push(30)?

for num in numbers {
    print_u32(num)
}
// Prints: 10 20 30
```

14.2.5 Memory and RAII

Vec<T> implements `Drop`, automatically freeing its memory when it goes out of scope:

```
fn process_data() -> Result<(), ExecError> {
    var data = Vec::<u32>::with_capacity(1000)?
```

```

    // ... use data ...

    return Result::Ok(())
} // data.drop() called automatically, memory freed

```

14.3 VecDeque<T> — Double-Ended Queue

`VecDeque<T>` is a double-ended queue implemented as a ring buffer. It provides $O(1)$ insertion and removal at both ends, making it ideal for queues, sliding windows, and work-stealing algorithms.

```

from std::collections::vecdeque import VecDeque

fn message_queue() -> Result<(), ExecError> {
    var queue = VecDeque::<Message>::with_capacity(32)?

    // Enqueue (add to back)
    queue.push_back(msg1)?
    queue.push_back(msg2)?

    // Urgent message to front
    queue.push_front(urgent_msg)?

    // Dequeue (remove from front)
    match queue.pop_front() {
        Option::Some(msg) => process(msg),
        Option::None => {} // empty
    }

    return Result::Ok(())
}

```

14.3.1 Ring Buffer Implementation

Internally, `VecDeque` uses a circular buffer with `head` and `len` indices:

```

// Logical view:      [A, B, C, D]
// Physical memory:  [C, D, _, _, A, B] (head=4, len=4)
//                  ^-- unused space

```

This allows both `push_front` and `push_back` to be $O(1)$ operations without shifting elements.

14.3.2 Core Operations

14.3.3 Use Cases

- **FIFO queues:** `push_back()` + `pop_front()`
- **LIFO stacks:** `push_back()` + `pop_back()`
- **Sliding windows:** Maintain a fixed-size window of recent values
- **Priority bypass:** Insert urgent items at front

Method	Returns	Description
<code>push_back(val)</code>	<code>Result<(), ExecError></code>	Add to back (enqueue)
<code>push_front(val)</code>	<code>Result<(), ExecError></code>	Add to front
<code>pop_back()</code>	<code>Option<T></code>	Remove from back
<code>pop_front()</code>	<code>Option<T></code>	Remove from front (dequeue)
<code>front()</code>	<code>Option<T></code>	Peek at front
<code>back()</code>	<code>Option<T></code>	Peek at back
<code>get(idx)</code>	<code>Option<T></code>	Access by logical index
<code>rotate_left(n)</code>	<code>()</code>	Rotate elements left
<code>rotate_right(n)</code>	<code>()</code>	Rotate elements right

Table 14.3: `VecDeque<T>` method summary

14.4 SmallVec<T, N> — Small Vector Optimization

`SmallVec<T, N>` is a vector that uses the `const` generic `N` as its initial capacity. This reduces allocation overhead when you know a typical size.

```
from std::collections::smallvec import SmallVec

// Create a SmallVec with initial capacity of 8
var coords = SmallVec::<Point, 8>::new()

// First push allocates N=8 capacity
coords.push(Point { x: 10, y: 20 })?

// Subsequent pushes use existing capacity
coords.push(Point { x: 30, y: 40 })?
```

Note

The current implementation allocates `N` elements on first push. A future version may store small arrays inline (true small-vector optimization) once Novus supports array types with `const` generic sizes.

14.4.1 When to Use SmallVec

Use `SmallVec` when:

- You know the typical size of your collection
- Most instances stay below a certain threshold
- You want to avoid the overhead of dynamic resizing

```
// A polygon usually has 3-8 vertices
var vertices = SmallVec::<Vertex, 8>::new()

// A function usually has 0-4 parameters
var params = SmallVec::<Param, 4>::new()
```

14.4.2 Core Operations

`SmallVec` provides the same interface as `Vec`:

Method	Returns	Description
<code>push(val)</code>	<code>Result<(), ExecError></code>	Append element
<code>pop()</code>	<code>Option<T></code>	Remove last element
<code>get(idx)</code>	<code>Option<T></code>	Read element by index
<code>get_mut(idx)</code>	<code>Option<*T></code>	Get mutable pointer
<code>len()</code>	<code>u32</code>	Element count
<code>capacity()</code>	<code>u32</code>	Current capacity
<code>is_empty()</code>	<code>bool</code>	Check if empty
<code>clear()</code>	<code>()</code>	Remove all elements

Table 14.4: `SmallVec<T, N>` method summary

14.5 `HashMap<K, V>` — Hash Map

`HashMap<K, V>` is a hash table with generic keys using open addressing and linear probing. Keys must implement the `Hash` and `Eq` traits.

```
from std::collections::hashmap import HashMap

fn symbol_table() -> Result<(), ExecError> {
    var symbols: HashMap<u32, u32> = HashMap::new()?

    // Insert key-value pairs
    symbols.insert(hash("foo"), 100)?
    symbols.insert(hash("bar"), 200)?

    // Lookup by key
    match symbols.get(&hash("foo")) {
        Option::Some(val_ptr) => print_u32(*val_ptr),
        Option::None => print("not found")
    }

    // Check existence
    if symbols.contains_key(&hash("foo")) {
        // key exists
    }

    // Remove
    match symbols.remove(&hash("bar")) {
        Option::Some(old_val) => {}, // was 200
        Option::None => {}
    }

    return Result::Ok(())
}
```

14.5.1 Implementation Details

`HashMap` uses several optimizations for 68k performance:

- **Open addressing:** All entries in a single array (better cache locality)
- **Linear probing:** Simple collision resolution
- **Tombstones:** Lazy deletion for correct probe sequences

- **Power-of-2 capacity:** Fast modulo via bitwise AND
- **70% load factor:** Automatic resize threshold
- **Cached entry size:** Avoids expensive multiplication

14.5.2 Core Operations

Method	Returns	Description
<code>insert(k, v)</code>	<code>Result<Option<V>, ExecError></code>	Insert/update
<code>get(&k)</code>	<code>Option<&V></code>	Lookup by key
<code>get_mut(&k)</code>	<code>Option<&var V></code>	Mutable lookup
<code>remove(&k)</code>	<code>Option<V></code>	Remove by key
<code>contains_key(&k)</code>	<code>bool</code>	Check existence
<code>len()</code>	<code>u32</code>	Number of pairs
<code>is_empty()</code>	<code>bool</code>	Check if empty
<code>clear()</code>	<code>()</code>	Remove all entries
<code>reserve(n)</code>	<code>Result<(), ExecError></code>	Reserve capacity
<code>compact()</code>	<code>Result<(), ExecError></code>	Clear tombstones
<code>iter()</code>	<code>HashMapIter<K, V></code>	Iterator

Table 14.5: HashMap<K, V> method summary

14.5.3 Iteration

The iterator returns `KeyValueRef` structs containing *pointers* to the keys and values:

```
// KeyValueRef structure (returned by iter().next())
struct KeyValueRef<K, V> {
    key: *K,      // Pointer to key
    value: *V,    // Pointer to value
}
```

Use pointer dereferencing (*) to access the actual values:

```
var map: HashMap<u32, u32> = HashMap::new()?
map.insert(1u32, 100u32)?
map.insert(2u32, 200u32)?

var iter = map.iter()
while true {
    match iter.next() {
        Option::Some(kv) => {
            // Note: kv.key and kv.value are pointers
            print_u32(*kv.key)    // Dereference to get key
            print_u32(*kv.value)  // Dereference to get value
        },
        Option::None => break
    }
}
```

Warning

Do not modify the map while iterating. This invalidates the iterator and may cause undefined behavior.

14.5.4 Key Requirements

Keys must implement `Hash` and `Eq`:

```
struct Point {
    x: i32,
    y: i32,
}

impl Hash for Point {
    fn hash(&self) -> u32 {
        var h: u32 = self.x.hash()
        h = h ^ (self.y.hash() * 2246822507u32)
        return h
    }
}

impl Eq for Point {
    fn eq(&self, other: &Point) -> bool {
        return self.x == other.x && self.y == other.y
    }
}

// Now Point can be used as a HashMap key
var point_map: HashMap<Point, u32> = HashMap::new()?
```

14.6 HashSet<T> — Hash Set

`HashSet<T>` is a set of unique values backed by `HashMap<T, ()>`. It provides $O(1)$ insert, remove, and contains operations.

```
from std::collections::hashset import HashSet

fn unique_ids() -> Result<(), ExecError> {
    var seen: HashSet<u32> = HashSet::new()?

    // Insert returns Ok(true) for new, Ok(false) for existing
    seen.insert(1u32)? // Ok(true)
    seen.insert(2u32)? // Ok(true)
    seen.insert(1u32)? // Ok(false) - already present

    // Check membership
    if seen.contains(&1u32) {
        // value is in set
    }

    // Remove
    seen.remove(&2u32) // returns true

    return Result::Ok(())
}
```



```
}
```

14.6.1 Set Operations

```
var a: HashSet<u32> = HashSet::new()?
var b: HashSet<u32> = HashSet::new()?

// Check subset relationship
if a.is_subset(&b) {
    // every element in a is also in b
}

// Check if sets have common elements
if a.intersects(&b) {
    // at least one element in common
}
```

14.7 BitSet — Compact Bit Storage

BitSet provides compact storage for boolean flags using 1 bit per element. This is 32x more memory-efficient than a bool array, essential for resource-constrained Amiga development.

```
from std::collections::bitset import BitSet

fn track_sprites() -> Result<(), ExecError> {
    // 64 bits = 8 bytes (vs 64 bytes for bool array)
    match BitSet::with_capacity(64) {
        Option::Some(var active_sprites) => {
            // Set bits
            active_sprites.set(0)    // sprite 0 is active
            active_sprites.set(5)    // sprite 5 is active

            // Check bits
            if active_sprites.get(0) {
                // sprite 0 is active
            }

            // Count active sprites
            let count = active_sprites.count_ones() // 2

            // Find first active sprite
            match active_sprites.find_first_set() {
                Option::Some(idx) => {}, // idx == 0
                Option::None => {}
            }
        },
        Option::None => {}
    }

    return Result::Ok(())
}
```

14.7.1 Bitwise Operations

```
var a = BitSet::with_capacity(32)?
var b = BitSet::with_capacity(32)?

// Set some bits
a.set(0)
a.set(2)
b.set(1)
b.set(2)

// Bitwise operations (mutate a)
a.and(&b)    // a = a & b (intersection)
a.or(&b)     // a = a | b (union)
a.xor(&b)    // a = a ^ b (symmetric difference)
a.not()     // a = ~a (complement)
```

14.7.2 Core Operations

Method	Returns	Description
set(idx)	bool	Set bit to true
clear(idx)	bool	Set bit to false
get(idx)	bool	Get bit value
toggle(idx)	bool	Flip bit
set_all()	()	Set all bits
clear_all()	()	Clear all bits
count_ones()	u32	Population count
count_zeros()	u32	Count cleared bits
find_first_set()	Option<u32>	First set bit
find_first_clear()	Option<u32>	First clear bit
all()	bool	All bits set?
any()	bool	Any bit set?
none()	bool	No bits set?

Table 14.6: BitSet method summary

Coming from C

BitSet uses Brian Kernighan’s algorithm for population count and binary search for first-set-bit. On 68020+, these could be replaced with native BFFFO (find first one) instructions for better performance.

14.8 RingBuffer<T> — Fixed-Size Circular Buffer

RingBuffer<T> is a fixed-capacity circular buffer ideal for streaming data, audio buffering, and interrupt-safe producer-consumer patterns.

```
from std::collections::ringbuffer import RingBuffer,
    RingBufferPow2

fn audio_buffer() {
    // Create a 256-sample buffer
```

```

match RingBuffer::<i16>::with_capacity(256) {
  Option::Some(var buffer) => {
    // Producer (e.g., interrupt handler)
    buffer.push_back(sample) // returns false if full

    // Consumer (e.g., main task)
    match buffer.pop_front() {
      Option::Some(sample) => play(sample),
      Option::None => {} // buffer empty
    }
  },
  Option::None => {}
}

```

14.8.1 Power-of-2 Optimization

For maximum performance, use `RingBufferPow2<T>` which uses bitwise AND instead of modulo for index wrapping:

```

// Request 100, get 128 (next power of 2)
match RingBufferPow2::<Sample>::with_capacity(100) {
  Option::Some(var buffer) => {
    // buffer.capacity() == 128
    // ~6x faster push/pop on 68000
  },
  Option::None => {}
}

```

Tip

Performance on 68000:

- `RingBuffer::push_back`: ~150 cycles (uses modulo)
- `RingBufferPow2::push_back`: ~25 cycles (uses bitmask)

That's approximately 6x faster! Always prefer `RingBufferPow2` for performance-critical code like audio streaming and interrupt handlers.

14.8.2 Interrupt Safety

Ring buffers are safe for single-producer/single-consumer patterns:

- **Producer:** Interrupt handler writes via `push_back()`
- **Consumer:** Main task reads via `pop_front()`

Warning

Ring buffers are NOT safe for multiple producers or multiple consumers without external synchronization (`Exec Forbid/Permit` or semaphores).

14.8.3 Amiga Use Cases

- Paula audio sample streaming (16-bit PCM at 28kHz)
- Serial port buffering (up to 115200 baud)
- Vertical blank event queues (50Hz PAL)
- Trackdisk sector buffering
- Interrupt-to-task message passing

14.9 IntrusiveList<T> — AmigaOS-Compatible Lists

`IntrusiveList<T>` is an intrusive doubly-linked list that matches AmigaOS Exec's `MinList/MinNode` pattern exactly. It uses native Exec list functions for maximum compatibility.

```
from std::collections::intrusive_list import IntrusiveList,
    init_min_node
from std::ffi::amiga_structs import MinNode

struct MyData {
    node: MinNode, // MUST be first field
    value: u32,
}

fn use_list() {
    var list = IntrusiveList::<MyData>::new()
    list.init() // MUST call after assignment

    // Create node (caller manages lifetime)
    var item = MyData {
        node: MinNode { mln_Succ: null, mln_Pred: null },
        value: 42
    }

    // Add to list
    list.add_tail(&item)

    // Iterate
    match list.first() {
        Option::Some(ptr) => {
            var current = ptr
            while true {
                print_u32((*current).value)
                match list.next(current) {
                    Option::Some(next) => current = next,
                    Option::None => break
                }
            }
        },
        Option::None => {}
    }

    // Remove before freeing
    list.remove(&item)
}
```

Warning

You **MUST** call `init()` immediately after creating an `IntrusiveList`. The list contains self-referential pointers that are invalid until `init()` is called. Using an uninitialized list will corrupt memory or crash.

14.9.1 Key Constraints

1. **Node must be first field:** The `MinNode` must be the first field of your struct
2. **Call `init()` after assignment:** The list has self-referential pointers that must be set up after the list is at its final location
3. **Caller manages node lifetime:** The list does NOT own the nodes
4. **One list per node:** A node can only be in one list at a time

14.9.2 PriorityList<T>

`PriorityList<T>` is similar but uses full `Node` structures with priority support:

```
from std::collections::intrusive_list import PriorityList,
    init_node
from std::ffi::amiga_structs import Node

struct Task {
    node: Node, // Full node with priority
    data: u32,
}

fn priority_queue() {
    var tasks = PriorityList::<Task>::new()
    tasks.init()

    var high = Task { node: @zeroed(Node), data: 1 }
    var low = Task { node: @zeroed(Node), data: 2 }

    // Set priorities
    init_node(&high.node, 10, null) // priority 10
    init_node(&low.node, 1, null)    // priority 1

    // Insert in priority order (higher = closer to head)
    tasks.enqueue(&high)
    tasks.enqueue(&low)

    // First will be the high-priority task
    match tasks.first() {
        Option::Some(ptr) => {
            // (*ptr).data == 1 (high priority)
        },
        Option::None => {}
    }
}
```

14.10 SlotMap<T> — Generational Arena

SlotMap<T> is a generational index arena that provides stable handles to objects. It detects use-after-free by tracking a generation counter for each slot.

```
from std::collections::slotmap import SlotMap, SlotKey

struct Entity {
    x: i32,
    y: i32,
    health: u32,
}

fn game_entities() {
    match SlotMap::<Entity>::new() {
        Option::Some(var entities) => {
            // Insert returns a SlotKey handle
            match entities.insert(Entity { x: 100, y: 50, health:
100 }) {
                Option::Some(player_key) => {
                    // Access via handle
                    match entities.get(player_key) {
                        Option::Some(entity) => {
                            print_i32((*entity).x)
                        },
                        Option::None => {}
                    }
                }

                // Remove the entity
                entities.remove(player_key)

                // Old key is now invalid!
                match entities.get(player_key) {
                    Option::Some(_) => {}, // Never reached
                    Option::None => {}      // Key is stale
                }
            },
            Option::None => {}
        }
    }
}
```

14.10.1 How Generations Work

Each slot has a generation counter. When an object is removed, the generation is incremented. Old keys have the old generation and fail the check:

```
// Slot 0, Generation 0 -> player inserted
let player_key = SlotKey { index: 0, generation: 0 }

// Slot 0, Generation 1 -> player removed
entities.remove(player_key)

// Slot 0 reused for enemy
let enemy_key = SlotKey { index: 0, generation: 1 }
```

```
// Old player_key fails - generation mismatch!
entities.get(player_key) // Returns None
```

14.10.2 SlotKey

SlotKey contains both index and generation:

```
struct SlotKey {
    index: u32,
    generation: u32,
}

// Pack into single u32 (16 bits each)
let packed: u32 = key.pack()

// Unpack
let key = SlotKey::unpack(packed)

// Sentinel/null key
let null_key = SlotKey::sentinel()
if key.is_sentinel() { /* invalid */ }
```

14.10.3 Use Cases

- **ECS game entities:** Players, enemies, projectiles, pickups
- **Resource handles:** Textures, sounds, fonts
- **Window/Screen handles:** Multiple windows with stable IDs
- **Network object sync:** Stable references across frames
- **Undo/redo systems:** References that survive modifications

14.10.4 SlotMapFixed<T>

SlotMapFixed<T> is a variant with fixed capacity (no growth). It uses separate arrays for generations, occupancy flags, and values, providing better cache behavior during iteration:

```
from std::collections::slotmap import SlotMapFixed

// Fixed capacity - never grows, but more cache-friendly
match SlotMapFixed::<Entity>::with_capacity(1024) {
    Option::Some(var entities) => {
        // Same API as SlotMap
        match entities.insert(player) {
            Option::Some(key) => { /* ... */ },
            Option::None => { /* capacity full */ }
        }
    },
    Option::None => {}
}
```

Note

Use `SlotMap<T>` when you need dynamic growth. Use `SlotMapFixed<T>` when you know the maximum count upfront and want better iteration performance.

14.11 ObjectPool<T> — Pool Allocator

`ObjectPool<T>` is a slab allocator for fixed-size objects with $O(1)$ allocation and deallocation. After initialization, no system calls are needed.

```
from std::collections::freelist import ObjectPool

struct Bullet {
    x: i32,
    y: i32,
    dx: i32,
    dy: i32,
}

fn bullet_system() {
    // Pre-allocate pool of 128 bullets
    match ObjectPool::<Bullet>::with_capacity(128) {
        Option::Some(var pool) => {
            // Allocate a bullet (O(1), no syscall)
            match pool.alloc() {
                Option::Some(bullet) => {
                    (*bullet).x = 100
                    (*bullet).y = 50
                    (*bullet).dx = 5
                    (*bullet).dy = 0

                    // ... use bullet ...

                    // Return to pool (O(1))
                    pool.free(bullet)
                },
                Option::None => {} // pool exhausted
            }
        },
        Option::None => {}
    }
}
```

14.11.1 Interrupt Safety

`ObjectPool` is safe for single-producer/single-consumer patterns:

- **Producer:** Interrupt handler allocates via `alloc()`
- **Consumer:** Main task frees via `free()`

Warning**Safety requirements:**

- Do NOT free objects not allocated from this pool

- Do NOT double-free objects
- Do NOT use objects after freeing them

14.11.2 ObjectPoolFixed<T>

`ObjectPoolFixed<T>` uses an index-stack approach instead of a linked free list. This provides additional capabilities:

- **Allocation tracking:** Knows exactly which slots are in use
- **Proper Drop support:** Calls `Drop` on all allocated objects when pool is dropped or reset
- **Index-based iteration:** Can iterate over allocated objects by index
- **Free validation:** Detects double-free and invalid free operations

```
from std::collections::freelist import ObjectPoolFixed

match ObjectPoolFixed::<Enemy>::with_capacity(32) {
    Option::Some(var pool) => {
        // Allocate
        match pool.alloc() {
            Option::Some(enemy) => {
                // Initialize...
            },
            Option::None => {}
        }

        // Iterate over all allocated enemies
        var i: u32 = 0
        while i < pool.capacity() {
            if pool.is_allocated(i) {
                match pool.get(i) {
                    Option::Some(enemy) => update(enemy),
                    Option::None => {}
                }
            }
            i = i + 1
        }

        // Reset calls Drop on all allocated objects
        pool.reset()
    },
    Option::None => {}
}
```

Note

Use `ObjectPool<T>` for maximum speed when you don't need iteration or `Drop` support. Use `ObjectPoolFixed<T>` when you need to iterate over allocated objects or when `T` owns resources that must be cleaned up.

14.12 Performance Summary

Collection	Insert	Remove	Access	Search	Memory
<code>Vec<T></code>	$O(1)^*$	$O(n)$	$O(1)$	$O(n)$	Contiguous
<code>VecDeque<T></code>	$O(1)$	$O(1)^{**}$	$O(1)$	$O(n)$	Ring buffer
<code>HashMap<K, V></code>	$O(1)^*$	$O(1)$	$O(1)$	$O(1)$	Hash table
<code>HashSet<T></code>	$O(1)^*$	$O(1)$	—	$O(1)$	Hash table
<code>BitSet</code>	$O(1)$	$O(1)$	$O(1)$	$O(n)$	1 bit/element
<code>RingBuffer<T></code>	$O(1)$	$O(1)$	$O(1)$	$O(n)$	Fixed size
<code>IntrusiveList<T></code>	$O(1)$	$O(1)$	$O(n)$	$O(n)$	No allocation
<code>SlotMap<T></code>	$O(1)^*$	$O(1)$	$O(1)$	—	Arena
<code>ObjectPool<T></code>	$O(1)$	$O(1)$	—	—	Fixed size

Table 14.7: Collection complexity summary. *Amortized. **Front/back only.

Coming from C

Unlike C's standard library, all Novus collections:

- Use AmigaOS `AllocMem/FreeMem` for memory management
- Are properly aligned for 68k DMA and cache efficiency
- Implement RAII via the `Drop` trait
- Return `Result` or `Option` instead of null pointers

Chapter 15

Strings

The `std::strings` module provides a comprehensive string handling system designed for Amiga development. It offers both borrowed views (zero-copy) and owned strings (heap-allocated), plus AmigaDOS-specific types like BCPL strings.

15.1 Overview

Novus provides multiple string types for different use cases:

Type	Size	Allocation	Use Case
<code>Str</code>	8 bytes	None (view)	Parameters, slices
<code>String</code>	12 bytes	Heap (Vec)	Building, ownership
<code>NameString</code>	36 bytes	Stack (32 chars)	Short names
<code>GadgetString</code>	132 bytes	Stack (128 chars)	UI text buffers
<code>PathString</code>	260 bytes	Stack (256 chars)	File paths
<code>StringBuilder</code>	12 bytes	Heap	Efficient concatenation
<code>BStr</code>	4 bytes	Heap (BCPL)	AmigaDOS interop

Table 15.1: String types and their characteristics

Note

All Novus strings are Latin-1 (ISO-8859-1) encoded, matching AmigaOS conventions. ASCII is a subset of Latin-1, so ASCII text works naturally. UTF-8 support may be added in future versions.

15.2 Str — Borrowed String Slice

`Str` is a “fat pointer” containing a pointer and length. It does NOT own the data—it’s just a view into existing memory. This makes it extremely efficient for passing strings to functions.

```
from std::strings::core import Str

fn greet(name: Str) {
    // name is just a view - no allocation, no copy
    print_str(name)
}

fn main() {
    // Create Str from C string literal
    match Str::from_cstr("Hello, World!") {
        Option::Some(s) => greet(s),
        Option::None => {}
    }
}
```

```
}

```

15.2.1 Creating Str Values

```
// From null-terminated C string (scans for length)
let s1 = Str::from_cstr("Hello")?

// From raw pointer and known length (unsafe - must be valid)
let s2 = Str::from_raw(ptr, 5)

```

Warning

`Str::from_raw()` does not validate the pointer. The caller must ensure `ptr` points to valid memory of at least `len` bytes.

15.2.2 Basic Operations

Method	Returns	Description
<code>len()</code>	<code>u32</code>	Length in bytes
<code>is_empty()</code>	<code>bool</code>	Check if length is zero
<code>as_ptr()</code>	<code>*u8</code>	Get raw pointer
<code>get(idx)</code>	<code>Option<u8></code>	Get byte at index
<code>as_cstr()</code>	<code>*u8</code>	Get C string pointer

Table 15.2: `Str` basic methods

15.2.3 Slicing

Create substrings without copying:

```
let text = Str::from_cstr("Hello, World!")?

// Slice with start and end indices
let hello = text.slice(0, 5)?    // "Hello"
let world = text.slice(7, 12)?   // "World"

// Slice from index to end
let rest = text.slice_from(7)?   // "World!"

// Slice from start to index
let prefix = text.slice_to(5)?   // "Hello"

```

Note

Slicing returns `Result<Str, StringError>` with bounds checking. The slice is a view into the original data—no allocation occurs.

15.2.4 Comparison

```

let a = Str::from_cstr("hello")?
let b = Str::from_cstr("HELLO")?
let c = Str::from_cstr("hello")?

// Exact equality
if a.equals(c) {
    // true - same bytes
}

// Case-insensitive equality
if a.eq_ignore_case(b) {
    // true - ignores ASCII case
}

// Lexicographic comparison (returns i32: <0, 0, or >0)
let cmp = a.compare(b) // negative (lowercase < uppercase in ASCII)

```

Coming from C

On 68020+, `equals()` uses optimized inline assembly with `cmpm.1` for ~4x speedup on strings ≥ 8 bytes. The compiler automatically uses byte-by-byte comparison on 68000/68010.

15.2.5 Searching

```

let text = Str::from_cstr("Hello, World!")?

// Check prefix/suffix
if text.starts_with(Str::from_cstr("Hello")?) {
    // true
}
if text.ends_with(Str::from_cstr("!")?) {
    // true
}

// Case-insensitive prefix check
if text.starts_with_ignore_case(Str::from_cstr("HELLO")?) {
    // true
}

// Find substring
match text.find(Str::from_cstr("World")?) {
    Option::Some(idx) => {}, // idx == 7
    Option::None => {}
}

// Check if contains substring
if text.contains(Str::from_cstr(",")?) {
    // true
}

// Find single byte
match text.find_byte($2C) { // ', '
    Option::Some(idx) => {}, // idx == 5

```

```
    Option::None => {}  
}  
  
// Find last occurrence of byte  
match text.rfind_byte($6F) { // 'o'  
    Option::Some(idx) => {}, // idx == 8 (second 'o')  
    Option::None => {}  
}
```

15.2.6 Trimming

```
let padded = Str::from_cstr("  hello  ")?  
  
let trimmed = padded.trim()           // "hello"  
let left = padded.trim_start()        // "hello "  
let right = padded.trim_end()         // "  hello"
```

Trimming removes ASCII whitespace: space (\$20), tab (\$09), LF (\$0A), and CR (\$0D).

15.3 String — Owned Heap String

`String` is a heap-allocated, growable, mutable string backed by `Vec<u8>`. Use it when you need to build strings dynamically or own string data.

```
from std::strings::core import String, Str  
  
fn build_greeting(name: Str) -> Result<String, StringError> {  
    var greeting = String::new()  
    greeting.push_str(Str::from_cstr("Hello, ")?)?  
    greeting.push_str(name)?  
    greeting.push_byte($21)? // '!'  
    return Result::Ok(greeting)  
}
```

15.3.1 Creating Strings

```
// Empty string (no allocation until first push)  
var s1 = String::new()  
  
// Pre-allocated capacity  
var s2 = String::with_capacity(256)?  
  
// From Str (allocates and copies)  
let s3 = String::new_from_str(Str::from_cstr("Hello"))??
```

15.3.2 Core Operations

15.3.3 Memory Management

`String` implements the `Drop` trait for automatic cleanup:

Method	Returns	Description
<code>push_str(s)</code>	<code>Result<(), StringError></code>	Append Str
<code>push_byte(b)</code>	<code>Result<(), StringError></code>	Append single byte
<code>len()</code>	<code>u32</code>	Current length
<code>capacity()</code>	<code>u32</code>	Current capacity
<code>is_empty()</code>	<code>bool</code>	Check if empty
<code>clear()</code>	<code>()</code>	Clear contents (keep capacity)
<code>reserve(n)</code>	<code>Result<(), StringError></code>	Reserve additional bytes
<code>as_str()</code>	<code>Str</code>	Borrow as Str slice
<code>as_ptr()</code>	<code>*u8</code>	Get raw pointer
<code>as_mut_ptr()</code>	<code>*u8</code>	Get mutable pointer

Table 15.3: String method summary

```
fn process() -> Result<(), StringError> {
    var message = String::with_capacity(100)?
    message.push_str(Str::from_cstr("Processing...")??)

    // ... use message ...

    return Result::Ok(())
} // message.drop() called automatically
```

15.3.4 Converting Between Types

```
// Str -> String (allocates)
let owned = str_slice.to_string()?

// String -> Str (borrows)
let borrowed = owned.as_str()

// Both implement AsRef<Str>
fn accepts_str<T: AsRef<Str>>(s: T) {
    let slice: Str = s.as_ref()
}
```

15.4 Fixed-Size Strings

For embedded systems work, heap allocation can be expensive or unpredictable. Novus provides stack-allocated fixed-size strings for common use cases.

15.4.1 NameString (32 bytes)

For short identifiers, variable names, and labels:

```
from std::strings::core import NameString, Str

var name = NameString::new()
name.push_str(Str::from_cstr("Player1")??)

if name.len() < name.capacity() {
```

```
    name.push_byte($31)? // '1'
}

let slice: Str = name.as_str()
```

15.4.2 GadgetString (128 bytes)

For Intuition gadget text buffers and UI strings:

```
from std::strings::core import GadgetString, Str

var button_text = GadgetString::new()
button_text.push_str(Str::from_cstr("Click Here"))??

// Use for gadget creation
let label_ptr: *u8 = button_text.as_mut_ptr()
```

15.4.3 PathString (256 bytes)

For AmigaDOS file paths (max path length is typically 256 bytes):

```
from std::strings::core import PathString, Str

var path = PathString::new()
path.push_str(Str::from_cstr("Work:Projects/"))??
path.push_str(Str::from_cstr("MyApp/main.novus"))??

// Get as Str for DOS functions
let path_slice: Str = path.as_str()
```

Tip

Fixed strings return `Result::Err(StringError::TooLong)` if you exceed their capacity. Always check the result or use `capacity()` to verify space.

15.5 StringBuilder — Efficient Construction

`StringBuilder` provides efficient string construction with built-in formatting methods for numbers:

```
from std::strings::core import StringBuilder, Str

fn format_status(x: i32, y: i32, health: u32) -> Result<String,
StringError> {
    var sb = StringBuilder::new()?

    sb.push_str(Str::from_cstr("Position: (")??)
    sb.push_i32(x)?
    sb.push_str(Str::from_cstr(", ")??)
    sb.push_i32(y)?
    sb.push_str(Str::from_cstr(") Health: ")??)
    sb.push_u32(health)?
```



```
    return Result::Ok(sb.build())
}
```

15.5.1 Formatting Methods

Method	Parameter	Description
push_str(s)	Str	Append string
push_byte(b)	u8	Append single byte
push_char(c)	u8	Append character (alias)
push_i32(n)	i32	Format signed integer
push_u32(n)	u32	Format unsigned integer
push_hex(n)	u32	Format as lowercase hex
push_hex_upper(n)	u32	Format as uppercase hex
push_bool(b)	bool	Format as “true”/“false”

Table 15.4: StringBuilder formatting methods

```
var sb = StringBuilder::new()?

sb.push_str(Str::from_cstr("Address: $")?)?
sb.push_hex_upper(0xDFF000)? // "DFF000"

sb.push_str(Str::from_cstr(" Active: ")?)?
sb.push_bool(true)? // "true"

let result: String = sb.build()
// "Address: $DFF000 Active: true"
```

Note

`StringBuilder::new()` pre-allocates 256 bytes. Use `with_capacity(n)` if you know you need more or less.

Warning

`StringBuilder::build()` **consumes** the builder—you cannot use the builder after calling `build()`. The `consuming` keyword in the signature enforces this at compile time.

15.6 BStr — BCPL Strings for AmigaDOS

AmigaDOS uses BCPL-style strings for many operations. These have a length byte prefix and use “BPTR” addressing (address divided by 4).

```
from std::strings::core import BStr, Str

fn dos_example() -> Result<(), StringError> {
    // Create BCPL string from Str
    let bstr =
        BStr::new_from_str(Str::from_cstr("Work:Projects")?)?

    // Get BPTR for DOS functions
```

```

let bptr: u32 = bstr.as_bptr()

// Get actual memory address
let ptr: *u8 = bstr.as_ptr()

// Get length
let len: u32 = bstr.len()

// View as Str (excludes length byte)
let slice: Str = bstr.as_str()

return Result::Ok(())
} // bstr automatically freed via Drop

```

Warning

BCPL strings have a maximum length of 255 bytes (stored in a single byte). `BStr::new_from_str()` returns `Err(StringError::TooLong)` if the input exceeds this limit.

15.6.1 BCPL String Format

```

// Memory layout of BStr for "Hello" (5 bytes)
// Address 0x1000: [05][H][e][l][l][o]
//                ^-- length byte
// BPTR = 0x1000 / 4 = 0x400

```

Coming from C

In C, you'd manage BPTR conversion manually:

```

BPTR bptr = MKBADDR(str); // str / 4
char *str = BADDR(bptr);  // bptr * 4

```

`BStr` handles this automatically and provides RAII cleanup.

Note

`BStr` has a `free()` method for explicit cleanup, but this is **deprecated**. Prefer letting the `Drop` trait handle cleanup automatically when the `BStr` goes out of scope.

15.7 String Parsing

`Str` provides methods to parse common types from string representations:

15.7.1 Integer Parsing

```

let decimal = Str::from_cstr("12345")?
let negative = Str::from_cstr("-42")?
let hex1 = Str::from_cstr("0xFF00")?
let hex2 = Str::from_cstr("ff00")?

```

```
// Parse unsigned
match decimal.parse_u32() {
    Result::Ok(n) => {},      // n == 12345
    Result::Err(e) => {}
}

// Parse signed (handles - and +)
match negative.parse_i32() {
    Result::Ok(n) => {},      // n == -42
    Result::Err(e) => {}
}

// Parse hex (case-insensitive, optional 0x prefix)
match hex1.parse_hex() {
    Result::Ok(n) => {},      // n == 0xFF00
    Result::Err(e) => {}
}
```

15.7.2 Boolean Parsing

```
// Accepts: "true", "false", "1", "0" (case-insensitive)
let t1 = Str::from_cstr("true")?.parse_bool()? // true
let t2 = Str::from_cstr("TRUE")?.parse_bool()? // true
let t3 = Str::from_cstr("1")?.parse_bool()?    // true
let f1 = Str::from_cstr("false")?.parse_bool()? // false
let f2 = Str::from_cstr("0")?.parse_bool()?    // false
```

15.7.3 Parse Errors

```
pub enum ParseError {
    Empty,                // Empty string
    InvalidChar(u32),     // Bad character at index
    Overflow,             // Value too large
    Underflow,           // Value too small (signed)
    InvalidBool,         // Not "true"/"false"/"1"/"0"
}
```

15.8 Splitting and Iteration

15.8.1 Splitting by Delimiter

```
let csv = Str::from_cstr("apple,banana,cherry")?

var iter = csv.split($2C) // ', ' = 0x2C
while true {
    match iter.next() {
        Option::Some(part) => {
            // part: "apple", "banana", "cherry"
            print_str(part)
        },
        Option::None => break
    }
}
```

```
    }  
}
```

15.8.2 Splitting into Lines

```
let text = Str::from_cstr("Line 1\nLine 2\r\nLine 3")?  
  
var lines = text.lines()  
while true {  
    match lines.next() {  
        Option::Some(line) => {  
            // Handles LF, CR, and CRLF  
            print_str(line)  
        },  
        Option::None => break  
    }  
}
```

Note

`SplitIter` and `LinesIter` implement the `Iterator<Str>` trait, returning `Str` slices into the original string. No allocation occurs during iteration.

15.9 String Transformations

15.9.1 Case Conversion

```
let text = Str::from_cstr("Hello, World!")?  
  
// Convert to uppercase (allocates new String)  
let upper: String = text.to_uppercase()?  
// "HELLO, WORLD!"  
  
// Convert to lowercase (allocates new String)  
let lower: String = text.to_lowercase()?  
// "hello, world!"
```

Case conversion handles both ASCII (A-Z, a-z) and Latin-1 extended characters (à-ö, À-Ö, etc.).

15.9.2 Search and Replace

```
let text = Str::from_cstr("foo bar foo baz foo")?  
  
// Replace all occurrences  
let result: String = text.replace(  
    Str::from_cstr("foo")?,  
    Str::from_cstr("qux")?  
)?  
// "qux bar qux baz qux"
```

15.9.3 Joining Strings

```
from std::strings::core import str_join, Str

// Join array of Str with separator
var parts: [Str; 3] = [
  Str::from_cstr("SYS"),
  Str::from_cstr("C"),
  Str::from_cstr("Dir")
]

let path: String = str_join((*Str)&parts, 3,
  Str::from_cstr("/")?)?
// "SYS/C/Dir"
```

15.10 AmigaDOS Path Manipulation

Novus provides AmigaDOS-aware path utilities:

```
from std::strings::core import (
  path_split, path_filename, path_dirname,
  path_is_absolute, path_join, PathComponents
)

let full_path = Str::from_cstr("Work:Projects/Novus/main.novus")?

// Split into components
let parts: PathComponents = path_split(full_path)
// parts.volume = Some("Work:")
// parts.directory = Some("Projects/Novus")
// parts.filename = Some("main.novus")

// Get just the filename
match path_filename(full_path) {
  Option::Some(name) => {}, // "main.novus"
  Option::None => {}
}

// Get the directory
match path_dirname(full_path) {
  Option::Some(dir) => {}, // "Work:Projects/Novus"
  Option::None => {}
}

// Check if absolute (has volume)
if path_is_absolute(full_path) {
  // true - has "Work:"
}

// Join paths (handles separators correctly)
let joined: String = path_join(
  Str::from_cstr("Work:Projects")?,
  Str::from_cstr("NewApp")?
)?
// "Work:Projects/NewApp"
```

Tip

AmigaDOS path conventions:

- Volume names end with colon: SYS:, Work:, RAM:
- Directories separated by forward slash: /
- Parent directory: / (just slash, unlike .. in Unix)
- Current directory: "" (empty string)

15.11 Display and Debug Traits

The `std::strings::format` module provides traits for formatted output:

```
from std::strings::format import Display, Debug, Formatter

// Display - user-facing output
pub trait Display {
    fn fmt(&self, f: &var Formatter) -> bool
}

// Debug - programmer-facing debug output
pub trait Debug {
    fn fmt(&self, f: &var Formatter) -> bool
}
```

15.11.1 Implementing Display

```
struct Point {
    x: i32,
    y: i32,
}

impl Display for Point {
    fn fmt(&self, f: &var Formatter) -> bool {
        if !f.write_char($28) { return false } // '('
        // ... write x ...
        if !f.write_str(Str::from_cstr(", ")) { return false }
        // ... write y ...
        if !f.write_char($29) { return false } // ')'
        return true
    }
}
```

15.11.2 Primitive Implementations

All primitive types implement `Display` and `Debug`:

- i8, i16, i32, i64 — decimal representation
- u8, u16, u32, u64 — decimal representation

- `bool` — “true” or “false”
- `Str`, `String` — string contents

15.11.3 StackFormatter

For zero-allocation formatting (e.g., in interrupt handlers), use `StackFormatter`:

```
from std::strings::format import StackFormatter

fn format_on_stack() {
    var fmt = StackFormatter::new() // 256-byte stack buffer

    fmt.write_str(Str::from_cstr("Value: ")?)
    // ... format values ...

    let result: Str = fmt.as_str()
    // WARNING: result only valid while fmt is in scope!
}
```

Warning

`StackFormatter::as_str()` returns a `Str` pointing into the formatter’s stack buffer. This reference becomes invalid when the formatter goes out of scope.

15.12 Error Types

15.12.1 StringError

```
pub enum StringError {
    OutOfBounds(u32, u32), // (index, length)
    InvalidUtf8(u32),      // position
    AllocationFailed,
    TooLong(u32),          // length (for BStr: max 255)
    InvalidFormat,
    NullPointer,
}
```

15.12.2 ParseError

```
pub enum ParseError {
    Empty, // Empty string
    InvalidChar(u32), // position of bad character
    Overflow, // Value too large
    Underflow, // Value too small (signed)
    InvalidBool, // Not true/false/1/0
}
```

15.13 C FFI Functions

For low-level operations and C library interop, standard C string functions are available:

```
extern fn strlen(str: *u8) -> u32
extern fn strcmp(s1: *u8, s2: *u8) -> i32
extern fn strncmp(s1: *u8, s2: *u8, n: u32) -> i32
extern fn strcpy(dest: *u8, src: *u8) -> *u8
extern fn strncpy(dest: *u8, src: *u8, n: u32) -> *u8
extern fn strcat(dest: *u8, src: *u8) -> *u8
extern fn strncat(dest: *u8, src: *u8, n: u32) -> *u8
extern fn strchr(str: *u8, ch: i32) -> *u8
extern fn strrchr(str: *u8, ch: i32) -> *u8
extern fn strstr(haystack: *u8, needle: *u8) -> *u8
extern fn memmove(dest: *u8, src: *u8, n: u32) -> *u8
extern fn memset(ptr: *u8, value: i32, n: u32) -> *u8
extern fn memcmp(s1: *u8, s2: *u8, n: u32) -> i32
```

Warning

These functions provide no bounds checking and can easily cause buffer overflows. Prefer the safe Novus string types whenever possible. These are provided for performance-critical code and C library interop.

15.14 Trait Implementations

String types implement several important traits for idiomatic usage:

15.14.1 Clone

String and StringBuilder implement Clone for deep copying:

```
let original = String::new_from_str(Str::from_cstr("Hello"))??

// Clone creates a complete copy (allocates new memory)
match original.clone() {
    Option::Some(copy) => {
        // copy is independent of original
    },
    Option::None => {
        // allocation failed
    }
}
```

15.14.2 Write

String, StringBuilder, and PathString implement the Write trait for generic output:

```
from std::core import Write

// Generic function that writes to any Write implementor
fn write_header<W: Write>(w: &var W) -> bool {
    return w.write_bytes("Header: ", 8)
}

var s = String::new()
write_header(&s) // Works with String
```



```

var sb = StringBuilder::new()?
write_header(&sb)  // Works with StringBuilder

var path = PathString::new()
write_header(&path)  // Works with PathString

```

15.14.3 Ord (Ordering)

Str implements Ord for lexicographic comparison:

```

let a = Str::from_cstr("apple")?
let b = Str::from_cstr("banana")?

// Compare using Ord trait
let cmp: i32 = a.cmp(&b)
if cmp < 0 {
    // a comes before b lexicographically
}

```

15.14.4 Iterator

SplitIter and LinesIter implement Iterator<Str>:

```

let csv = Str::from_cstr("a,b,c")?

// Iterators implement Iterator<Str> trait
var iter = csv.split($2C)  // ','

// Can use in while loops with next()
while true {
    match iter.next() {
        Option::Some(part) => process(part),
        Option::None => break
    }
}

```

15.14.5 AsRef

Multiple types implement AsRef<Str> for unified string access:

```

// These all implement AsRef<Str>:
// - Str (identity)
// - String
// - StringBuilder
// - PathString

fn print_anything<T: AsRef<Str>>(s: T) {
    let slice: Str = s.as_ref()
    // Now use slice...
}

let owned = String::new_from_str(Str::from_cstr("Hello"))?
print_anything(owned)  // Works!

```

```
let path = PathString::new()
print_anything(path)    // Also works!
```

15.15 Performance Summary

Operation	Str	String	FixedString
Create empty	—	O(1)	O(1)
Create from C str	O(n)	O(n)	O(n)
Length	O(1)	O(1)	O(1)
Index byte	O(1)	O(1)	O(1)
Slice	O(1)	—	—
Push byte	—	O(1)*	O(1)
Push string	—	O(n)*	O(n)
Equals	O(n)	O(n)	O(n)
Find substring	O(nm)	O(nm)	O(nm)

Table 15.5: String operation complexity. *Amortized for String (may reallocate).

Coming from C

Key differences from C strings:

- **Length stored:** No need to scan for null terminator
- **Bounds checking:** Index operations return `Option`
- **RAII cleanup:** `String` and `BStr` free memory automatically
- **Safe slicing:** No pointer arithmetic errors
- **Type safety:** `Str` vs `String` distinguish borrowed vs owned

Chapter 16

Memory Management

Memory management on the Amiga is both powerful and perilous. Unlike modern systems with virtual memory and memory protection, the Amiga has a flat address space where any program can read from or write to any memory location. One wrong pointer dereferences the wrong address, and the entire system crashes—often spectacularly.

Novus addresses this challenge head-on with a memory system designed around three principles:

1. **RAII everywhere:** Memory is automatically freed when variables go out of scope
2. **Size tracking:** No more passing size parameters that can get out of sync
3. **Result/Option returns:** Allocation failures are explicit, not hidden null pointers

Warning

There is no memory protection on the Amiga. A null pointer dereference, buffer overflow, or use-after-free doesn't cause a clean error message—it corrupts memory, hangs the system, or produces bizarre behavior. The patterns in this chapter exist to prevent these disasters.

16.1 Amiga Memory Architecture

Understanding Amiga memory types is essential for effective programming:

16.1.1 Chip RAM vs Fast RAM

Type	Access	Use Cases
Chip RAM	CPU + DMA	Graphics, audio, blitter operations
Fast RAM	CPU only	Code, data, general computation

Table 16.1: Amiga memory types

Chip RAM (up to 2MB) is accessible by both the CPU and the custom chips (Paula, Blitter, Copper). All graphics data, audio samples, and blitter operations *must* use chip RAM.

Fast RAM (expansion memory) is accessible only by the CPU but runs at full speed without DMA contention. Use fast RAM for code, data structures, and anything that doesn't need custom chip access.

Coming from C

In C, you'd use `AllocMem()` with `MEMF_CHIP` or `MEMF_FAST` flags and manually track the size for `FreeMem()`. Forget the size or use the wrong flag, and you're debugging a guru meditation.

16.1.2 Memory Flags

```
from std::ffi::amiga_consts import MEMF_PUBLIC, MEMF_CHIP,
    MEMF_FAST, MEMF_ANY, MEMF_CLEAR
```

Flag	Purpose
MEMF_PUBLIC	Accessible by other tasks
MEMF_CHIP	Chip RAM (DMA accessible)
MEMF_FAST	Fast RAM (CPU only)
MEMF_ANY	Either chip or fast
MEMF_CLEAR	Zero-fill on allocation

Common flag combinations:

- MEMF_CHIP | MEMF_CLEAR: Graphics buffers, audio samples
- MEMF_FAST | MEMF_PUBLIC: General data structures
- MEMF_ANY | MEMF_PUBLIC: Don't care, just give me memory

16.2 Core Memory Types

The `std::memory` module provides a hierarchy of memory types, from low-level to high-level:

Type	Level	Size Tracking	Use Case
MemoryBlock	Low	Manual	FFI buffers, raw bytes
Allocation<T>	Mid	Automatic	Typed arrays
Box<T>	High	Automatic	Single heap values
MemHandle	High	Automatic	Simple RAII wrapper

Table 16.2: Memory type hierarchy

16.3 MemoryBlock — Raw Memory with Size Tracking

`MemoryBlock` is the fundamental building block. It wraps a raw pointer and its size, ensuring you can never “forget” how big an allocation is.

```
from std::memory::block import MemoryBlock
from std::ffi::amiga_consts import MEMF_FAST

fn allocate_buffer() {
    // Allocate 1024 bytes of fast RAM
    match MemoryBlock::alloc(1024, MEMF_FAST) {
        Option::Some(var block) => {
            // Use the block
            let ptr: *u8 = block.ptr()
            let size: u32 = block.size()

            // Write data
            ptr[0] = 42u8
            ptr[1] = 43u8

            // Memory automatically freed here!
```

```

    },
    Option::None => {
        // Allocation failed - handle gracefully
    }
}
}

```

16.3.1 Key Points

- `alloc()` returns `Option<MemoryBlock>`—no null pointer surprises
- `size()` always returns the correct size—no tracking bugs
- Implements `Drop`—memory freed automatically when block goes out of scope
- `free()` can be called explicitly if needed (idempotent)

16.3.2 The `try_alloc` Variant

For code using the `?` operator, use `try_alloc()`:

```

fn process_data() -> Result<(), ExecError> {
    var block = MemoryBlock::try_alloc(1024, MEMF_FAST)?

    // Use block...

    return Result::Ok(())
} // block freed automatically

```

16.3.3 Resizing Memory

```

match MemoryBlock::alloc(256, MEMF_FAST) {
    Option::Some(var block) => {
        // Need more space
        if block.resize(512, MEMF_FAST) {
            // Success - block now holds 512 bytes
            // Old data is preserved (up to min(old_size,
            //                               new_size))
        } else {
            // Resize failed - block unchanged
        }
    },
    Option::None => {}
}

```

Warning

Unlike C's `realloc()`, Novus always allocates new memory and copies—AmigaOS doesn't support in-place growth. This means resizing is $O(n)$. Pre-allocate when possible!

16.3.4 Taking Ownership with `into_raw()`

Sometimes you need to transfer ownership to external code:

```

match MemoryBlock::alloc(1024, MEMF_CHIP) {
  Option::Some(var block) => {
    // Take ownership - block won't free on drop
    let (ptr, size) = block.into_raw()

    // ptr and size are now YOUR responsibility!
    // You MUST call FreeMem(ptr, size) eventually
  },
  Option::None => {}
}

```

Warning

into_raw() transfers ownership to **YOU**. The returned pointer and size **MUST** be manually freed with `FreeMem(ptr, size)`. Forgetting this causes a memory leak. On the Amiga, **leaked memory is lost until reboot**—there is no garbage collector, no automatic cleanup, nothing. This function exists for FFI interop; avoid it unless absolutely necessary.

16.4 Allocation<T> — Type-Safe Bulk Allocation

`Allocation<T>` provides type-safe allocation for arrays of a specific type. It automatically calculates the byte size from the element count.

```

from std::memory::block import Allocation
from std::ffi::amiga_consts import MEMF_FAST

struct Vertex {
  x: i32,
  y: i32,
  z: i32,
}

fn allocate_vertices() {
  // Allocate space for 100 vertices
  match Allocation::<Vertex>::new(100, MEMF_FAST) {
    Option::Some(var vertices) => {
      let count: u32 = vertices.count()           // 100
      let bytes: u32 = vertices.size_bytes()      // 100 * 12
                                              = 1200

      // Get typed pointer
      let ptr: *Vertex = vertices.as_mut_ptr()
      ptr[0] = Vertex { x: 10, y: 20, z: 30 }

      // Resize to 200 vertices
      if vertices.resize(200, MEMF_FAST) {
        // Now holds 200 vertices
      }
    },
    Option::None => {}
  }
} // Memory freed automatically

```

16.4.1 Overflow Protection

Allocation checks for integer overflow when calculating total bytes:

```
// This would overflow: $FFFFFFFF * 4 > u32::MAX
match Allocation::<u32>::new($FFFFFFFFu32, MEMF_FAST) {
    Option::Some(_) => {}, // Never reached
    Option::None => {}      // Returns None - overflow detected
}
```

16.5 Box<T> — Single Heap Values

Box<T> is the simplest way to put a single value on the heap. Use it for large structs or recursive data structures.

```
from std::memory::block import Box
from std::ffi::amiga_consts import MEMF_CHIP

struct LargeBuffer {
    data: [u8; 4096],
}

fn use_large_struct() {
    // Too big for stack? Put it on the heap
    let initial = LargeBuffer { data: [0u8; 4096] }

    match Box::new(initial) {
        Option::Some(var boxed) => {
            // Access via get() / get_mut()
            let ptr: *LargeBuffer = boxed.get_mut()
            (*ptr).data[0] = 42u8
        },
        Option::None => {}
    }
} // Memory freed automatically

// Or with specific memory flags:
fn chip_ram_struct() {
    match Box::new_in(my_value, MEMF_CHIP) {
        Option::Some(var boxed) => {
            // boxed.get() is in chip RAM - DMA accessible
        },
        Option::None => {}
    }
}
```

16.6 MemHandle — Simple RAI Wrapper

MemHandle provides convenient factory methods for common allocation patterns using AllocVec/FreeVec:

```
from std::memory::amiga import MemHandle

fn allocation_examples() {
    // Chip RAM (DMA accessible)
```

```

match MemHandle::new_chip(1024) {
    Option::Some(var chip_mem) => {
        let ptr: *u8 = chip_mem.ptr()
        // Use for graphics, audio...
    },
    Option::None => {}
}

// Fast RAM (CPU only, faster)
match MemHandle::new_fast(2048) {
    Option::Some(var fast_mem) => {
        let ptr: *u8 = fast_mem.ptr()
        // Use for data structures...
    },
    Option::None => {}
}

// Any available memory
match MemHandle::new_any(512) {
    Option::Some(var any_mem) => {
        // Don't care where it comes from
    },
    Option::None => {}
}

// Custom flags
match MemHandle::new_with_flags(1024, MEMF_CHIP | MEMF_CLEAR)
{
    Option::Some(var mem) => {
        // Chip RAM, zero-filled
    },
    Option::None => {}
}
} // All memory freed automatically

```

16.7 Slice<T> — Fat Pointers for Safety

Slice<T> is a “fat pointer” containing both a pointer and length. This eliminates the classic C bug of passing arrays without their sizes.

```

from std::memory::slice import Slice

fn sum_values(numbers: &Slice<i32>) -> i32 {
    var total: i32 = 0
    var i: u32 = 0
    while i < numbers.len() {
        // Safe access with bounds checking
        match numbers.get(i) {
            Option::Some(val) => total = total + val,
            Option::None => {}
        }
        i = i + 1
    }
    return total
}

```



```
fn use_slice() {
    var arr: [i32; 5] = [1, 2, 3, 4, 5]
    let slice = Slice::new((*i32)&arr[0], 5)

    let result = sum_values(&slice) // No size parameter needed!
}
```

Tip

Slices don't own their data—they're just views. The underlying memory must remain valid for the slice's lifetime.

16.8 Chip RAM Management

Chip RAM is precious on the Amiga (typically 512KB to 2MB). Novus provides specialized tools for managing it efficiently.

16.8.1 ChipBuffer — Direct Allocation

For simple chip RAM needs:

```
from std::memory::chip import ChipBuffer

fn graphics_buffer() {
    match ChipBuffer::alloc(32000) {
        Option::Some(var buffer) => {
            // Copy data with automatic cache flush (68040/060
            // safe)
            buffer.copy_from(source_ptr, 32000)

            // Get pointer for DMA operations
            let dma_ptr: *u8 = buffer.as_ptr()
        },
        Option::None => {
            // Not enough chip RAM
        }
    }
} // Chip RAM freed automatically
```

16.8.2 Checking Memory Type

```
from std::memory::chip import is_chip_ram

fn validate_dma_pointer(ptr: *u8) -> bool {
    if !is_chip_ram(ptr) {
        // ERROR: This pointer can't be used for DMA!
        return false
    }
    return true
}
```

16.8.3 Cache-Safe Copying (68040/68060)

On 68040 and 68060 systems, the CPU cache can contain data that DMA hardware hasn't seen. Always use cache-safe functions:

```
from std::memory::chip import copy_to_chip

fn safe_copy(dst_chip: *u8, src_fast: *u8, size: u32) {
    // Copies data AND flushes CPU cache for this range
    copy_to_chip(dst_chip, src_fast, size)
    // Now safe for Blitter/Paula to read
}
```

Warning

Never use raw CopyMem() for chip RAM on 68040/060! The CPU cache may contain data that DMA hardware can't see. Always use copy_to_chip() or manually call CacheClearE().

16.9 ChipPool — Pooled Allocation

Frequent small allocations fragment memory. ChipPool uses AmigaOS's pool functions (V39+) to allocate from pre-allocated “puddles,” reducing fragmentation.

```
from std::memory::chip_pool import ChipPool, PoolSizeClass

fn pooled_allocation() {
    // Initialize pool (once at startup)
    ChipPool::init()

    // Allocate from pool (returns pointer to size-class block)
    match ChipPool::alloc(3000) {
        Option::Some(ptr) => {
            // Got 4KB block (smallest class >= 3000)

            // Use for sprites, samples, etc...

            // Return to pool
            ChipPool::free(ptr)
        },
        Option::None => {}
    }

    // Shutdown frees all pool memory
    ChipPool::shutdown()
}
```

16.9.1 Size Classes

Requests are rounded up to the next size class. For allocations > 32KB, use direct AllocMem().

Class	Size	Puddle	Use Case
Small	4KB	32KB	Sprites, small samples
Medium	8KB	64KB	Copper lists, medium samples
Large	16KB	64KB	Streaming buffers
Huge	32KB	128KB	Large streaming buffers

Table 16.3: ChipPool size classes

16.10 ChipCache — LRU Caching

ChipCache provides automatic caching of assets in chip RAM with priority-based LRU eviction. Assets in fast RAM are transparently copied to chip RAM when needed.

```

from std::memory::chip_cache import ChipCache
from std::memory::chip import Priority

fn cached_asset() {
    // Initialize cache (once at startup)
    ChipCache::init()

    // Get asset into chip RAM (copies if not already there)
    let asset_id: u32 = (u32)my_asset_ptr // Unique ID
    match ChipCache::get(
        asset_id,
        my_asset_ptr, // Source (may be fast RAM)
        asset_size,
        Priority::Normal,
        "player_sprite" // Debug name
    ) {
        Result::Ok(chip_ptr) => {
            // chip_ptr is guaranteed to be in chip RAM
            // Use for DMA operations
        },
        Result::Err(e) => {
            // Not enough chip RAM even after eviction
        }
    }
}

```

16.10.1 Pinning Assets

During DMA operations, you must prevent eviction:

```

// Before DMA
ChipCache::pin(asset_id)

// ... DMA operation (Blitter, audio, etc.) ...

// After DMA complete
ChipCache::unpin(asset_id)

```

Warning

Always pin assets during DMA! If an asset is evicted while the Blitter or Paula is reading it, you'll get corrupted graphics or audio—or worse, the system reads freed memory and crashes.

16.10.2 Priority Levels

```
pub enum Priority {
    Low,      // Evict first (large, infrequent assets)
    Normal,   // Default
    High,     // Evict last (small, frequently used)
}
```

The cache evicts in this order:

1. Lowest priority first
2. Within same priority, oldest access first (LRU)
3. Pinned assets are never evicted

16.10.3 Cache Statistics

```
let stats = ChipCache::stats()
// stats.entries      - Cached asset count
// stats.pinned_entries - Pinned asset count
// stats.total_cached  - Total bytes cached
// stats.cache_hits    - Hit count
// stats.cache_misses  - Miss count
// stats.evictions     - Eviction count
// stats.chip_available - Available chip RAM
```

The CacheStats structure:

```
pub struct CacheStats {
    entries: u32,      // Total cached entries
    pinned_entries: u32, // Pinned (non-evictable) entries
    total_cached: u32,  // Bytes in cache
    pinned_bytes: u32,  // Bytes in pinned entries
    chip_available: u32, // Free chip RAM (AvailMem)
    chip_largest: u32,  // Largest contiguous block
    cache_hits: u32,    // Successful lookups
    cache_misses: u32,  // Lookups requiring load
    evictions: u32,     // Times entries were evicted
}
```

16.10.4 Advanced Cache Operations

```
// Update LRU timestamp (call once per frame)
ChipCache::tick()

// Manually evict a specific asset
```

```
ChipCache::evict(asset_id)

// Evict all unpinned assets (emergency memory recovery)
ChipCache::clear_unpinned()
```

Tip

Call `ChipCache::tick()` once per frame in your main loop. This updates the LRU timestamps used for eviction decisions. Without it, the cache can't distinguish recently-used assets from stale ones.

16.11 ResourceBank — Proactive Management

While `ChipCache` is reactive (caching on use), `ResourceBank` is proactive—you explicitly load groups of related assets when entering a level or scene.

```
from std::memory::resource_bank import ResourceBank, BankBuilder
from std::memory::chip import Priority

fn level_loading() -> Result<(), BankError> {
    // Build a bank for level 1 assets
    // Methods marked @chain return &var self for fluent chaining
    var level1 = ResourceBank::builder("level1")
        .add_asset(sprite_id, sprite_ptr, sprite_size,
            Priority::High)
        .add_asset(tileset_id, tileset_ptr, tileset_size,
            Priority::Normal)
        .add_asset_named(music_id, music_ptr, music_size,
            Priority::Low, "level_music")
        .persistent(false) // Unload when done
        .build()?

    // Load all assets atomically
    level1.load()?

    // Get chip RAM pointer for an asset
    match level1.get_by_id(sprite_id) {
        Option::Some(chip_ptr) => {
            // Use for DMA...
        },
        Option::None => {}
    }

    // When leaving level, unload everything
    level1.unload()

    return Result::Ok(())
}
```

16.11.1 Bank Features

- **Atomic loading:** All assets load or none do (rollback on failure)
- **Reference counting:** Shared assets across banks are pinned correctly

- **will_fit() check:** Verify before loading to avoid partial failures
- **Persistent banks:** Mark UI/font banks to never unload

16.11.2 Pinning Banks for DMA Safety

For multi-frame DMA operations (e.g., a long Blitter chain), pin the entire bank:

```
// Before multi-frame DMA
level1.pin() // Increment pin count on all assets

// ... several frames of DMA operations ...

// When done
level1.unpin() // Decrement pin counts
```

16.11.3 Bank State Checking

```
// Check if bank is loaded
if level1.is_loaded() {
    let ptr = level1.get_by_id(sprite_id)
}

// Get detailed state
match level1.state() {
    BankState::Unloaded => { /* Not in chip RAM */ },
    BankState::Loading => { /* Async load in progress */ },
    BankState::Loaded => { /* Ready for use */ },
    BankState::Failed => { /* Load failed */ },
}

// Statistics
let count = level1.entry_count() // Number of assets
let bytes = level1.total_size() // Total bytes
let name = level1.name() // Bank name
```

16.11.4 Memory Budgeting

```
from std::memory::resource_bank import ChipBudget

fn plan_memory() {
    let available_chip = AvailMem(MEMF_CHIP)

    // Typical budget: 25% persistent, 50% level, 25% streaming
    let budget = ChipBudget::typical(available_chip)

    // Check if level bank fits
    if budget.level_will_fit(&level1_bank) {
        level1_bank.load()?
    } else {
        // Need to unload something first
    }
}
```

16.12 Streaming Large Assets

For assets too large to fit in chip RAM (music modules, video), use streaming with double or triple buffering.

16.12.1 The Streamable Trait

```
from std::memory::streamable import Streamable

impl Streamable for MyAudioAsset {
    fn stream_data(&self) -> *u8 {
        return self.data_ptr
    }

    fn stream_size(&self) -> u32 {
        return self.data_size
    }

    fn suggested_buffer_size(&self) -> u32 {
        return 16384 // 16KB per buffer
    }

    fn suggested_buffer_count(&self) -> u8 {
        return 3 // Triple-buffered
    }
}
```

16.12.2 ChipStream — High-Level Streaming

ChipStream manages streaming buffers for you. The constructor:

```
pub fn new(source: *u8, size: u32, buffer_count: u8, buffer_size:
u32)
    -> Result<ChipStream, ChipCacheError>
```

Usage:

```
from std::memory::streamable import ChipStream

fn stream_audio(audio: &MyAudioAsset) -> Result<(),
ChipCacheError> {
    // Create stream with triple buffering
    var stream = ChipStream::new(
        audio.stream_data(),
        audio.stream_size(),
        3, // buffer count
        16384 // buffer size
    )?

    // Prime all buffers
    stream.fill_all()?

    // Playback loop
    while !stream.is_finished() {
        // Get current buffer for Paula
```

```

    let chip_ptr = stream.current_buffer()
    let filled = stream.current_buffer_filled()

    // Start Paula playing this buffer...
    // ... wait for buffer to finish ...

    // Advance to next buffer (refills the one just played)
    stream.advance()?
}

return Result::Ok(())
}

```

16.12.3 Buffer Strategy

- **Double buffering (2):** One plays while one fills. Tight timing.
- **Triple buffering (3):** Extra buffer absorbs timing jitter. Recommended.

16.13 DMA Cache Management (68040/68060)

On 68040 and 68060 systems with data caches, you must synchronize between CPU and DMA:

```

from std::memory::chip import (
    prepare_for_dma_read,
    complete_dma_read,
    prepare_for_dma_write,
    complete_dma_write
)

fn blitter_copy(src: *u8, dst: *u8, size: u32) {
    // Prepare source (flush cache so Blitter sees correct data)
    let src_dma = prepare_for_dma_read(src, size)

    // Prepare destination (invalidate cache)
    let dst_dma = prepare_for_dma_write(dst, size)

    // ... perform Blitter operation ...

    // Complete (invalidate destination cache so CPU sees DMA
    // result)
    complete_dma_read(src, size)
    complete_dma_write(dst, size)
}

```

Note

These functions are no-ops on 68000/020/030. They're always safe to call—the cost is negligible on systems that don't need them.

Warning

The function names describe the DMA hardware's perspective, not the CPU's:

- `prepare_for_dma_read`: DMA will **READ** from this buffer (e.g., Blitter source, Paula sample). Flushes CPU cache so DMA sees correct data.
- `prepare_for_dma_write`: DMA will **WRITE** to this buffer (e.g., Blitter destination). Invalidates CPU cache.

This naming convention matches the AmigaOS `CacheClearE()` flags.

16.13.1 Critical: `WaitBlit()` Before Freeing

When freeing chip RAM that the Blitter may have been using, you **must** wait for the Blitter to finish first:

```
from std::ffi::graphics import WaitBlit

fn safe_free_chip_buffer(buffer: &var MemoryBlock) {
    // CRITICAL: Blitter might still be reading/writing this
    // memory!
    WaitBlit() // Block until all Blitter DMA completes

    // Now safe to free
    buffer.free()
}
```

Warning

Never free chip RAM without calling `WaitBlit()` first! The Blitter operates asynchronously. If you free memory while the Blitter is still using it, the Blitter will read or write freed (or reallocated) memory, causing data corruption, graphical glitches, or system crashes.

This rule applies to:

- Manual `FreeMem()/FreeVec()` calls
- RAII `Drop` implementations (Novus's built-in types handle this)
- Pool returns via `ChipPool::free()`

16.14 Common Patterns

16.14.1 Pattern 1: Load-Use-Free

The most common pattern—allocate, use, automatically free:

```
fn process() -> Result<(), ExecError> {
    var buffer = MemoryBlock::try_alloc(1024, MEMF_FAST)?

    // Use buffer...
    let ptr = buffer.ptr()
    // ... work with ptr ...

    return Result::Ok(())
} // buffer.drop() called automatically
```

16.14.2 Pattern 2: defer for Complex Logic

When you need cleanup at multiple exit points:

```
fn complex_operation() -> Result<(), Error> {
    match MemoryBlock::alloc(1024, MEMF_FAST) {
        Option::Some(var buffer) => {
            defer {
                buffer.free() // Always runs on scope exit
            }

            if some_condition() {
                return Result::Err(Error::Something)
                // buffer freed here
            }

            if other_condition() {
                return Result::Err(Error::Other)
                // buffer freed here too
            }

            return Result::Ok(())
            // buffer freed here as well
        },
        Option::None => return Result::Err(Error::NoMem)
    }
}
```

16.14.3 Pattern 3: Pre-allocation for Performance

Avoid allocation in hot paths:

```
struct GameState {
    sprite_buffer: MemoryBlock,
    tile_buffer: MemoryBlock,
}

fn init_game() -> Option<GameState> {
    let sprites = MemoryBlock::alloc(64000, MEMF_CHIP)?
    let tiles = MemoryBlock::alloc(128000, MEMF_CHIP)?

    return Option::Some(GameState {
        sprite_buffer: sprites,
        tile_buffer: tiles,
    })
}

fn game_loop(state: &var GameState) {
    // No allocation here - use pre-allocated buffers
    let sprite_ptr = state.sprite_buffer.ptr()
    // ... fast path, no allocations ...
}
```

16.14.4 Pattern 4: Graceful Degradation

When memory is tight, fall back to smaller allocations:

```

fn allocate_with_fallback() -> Option<MemoryBlock> {
    // Try ideal size first
    match MemoryBlock::alloc(65536, MEMF_CHIP) {
        Option::Some(block) => return Option::Some(block),
        Option::None => {}
    }

    // Fall back to smaller
    match MemoryBlock::alloc(32768, MEMF_CHIP) {
        Option::Some(block) => return Option::Some(block),
        Option::None => {}
    }

    // Minimum viable
    return MemoryBlock::alloc(16384, MEMF_CHIP)
}

```

16.15 Memory Safety Checklist

Follow these rules to avoid memory disasters:

1. **Always check allocation results**—never assume success
2. **Let RAII handle cleanup**—don't call `free()` unless necessary
3. **Pin assets during DMA**—unpinned assets can be evicted
4. **Use cache-safe copies on 68040/060**—raw `CopyMem` causes corruption
5. **Pre-allocate in init, not in loops**—avoid fragmentation
6. **Check `is_chip_ram()` before DMA**—fast RAM won't work
7. **Match allocation and deallocation**—`AllocMem/FreeMem`, `AllocVec/FreeVec`
8. **Never use freed memory**—the RAII types make this hard to do wrong
9. **Track sizes with types, not variables**—use `MemoryBlock`, not raw pointers
10. **Call `WaitBlit()` before freeing chip RAM**—the Blitter runs asynchronously

Warning

The Amiga has no safety net. Modern operating systems kill your process when you access invalid memory. The Amiga doesn't—it corrupts data, hangs, or produces bizarre effects. The discipline enforced by these patterns isn't optional; it's survival.

Operation	Cost	Notes
<code>MemoryBlock::alloc()</code>	$O(1)$	AmigaOS first-fit
<code>MemoryBlock::free()</code>	$O(1)$	Returns to free list
<code>MemoryBlock::resize()</code>	$O(n)$	Always copies
<code>ChipPool::alloc()</code>	$O(1)$	Pool allocation
<code>ChipCache::get()</code> (hit)	$O(n)$	Linear search
<code>ChipCache::get()</code> (miss)	$O(n)$	Search + alloc + copy
<code>is_chip_ram()</code>	$O(1)$	<code>TypeOfMem</code> syscall

Table 16.4: Memory operation costs

16.16 Performance Considerations

Tip

Allocation is not free. Every `AllocMem` call takes time. For real-time code (games, audio), pre-allocate everything during initialization and reuse buffers during gameplay.

Chapter 17

Input/Output

Input and output on the Amiga spans two distinct worlds: the CLI (Command Line Interface) where text-based programs read arguments and write output, and Workbench where programs are launched from icons with tooltypes. Novus provides unified support for both, with RAII wrappers that prevent the resource leaks that plagued classic Amiga development.

17.1 Console Output

The `std::io::file` module provides basic console I/O functions:

```
from std::io::file import write, print, println

fn hello() {
    // Simple string output
    print("Hello, ")
    println("Amiga!") // Adds newline

    // Printf-style formatting (AmigaOS RawDoFmt syntax)
    let count: i32 = 42
    let name: *u8 = "Player"
    write("Score: %ld for %s\n", count, name)
}
```

Warning

Use `%ld` for 32-bit integers, not `%d`! AmigaOS `RawDoFmt` expects 16-bit values for `%d`. Since Novus passes all integers as 32-bit, always use the `l` (long) prefix: `%ld`, `%lu`, `%lx`.

17.1.1 Format Specifiers

Specifier	Meaning
<code>%ld</code>	Signed 32-bit decimal
<code>%lu</code>	Unsigned 32-bit decimal
<code>%lx</code>	Hexadecimal (lowercase)
<code>%lX</code>	Hexadecimal (uppercase)
<code>%s</code>	Null-terminated string
<code>%c</code>	Single character
<code>%b</code>	BSTR (BPTR to byte-counted string)

Table 17.1: AmigaOS format specifiers

17.1.2 Array-Based Formatting

When building arguments programmatically:

```
from std::io::file import write_array

fn formatted_output(major: i32, minor: i32, patch: i32) {
    var args: [i32; 3] = [major, minor, patch]
    write_array("Version: %ld.%ld.%ld\n", &args[0])
}
```

17.2 Console Input

```
from std::io::file import wait_for_key, key_pressed, read_key

fn input_examples() {
    // Block until any key pressed
    print("Press any key to continue...")
    wait_for_key()
    println("")

    // Non-blocking check
    if key_pressed() {
        let ch = read_key()
        write("You pressed: %c\n", (i32)ch)
    }

    // Blocking read of single key
    print("Enter a character: ")
    let input = read_key() // Returns u8
    write("Got: %c\n", (i32)input) // Cast to i32 for format
                                string
}
```

Note

`read_key()` returns `u8`. The cast to `i32` is required because AmigaOS `%c` format expects a 32-bit value on the stack.

17.3 ANSI Console Control

The Amiga console supports ANSI X3.64 escape sequences for cursor control, colors, and text styling. The `std::io::ansi` module provides type-safe access to these sequences.

Note

Return types: Simple functions like `reset()`, `bold()`, `clear_screen()` return `Str` (static escape sequences). Parameterized functions like `move_to()`, `fg()`, `cursor_up()` return `String` (dynamically built). Both have `as_cstr()` for use with `write()`.

17.3.1 Basic Functions

```

from std::io::ansi import clear_screen, home, reset, bold,
    underline, reverse
from std::io::file import write

fn styled_output() {
    // Clear screen and move to top-left
    write("%s%s", clear_screen().as_cstr(), home().as_cstr())

    // Styled text
    write("%sThis is bold%s\n", bold().as_cstr(),
        reset().as_cstr())
    write("%sThis is underlined%s\n", underline().as_cstr(),
        reset().as_cstr())
    write("%sThis is reversed%s\n", reverse().as_cstr(),
        reset().as_cstr())
}

```

17.3.2 Pen Colors

The Amiga console uses *pen numbers*, not fixed colors. The actual color depends on the Workbench palette:

```

from std::io::ansi import Pen, fg, bg, reset
from std::io::file import write

fn colored_output() {
    // Foreground colors
    write("%sText in Pen 1%s\n", fg(Pen::Pen1).as_cstr(),
        reset().as_cstr())
    write("%sText in Pen 2%s\n", fg(Pen::Pen2).as_cstr(),
        reset().as_cstr())
    write("%sText in Pen 3%s\n", fg(Pen::Pen3).as_cstr(),
        reset().as_cstr())

    // Background colors
    write("%sHighlighted%s\n", bg(Pen::Pen3).as_cstr(),
        reset().as_cstr())
}

```

Pen	Workbench 2.0+ Default	ANSI Code
Pen0	Grey (background)	30/40
Pen1	Black (text)	31/41
Pen2	White (highlight)	32/42
Pen3	Blue (shine/fill)	33/43
Pen4–7	Only on 8+ color screens	34–37/44–47

Table 17.2: Amiga console pen mapping

Note

On 4-color screens, pens 4–7 are remapped: Pen4 becomes invisible (same as background), Pen5 maps to Pen1, Pen6 to Pen2, Pen7 to Pen3.

17.3.3 Cursor Movement

```
from std::io::ansi import move_to, cursor_up, cursor_down,
    cursor_left, cursor_right
from std::io::file import write

fn cursor_demo() {
    // Absolute positioning (1-indexed, row then column)
    write("%s", move_to(10, 20).as_cstr())
    write("At row 10, column 20")

    // Relative movement
    write("%s", cursor_up(5).as_cstr()) // Move up 5 rows
    write("%s", cursor_right(10).as_cstr()) // Move right 10
        columns
}
```

17.3.4 Screen and Line Clearing

```
from std::io::ansi import clear_screen, clear_line, clear_to_eol,
    clear_to_eos
from std::io::ansi import ClearMode, clear
from std::io::file import write

fn clearing_demo() {
    // Simple functions
    write("%s", clear_screen().as_cstr()) // Clear entire screen
    write("%s", clear_line().as_cstr()) // Clear current line
    write("%s", clear_to_eol().as_cstr()) // Clear to end of line
    write("%s", clear_to_eos().as_cstr()) // Clear to end of
        screen

    // Parameterized clearing
    write("%s", clear(ClearMode::ToEnd).as_cstr()) // From
        cursor to end
    write("%s", clear(ClearMode::ToStart).as_cstr()) // From
        start to cursor
    write("%s", clear(ClearMode::All).as_cstr()) // Entire
        screen
}
```

17.3.5 The Ansi Builder

For complex styling, use the builder pattern:

```
from std::io::ansi import Ansi, Pen, reset
from std::io::file import write

fn builder_demo() {
    // Build complex style
    var style = Ansi::new()
    style.bold()
    style.fg(Pen::Pen2)
    style.bg(Pen::Pen1)
```



```

    // Apply and reset
    let seq = style.build()
    write("%sStyled Text%s\n", seq.as_cstr(), reset().as_cstr())
}

// Method chaining
fn chained_demo() {
    var style = Ansi::new()
    style.bold().underline().fg(Pen::Pen3)

    write("%sChained%s\n", style.build().as_cstr(),
        reset().as_cstr())
}

```

17.4 File I/O

The `std::os::dos` module provides RAII wrappers for AmigaOS file operations.

17.4.1 OwnedFileHandle — RAII File Access

```

from std::os::dos import OwnedFileHandle
from std::ffi::amiga_consts import MODE_OLDFILE, MODE_NEWFILE

fn read_file_example() -> Result<(), DosError> {
    // Open file - automatically closed when handle goes out of
    // scope
    let file = OwnedFileHandle::open("S:Startup-Sequence",
        MODE_OLDFILE)?

    // Read content
    var buffer: [u8; 1024] = [0u8; 1024]
    let bytes_read = file.read((*u8)&buffer[0], 1024)

    if bytes_read > 0 {
        // Process buffer...
    }

    return Result::Ok(())
} // file automatically closed here!

```

Coming from C

In C, you'd use `Open()` and must remember to call `Close()` on every code path—including error paths. With `OwnedFileHandle`, the file is closed automatically when the handle goes out of scope, even if an error occurs.

17.4.2 File Open Modes

17.4.3 Reading and Writing

```

from std::os::dos import OwnedFileHandle
from std::ffi::amiga_consts import MODE_NEWFILE

```

Mode	Value	Description
MODE_OLDFILE	1005	Open existing file (read/write)
MODE_NEWFILE	1006	Create new file (truncate if exists)
MODE_READWRITE	1004	Open with shared lock

Table 17.3: DOS file open modes

```

fn write_file_example() -> Result<(), DosError> {
    let file = OwnedFileHandle::open("RAM:output.txt",
        MODE_NEWFILE)?

    // Write bytes
    let data: *u8 = "Hello, Amiga!\n"
    let written = file.write(data, 14)

    // Write all bytes (loops until complete)
    file.write_all(data, 14)?

    return Result::Ok(())
}

fn read_all_example() -> Result<(), DosError> {
    let file = OwnedFileHandle::open("RAM:input.txt",
        MODE_OLDFILE)?

    var buffer: [u8; 4096] = [0u8; 4096]

    // Read up to buffer size
    let total = file.read_all((*u8)&buffer[0], 4096)?

    write("Read %ld bytes\n", (i32)total)

    return Result::Ok(())
}

```

17.4.4 Reading Lines

```

from std::os::dos import OwnedFileHandle
from std::ffi::amiga_consts import MODE_OLDFILE

fn read_lines_example() -> Result<(), DosError> {
    let file = OwnedFileHandle::open("S:Startup-Sequence",
        MODE_OLDFILE)?

    var line_buffer: [u8; 256] = [0u8; 256]
    var line_num: i32 = 1

    loop {
        match file.read_line((*u8)&line_buffer[0], 256) {
            Result::Ok(len) => {
                if len == 0 {
                    break // EOF
                }
                write("%ld: %s", line_num, (*u8)&line_buffer[0])
            }
        }
    }
}

```

```

        line_num++
    },
    Result::Err(_) => break,
}
}

return Result::Ok(())
}

```

17.4.5 Seeking

```

from std::os::dos import OwnedFileHandle
from std::ffi::amiga_consts import MODE_OLDFILE,
    OFFSET_BEGINNING, OFFSET_CURRENT, OFFSET_END

fn seek_example() -> Result<(), DosError> {
    let file = OwnedFileHandle::open("RAM:data.bin",
        MODE_OLDFILE)?

    // Seek to beginning
    file.seek(0, OFFSET_BEGINNING)

    // Seek forward 100 bytes from current position
    file.seek(100, OFFSET_CURRENT)

    // Get file size by seeking to end
    let size = file.seek(0, OFFSET_END)
    write("File size: %ld bytes\n", size)

    // Seek back to start
    file.seek(0, OFFSET_BEGINNING)

    return Result::Ok(())
}

```

Constant	Value	Meaning
OFFSET_BEGINNING	-1	From start of file
OFFSET_CURRENT	0	From current position
OFFSET_END	1	From end of file

Table 17.4: DOS seek offset constants

17.4.6 Ownership Transfer

When you need to pass a file handle to external code:

```

fn ownership_example() -> Result<i32, DosError> {
    let file = OwnedFileHandle::open("RAM:test.txt",
        MODE_OLDFILE)?

    // Transfer ownership - prevents automatic close
    let raw_handle: i32 = file.into_raw()

    // raw_handle is now YOUR responsibility!
}

```

```

    // You must call Close(raw_handle) when done

    return Result::Ok(raw_handle)
}

fn from_raw_example(handle: i32) {
    // Take ownership of an existing handle
    let file = OwnedFileHandle::from_raw(handle)

    // Use file...

} // Handle closed automatically

```

Warning

Never use `from_raw()` on system handles! Handles like `Input()`, `Output()`, and `Error()` are owned by the system and must not be closed. Only use `from_raw()` on handles you own.

17.5 Directory Locks

AmigaOS uses “locks” to access directories. `DirLockHandle` provides RAII management:

```

from std::os::dos import DirLockHandle
from std::ffi::amiga_consts import ACCESS_READ, ACCESS_WRITE

fn lock_example() -> Result<(), DosError> {
    // Lock directory for reading
    let lock = DirLockHandle::lock("SYS:", ACCESS_READ)?

    // Get raw lock for FFI
    let raw_lock = lock.handle()

    // Use with Examine(), ExNext(), etc...

    return Result::Ok(())
} // Lock automatically released

```

Warning

Lock leaks prevent floppy ejection! If a lock on a floppy disk isn’t released, the user cannot eject the disk and sees “Volume is in use.” Always use `DirLockHandle` to ensure locks are released.

17.6 File System Helpers

```

from std::os::dos import file_exists, delete_file, create_dir,
    file_info

fn fs_examples() -> Result<(), DosError> {
    // Check existence
    if file_exists("RAM:test.txt") {

```

```

        println("File exists!")
    }

    // Delete file
    delete_file("RAM:temp.txt")?

    // Create directory
    create_dir("RAM:NewFolder")?

    // Get file information
    match file_info("SYS:C/Dir") {
        Result::Ok(info) => {
            if info.is_file {
                write("Size: %ld bytes\n", info.size)
            } else if info.is_directory {
                println("It's a directory")
            }
        },
        Result::Err(_) => println("Cannot get info"),
    }

    return Result::Ok(())
}

```

The FileInfo structure:

```

pub struct FileInfo {
    is_file: bool,      // True if regular file
    is_directory: bool, // True if directory
    size: i32,          // File size in bytes
    protection: i32,    // Protection bits (FIBF_*)
}

```

17.7 Command-Line Arguments

AmigaOS uses `ReadArgs()` for parsing command-line arguments with a template syntax. The `std::args` module provides a type-safe wrapper.

17.7.1 Basic Argument Parsing

```

from std::args import Args

pub fn main() -> i32 {
    match Args::parse("FROM/A,TO/A,VERBOSE/S") {
        Result::Ok(args) => {
            // Get required string arguments
            let from = args.get_str("FROM") or { return 10 }
            let to = args.get_str("TO") or { return 10 }

            // Get optional switch
            let verbose = args.get_switch("VERBOSE")

            if verbose {
                write("Copying %s to %s\n", from.as_cstr(),
                    to.as_cstr())
            }
        }
    }
}

```

```

    }

    // ... do work ...

    return 0
},
Result::Err(e) => {
    write("Error: %s\n", e.message())
    return 10
}
}
}

```

17.7.2 Template Syntax

Modifier	Meaning	Example
(none)	Optional positional	FILE
/A	Required argument	FILE/A
/S	Switch (boolean flag)	VERBOSE/S
/K	Keyword (NAME=value)	OUTPUT/K
/N	Numeric argument	COUNT/N
/M	Multiple values	FILES/M
/F	Rest of line	COMMAND/F

Table 17.5: ReadArgs template modifiers

Modifiers can be combined: COUNT/N/A means required number, OUTPUT/K/A means required keyword.

17.7.3 Argument Types

```

from std::args import Args

fn argument_types() {
    match Args::parse("FILE,COUNT/N,VERBOSE/S,FILES/M") {
        Result::Ok(args) => {
            // String argument
            match args.get_str("FILE") {
                Option::Some(f) => write("File: %s\n",
                    f.as_cstr()),
                Option::None => println("No file specified"),
            }

            // Numeric argument
            let count = args.get_number("COUNT") or { 10 } //
                Default 10
            write("Count: %ld\n", count)

            // Switch (boolean)
            if args.get_switch("VERBOSE") {
                println("Verbose mode enabled")
            }

            // Multiple arguments

```

```

        match args.get_multi("FILES") {
            Option::Some(var files) => {
                while var i < files.len {
                    match files.get(i) {
                        Option::Some(f) => write("  %s\n",
                                                f.as_cstr()),
                        Option::None => {},
                    }
                    i++
                }
            },
            Option::None => {},
        }
    },
    Result::Err(_) => {},
}
}
}

```

17.7.4 ArgParser Builder

For programmatic template construction:

```

from std::args import ArgParser, ArgType

fn builder_example() -> Result<(), ArgsError> {
    var parser = ArgParser::new()
    parser.add("FROM", ArgType::Required)
    parser.add("TO", ArgType::Required)
    parser.add("VERBOSE", ArgType::Switch)
    parser.add("BUFFER", ArgType::Number)

    let args = parser.parse()?

    // Retrieve values just like with Args::parse()
    let from = args.get_str("FROM") or { return
        Result::Err(ArgsError::MissingRequired) }
    let to = args.get_str("TO") or { return
        Result::Err(ArgsError::MissingRequired) }
    let verbose = args.get_switch("VERBOSE")
    let buffer_size = args.get_number("BUFFER") or { 4096 } //
        Default

    if verbose {
        write("Copying %s to %s (buffer: %ld)\n", from.as_cstr(),
            to.as_cstr(), buffer_size)
    }

    return Result::Ok(())
}

```

17.7.5 Argument Errors

ArgType	Template Equivalent
Optional	(none)
Required	/A
Switch	/S
Keyword	/K
Number	/N
Multi	/M
RestOfLine	/F
RequiredNumber	/N/A
RequiredKeyword	/K/A

Table 17.6: ArgType enum values

ArgsError	Meaning
Dos(DosError)	Underlying DOS error (wrapped)
InvalidTemplate	Template syntax error
MissingRequired	Required argument not provided
InvalidValue	Argument value is invalid
TooManyArgs	Template has > 32 arguments
AllocationFailed	Memory allocation failed
ArgNotFound	Argument name not in template

Table 17.7: ArgsError enum variants

17.8 Workbench Support

Programs can be launched from either the CLI or Workbench. The two modes require different handling.

17.8.1 Detecting Launch Mode

```
from std::args import is_workbench, Args, get_workbench_startup

pub fn main() -> i32 {
  if is_workbench() {
    // Launched from icon - no stdin/stdout
    return handle_workbench()
  } else {
    // Launched from CLI
    return handle_cli()
  }
}

fn handle_cli() -> i32 {
  match Args::parse("FILE/A") {
    Result::Ok(args) => {
      let file = args.get_str("FILE") or { return 10 }
      // Process file...
      return 0
    },
    Result::Err(_) => return 10,
  }
}
```



```

fn handle_workbench() -> i32 {
    match get_workbench_startup() {
        Option::Some(wb) => {
            // Process Workbench arguments
            if wb.num_args() > 1 {
                // User shift-clicked files
                let arg = wb.arg(1) or { return 20 }
                // Process first selected file...
            }
            return 0
        },
        Option::None => return 20,
    }
}

```

17.8.2 Workbench Startup Message

When launched from Workbench, the program receives a `WBStartup` message containing information about the icon clicked and any files selected:

Note

Message reply is automatic. The Novus runtime startup code retrieves the `WBStartup` message and handles replying to it when the program exits. You do **not** need to manually reply to the message—the runtime handles this for you.

```

from std::args import get_workbench_startup, WorkbenchStartup

fn workbench_example() {
    let wb = get_workbench_startup() or { return }

    // Number of arguments (index 0 = program icon)
    let count = wb.num_args()
    write("Got %ld arguments\n", (i32)count)

    // Get program's own icon (always index 0)
    match wb.arg(0) {
        Option::Some(prog) => {
            match prog.name() {
                Option::Some(name) => {
                    write("Program: %s\n", name.as_cstr())
                },
                Option::None => {},
            }
        },
        Option::None => {},
    }

    // Process any selected files (indices 1..n)
    while var i: u32 = 1 {
        if i >= count { break }

        match wb.arg(i) {
            Option::Some(arg) => {
                match arg.full_path() {
                    Result::Ok(path) => {

```

```

        write("File: %s\n", path.as_cstr())
    },
    Result::Err(_) => {},
}
},
Option::None => {},
}
i++
}
}
}

```

17.8.3 WBAArgument Structure

Each Workbench argument provides:

```

pub struct WBAArgument {
    lock: i32,      // Lock on containing directory
    name_ptr: *u8,  // Filename (not full path)
}

impl WBAArgument {
    /// Get directory lock (for CurrentDir)
    pub fn lock(&self) -> i32

    /// Get filename only
    pub fn name(&self) -> Option<Str>

    /// Get full path (returns String, auto-freed on drop)
    pub fn full_path(&self) -> Result<String, DosError>
}

```

Tip

The `lock` field is a DOS lock on the *directory* containing the file, not the file itself. Use it with `CurrentDir()` to access the file by name, or call `full_path()` to get the complete path. The `String` returned by `full_path()` is automatically freed when it goes out of scope.

17.8.4 Tool Window

If the program's icon has a `TOOLWINDOW` tooltype, you can get the window specification:

```

fn check_tool_window(wb: &WorkbenchStartup) {
    match wb.tool_window() {
        Option::Some(window) => {
            // window is something like "CON:0/0/640/200/Output"
            write("Tool window: %s\n", window.as_cstr())
        },
        Option::None => {
            // No TOOLWINDOW specified
        },
    }
}
}

```

17.9 Error Handling

I/O operations return `Result` types for proper error handling:

```
from std::os::dos import OwnedFileHandle, DosError
from std::ffi::dos import IoErr

fn robust_io() -> Result<(), DosError> {
    // The ? operator propagates errors
    let file = OwnedFileHandle::open("RAM:data.txt",
        MODE_OLDFILE)?

    var buffer: [u8; 1024] = [0u8; 1024]
    let result = file.read_all((*u8)&buffer[0], 1024)?

    return Result::Ok(())
}

fn detailed_error_handling() {
    match OwnedFileHandle::open("NONEXISTENT:", MODE_OLDFILE) {
        Result::Ok(file) => {
            // Use file...
        },
        Result::Err(e) => {
            write("Error: %s\n", e.message())

            // Get detailed DOS error code
            let code = IoErr()
            write("DOS error code: %ld\n", code)
        },
    }
}
```

17.9.1 Common DOS Errors

Error Code	Constant	Meaning
103	ERROR_NO_FREE_STORE	Out of memory
205	ERROR_OBJECT_NOT_FOUND	File not found
212	ERROR_OBJECT_IN_USE	File locked by another
214	ERROR_DISK_FULL	No space on disk
221	ERROR_DISK_WRITE_PROTECTED	Write-protected
225	ERROR_SEEK_ERROR	Invalid seek position
232	ERROR_NO_MORE_ENTRIES	Directory scan complete

Table 17.8: Common DOS error codes

17.10 I/O Patterns

17.10.1 Pattern 1: Read Entire File

```
from std::os::dos import OwnedFileHandle, file_info
from std::memory::block import MemoryBlock
```

```
from std::ffi::amiga_consts import MODE_OLDFILE, MEMF_ANY,
    OFFSET_BEGINNING

fn read_entire_file(path: Str) -> Result<MemoryBlock, DosError> {
    // Get file size
    let info = file_info(path)?
    let size = (u32)info.size

    // Allocate buffer
    let buffer = MemoryBlock::try_alloc(size, MEMF_ANY)?

    // Read file
    let file = OwnedFileHandle::open(path, MODE_OLDFILE)?
    let read = file.read_all(buffer.ptr(), size)?

    return Result::Ok(buffer)
}
```

17.10.2 Pattern 2: Write with Backup

```
from std::os::dos import OwnedFileHandle, file_exists, delete_file
from std::ffi::amiga_consts import MODE_NEWFILE

fn write_with_backup(path: Str, data: *u8, len: u32) ->
    Result<(), DosError> {
    // Create backup filename
    var backup_path = String::from(path)
    backup_path.push_str(".bak")

    // Delete old backup
    if file_exists(backup_path.as_str()) {
        let _ = delete_file(backup_path.as_str())
    }

    // Rename current to backup (if exists)
    if file_exists(path) {
        // Use Rename() from FFI...
    }

    // Write new file
    let file = OwnedFileHandle::open(path, MODE_NEWFILE)?
    file.write_all(data, len)?

    return Result::Ok(())
}
```

17.10.3 Pattern 3: Process File Line by Line

```
fn process_lines(path: Str, handler: fn(*u8, u32) -> bool) ->
    Result<u32, DosError> {
    let file = OwnedFileHandle::open(path, MODE_OLDFILE)?

    var buffer: [u8; 1024] = [0u8; 1024]
    var line_count: u32 = 0
```

```

loop {
    match file.read_line((*u8)&buffer[0], 1024) {
        Result::Ok(len) => {
            if len == 0 { break } // EOF

            // Call handler, stop if it returns false
            if !handler((*u8)&buffer[0], len) {
                break
            }
            line_count++
        },
        Result::Err(_) => break,
    }
}

return Result::Ok(line_count)
}

```

17.11 I/O Best Practices

1. **Always use RAII wrappers**—`OwnedFileHandle` and `DirLockHandle` prevent leaks
2. **Check for Workbench vs CLI**—handle both launch modes
3. **Use %ld for integers**—AmigaOS expects 32-bit values
4. **Handle errors explicitly**—use `Result` and the `?` operator
5. **Buffer reads and writes**—minimize system calls for performance
6. **Close files promptly**—don't hold files open longer than necessary
7. **Check disk space**—verify writes succeeded (`ERROR_DISK_FULL`)

Warning

DOS has limited file handles! Each process can have at most 20 files open simultaneously. If you open many files, close them promptly or batch your operations.

Appendix A

Guide for C Programmers

This appendix consolidates the key differences between C and Novus for experienced C programmers transitioning to the language. If you've written AmigaOS code in C, this chapter will help you map familiar concepts to their Novus equivalents.

A.1 Philosophy Differences

Novus and C share a systems programming heritage, but differ in philosophy:

Aspect	C	Novus
Safety	Trust the programmer	Verify at compile time
Defaults	Mutable, public	Immutable, private
Errors	Return codes, <code>errno</code>	<code>Result<T, E></code> types
Memory	Manual <code>malloc/free</code>	RAII with <code>Drop</code> trait
Null	Implicit in all pointers	Explicit <code>*T</code> vs non-null <code>&T</code>

A.2 Variable Declarations

C declares type first, Novus declares intent first:

```
// C
int x = 10;           // Mutable by default
const int Y = 20;     // Immutable
static int z = 30;    // File-local
```

```
// ... (Novus equivalent - illustrative)
var x = 10            // Mutable (explicit)
let y = 20            // Immutable (explicit)
const Z: i32 = 30     // Compile-time constant
```

Key differences:

- Type inference: `var x = 10` infers `i32`
- No semicolons required (optional)
- `let` bindings cannot be reassigned
- `const` must be computable at compile time

A.3 Type System

A.3.1 Primitive Types

C	Novus	C	Novus
char	i8 or u8	unsigned char	u8
short	i16	unsigned short	u16
int / long	i32	unsigned int	u32
long long	i64	unsigned long long	u64
float	f32	double	f64
void*	*u8	NULL	null

Novus also provides fixed-point types not available in standard C:

- `fixed16` — 8.8 fixed-point (useful for Amiga graphics)
- `fixed32` — 16.16 fixed-point

A.3.2 Booleans

```
// C: integers are implicitly boolean
if (ptr) { ... }           // Non-zero = true
if (count) { ... }        // Non-zero = true
int flag = 1;              // "true"
```

```
// Novus: explicit bool type, no implicit conversion
if ptr != null { ... }    // Must compare explicitly
if count > 0 { ... }      // Must compare explicitly
let flag: bool = true     // Dedicated bool type
```

A.3.3 Structs

```
// C: typedef dance required
typedef struct {
    int x;
    int y;
} Point;

Point p = {10, 20};
Point *ptr = &p;
int x = ptr->x;           // Arrow for pointer access
```

```
struct Point {
    x: i32,
    y: i32,
}
```

Usage:

```
// ... (with Point defined above)
let p = Point { x: 10, y: 20 }
let ptr = &p
let x = ptr.x             // Dot works for everything
```


Key differences:

- No `typedef` needed—struct name is the type
- Field syntax is `name: type` (reversed from C)
- No `->` operator—dot (`.`) auto-dereferences
- Struct literals use `Type { field: value }` syntax

A.4 Pointers and References

C has one pointer type; Novus distinguishes three:

Novus	C Equivalent	Properties
<code>&T</code>	<code>const T*</code>	Non-null, immutable
<code>&var T</code>	<code>T*</code>	Non-null, mutable
<code>*T</code>	<code>T*</code>	Nullable, requires <code>unsafe</code>

```
// C: All pointers can be null
void process(Player *p);           // Might modify, might be null
void print_info(const Player *p);  // Won't modify, might be null
```

```
// Novus: Intent is explicit
fn process(p: &var Player)         // Will modify, never null
fn print_info(p: &Player)          // Won't modify, never null
fn maybe_process(p: *Player)       // Might be null (use in unsafe)
```

The call site makes mutability visible:

```
process(&var player)    // Clearly shows modification
print_info(&player)     // Clearly shows read-only
```

A.5 Functions

A.5.1 Declaration Syntax

```
// C: return type first
int add(int a, int b) {
    return a + b;
}

static int helper(void) { return 42; } // File-local
```

```
// Novus: fn keyword, return type last
fn add(a: i32, b: i32) -> i32 {
    return a + b
}

fn helper() -> i32 { return 42 } // Private by default
pub fn api() -> i32 { return 0 } // Explicitly public
```

A.5.2 Methods

C doesn't have methods; Novus does:

```
// C: Pass struct explicitly
int rect_area(const Rectangle *r) {
    return r->width * r->height;
}

void rect_scale(Rectangle *r, int factor) {
    r->width *= factor;
    r->height *= factor;
}

// Usage
int a = rect_area(&rect);
rect_scale(&rect, 2);
```

```
// ... (with Rectangle defined elsewhere - illustrative)
impl Rectangle {
    fn area(&self) -> i32 {
        return self.width * self.height
    }
    fn scale(&var self, factor: i32) {
        self.width = self.width * factor
        self.height = self.height * factor
    }
}
```

Usage:

```
// ... (with Rectangle defined above)
let a = rect.area()
rect.scale(2)
```

A.6 Control Flow

A.6.1 If Statements

```
// C: Parentheses required, braces optional
if (x > 0)
    return 1;
else if (x < 0)
    return -1;
else
    return 0;
```

```
// Novus: No parentheses, braces required
if x > 0 {
    return 1
} else if x < 0 {
    return -1
} else {
    return 0
}
```

```
}
```

A.6.2 Switch vs Match

```
// C: switch with fall-through, break required
switch (cmd) {
    case CMD_READ:
        handle_read();
        break;           // Easy to forget!
    case CMD_WRITE:
    case CMD_UPDATE:     // Fall-through
        handle_write();
        break;
    default:
        handle_unknown();
}
```

```
// Novus: match with no fall-through
match cmd {
    CMD_READ => handle_read(),
    CMD_WRITE | CMD_UPDATE => handle_write(), // Explicit OR
    _ => handle_unknown(),
}
```

Key differences:

- No fall-through—each arm is independent
- No `break` needed
- Use `|` for multiple patterns
- `_` is the wildcard (like `default`)
- `match` can be an expression returning a value

A.6.3 Loops

```
// C: Traditional for loop
for (int i = 0; i < 10; i++) {
    process(i);
}
```

```
// Novus: Range-based for (preferred)
for i in 0..10 {
    process(i)
}

// C-style also supported
for var i = 0; i < 10; i++ {
    process(i)
}
```

A.7 Memory Management

A.7.1 The goto cleanup Pattern

```
// C: Manual cleanup with goto
int do_work(void) {
    void *a = NULL, *b = NULL;
    int result = -1;

    a = AllocMem(100, MEMF_ANY);
    if (!a) goto cleanup;

    b = AllocMem(200, MEMF_ANY);
    if (!b) goto cleanup;

    // ... actual work ...
    result = 0;

cleanup:
    if (b) FreeMem(b, 200); // Must remember size!
    if (a) FreeMem(a, 100);
    return result;
}
```

```
// Novus: RAII handles cleanup automatically
fn do_work() -> Result<(), ExecError> {
    var a = MemoryBlock::try_alloc(100, MEMF_ANY)?
    var b = MemoryBlock::try_alloc(200, MEMF_ANY)?

    // ... actual work ...

    return Result::Ok(())
} // Both freed automatically, correct order, size tracked
```

A.7.2 RAII vs Manual Management

Aspect	C	Novus
Allocation	AllocMem(size, flags)	MemoryBlock::alloc(size, flags)
Deallocation	FreeMem(ptr, size)	Automatic via Drop
Size tracking	Manual (error-prone)	Automatic
Cleanup order	Manual (LIFO required)	Automatic (guaranteed LIFO)
Early return	Requires goto	Just return or ?

A.8 Error Handling

```
// C: Return codes and errno
BPTR file = Open("data.txt", MODE_OLDFILE);
if (!file) {
    LONG err = IoErr();
    printf("Error: %ld\n", err);
    return -1;
}
```

```

}
// ... use file ...
Close(file);

```

```

// Novus: Result type with ? operator
fn read_data() -> Result<String, DosError> {
    let file = open("data.txt", MODE_OLDFILE)? // Returns on
        error
    let data = read_all(&file)?
    return Result::Ok(data)
} // file closed automatically

```

The ? operator:

1. Checks if result is `Err`
2. If error, returns immediately with that error
3. If success, unwraps the value and continues

A.9 Headers and Modules

```

// C: Textual inclusion
#ifndef MYHEADER_H
#define MYHEADER_H

#include <exec/types.h>
#include <proto/exec.h> // Order can matter!

extern int global_count;
void my_function(void);

#endif

```

```

// Novus: Semantic imports
from std::ffi::exec import AllocMem, FreeMem
from std::ffi::amiga_consts import MEMF_FAST

pub fn my_function() {
    // ...
}

```

Key differences:

- No include guards needed
- Import order doesn't matter
- Only imported names are in scope
- No separate header files
- `pub` explicitly exports (vs C's implicit visibility)

A.10 Visibility

C	Novus
<code>static</code> (file-local)	(default, no modifier)
Non-static + header	<code>pub</code>
<code>extern</code> declaration	<code>extern</code> (for FFI)

Note: Novus inverts C's default. In C, symbols are public unless marked `static`. In Novus, symbols are private unless marked `pub`.

A.11 Common Intrinsics

C	Novus
<code>sizeof(T)</code>	<code>@sizeof(T)</code>
<code>_Alignof(T)</code>	<code>alignof(T)</code> (in inline asm only)
<code>memset(&x, 0, sizeof(x))</code>	<code>@zeroed(T)</code>

A.12 AmigaOS NDK Access

The entire AmigaOS NDK is available from Novus. Every function, structure, and constant from the 3.9 NDK has been imported and can be called directly. You don't need wrappers—you can use the raw API just like you would in C.

A.12.1 FFI Modules

The NDK bindings are organized by library in `std::ffi`:

```
// ... (FFI import examples - illustrative)
from std::ffi::exec import AllocMem, FreeMem, FindTask, Wait,
    Signal
from std::ffi::dos import Open, Close, Read, Write, IoErr, Lock,
    UnLock
from std::ffi::graphics import BltBitMap, SetAPen, RectFill, Text
from std::ffi::intuition import OpenWindow, CloseWindow,
    RefreshGList
from std::ffi::gadtools import CreateMenus, FreeMenus, LayoutMenus
from std::ffi::layers import CreateUpfrontLayer, DeleteLayer
from std::ffi::utility import AllocateTagItems, FreeTagItems
```

Structures and constants:

```
// ... (structure and constant imports - illustrative)
from std::ffi::amiga_structs import Window, Screen, BitMap
from std::ffi::amiga_consts import MEMF_CHIP, MEMF_FAST,
    MEMF_CLEAR
```

A.12.2 Calling Raw NDK Functions

Raw AmigaOS calls require `unsafe` blocks because they bypass Novus's safety guarantees:

```
from std::ffi::exec import AllocMem, FreeMem
from std::ffi::amiga_consts import MEMF_CHIP, MEMF_CLEAR
```

```

fn allocate_chip_buffer(size: u32) -> *u8 {
    unsafe {
        return AllocMem(size, MEMF_CHIP | MEMF_CLEAR)
    }
}

fn free_chip_buffer(ptr: *u8, size: u32) {
    if ptr != null {
        unsafe {
            FreeMem(ptr, size)
        }
    }
}

```

This is identical to what you’d write in C, just with explicit `unsafe` markers.

A.12.3 When to Use Raw vs Wrappers

Use	Raw NDK	Stdlib Wrappers
When	Porting C code, one-off calls, no wrapper exists	New code, repeated patterns
Safety	Manual (like C)	Automatic (RAII, Result)
Cleanup	Manual <code>Free*</code> calls	Automatic via <code>Drop</code>
Errors	Check return values	<code>Result<T, E></code>

Both approaches are valid. The wrappers exist for convenience and safety, but they’re built on the same raw FFI you have access to. If a wrapper doesn’t exist for what you need, use the raw API directly.

A.12.4 Example: Raw Intuition Window

Here’s opening a window the “C way” in Novus:

```

from std::ffi::intuition import OpenWindowTags, CloseWindow
from std::ffi::amiga_consts import WA_Width, WA_Height, WA_Title,
    WA_IDCMP, WA_Flags, WFLG_CLOSEGADGET, WFLG_DRAGBAR,
    IDCMP_CLOSEWINDOW, TAG_DONE

fn main() -> i32 {
    var window: *Window = null

    unsafe {
        window = OpenWindowTags(null,
            WA_Width, 640,
            WA_Height, 480,
            WA_Title, "Raw NDK Window",
            WA_Flags, WFLG_CLOSEGADGET | WFLG_DRAGBAR,
            WA_IDCMP, IDCMP_CLOSEWINDOW,
            TAG_DONE)
    }

    if window == null {
        return 1
    }

    // ... use window ...

```

```
unsafe {  
    CloseWindow(window)  
}  
  
return 0  
}
```

This is nearly identical to C code—the only differences are:

- `unsafe {}` blocks around library calls
- `null` instead of `NULL`
- Novus variable declaration syntax

A.12.5 BCPL Pointers (DOS Library)

DOS library functions use BCPL pointers (BPTR/BSTR), which are `i32` values, not actual pointers:

```
from std::ffi::dos import Open, Close, Read, Write  
from std::ffi::amiga_consts import MODE_OLDFILE  
  
fn read_file_raw(path: *u8) -> i32 {  
    // DOS file handle is i32 (BPTR), not *File  
    var file: i32 = 0  
  
    unsafe {  
        file = Open(path, MODE_OLDFILE)  
    }  
  
    if file == 0 {  
        return -1  
    }  
  
    var buffer = [0u8; 1024]  
    var bytes_read: i32 = 0  
  
    unsafe {  
        bytes_read = Read(file, &buffer as *u8, 1024)  
        Close(file)  
    }  
  
    return bytes_read  
}
```

A.12.6 Comparison: Raw vs Wrapper

For comparison, here's the same operation using the `stdlib` wrapper:

```
// C: Direct library calls  
struct Window *win = OpenWindowTags(NULL,  
    WA_Width, 640,  
    WA_Height, 480,  
    WA_Title, "My Window",  
    TAG_DONE);  
if (win) {
```



```
// ... use window ...
CloseWindow(win);
}
```

```
// Novus: RAII wrappers preferred
match WindowBuilder::new()
    .size(640, 480)
    .title("My Window")
    .build()
{
    Result::Ok(window) => {
        // ... use window ...
    }, // Closed automatically
    Result::Err(e) => {
        println(e.message())
    }
}
```

A.12.7 Hardware Access

Both languages support direct hardware access, but Novus requires `unsafe`:

```
// C: Direct hardware poke
volatile UWORD *custom = (UWORD *)0xdff000;
custom[0x96/2] = 0x8020; // DMACON
```

```
// ... (Novus unsafe block - illustrative)
unsafe {
    let custom = 0xdff000 as *u16
    *(custom + 0x96 / 2) = 0x8020
}
```

A.13 Quick Reference Table

C	Novus
<code>int x = 10;</code>	<code>var x = 10</code>
<code>const int X = 10;</code>	<code>const X: i32 = 10</code>
<code>int arr[3] = {1,2,3};</code>	<code>let arr = [1, 2, 3]</code>
<code>if (ptr)</code>	<code>if ptr != null</code>
<code>switch/case/break</code>	<code>match { pat => }</code>
<code>for(i=0;i<n;i++)</code>	<code>for i in 0..n</code>
<code>ptr->field</code>	<code>ptr.field</code>
<code>struct S {...};</code>	<code>struct S {...}</code>
<code>typedef struct</code>	(not needed)
<code>malloc/free</code>	MemoryBlock, Vec
<code>sizeof(T)</code>	<code>@sizeof(T)</code>
<code>NULL</code>	<code>null</code>
<code>errno</code>	<code>Result::Err(...)</code>
<code>static</code>	(default)
<code>#include</code>	<code>from ... import</code>

A.14 Migration Tips

1. **Start with `var`** — You can always tighten to `let` later
2. **Use `stdlib` wrappers** — `MemoryBlock`, `Vec`, `WindowHandle` instead of raw FFI
3. **Embrace `Result`** — The `?` operator makes error handling concise
4. **Trust the `Drop`** — Let RAII handle cleanup instead of manual `goto cleanup`
5. **Check explicit `null`** — `if ptr != null` instead of `if (ptr)`
6. **Use `match`** — More powerful and safer than `switch`
7. **Read the sidebars** — “Coming from C” boxes throughout the guide explain specific differences

A.15 What Stays the Same

Not everything is different! These concepts transfer directly:

- Structs are laid out the same way (C ABI compatible)
- Pointer arithmetic works the same
- Bitwise operations are identical
- AmigaOS structures and constants are the same
- The calling convention matches (`d0/d1/a0/a1` for args)
- You can still write inline assembly when needed

Novus is designed to feel familiar to C programmers while adding safety. Your AmigaOS knowledge transfers directly—you’re just expressing it in a slightly different syntax with more compile-time guarantees.

Appendix B

Quick Reference

This appendix provides a compact reference for Novus syntax and common patterns.

B.1 Variables and Constants

```
// ... (variable examples - illustrative)
var x = 10
x = 20

// Immutable binding
let y = 42

// Compile-time constant
pub const MAX_SIZE: u32 = 1024

// Type annotation (when needed)
let z: i32 = 100
var buffer = [0u8; 256] // Size inferred
```

B.2 Primitive Types

Type	Size	Range/Description
i8	8-bit	−128 to 127
i16	16-bit	−32,768 to 32,767
i32	32-bit	±2 billion (default)
i64	64-bit	Very large signed
u8	8-bit	0 to 255
u16	16-bit	0 to 65,535
u32	32-bit	0 to 4 billion
u64	64-bit	Very large unsigned
f32	32-bit	Single-precision float
f64	64-bit	Double-precision float
fixed16	16-bit	8.8 fixed-point
fixed32	32-bit	16.16 fixed-point
bool	8-bit	true or false
*T	32-bit	Raw pointer (nullable)
&T	32-bit	Reference (non-null)
&var T	32-bit	Mutable reference

B.3 Literals

```
// ... (literal examples - illustrative)
42          // i32 (default)
42u32       // u32
42i8        // i8
0xFF        // Hexadecimal (C-style)
$FF         // Hexadecimal (Amiga-style)
%1010       // Binary
1_000_000   // Underscores for readability

// Floats
3.14        // f64 (default)
3.14f32     // f32

// Strings
"Hello"      // String literal
"Line1\nLine2" // With escape
f"Value: {x}" // F-string (interpolation)

// Characters
'A'          // u8 (65)
'\n'         // Newline
'\x1B'       // Hex byte (ESC)

// Other
true, false  // Booleans
null         // Null pointer
```

B.4 Operators

Op	Meaning	Op	Meaning
+	Add	&&	Logical AND
-	Subtract		Logical OR
*	Multiply	!	Logical NOT
/	Divide	&	Bitwise AND
%	Modulo		Bitwise OR
==	Equal	^	Bitwise XOR
!=	Not equal	~	Bitwise NOT
<	Less than	«	Left shift
>	Greater than	»	Right shift
<=	Less or equal	..	Range (exclusive)
>=	Greater or equal	..=	Range (inclusive)

Compound: +=, -=, *=, /=, %=, &=, |=, ^=, <=, >=

Increment/Decrement: ++x, x++, -x, x-

B.5 Control Flow

```
// If-else
if condition {
    // ...
} else if other {
    // ...
}
```

```

} else {
    // ...
}

// If as expression
let max = if a > b { a } else { b }

// Match
match value {
    1 => handle_one(),
    2 | 3 => handle_two_or_three(),
    4..10 => handle_range(),
    _ => handle_default(),
}

// While loop
while condition {
    // ...
}

// For loop
for i in 0..10 {
    // i goes 0 to 9
}

for item in collection {
    // iterate items
}

// Loop (infinite)
loop {
    if done { break }
}

// Break and continue
break           // Exit loop
continue        // Next iteration
break value     // Return value from loop

```

B.6 Functions

```

// Basic function
fn add(a: i32, b: i32) -> i32 {
    return a + b
}

// No return value
fn print_hello() {
    println("Hello")
}

// Public function
pub fn api_function() -> i32 {
    return 42
}

```

```
// Generic function
fn swap<T>(a: &var T, b: &var T) {
    let tmp = *a
    *a = *b
    *b = tmp
}

// With trait bounds
fn print_all<T>(items: &[T]) where T: Display {
    for item in items {
        println(item.to_string())
    }
}
```

B.7 Structs and Enums

```
// Struct
struct Point {
    x: i32,
    y: i32,
}

// Create instance
let p = Point { x: 10, y: 20 }

// Impl block
impl Point {
    pub fn new(x: i32, y: i32) -> Point {
        return Point { x: x, y: y }
    }

    pub fn distance(&self, other: &Point) -> f32 {
        // ...
    }
}

// Enum
enum Direction {
    Up,
    Down,
    Left,
    Right,
}

// Enum with data
enum Message {
    Quit,
    Move(i32, i32),
    Write(String),
}

// Match enum
match msg {
    Message::Quit => exit(),
```

```

    Message::Move(x, y) => move_to(x, y),
    Message::Write(s) => println(s),
}

```

B.8 Option and Result

```

// ... (Option/Result patterns - illustrative)
let maybe: Option<i32> = Option::Some(42)
let none: Option<i32> = Option::None

match maybe {
    Option::Some(x) => use_value(x),
    Option::None => handle_missing(),
}

let value = maybe.unwrap_or(0)

// Result
let result: Result<i32, Error> = Result::Ok(42)
let error: Result<i32, Error> = Result::Err(Error::NotFound)

match result {
    Result::Ok(x) => use_value(x),
    Result::Err(e) => handle_error(e),
}

// ? operator (propagate error)
fn process() -> Result<i32, Error> {
    let file = open(path)? // Returns Err early
    let data = read(&file)?
    Result::Ok(parse(data))
}

```

B.9 Memory and References

```

// Immutable reference
let r: &i32 = &x

// Mutable reference
let r: &var i32 = &var x
*r = 10 // Modify through reference

// Raw pointer
let p: *u8 = (*u8)address

// Dereference
let value = *r
let byte = *p

// Null check
if ptr != null {
    // safe to use
}

```

```
}

// defer (cleanup)
fn process() {
    let resource = acquire()
    defer { release(resource) }
    // ... use resource
} // release() called here

// Drop trait (RAII)
impl Drop for MyHandle {
    fn drop(&var self) {
        // Automatic cleanup
    }
}
```

B.10 RAII Patterns

Pattern	Use
Handle	Wrap OS resources (WindowHandle, FileHandle)
Guard	Temporary locks (SemaphoreGuard, BlitterGuard)
Builder	Accumulate config (MUIAppBuilder)

```
// Handle: wrap resource, auto-cleanup
var window = WindowBuilder::new().build()?
// window closed automatically on drop

// Guard: temporary exclusive access
var lock = semaphore.obtain()
// ReleaseSemaphore() called on drop

// Builder: configure then create
var app = MUIAppBuilder::new()
    .title("App")
    .build()?
```

B.11 Arrays and Collections

```
// Fixed array (size inferred)
let arr = [1, 2, 3]
let zeros = [0; 100] // 100 zeros
let first = arr[0u32]

// Vec (dynamic array)
var v: Vec<i32> = Vec::new()
v.push(1)
v.push(2)
let len = v.len()
let item = v.get(0u32) // Option<T>
```



```
// Tuple
let point: (i32, i32) = (10, 20)
let (x, y) = point // Destructure
```

B.12 Modules and Imports

```
// ... (import examples - illustrative)
from std::collections::vec import Vec
from std::io::file import println, read_file

// Multiple imports
from std::core import Option, Result, Drop

// Use in code
let v: Vec<i32> = Vec::new()
```

B.13 Traits

```
trait Printable {
    fn print(&self)
}
```

Implementation:

```
// ... (with Point and Printable defined above)
impl Printable for Point {
    fn print(&self) {
        println(f"({self.x}, {self.y})")
    }
}
```

Common traits: Drop (destructor), Clone, Eq, Hash.

B.14 Common Attributes

```
@test                // Test function
@test(skip = "reason") // Skipped test
@benchmark           // Benchmark function
@inline              // Hint: inline this function
@noinline            // Prevent inlining
@packed              // No struct padding
@export              // Keep function, don't mangle name
@atomic              // Wrap in Forbid/Permit
#[derive(Eq, Hash)]   // Auto-implement traits
#[stack_size(65536)] // Set program stack size
```

B.15 Ininsics

```
@sizeof(T)      // Size of type in bytes
@zeroed(T)      // Zero-initialized value
```

Note: `alignof()` is only available in inline assembly blocks.

B.16 Compiler Commands

```
# Compile single file
novus compile file.novus -o output

# Build project
novus build

# Run tests
novus test file.novus

# Format code
novus fmt

# With options
novus compile file.novus -o out --release --cpu 68040
```

B.17 C to Novus Quick Reference

C	Novus
<code>int x = 10;</code>	<code>var x = 10</code>
<code>const int X = 10;</code>	<code>const X: i32 = 10</code>
<code>if (ptr)</code>	<code>if ptr != null</code>
<code>switch/case/break</code>	<code>match { pat => }</code>
<code>for(i=0;i<n;i++)</code>	<code>for i in 0..n</code>
<code>NULL</code>	<code>null</code>
<code>struct S {...};</code>	<code>struct S {...}</code>
<code>typedef struct</code>	(not needed)
<code>ptr->field</code>	<code>ptr.field</code> (auto-deref)
<code>malloc/free</code>	Use <code>Vec</code> , <code>MemoryBlock</code> , or <code>allocator</code>
<code>sizeof(T)</code>	<code>@sizeof(T)</code>
<code>errno</code>	<code>Result::Err(DosError::...)</code>

B.18 Escape Sequences

Sequence	Meaning
<code>\n</code>	Newline
<code>\r</code>	Carriage return
<code>\t</code>	Tab
<code>\0</code>	Null byte
<code>\\</code>	Backslash
<code>\"</code>	Double quote
<code>\'</code>	Single quote
<code>\xNN</code>	Hex byte